

Discussion 9 (week 10)

Concurrent Programming and Final Review

Brandon Wu

Reminders

- Final is Tuesday 12/16 8:00AM – 11:00AM
- Super office hour Monday 10AM-4PM in 4670 Boelter Hall
- Course Evaluations – please evaluate me!

Concurrent Processes

- Process is an abstraction of a running program
- A process **context** describes its state
 - Registers
 - Virtual address space
- Fork spawns a child process
 - Duplicates the context

Processes are Expensive!

- OS overhead in creating a process
- Sharing state information is tough
 - Process have separate Virtual Address space
 - Must communicate via OS IPC (interprocess communication)
- On the other hand, processes are hard to screw up, cannot accidentally step on memory of another process

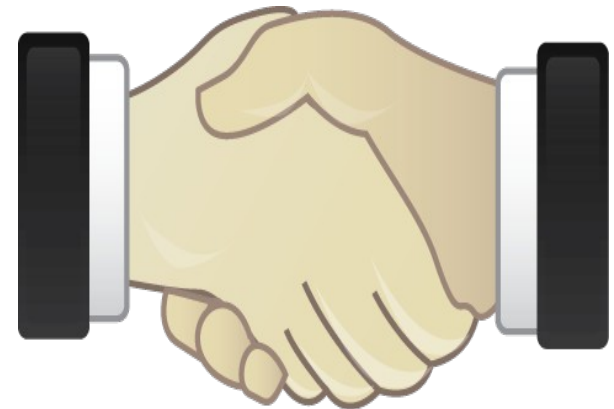
Threads

- **Threads** are logical flow that share context of program
- Has its own “agenda”
 - Can execute independently of its peers
 - Has private set of registers and stack
- Shares process context with all other threads within process
 - Share the same addr space



Thread context and Sharing

- Threads have their own context
 - Stack
 - General purpose registers
- Can share global variables with other threads within the process



OS does not guarantee order of execution across threads!

OS also does not guarantee exclusion


```
volatile int count = 0;
int main() {
    ...
    pthread_create(&tid1, NULL, thread, &arg);
    pthread_create(&tid2, NULL, thread, &arg);
    ...
    return 0;
}

void* thread(void *arg) {
    num = *(int* ) arg);
    for (int i=0; i<num; i++)
        count++;
}
```

Thread concurrency

- After tid1 and tid2 terminate, we expect $\text{count} = 2 * \text{arg}$ but it doesn't
- Threads can execute concurrently, even in parallel
- We don't guarantee exclusion of shared resources
- Unguarded modification of shared data results in **race conditions**

Mutual Exclusion: Semaphores

- A global variable used only to control access to other global variables
- Two operations on a semaphore s :
 - **Down**: $P(s)$ return $(s > 0) ? -- s : \text{RETRY}$
 - **Up**: $V(s)$ performs $++s$
- In C library:

```
int sem_init(sem_t* sem, 0, unsigned val);
```

```
int sem_wait(sem_t* s); /* P(s) */
```

```
int sem_post(sem_t* s); /* V(s) */
```

Up and Down operations must be atomic.
WHY?

Semaphores

- Semaphores with value = 1 are **Locks**
 - $s == 1 \rightarrow$ unlocked
 - $s == 0 \rightarrow$ locked
 - $P(s) \rightarrow$ attempt to acquire lock s
 - $V(s) \rightarrow$ unlock s
-
- Binary semaphores sometimes called **Mutexes**



Fixing the multi-threaded Counting Problem

```
volatile int count = 0;
```

```
sem_t mutex;
```

```
...
```

```
void* thread(void* arg) {
```

```
    sem_wait(&mutex);    /* P(&mutex) */
```

```
    count++;
```

```
    sem_post(&mutex);    /* V(&mutex) */
```

```
}
```

Readers-Writers Problem

- Two types of players: **readers** and **writers**
- Many readers can read the value at the same time
- Modifying (writing) the value requires exclusive access

Readers-Writers Problem

```
/* GLOBAL VARIABLES */
```

```
int readcnt = 0;
```

```
sem_t rc_mutex, wlock; /* initially 1 */
```


The Reader

```
void reader( ) {  
    sem_wait(&rc_mutex); // P(&rc_mutex)  
    readcnt++;  
    if (readcnt == 1)  
        sem_wait(&wlock);           // P (&wlock)  
    sem_post(&rc_mutex); // V(&rc_mutex)  
    // Reading happens...  
    sem_wait(&rc_mutex); //R(&rc_mutex);  
    readcnt --;  
    if (readcnt == 0)  
        sem_post(&wlock);           // V(&wlock);  
    sem_post(&rc_mutex); //V(&rc_mutex);  
}
```

The Reader

- `rc_mutex` protects access to `readcnt`
 - Readers modify this
- `wlock` is the lock required to obtain exclusive access to the data
 - Writer cannot have this while there are active readers

The Writer

```
void writer( ) {  
    sem_wait(&wlock); // Acquire write lock  
    ... do some writing  
    sem_post(&wlock); // Release write lock  
}
```



First Readers-Writers Problem

- This implementation favors readers
- While there is at least one reader, writer must wait
- If readers keep coming, writer never executes
 - This is an example of starvation

CHAPTER 2: DATA REPRESENTATION

Floating Point Conversions

IEEE Floating Point Recap

- Decimal value $x = (-1)^s * 2^{(E-BIAS)} * M$



- Where $BIAS = 2^{(k-1)} - 1$
- Normalized FP:
 - When $E \neq 0$
- Denormalized
 - $E == 0$, exponent = $1 - BIAS$

Floating Point Example

- Consider 9 bit floating point
- 3 bit exponent, 5 bit mantissa
- What is the bias
- What is the smallest possible range between two representable values?
- What is the range of numbers where the difference between any two representable numbers is $> 2^{-5}$

Example from Uen-Tao

Consider the two floating point representations:

A: 0 000 00000 (three exp, five fraction)

B: 0 00000 000 (five exp, three fraction)

What are the closest Format B approximations for the Format A value:

0 000 10101

Floating Point Example

- Format A:
- 0 000 10101
- Bias = 3
- $E = 1 - \text{Bias} = -2$ (Denormalized)
- $V = 2^{-2} * (1/2 + 1/2^3 + 1/2^5) = 2^{-3} * (1 + 1/2^2 + 1/2^4) = .1640625$
- Format B:
- Format B has enough range in the exponent bits to represent the full exponent range of format A.
- Bias = 15
- $E = -3 = e - \text{Bias}$
- $e = 12$
- We want this V to equal the V from format A ($2^{-3} * (1 + 1/2^2 + 1/2^4)$)

Floating Point Example

However, it is impossible to represent this in format B

$$V = 2^{-3} * (1 + 1/22 + 1/24)$$

Format B only has three fractional bits and cannot represent a fractional contribution of $1/24$

Rounding down solution:

$$V = 2^{-3} * (1 + 1/22) = .15625$$

0 01100 010

Rounding up solution:

$$V = 2^{-3} * (1 + 1/22 + 1/23) = .171875$$

0 01100 011

CHAPTER 3: MACHINE CODE

Reverse Engineering and Stack Discipline

Inferring Operation Suffix

- If you can, look to the destination size

`mov_ $0x7, %ebx`

`add_ %ecx, %edx`

`sar_ %al, %edi`

`sub_ 0x63(%eax), %esi`

- When destination is memory, it can be ambiguous

`mov_ $0xfed3, $0xc(%ebp)`

Registers in x86

63	31	15	8	7	0	
%rax	%eax	%ax	%ah	%al		Return value
%rbx	%ebx	%bx	%bh	%bl		Callee saved
%rcx	%ecx	%cx	%ch	%cl		4th argument
%rdx	%edx	%dx	%dh	%dl		3rd argument
%rsi	%esi	%si		%sil		2nd argument
%rdi	%edi	%di		%dil		1st argument
%rbp	%ebp	%bp		%bpl		Callee saved
%rsp	%esp	%sp		%spl		Stack pointer

More Registers

63	31	15	
%r8	%r8d %r8w		%r8b 5th argument
%r9	%r9d %r9w		%r9b 6th argument
%r10	%r10d %r10w		%r10b Caller saved
%r11	%r11d %r11w		%r11b Caller saved
%r12	%r12d %r12w		%r12b Callee saved
%r13	%r13d %r13w		%r13b Callee saved
%r14	%r14d %r14w		%r14b Callee saved
%r15	%r15d %r15w		%r15b Callee saved

x86 Reverse Engineering Problems

- Unless it specifies otherwise, assume function parameters make use of registers (for x86 problems)
 - i.e. first parameter stored in %rdi
 - Second parameter stored in %rsi
 - Similar to %ebp+8 and %ebp+12 convention in IA32

Reverse Engineering: Problem 3.59

```
int switch_prob(int x, int n) {  
    int result = x;  
    switch(n) {  
        /* Fill in the code here */  
    }  
    return result;  
}
```

```
(gdb) x/6w 0x80485d0
```

```
0x80485d0: 0x8048438  0x8048448  0x8048438  0x804843d
```

```
0x80485e0: 0x8048442  0x8048455
```


1 8048420 <switch_prob>:

2 8048420: 55

3 8048421: 89 e5

4 8048423: 8b 45 08

5 8048426: 8b 55 0c

6 8048429: 83 ea 32

7 804842c: 83 fa 05

8 804842f: 77 17

9 8048431: ff 24 95 d0 85 04 08

10 8048438: c1 e0 02

11 804843b: eb 0e

12 804843d: c1 f8 02

13 8048440: eb 09

14 8048442: 8d 04 40

15 8048445: 0f af c0

16 8048448: 83 c0 0a

17 804844b: 5d

18 804844c: c3

push %ebp

mov %esp,%ebp

mov 0x8(%ebp),%eax

mov 0xc(%ebp),%edx

sub \$0x32,%edx

cmp \$0x5,%edx

ja 8048448 <switch_prob+0x28>

jmp *0x80485d0(,%edx,4)

shl \$0x2,%eax

jmp 804844b <switch_prob+0x2b>

sar \$0x2,%eax

jmp 804844b <switch_prob+0x2b>

lea (%eax,%eax,2),%eax

imul %eax,%eax

add \$0xa,%eax

pop %ebp

ret

1 8048420 <switch_prob>:

2 8048420: 55

3 8048421: 89 e5

4 8048423: 8b 45 08

5 8048426: 8b 55 0c

6 8048429: 83 ea 32

7 804842c: 83 fa 05

8 804842f: 77 17

9 8048431: ff 24 95 d0 85 04 08

JTAB[0] JTAB[2] 10 8048438: c1 e0 02

11 804843b: eb 0e

JTAB[3] 12 804843d: c1 f8 02

13 8048440: eb 09

JTAB[4] 14 8048442: 8d 04 40

15 8048445: 0f af c0

JTAB[1] 16 8048448: 83 c0 0a

17 804844b: 5d

18 804844c: c3

push %ebp

mov %esp,%ebp

mov 0x8(%ebp),%eax

mov 0xc(%ebp),%edx

sub \$0x32,%edx

cmp \$0x5,%edx

ja 8048448 <switch_prob+0x28>

jmp *0x80485d0(,%edx,4)

shl \$0x2,%eax

jmp 804844b <switch_prob+0x2b>

sar \$0x2,%eax

jmp 804844b <switch_prob+0x2b>

lea (%eax,%eax,2),%eax

imul %eax,%eax

add \$0xa,%eax

pop %ebp

ret

```
int switch_prob(int x, int n) {  
    int result = x;  
    switch(n) {  
        case 50:  
        case 52:  
            result = x << 2;  
            break;  
        case 53:  
            result = x >> 2;  
            break;  
        case 54:  
            result = 3*x;  
            result *= result;  
        default:  
            result += 10;  
    }  
    return result;  
}
```

Problem 3.58

```
Int switch3(int *p1, int *p2, mode_t action) {  
    int result = 0;  
    switch(action) {  
        case MODE_A:  
        case MODE_B:  
        case MODE_C:  
        case MODE_D:  
        default:  
    }  
    return result;  
}
```

1	.L17:	<i>MODE_E</i>	21	.L15:	<i>MODE_C</i>
2	movl \$17, %edx		22	movl 12(%ebp), %edx	
3	jmp .L19		23	movl \$15, (%edx)	
4	.L13:	<i>MODE_A</i>	24	movl 8(%ebp), %ecx	
5	movl 8(%ebp), %eax		25	movl (%ecx), %edx	
6	movl (%eax), %edx		26	jmp .L19	
7	movl 12(%ebp), %ecx		27	.L16:	<i>MODE_D</i>
8	movl (%ecx), %eax		28	movl 8(%ebp), %edx	
9	movl 8(%ebp), %ecx		29	movl (%edx), %eax	
10	movl %eax, (%ecx)		30	movl 12(%ebp), %ecx	
11	jmp .L19		31	movl %eax, (%ecx)	
12	.L14:	<i>MODE_B</i>	32	movl \$17, %edx	
13	movl 12(%ebp), %edx		33	.L19:	<i>default</i>
14	movl (%edx), %eax		34	movl %edx, %eax	<i>Set return value</i>
15	movl %eax, %edx				
16	movl 8(%ebp), %ecx				
17	addl (%ecx), %edx				
18	movl 12(%ebp), %eax				
19	movl %edx, (%eax)				
20	jmp .L19				

Problem 3.58

```
result = -1;
switch(action) {
case MODE_A:
    result = *p1;
    *p1 = *p2;
    break;
case MODE_B:
    result = *p2;
    *p2 = *p1;
    break;
case MODE_C:
    *p2 = 15;
    result = *p1;
    break;
```

```
case MODE_D:
    *p2 = *p1;
    result = 17;
    break;
case MODE_E:
    result = 17
default;;}
return result;
```

Problem 3.65

```
typedef struct {  
    short x[A][B];  
    int y;  
} str 1;  
  
typedef struct {  
    char array[A];  
    int t;  
    short s[B];  
    int u;  
} str2;
```

```
void setVal(str1 *p, str2 *q) {  
    int v1 = q->t;  
    int v2 = q->u;  
    p->y = v1+v2;  
}
```

Problem 3.65

```
movl 12(%ebp), %eax  
movl 36(%eax), %edx  
addl 12(%eax), %edx  
movl 8(%ebp), %eax  
movl %edx, 92(%eax)
```

What is the Value of A and B?

The solution is unique

Problem 3.69

```
typedef struct ELE *tree_ptr;
struct ELE {
    long val;
    tree_ptr left;
    tree_ptr right;
}
long trace(tree_ptr tp);    // What does this do?
```

trace:

movl \$0, %eax

testq %rdi, %rdi

je .L3

.L5:

movq (%rdi), %rax

movq 16(%rdi), %rdi

testq %rdi, %rdi

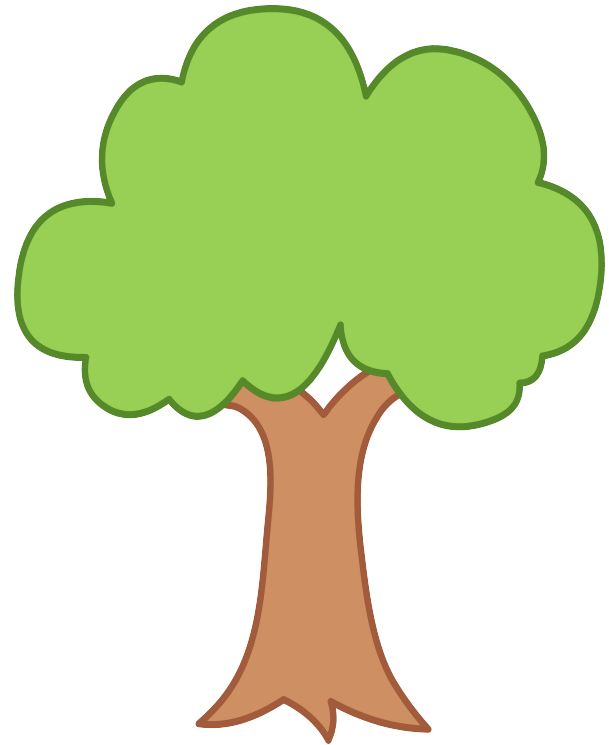
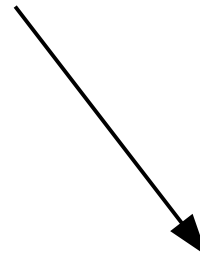
jne .L5

.L3

rep

ret

```
long trace(tree_ptr tp) {  
    long eax = 0;  
    while (tp != 0) {  
        eax = tp->val;  
        tp = tp->right;  
    }  
    return eax;  
}
```



Stack Discipline

Suppose $ebp = 4$, what is return address?

0				4				8				12			
A0	D3	7F	00	10	00	00	00	36	D3	77	E4	3D	D3	FF	B4
16				20				24				28			
BB	88	7F	00	FF	FF	39	24	D9	FD	11	00	00	00	FF	4B
32				36				40				44			
9C	11	3F	10	48	D0	99	24	89	D9	43	58	80	00	88	44

What is first parameter, second parameter if they are both ints?

Stack Discipline

ebp = 4 → return address = 0xE477D336

0				4				8				12			
A0	D3	7F	00	20	00	00	00	36	D3	77	E4	3D	D3	FF	B4
16				20				24				28			
BB	88	7F	00	FF	FF	39	24	D9	FD	11	00	00	00	FF	4B
32				36				40				44			
9C	11	3F	10	48	D0	99	24	89	D9	43	58	80	00	88	44

Param1 = 0xB4FFD33D Param 2 = 0x7F88BB

Stack Discipline

ebp = 4. What is first parameter of the caller?

0				4				8				12			
A0	D3	7F	00	20	00	00	00	36	D3	77	E4	3D	D3	FF	B4
16				20				24				28			
BB	88	7F	00	FF	FF	39	24	D9	FD	11	00	00	00	FF	4B
32				36				40				44			
9C	11	3F	10	48	D0	99	24	89	D9	43	58	80	00	88	44

Stack Discipline

ebp = 4. What is first parameter of the caller?

0				4				8				12			
A0	D3	7F	00	20	00	00	00	36	D3	77	E4	3D	D3	FF	B4
16				20				24				28			
BB	88	7F	00	FF	FF	39	24	D9	FD	11	00	00	00	FF	4B
32				36				40				44			
9C	11	3F	10	48	D0	99	24	89	D9	43	58	80	00	88	44

Old ebp = 0x00000020 = 32

1st param = 0x5843D989

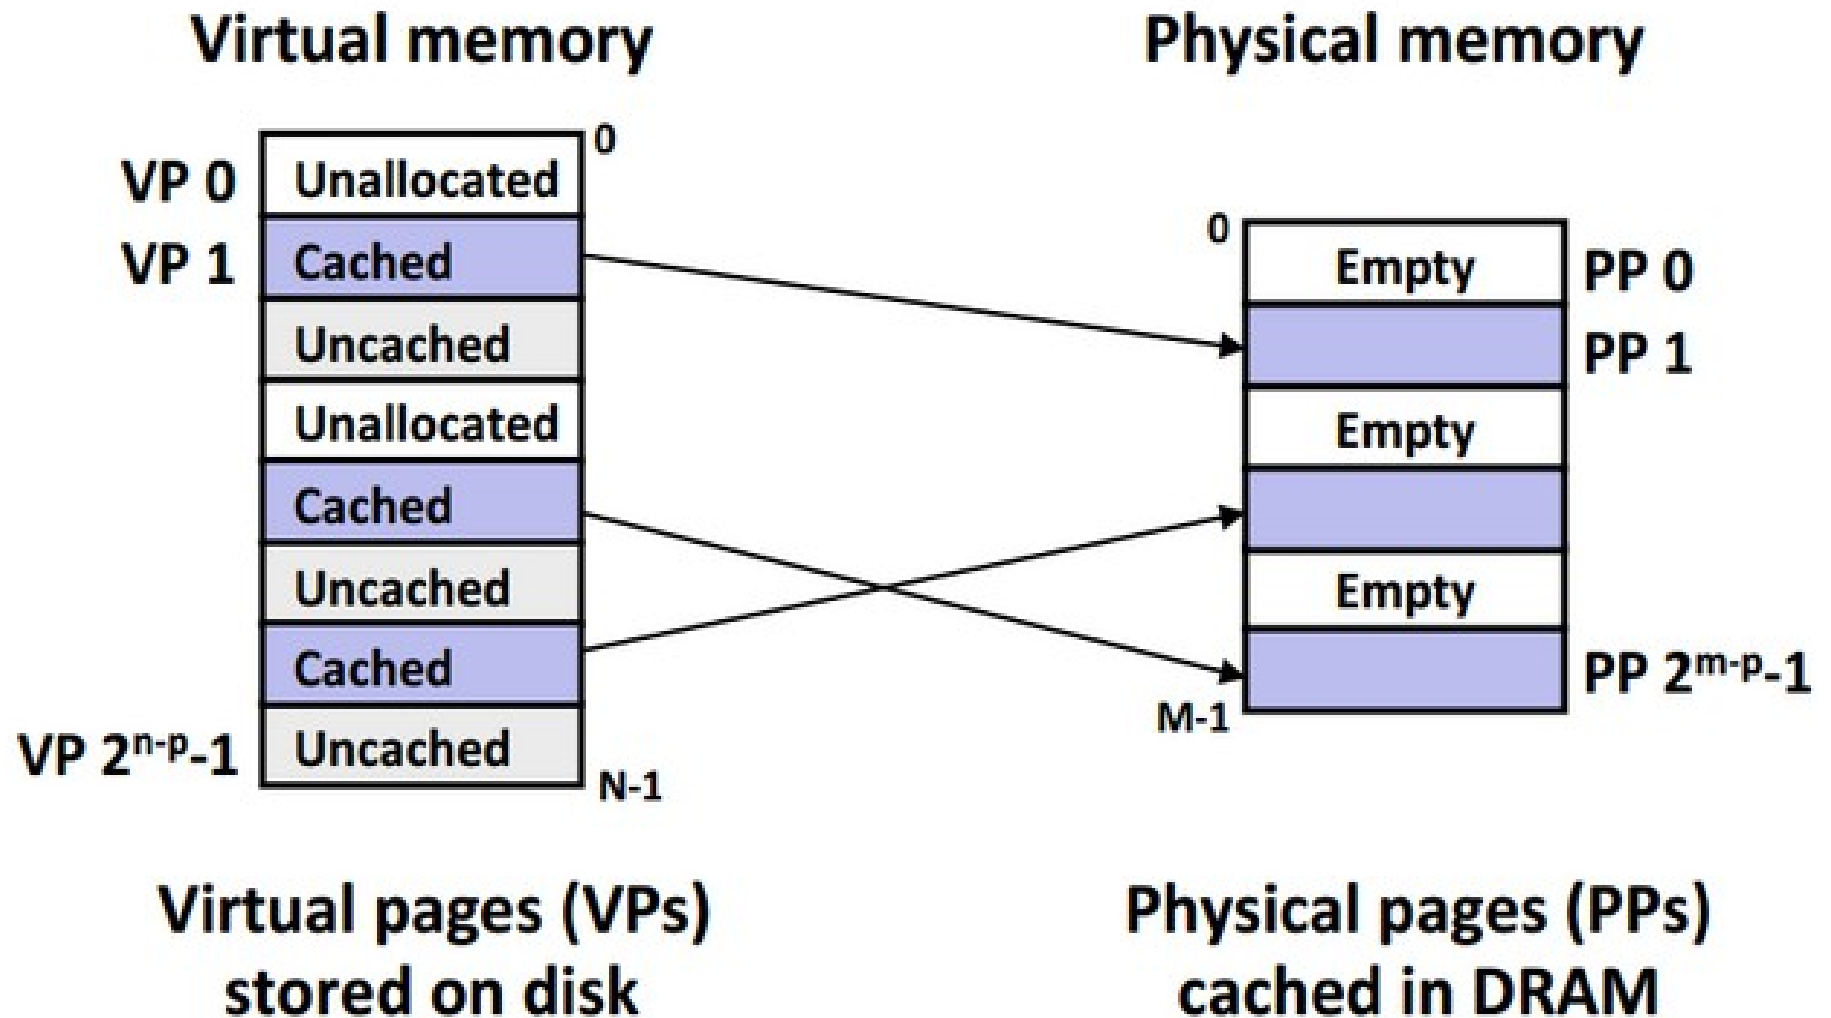
Other Misc Notes

- Linux alignment policy: all 2B data types (short) are aligned on multiples of 2B
- Any primitive larger than 2B is aligned on multiples of 4B
 - Even 8B primitives like doubles and long longs

CHAPTER 6: Virtual Memory

Page Table and TLB

Remember how VM works



VM basics

- Virtual Address (VA) space is 8GB
 - How long is a virtual address?
- Page size is 4KB
 - How many bits needed to represent page offset
- How many bits to represent the Virtual page number?
- What is the virtual page number for this address:
 - `&x = 0x13477fad7`

VM basics

- Virtual Address (VA) space is 8GB = $(2^3) * (2^{30})B$
 - 33 bits
- Page size is 4KB = $(2^2) * (2^{10}) B$
 - 12 bits for the page offset
- Virtual page number is $33 - 12 = 21$ bits
- What is the virtual page number for this address:
 - $\&x = 0x13477fad7$
 - Page number is 0x13477f

VM basics

- VA space is 8GB, page size is 4KB
- How many pages in the PT?

VM basics

- VA space is 8GB, page size is 4KB
- How many pages in the PT?
 - 2^{21} possible virtual pages == 2^{21} page table entries

Assume 12 bit Virtual address, 10 bit Physical Address, Page size 16 B.

Virtual Address:

[illegible]Page
Table:

V	Physical Page Number
0	0x2C
1	0x37
1	0x2B
0	null
1	0x1F
0	0x08
1	0x1A
	...

bits VPO = ?

bits PPO = ?

bits VPN = ?

bits PPN = ?

Physical Address:

[illegible]

Assume 12 bit Virtual address, 10 bit Physical Address, Page size 16 B.

Virtual Address:

Virtual Page Number					Virtual Page Offset				
0	1	2	3	4	0	1	2	3	4
1	2	3	4	5	5	6	7	8	9
2	3	4	5	6	10	11	12	13	14
3	4	5	6	7	15	16	17	18	19
4	5	6	7	8	20	21	22	23	24
5	6	7	8	9	25	26	27	28	29
6	7	8	9	10	30	31	32	33	34
7	8	9	10	11	35	36	37	38	39
8	9	10	11	12	40	41	42	43	44
9	10	11	12	13	45	46	47	48	49
10	11	12	13	14	50	51	52	53	54
11	12	13	14	15	55	56	57	58	59
12	13	14	15	16	60	61	62	63	64
13	14	15	16	17	65	66	67	68	69
14	15	16	17	18	70	71	72	73	74
15	16	17	18	19	75	76	77	78	79
16	17	18	19	20	80	81	82	83	84
17	18	19	20	21	85	86	87	88	89
18	19	20	21	22	90	91	92	93	94
19	20	21	22	23	95	96	97	98	99

V Physical Page Number

Page
Table:

0	0x2C
1	0x37
1	0x2B
0	null
1	0x1F
0	0x08
1	0x1A
	...

bits VPO = 4

bits PPO = 4

bits VPN = 8

bits PPN = 6

Physical Address:

Physical Page Number	Physical Page Offset
0	0
1	1
2	2
3	3
4	4
5	5
6	6
7	7
8	8
9	9
10	10
11	11
12	12
13	13
14	14
15	15
16	16
17	17
18	18
19	19
20	20
21	21
22	22
23	23
24	24
25	25
26	26
27	27
28	28
29	29
30	30
31	31
32	32
33	33
34	34
35	35
36	36
37	37
38	38
39	39
40	40
41	41
42	42
43	43
44	44
45	45
46	46
47	47
48	48
49	49
50	50
51	51
52	52
53	53
54	54
55	55
56	56
57	57
58	58
59	59
60	60
61	61
62	62
63	63
64	64
65	65
66	66
67	67
68	68
69	69
70	70
71	71
72	72
73	73
74	74
75	75
76	76
77	77
78	78
79	79
80	80
81	81
82	82
83	83
84	84
85	85
86	86
87	87
88	88
89	89
90	90
91	91
92	92
93	93
94	94
95	95
96	96
97	97
98	98
99	99

Assume 12 bit Virtual address, 10 bit Physical Address, Page size 16 B.

Virtual Address: 0x03A

0	0	0	0	0	0	1	1	1	0	1	0
---	---	---	---	---	---	---	---	---	---	---	---

V Physical Page Number

Page
Table:

0	0x2C
1	0x37
1	0x2B
0	null
1	0x1F
0	0x08
1	0x1A
	...

bits VPO = 4

bits PPO = 4

bits VPN = 8

bits PPN = 6

Physical Address:

[illegible]

Assume 12 bit Virtual address, 10 bit Physical Address, Page size 16 B.

Virtual Address: 0x03A

0	0	0	0	0	0	1	1	1	0	1	0
---	---	---	---	---	---	---	---	---	---	---	---

Page
Table:

V Physical Page Number

0	0x2C
1	0x37
1	0x2B
0	null
1	0x1F
0	0x08
1	0x1A
	...

bits VPO = 4

bits PPO = 4

bits VPN = 8

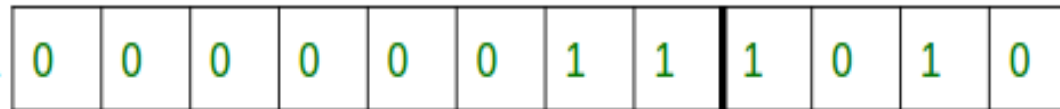
bits PPN = 6

Physical Address:

[illegible]

Assume 12 bit Virtual address, 10 bit Physical Address, Page size 16 B.

Virtual Address: 0x03A



Page
Table:

Page Hit!

V Physical Page Number

0	0x2C
1	0x37
1	0x2B
0	null
1	0x1F
0	0x08
1	0x1A
...	

bits VPO = 4

bits PPO = 4

bits VPN = 8

bits PPN = 6

Physical Address: 0x2BA



Assume 12 bit Virtual address, 10 bit Physical Address, Page size 16 B.

Virtual Address: 0x047

0	0	0	0	0	1	0	0	0	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---

V Physical Page Number

Page
Table:

0	0x2C
1	0x37
1	0x2B
0	null
1	0x1F
0	0x08
1	0x1A
	...

bits VPO = 4

bits PPO = 4

bits VPN = 8

bits PPN = 6

Physical Address:

[illegible]

Assume 12 bit Virtual address, 10 bit Physical Address, Page size 16 B.

Virtual Address: 0x047

0	0	0	0	0	1	0	0	0	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---

Page
Table:

V	Physical Page Number
0	0x2C
1	0x37
1	0x2B
0	null
1	0x1F
0	0x08
1	0x1A
	...

bits VPO = 4

bits PPO = 4

bits VPN = 8

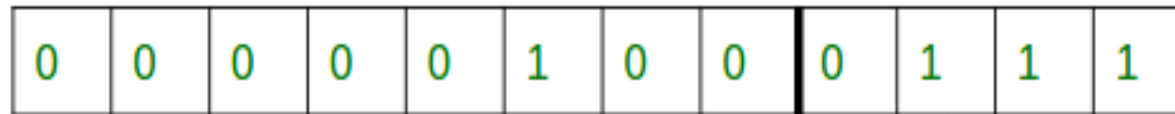
bits PPN = 6

Physical Address:

[illegible]

Assume 12 bit Virtual address, 10 bit Physical Address, Page size 16 B.

Virtual Address: 0x047



Page
Table:

Page Hit!

V Physical Page Number

0	0x2C
1	0x37
1	0x2B
0	null
1	0x1F
0	0x08
1	0x1A

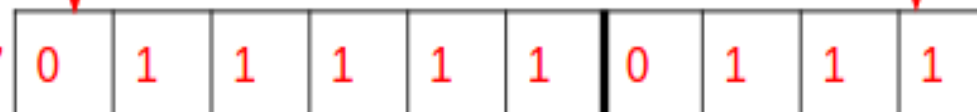
bits VPO = 4

bits PPO = 4

bits VPN = 8

bits PPN = 6

Physical Address: 0x1F7



Assume 12 bit Virtual address, 10 bit Physical Address, Page size 16 B.

Virtual Address: 0x05F

0	0	0	0	0	1	0	1	1	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---

V Physical Page Number

Page
Table:

0	0x2C
1	0x37
1	0x2B
0	null
1	0x1F
0	0x08
1	0x1A
	...

bits VPO = 4

bits PPO = 4

bits VPN = 8

bits PPN = 6

Physical Address:

[illegible]

Assume 12 bit Virtual address, 10 bit Physical Address, Page size 16 B.

Virtual Address: 0x05F

0	0	0	0	0	1	0	1	1	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---

Page
Table:

Page Fault!
=(
Page is not
resident in
Memory

V	Physical Page Number
0	0x2C
1	0x37
1	0x2B
0	null
1	0x1F
0	0x08
1	0x1A
	...

bits VPO = 4

bits PPO = 4

bits VPN = 8

bits PPN = 6

Physical Address:

[illegible]

Assume 12 bit Virtual address, 10 bit Physical Address, Page size 16 B.

Virtual Address: 0x05F

0	0	0	0	0	1	0	1	1	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---

Page
Table:

Page Fault!
=(
Page is not
resident in
Memory

What do we do?

Physical Address:

V	Physical Page Number
0	0x2C
1	0x37
1	0x2B
0	null
1	0x1F
1	0x2E
1	0x1A
	...

bits VPO = 4

bits PPO = 4

bits VPN = 8

bits PPN = 6

- Let page fault exception handler take care of bringing the page into memory, so we can use it.

[illegible]

Assume 12 bit Virtual address, 10 bit Physical Address, Page size 16 B.

Virtual Address: 0x05F

0	0	0	0	0	1	0	1	1	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---

Page
Table:

Page Fault!
=(
Page is not
resident in
Memory

What do we
do?

V Physical Page Number

0	0x2C
1	0x37
1	0x2B
0	null
1	0x1F
1	0x2E
1	0x1A
...	

bits VPO = 4

bits PPO = 4

bits VPN = 8

bits PPN = 6

Let page fault exception handler
take care of bringing the page in
from the disc so we can use it.

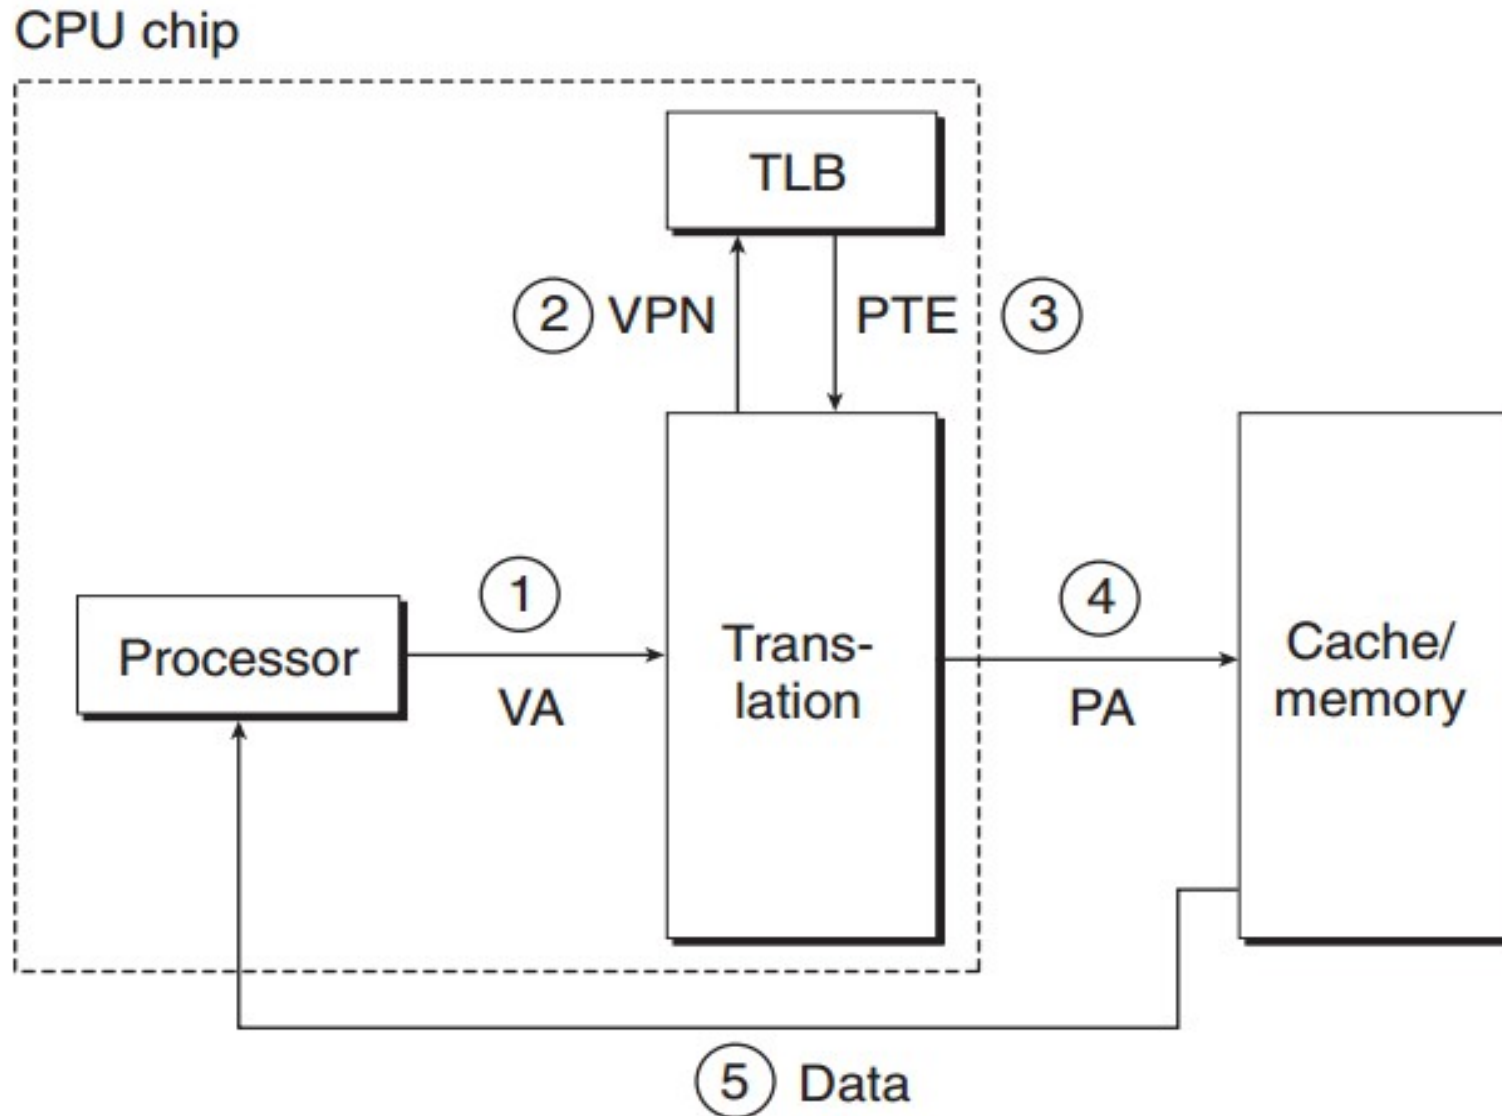
Physical Address: 0x2EF

1	0	1	1	1	0	1	1	1	1
---	---	---	---	---	---	---	---	---	---

Improving Performance of Cache

- Last time we noticed that using a PT could be expensive
- Need one I/O to read the PT and retrieve addr translation
- Need one more I/O to fetch the data at the correct PA

Using TLB to Cache the PT



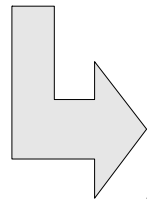
TLB is Just like a Cache

- Caches Page Table entries (e.g. caches Virtual
→ page number translation mappings)

TLB is Just like a Cache

- Caches Page Table entries (e.g. caches Virtual → page number translation mappings)
- One Cache line contains one translation

TLB Index



- And TLB is typically highly associative

Determining TLB Index and Tag

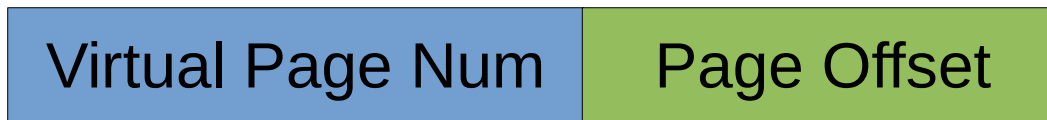
- Before, we chopped our VA into 2 parts (VPN and VPO)

Virtual Page Num

Page Offset

Determining TLB Index and Tag

- Before, we chopped our VA into 2 parts (VPN and VPO)



- Now, divide VPN into TLB tag and index



TLB Example

- $|VA| = 20$ bits, $|PA| = 18$ bits
- Page Size = 4KB
- TLB is direct mapped, with 8 sets
 - How many bits do I need for:
 - Virtual Page Number (VPN)
 - Physical Page Number (PPN)
 - TLB Tag
 - TLB index

TLB Example

- $|VA| = 20$ bits, $|PA| = 18$ bits
- Page Size = 4KB
- TLB is direct mapped, with 8 sets
 - How many bits do I need for:
 - Virtual Page Number (VPN) = 8 bits
 - Physical Page Number (PPN) = 6 bits
 - TLB Tag = 5 bits
 - TLB index = 3 bits

TLB Example

What is the
addresss
translation for
VA = 0xEE000?

	VALID	1	PPN Mapping
0	1	0x0A	b100100
1	1	0x17	b101010
2	0	0x11	b000001
3	1	0x0F	b011111
4	1	0x03	b111000
5	1	0x01	b010001
6	1	0x1D	b110011
0	1	0x13	b111000

TLB Example

What is the
address
translation for
VA = 0xEE000?

	VALID	1	PPN Mapping
0	1	0x0A	b100100
1	1	0x17	b101010
2	0	0x11	b000001
3	1	0x0F	b011111
4	1	0x03	b111000
5	1	0x01	b010001
6	1	0x1D	b110011
0	1	0x13	b111000

PPN = b110011
→ **PA = 0x33000**

TLB Lookup Procedure

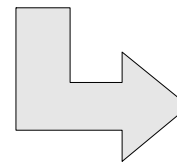
1. Given Virtual address, chop off lower bits corresponding to Page Offset
2. The remaining bits specify Virtual Page Number. Separate into TLB Tag and TLB index (same procedure as in caching)
3. Use TLB set-index as an index lookup into the TLB
4. If there is a valid entry in the set with matching tags, take the contents of the block (the Physical Page Number)

TLB and Page Table

Small TLB

VALID	1	PPN Mapping
1	0x0A	b100100
1	0x17	b101010
0	0x11	b000001
1	0x0F	b011111
1	0x03	b111000
1	0x01	b010001
1	0x1D	b110011
1	0x13	b111000

PT entry
0xEE



VALID	PPN Mapping
	...
1	b110011
	...

BIG Page Table

VA = 0xEE000

→ **VPN = 0xEE**

→ **TLB tag = 0x1D index = 6**

TLB and Page Table

- Valid bits **do not** mean the same thing in PT and TLB
 - In Page Table, valid indicates whether that page is in memory
 - If not generate a **PAGE FAULT**
 - In TLB, valid indicates whether that cache line is valid
- e.g. a line in the TLB might be invalid, but PT lookup may yield a valid Physical Page translation
 - i.e. can be a **TLB miss** without generating a **Page Fault**
- If entry is not valid in PT, should **not** be valid in TLB

TLB Miss

- If the TLB lookup results in a miss, must fetch the underlying translation from the PT (just like a cache miss)
 - If the PT entry is invalid, must also generate a **PAGE FAULT**

TLB Miss

Small TLB

VALID	1	PPN Mapping
1	0x0A	b100100
1	0x17	b101010
0	0x11	b000001
1	0x0F	b011111
1	0x03	b111000
1	0x01	b010001
1	0x1D	b110011
1	0x13	b111000

Now Consider:

VA = 0xED000

→ **VPN = 0xED**

→ **TLB tag = 0x1D index = 5**

VALID	PPN Mapping
1	b000010
1	b110011
	...

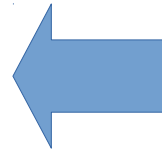
BIG Page Table

TLB Miss

Small TLB

VALID	1	PPN Mapping
1	0x0A	b100100
1	0x17	b101010
0	0x11	b000001
1	0x0F	b011111
1	0x03	b1111000
1	0x01	b010001
1	0x1D	b110011
1	0x13	b1111000

TLB MISS



VA = 0xED000

→ **VPN = 0xED**

→ **TLB tag = 0x1D index = 5**

VALID	PPN Mapping
1	b000010
1	b110011
	...

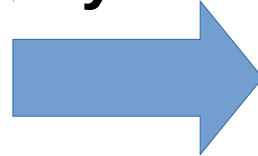
BIG Page Table

TLB Miss

Small TLB

VALID	1	PPN Mapping
1	0x0A	b100100
1	0x17	b101010
0	0x11	b000001
1	0x0F	b011111
1	0x03	b111000
1	0x01	b010001
1	0x1D	b110011
1	0x13	b111000

PT lookup
Entry 0xED



VALID	PPN Mapping
1	b000010
1	b110011
	...

VA = 0xED000

→ **VPN = 0xED**

→ **TLB tag = 0x1D index = 5**

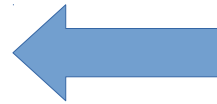
BIG Page Table

TLB Miss

Small TLB

VALID	1	PPN Mapping
1	0x0A	b100100
1	0x17	b101010
0	0x11	b000001
1	0x0F	b011111
1	0x03	b111000
1	0x1D	b000010
1	0x1D	b110011
1	0x13	b111000

Block Move



VALID	PPN Mapping
1	b000010
1	b110011
	...

VA = 0xED000

→ **VPN = 0xED**

→ **TLB tag = 0x1D index = 5**

BIG Page Table

THE END

Good luck on Finals!!