

CS 33 – Discussion Week 7

Agenda:

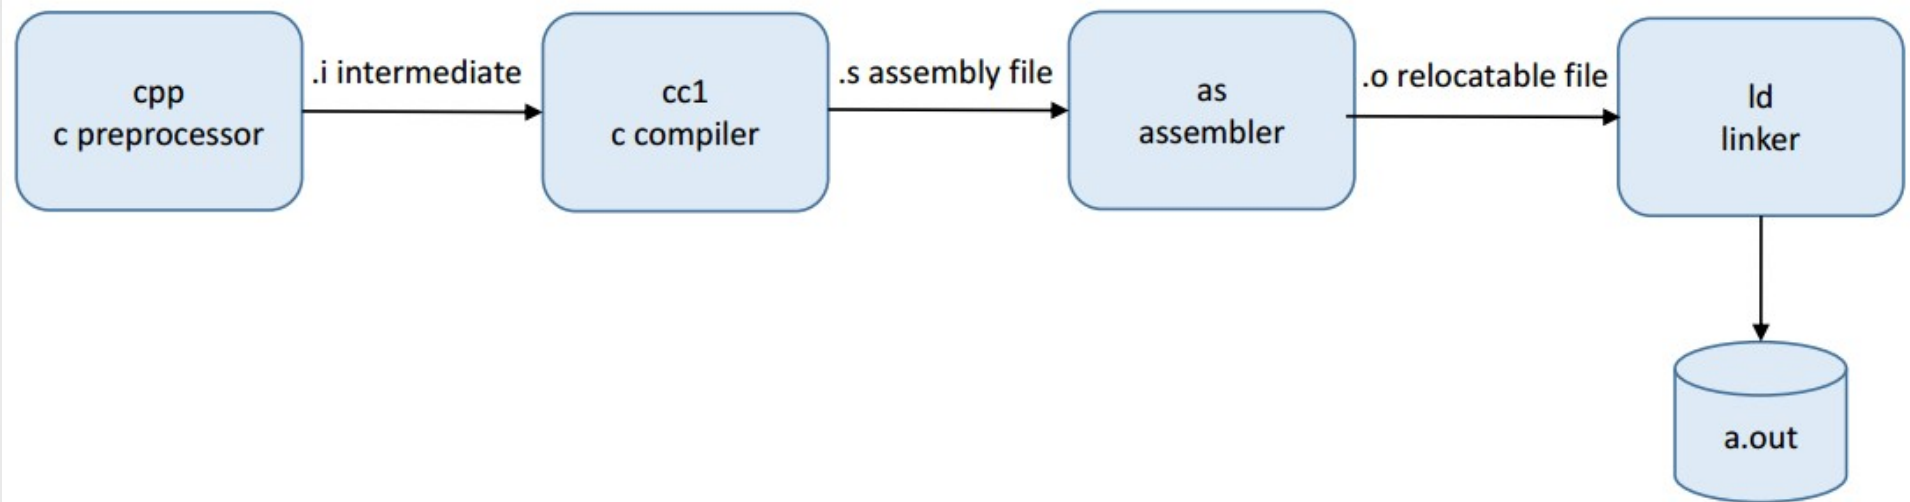
- Midterm
- Lab 3
- Linking
- Exception Control Flow

Lab 3

- Write a cache simulator
- Keep track of read and write hits/misses.
- See how the hit rate changes when we perform a matrix transposition in several different ways.
 - Row major order
 - Column major order
 - Blocking
- LRU and write-back
- Need to use C for this lab
- Compile with `gcc cachelab.c -lm`
- From now on, avoid using extra compiler flags

Linking

- Linking is the last step in compiling.



- Linker inputs: Relocatable object files (e.g. `main.o`, `swap.o`) that are compiled separately.
- Linker output: Fully linked executable object file (e.g. `a.out`) ready to be copied into memory and begin execution.

Linking

- Two tasks performed by linker:
 - Symbol Resolution – associate symbols with exactly one symbol definition. What exactly are symbols?
 - Relocation – Merges code and data into single sections. Relocates symbols from their locations in each “.o” to their final memory locations
- We will go into more detail about these 2 tasks.

What are Relocatable Object Files?

- .o files (e.g. swap.o)
- Piece of compiled code and data.
- Not yet executable without linking.
- What does it look like?
- Follows ELF: Unix Executable and Linkable Format
- pp.658-59 of textbook go into detail about each of these sections

Typical ELF Relocatable Object File

ELF Header
.text: machine code
.rodata: read only data
.data: initialized globals
.bss: uninitialized globals (description)
.symtab: symbol table (globals and external function info)
.rel.text: relocation information for externals
.rel.data: relocation information for cross referenced data
.debug: -g symbols for gdb
.line: -g line numbers for gdb
.strtab: descriptive strings for .symtab
Section header table: which sections are in the table

Symbols and Symbol Tables

- Every Relocatable Object File has its own symbol table, with information about its own symbols. This lives in `.symtab`
- 3 types of symbols:
 - Global – All *nonstatic* C functions and *nonstatic* global variables defined in this module.
 - External – C functions and variables referenced in current module but defined in other modules.
 - Local – only used within this object file. e.g. global variables with `static` attribute. NOTE: These are **not** the same as local variables!

What symbols do we have here?

```
1 // Code for main.c
2 int buf[2] = {1, 2};
3
4 int main()
5 {
6     swap();
7     return 0;
8 }
```

Symbols:

- Global: buf, main()
- Extern: swap()

```
1 // Code for swap.c
2 extern int buf[];
3
4 int *bufp0 = &buf[0];
5 static int *bufp1;
6 void swap()
7 {
8     int temp;
9     bufp1 = &buf[1];
10    temp = *bufp0;
11    *bufp0 = *bufp1;
12    *bufp1 = temp;
13 } |
```

Symbols:

- Global: bufp0, swap()
- Extern: buf[]
- Local: bufp1

Linker – Symbol Resolution

- Match up identically named external (global) symbols. Some defined within the module, some outside.
- How is this done? Program symbols are put into two categories:
 - Strong – functions and initialized globals
 - Weak – uninitialized globals

What symbols do we have here?

```
1 // Code for main.c
2 int buf[2] = {1, 2};
3
4 int main()
5 {
6     swap();
7     return 0;
8 }
```

Symbols:

- Global: buf^s, main(^s)
- Extern: swap(^w)

```
1 // Code for swap.c
2 extern int buf[];
3
4 int *bufp0 = &buf[0];
5 static int *bufp1;
6 void swap()
7 {
8     int temp;
9     bufp1 = &buf[1];
10    temp = *bufp0;
11    *bufp0 = *bufp1;
12    *bufp1 = temp;
13 } |
```

Symbols:

- Global: bufp0^s, swap(^s)
- Extern: buf[]^w
- Local: bufp1^w

Symbol Resolution

- **Rule 1: Multiple strong symbols are not allowed**
 - Each item can be defined only once
 - Otherwise: Linker error
- **Rule 2: Given a strong symbol and multiple weak symbol, choose the strong symbol**
 - References to the weak symbol resolve to the strong symbol
- **Rule 3: If there are multiple weak symbols, pick an arbitrary one**
 - Can override this with `gcc -fno-common`

Examples

```
1 // foo1.c
2 int main() {
3     return 0;
4 }
```

```
1 // foo2.c
2 int main() {
3     return 5;
4 }
```

What happens when I try to compile this?

```
$ gcc foo1.c foo2.c
```

Examples

```
1 // foo3.c
2 #include <stdio.h>
3 void f(void);
4
5 int x = 15;
6
7 int main() {
8     f();
9     printf("x = %d\n", x);
10    return 0;
11 }
```

```
1 // bar3.c
2 int x;
3
4 void f() {
5     x = 25;
6 }
```

What happens when I try to compile this?

\$ gcc foo3.c bar3.c

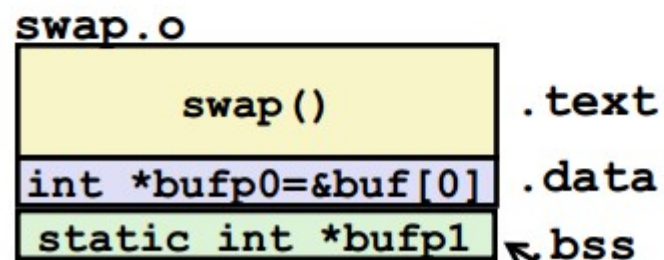

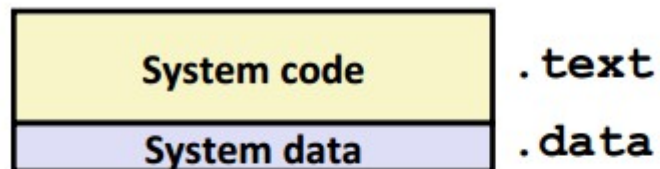
What happens when I run it?

Notice

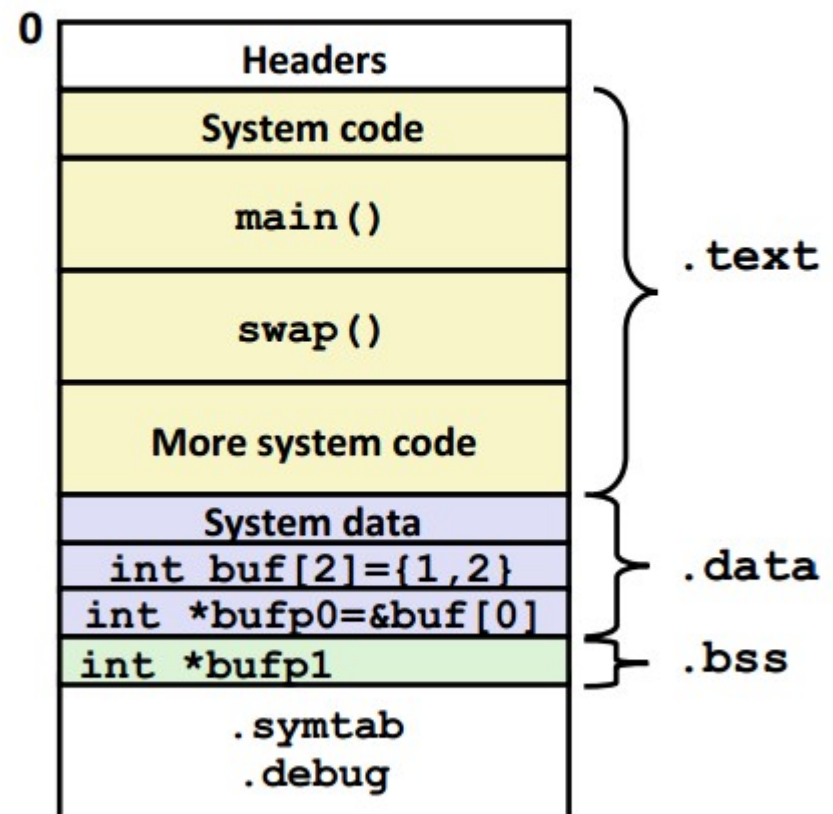
- These previous slides should illustrate why using global variables can be very dangerous, and is generally avoided in good coding practice.

Linking - Relocation

Relocatable Object Files



Executable Object File



Even though private to swap, requires allocation in .bss

Relocation

- Aggregates sections of the same type from each of the different relocatable object files.
- Changes any symbol reference to point to the correct run-time address.

Example

main.c

```
int buf[2] =
    {1,2};

int main()
{
    swap();
    return 0;
}
```

main.o

```
00000000 <main>:
 0: 8d 4c 24 04      lea    0x4(%esp),%ecx
 4: 83 e4 f0        and    $0xffffffff0,%esp
 7: ff 71 fc        pushl  0xffffffffc(%ecx)
 a: 55             push   %ebp
 b: 89 e5          mov    %esp,%ebp
 d: 51             push   %ecx
 e: 83 ec 04        sub    $0x4,%esp
11: e8 fc ff ff ff   call   12 <main+0x12>
16: b8 00 00 00 00   mov    $0x0,%eax
1b: 83 c4 04        add    $0x4,%esp
1e: 59             pop     %ecx
1f: 5d             pop     %ebp
20: 8d 61 fc        lea    0xffffffffc(%ecx),%esp
23: c3             ret
```

12: R_386_PC32 swap

-4

Disassembly of section .data:

```
00000000 <buf>:
 0: 01 00 00 00 02 00 00 00
```

Source: objdump -r -d main.o

Source: objdump -j .data -d main.o 20

Continued

```
00000000 <main>:
...
   e:   83 ec 04          sub    $0x4,%esp
  11:   e8 fc ff ff ff    call   12 <main+0x12>
                                12: R_386_PC32 swap
  16:   b8 00 00 00 00    mov    $0x0,%eax
...
```

Link time:

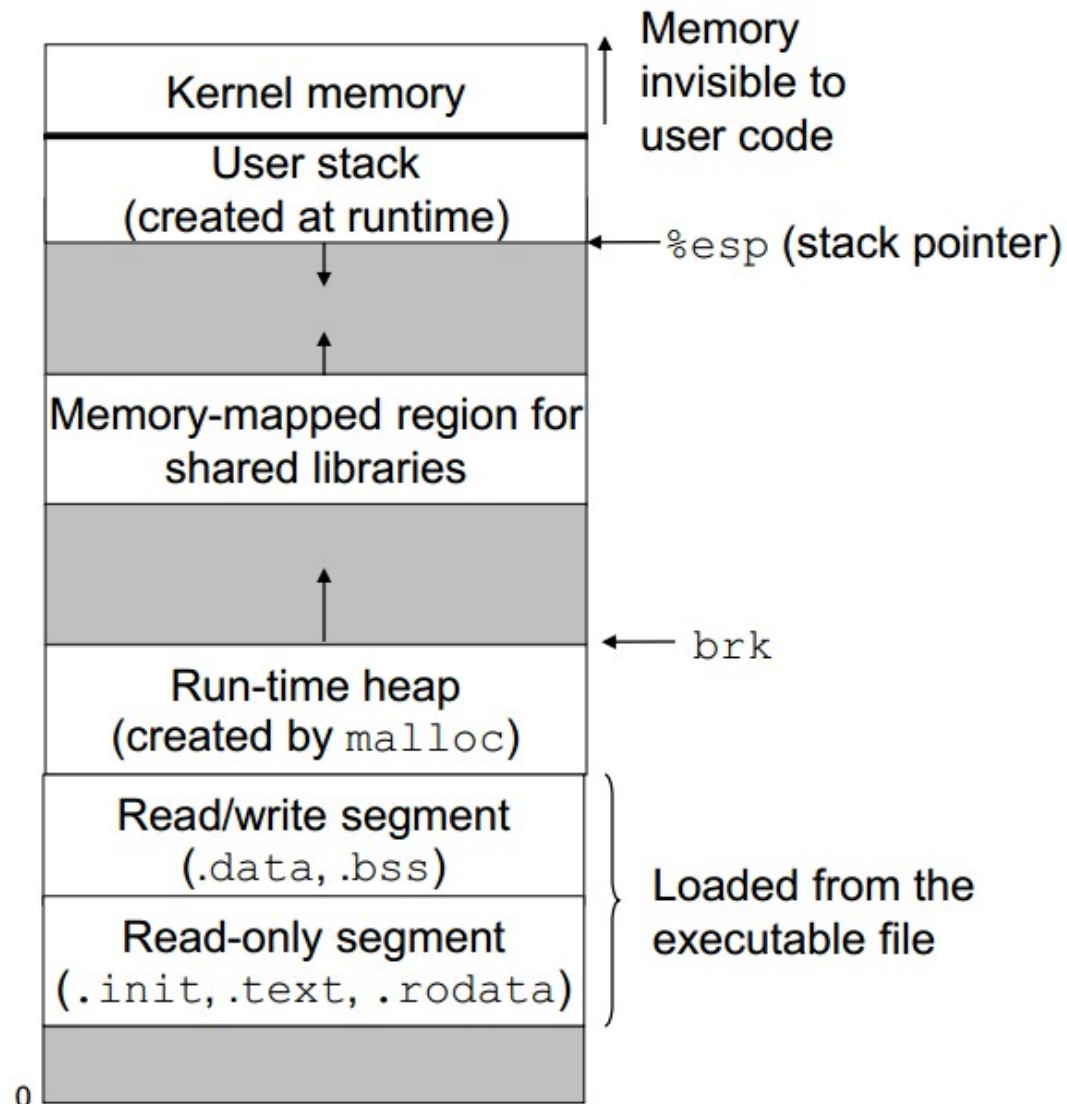
$0x8048398 + (-4)$
 $- 0x8048386 = 0xe$

Runtime:

$0x804838a + 0xe$
 $= 0x8048398$

```
08048374 <main>:
8048374:      8d 4c 24 04          lea    0x4(%esp),%ecx
8048378:      83 e4 f0            and    $0xfffffffff0,%esp
804837b:      ff 71 fc            pushl  0xfffffffffc(%ecx)
804837e:      55                  push   %ebp
804837f:      89 e5                mov    %esp,%ebp
8048381:      51                  push   %ecx
8048382:      83 ec 04            sub    $0x4,%esp
8048385:      e8 0e 00 00 00      call   8048398 <swap>
804838a:      b8 00 00 00 00      mov    $0x0,%eax
804838f:      83 c4 04            add    $0x4,%esp
8048392:      59                  pop    %ecx
8048393:      5d                  pop    %ebp
8048394:      8d 61 fc            lea    0xfffffffffc(%ecx),%esp
8048397:      c3                  ret
```

How is the executable loaded in memory?



How do we see Program Execution?

- When we execute a program, %rip (or %eip) points to the first instruction.
- Each instruction in the program is executed step by step (or jumping as the case may be), using the CPU, RAM, and the CPU registers.
- Life is good.
- But let's look at the bigger picture for a moment.

The Bigger Picture

- Programs and applications are run on top of the operating system right.
- For one thing, what happens the program finishes? Presumably, `%rip` points to the next program to execute?
- But wait, how can we run multiple programs at once right? There's only one `%rip` per CPU right? Do they also share the stack?
- Can we go back to not thinking about this?

Nope

- For now, focus on one concern:
 - What happens when something unusual (one could even say... exceptional) occurs.
- What do I mean?
 - Divide by zero
 - Invalid operation
 - OS needs to interrupt or halt program execution.

Exceptions

- “An abrupt change in the control flow in response to some change in the processor's state”
- Come in four flavors:
 - Interrupts
 - Traps
 - Faults
 - Aborts

Interrupts

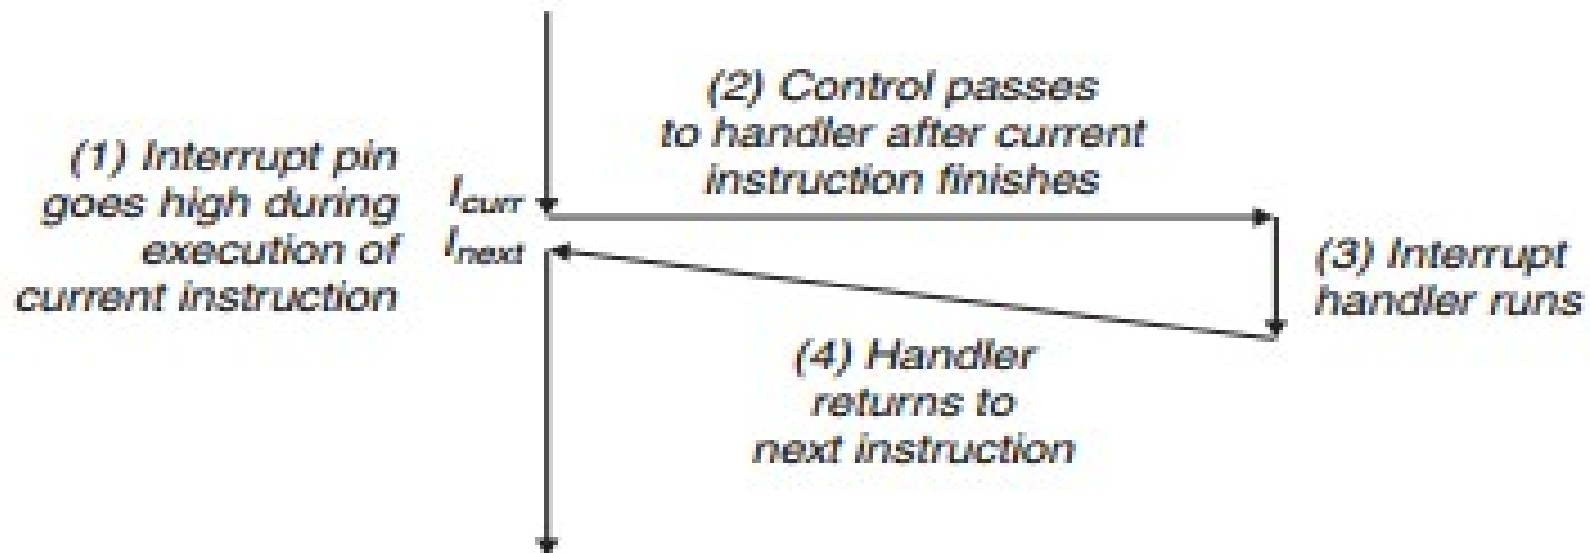
- Most commonly signals from I/O devices.
 - Keyboard key presses.
 - Mouse movement
 - Network adapter activity
 - Etc.
- Asynchronous
 - Occurs independently of currently executing program

Interrupt Handling

- I/O device triggers the “interrupt pin”
- After current instruction, stop executing current program and “control switches to interrupt handler”.
 - What does “control flow” and “passing control” mean?
 - High level: control flow is the execution of a single program and switching control means to allow another program to use the CPU resources to execute.
 - But more on that later

Interrupt Handling

- Interrupt handler handles interrupt.
- Control is given back to previously executing program.
- Previous program executes the next instruction.

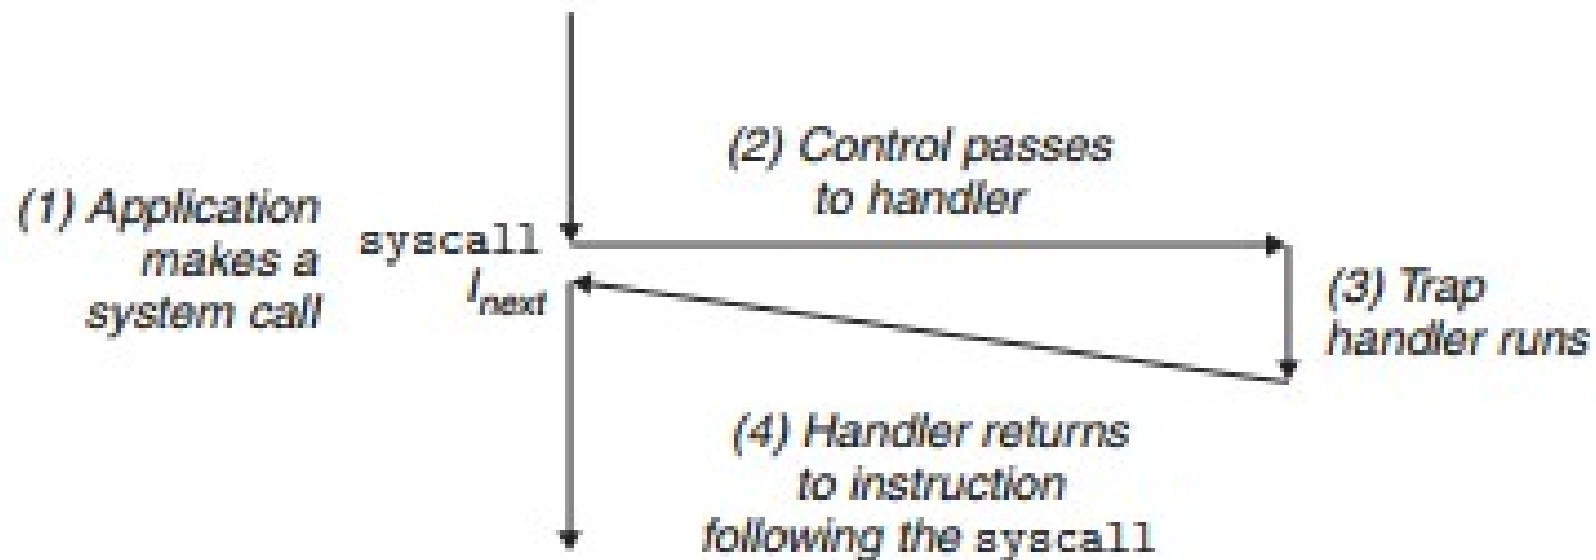


Trap

- An intentional exception triggered by user. What for?
- Sometimes we need to do things that are not within the scope of what the program alone can do.
 - Read a file
 - Create a new process
 - Load a new program
- Synchronous: occurs as a result of program instruction.

Trap handling

- Same as interrupt handling, except caused by an explicit instruction.

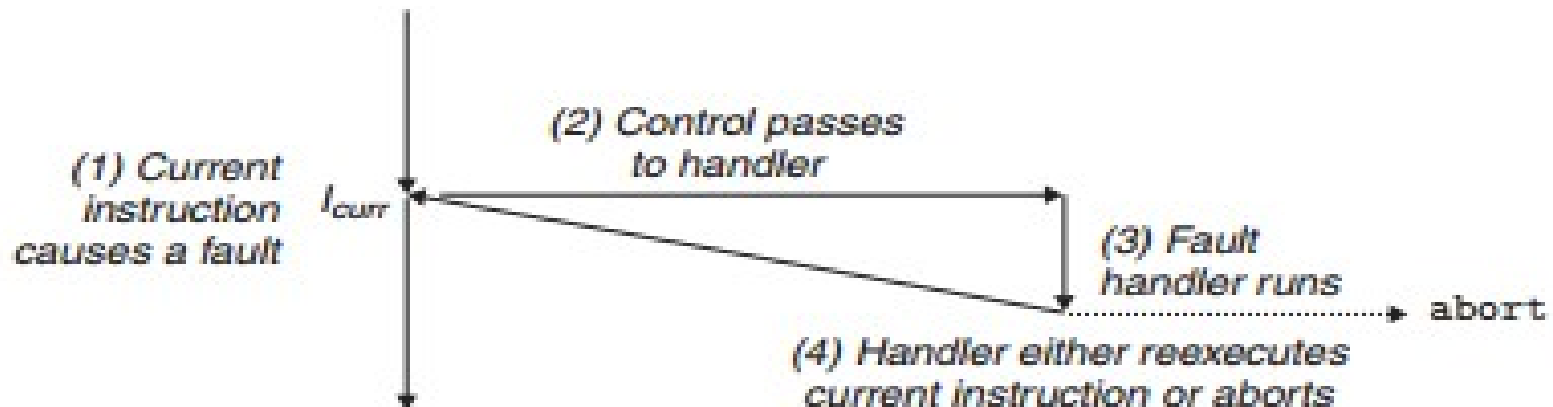


Fault

- Caused by a potentially recoverable, but unexpected error.
 - Divide by zero (in Linux, won't recover)
 - Invalid memory access (usually won't recover)
 - Page faults (must recover)
 - Like cache misses but oh so much worse.
 - But more on that later.
- Synchronous

Fault Handling

- Control passes to fault handler.
- Fault handler executes. If recovery is possible, return to instruction that caused fault. Else, halt.
 - Execute the instruction that caused the fault again?
 - If recoverable, whatever caused fault will be fixed and the instruction can be run without error.



Abort

- Unrecoverable, fatal error.
 - Corrupted memory
 - Fatal hardware error
- Abort handling
 - Abort with no chance of recovery.

Trap: syscall

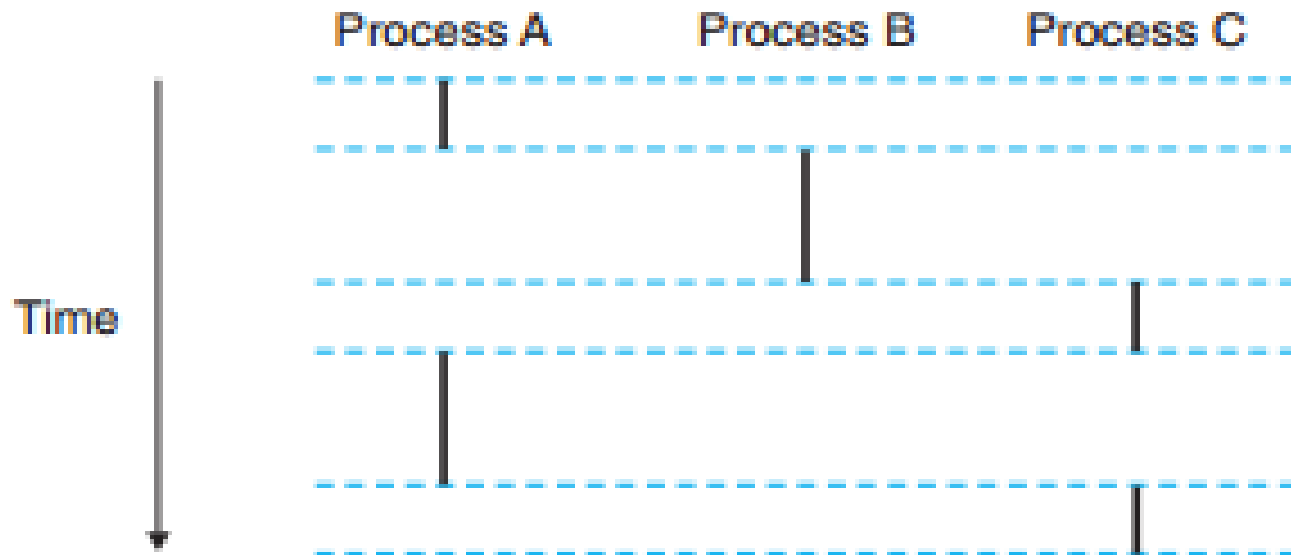
- Linux provides “system calls” which provides services from the OS to an executing program.
- In C, can use syscall function, but can more simply use wrapper functions.
 - read, write, open, close, execve, exit, fork, etc.
- These syscalls will cause a trap.

But hold on...

- How exactly is a program “paused” so that an exception handler can execute?
- For that matter, how can multiple programs run at the same time? There's only one %rip, %rbp, etc.

The dark secret

- Programs on a single CPU do not truly run simultaneously.
- A program generally corresponds to a single process and processes share...



Processes

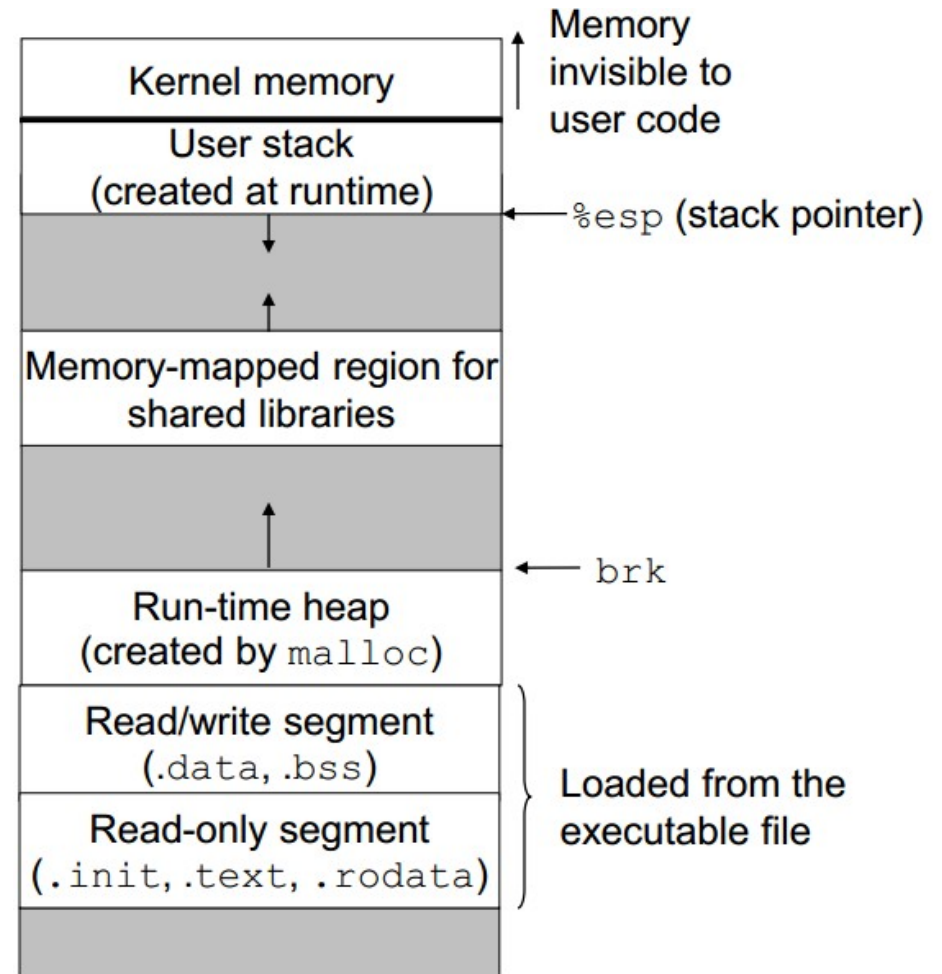
- Programs run atop a process, which appears to provide:
 - Control flow, or exclusive use of the processor to execute instructions
 - Its own memory.
- Every process is special... just like every other process.
- In reality, multiple processes take turn using the processor.

Context Switching

- When the CPU needs to switch to another process to execute, the current process' state (registers, memory) must be stored.
- The state of the next process to execute is restored and the next process runs.
- The previous process is none the wiser
 - This is context switching.
- This is what happens when switching to exception handlers.

What about memory?

- Do we need to save the entire addressable space?
- This is only what the process thinks it has.
- But more on this later (Virtual Memory).



Processes in C

- A program usually corresponds to a single process.
- But, you can actually refer to and create a process from within a program.
- Processes are referred to by a number id or in C, the data type “pid_t”.
- The syscall (wrapper) `fork()` will create a child process.

Processes in C

```
#include "csapp.h"
int main()
{
    pid_t pid;
    int x = 1;
    pid = fork();

    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        exit(0);
    }

    /* Parent */
    printf("parent: x=%d\n", --x);
    exit(0);
}
```

- `pid = fork()`
- As soon as `fork` is called a child process with an identical duplicate of the parents memory is made, which one exception.
- `pid` (parent process) = child process' id
- `pid` (child process) = 0
- `fork()` returns 0 for the child
.

Processes in C Major

```
#include "csapp.h"
int main()
{
    pid_t pid;
    int x = 1;
    pid = fork();

    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        exit(0);
    }

    /* Parent */
    printf("parent: x=%d\n", --x);
    exit(0);
}
```

- Both child and process will run the same code in parallel, but now the difference in pid will yield different behavior.