

# CS33 : Introduction to Computer Organization

Jochen Haber  
MS4000A

[jhaber@cs.ucla.edu](mailto:jhaber@cs.ucla.edu)

Lecture: MW 4PM-6PM

Office Hours:  
MW: 6PM-7PM  
BH4532B

Discussion as Enrolled:

Sec 1A PAB1749: F2PM-4PM Garrett Johnston ([gjohnston@cs.ucla.edu](mailto:gjohnston@cs.ucla.edu))  
TuTh: 9:30-10:30

Sec 1B BH5272: F4PM-6PM Peng Wei ([pengweiprc@ucla.edu](mailto:pengweiprc@ucla.edu))  
TuTh: 3:30-4:30

Sec 1C BH3400: F4PM-6PM Brandon Wu ([brandonwu@cs.ucla.edu](mailto:brandonwu@cs.ucla.edu))  
MTh: 9:30AM-10:30AM

Sec 1D BH5252: F4PM-6PM Uen-Tao Wang ([cerberiga@ucla.edu](mailto:cerberiga@ucla.edu))  
M1:00PM-3:00PM

Sec 1E BH5273: F4PM-6PM Alex Wood ([alex.wood@cs.ucla.edu](mailto:alex.wood@cs.ucla.edu))  
MW 11:30AM-12:30AM

Text:

Randal E. Bryant and David R. O'Hallaron  
Computer Systems: A Programmer's Perspective (CSPP)  
2nd (North American!) Edition, Prentice Hall 2010.

## Computer organization:

How does a computer really work?

From a hardware perspective but exposed by software

## Course focus:

chemicals and manufacturing:	NO
imbedded architecture	NO, but just an exposure
general architecture	YES, mostly about software
heavy on memory architecture	YES
compilers	YES
enabling software (systems programs)	NO

## Outline of course : (CSPP page xxviii Figure 2 “Course” ICS plus concurrent programming):

<u>Topic</u>	<u>CSPP</u>
Introduction	(1)
Tour of systems	(1)
Data representation	(2)
Machine language	(3)
Code optimization	(5)
Memory hierarchy	(6)
Linking	(7)
Exception control flow	(8)
Virtual memory	(9)
Concurrent programming	(12)

## Course work

18 lectures	36 hours		
9 discussions	18 hours		
2 mid-term exams	open notes: 15 points (each)	30	
final exam	open notes: 30 points	30	

## Outside work

90 hours

4 labs	15 points (each)	60	
? homework	only passing grade +1 point for pass, extra credit		

**Perfect grade**

**120 plus EC homeworks**

## Labs/homeworks

Will be tested on SEAS Linux machines.

Collaboration is encouraged but should be symmetrically acknowledged.

Late submissions not accepted.

## Exams:

Postponed exams by prior notice only

## Course Enrollment Issues

## Spreadsheets?

# Tentative Lecture Schedule

Lecture	Date	Topic
1	10/6/2014	Introduction/Tour of Systems
2	10/8/2014	Data Representation
3	10/13/2014	Data Representation
4	10/15/2014	Machine Language
5	10/20/2014	Assembly Language
6	10/22/2014	Midterm*
7	10/27/2014	Code Optimization
8	10/29/2014	Code Optimization
9	11/3/2014	Memory Hierarchy
10	11/5/2014	Caching
11	11/10/2014	Caching
12	11/12/2014	Midterm*
13	11/17/2014	Linking
14	11/19/2014	Exception Control Flow
15	11/24/2014	Virtual Memory
16	11/26/2014	Virtual Memory
17	12/1/2014	Concurrent Programming
18	12/3/2014	Concurrent Programming
19	12/8/2014	Review
20	12/10/2014	Review
	12/16/2014	Final Exam 8AM-11AM
	11/28/2014	No Discussion - Thanksgiving Holiday
		* No office hour

# What is a Computer?

Abacus?  
Cash Register?  
Thermometer?  
Stereo Amplifier?

## **Analog** (continuous, real)

Transducers  
Problem Solvers

## **Digital** (discrete, rational)

Mainframe  
Mini  
PC  
Imbedded

## **Theoretical** (set theory based)


FSM	regular languages
PDA	context free
LBA	context sensitive
TM	recursive

TMHP, Goedel's Incompleteness, Russell's paradox  
Relationship to grammars  
Chomsky, Greibach

## **Non-programmable versus programmable**

Stored program vs fixed.  
Modify its stored program

# Abacus



# What is a Computer?

Abacus?  
Cash Register?  
Thermometer?  
Stereo Amplifier?

## **Analog** (continuous, real)

Transducers  
Problem Solvers

## **Digital** (discrete, rational)

Mainframe  
Mini  
PC  
Imbedded

## **Theoretical** (set theory based)

FSM	regular languages
PDA	context free
LBA	context sensitive
TM	recursive


TMHP, Goedel's Incompleteness, Russell's paradox  
Relationship to grammars  
Chomsky, Greibach

## **Non-programmable versus programmable**

Stored program vs fixed.  
Modify its stored program

# Short Digression on Theoretical Computers

Input Tape




Valid Input Characters:  $A = \{a, b, c, \dots\}$  = alphabet

$A^*$  = set of all possible strings made up of alphabet

Any subset of  $A^*$  is a “language”

State Diagram



States:  $S = \{V, W, X, \dots\}$  Initial and Final State

Arrows are transitions from state to state depending of what is under the read head of the tape.

Machine starts in initial state with a string on the tape, stops when no further moves are possible.



# What is a Computer?

Abacus?  
Cash Register?  
Thermometer?  
Stereo Amplifier?

## **Analog** (continuous, real)

Transducers  
Problem Solvers

## **Digital** (discrete, rational)

Mainframe  
Mini  
PC  
Imbedded

## **Theoretical** (set theory based)

FSM	regular languages
PDA	context free
LBA	context sensitive
TM	recursive

TMHP, Goedel's Incompleteness, Russell's paradox  
Relationship to grammars  
Chomsky, Greibach

## **Non-programmable versus programmable**

Stored program vs fixed.  
Modify its stored program

## A Short History

2400 BC	Greek - Abacus
1100 BC	Chinese - Abacus
800 AD	Persians - Analog Machines/Astronomical Calculators
1200	Europe – Logic Machines
1850	Charles Babbage – Differential Calculator
1890	Herman Hollerith – Census Tabulator
1936	Alan Turing – Article on Turing Machines
1936	Konrad Zuse – Programmable Mechanical Machine (Z1)
1939	Hewlett Packard founded
1944	IBM – Howard Mark I
1945	John von Neumann – report on programmable computers
1946	ENIAC – U Penn
1954	WEIZAC – Weizmann Institute
1954	IBM 650 – commercially produced computer
1957	Fortran programming language
1959	COBOL programming language
1964	PL/1 programming language
1964	IBM System 360
1965	DEC PDP-8 mini computer
1965	Packet Switching – basis of Internet UCLA/Stanford
1970	UNIX
1975	Microsoft founded

## Punched Cards



Figure 4. Card Codes and Graphics for 64-Character Set

## A Short History


2400 BC	Greek - Abacus
1100 BC	Chinese - Abacus
800 AD	Persians - Analog Machines/Astronomical Calculators
1200	Europe – Logic Machines
1850	Charles Babbage – Differential Calculator
1890	Herman Hollerith – Census Tabulator
1936	Alan Turing – Article on Turing Machines
1936	Konrad Zuse – Programmable Mechanical Machine (Z1)
1939	Hewlett Packard founded
1944	IBM – Howard Mark I
1945	John von Neumann – report on programmable computers
1946	ENIAC – U Penn
1954	WEIZAC – Weizmann Institute
1954	IBM 650 – commercially produced computer
1957	Fortran programming language
1959	COBOL programming language
1964	PL/1 programming language
1964	IBM System 360
1965	DEC PDP-8 mini computer
1965	Packet Switching – basis of Internet UCLA/Stanford
1970	UNIX
1975	Microsoft founded



## A Short History

2400 BC	Greek - Abacus
1100 BC	Chinese - Abacus
800 AD	Persians - Analog Machines/Astronomical Calculators
1200	Europe – Logic Machines
1850	Charles Babbage – Differential Calculator
1890	Herman Hollerith – Census Tabulator
1936	Alan Turing – Article on Turing Machines
1936	Konrad Zuse – Programmable Mechanical Machine (Z1)
1939	Hewlett Packard founded
1944	IBM – Howard Mark I
1945	John von Neumann – report on programmable computers
1946	ENIAC – U Penn
1954	WEIZAC – Weizmann Institute
1954	IBM 650 – commercially produced computer
1957	Fortran programming language
1959	COBOL programming language
1964	PL/1 programming language
1964	IBM System 360
1965	DEC PDP-8 mini computer
1965	Packet Switching – basis of Internet UCLA/Stanford
1970	UNIX
1975	Microsoft founded

# Babbage Difference Engine




## A Short History

2400 BC	Greek - Abacus
1100 BC	Chinese - Abacus
800 AD	Persians - Analog Machines/Astronomical Calculators
1200	Europe – Logic Machines
1850	Charles Babbage – Differential Calculator
1890	Herman Hollerith – Census Tabulator
1936	Alan Turing – Article on Turing Machines
1936	Konrad Zuse – Programmable Mechanical Machine (Z1)
1939	Hewlett Packard founded
1944	IBM – Howard Mark I
1945	John von Neumann – report on programmable computers
1946	ENIAC – U Penn
1954	WEIZAC – Weizmann Institute
1954	IBM 650 – commercially produced computer
1957	Fortran programming language
1959	COBOL programming language
1964	PL/1 programming language
1964	IBM System 360
1965	DEC PDP-8 mini computer
1965	Packet Switching – basis of Internet UCLA/Stanford
1970	UNIX
1975	Microsoft founded



## And Gate




		Control	
		0	1
Input	0	0	0
	1	0	1


**Truth Table**

## Or Gate

Control = 0




Control = 1



		Control	
		0	1
Input	0	0	1
	1	1	1

**Truth Table**

## Not Gate



		Control	
		0	1
Input	0	1	
	1		0

**Truth Table**

## Mainframes

IBM (605, 1401, 7040, 7094, 360, Model 91, 370, z Series)  
1954 1959 1961 1962 1964 1968 1970, 1990  
NCR, CDC, Univac  
Cray, Illiac

## Mini Computers

1965 Digital Equipment PDP-8  
Hewlett Packard

## PCs

1971 Intel  
1975 ALTAIR  
1976 Apple  
1977 Radio Shack TRS-80, Commodore  
1981 IBM DOS/IBM PCs  
Compaq  
Sun  
Dell  
HP

## **Handheld**

- Digital Watches
- Simple calculators
- HP55
- Blackberry
- Smartphones

## **Built-in**

- Cars
- Air Navigation
- Appliances

Enrollment situation – new section

Auditors – invite

Kiran Sivakumar

Collaboration -- more to say

This lecture:

- Computer overview

- Memories

- Main memory

- Binary/hex

- Binary integers

- Integer addition


- Integer multiplication


- Fixed point integer

- Floating point representation

- Boolean Logic/arithmetic

# Overview of a Digital Computer





## Memory Combinations:

Registers/Main Memory

ALU/Main Memory

Main Memory/Disk Controller



# Main Memory

Huge array of bytes (bits) Each byte has an address. Lets say that x is an int. x has address 0x78, y is a char. Its address is 0x95. z is a short, its address is 0xb4. Memory contains aggregates.

	<u>0</u>	<u>1</u>	<u>2</u>	<u>3</u>	<u>4</u>	<u>5</u>	<u>6</u>	<u>7</u>	<u>8</u>	<u>9</u>	<u>A</u>	<u>B</u>	<u>C</u>	<u>D</u>	<u>E</u>	<u>F</u>
<u>0</u>																
<u>1</u>																
<u>2</u>																
<u>3</u>																
<u>4</u>																
<u>5</u>																
<u>6</u>																
<u>7</u>									<u>x<sub>3</sub></u>	<u>x<sub>2</sub></u>	<u>x<sub>1</sub></u>	<u>x<sub>0</sub></u>				
<u>8</u>																
<u>9</u>						<u>y</u>										
<u>A</u>																
<u>B</u>					<u>z<sub>1</sub></u>	<u>z<sub>0</sub></u>										
<u>C</u>																
<u>D</u>																
<u>E</u>																
<u>F</u>																

IA32:  $2^{32}-1$  is maximum memory address

X86-64  $2^{64}-1$  is maximum.

## **Characters: (one byte)**

ASCII (Teletype)

EBCDIC

## **Numbers many different uses, sizes**

Integer

Floating point

Decimal?

Pointers

## **Bit matrices**

Black and white image

Color images: RGB

## **Machine instructions**

Op code, operands

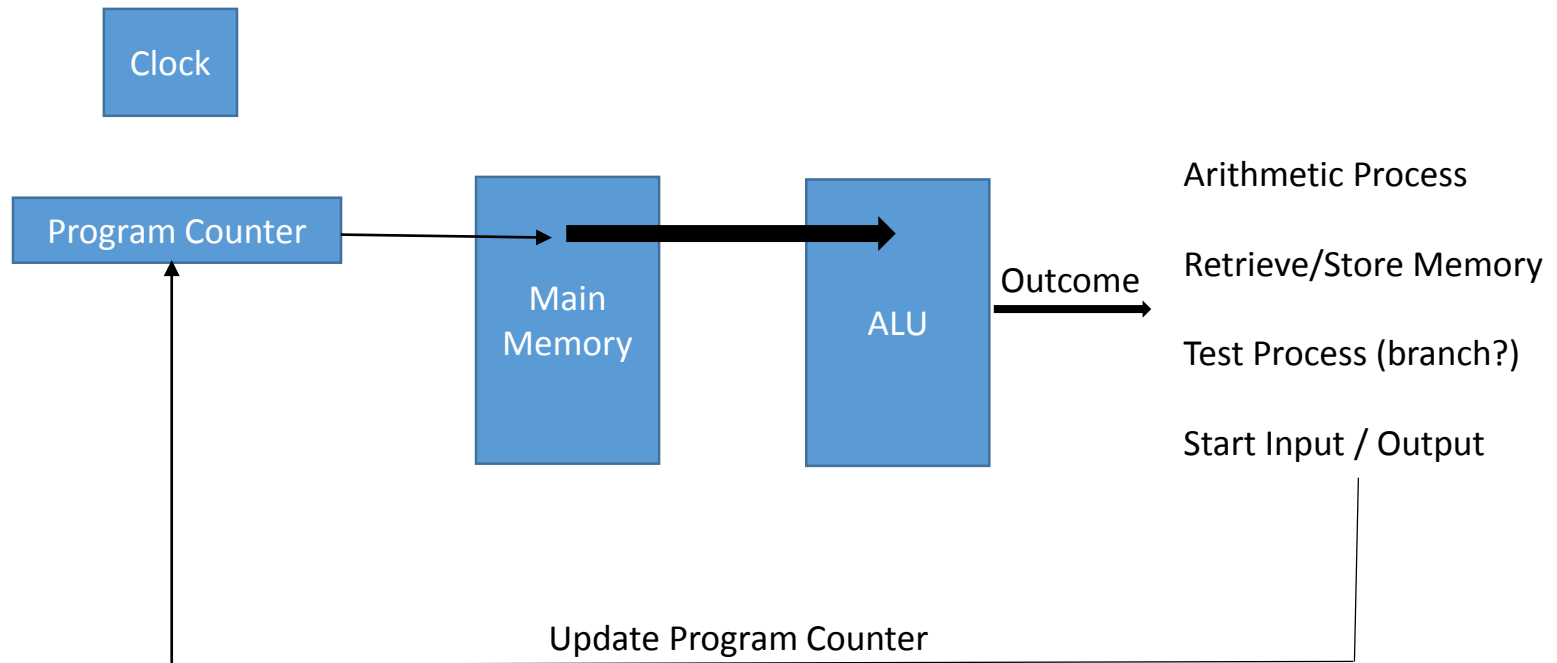
# ASCII Code Table

Dec	Hx	Oct	Char	Dec	Hx	Oct	Html	Chr	Dec	Hx	Oct	Html	Chr	Dec	Hx	Oct	Html	Chr
0	0	000	<b>NUL</b> (null)	32	20	040	&#32;	<b>Space</b>	64	40	100	&#64;	<b>@</b>	96	60	140	&#96;	<b>`</b>
1	1	001	<b>SOH</b> (start of heading)	33	21	041	&#33;	<b>!</b>	65	41	101	&#65;	<b>A</b>	97	61	141	&#97;	<b>a</b>
2	2	002	<b>STX</b> (start of text)	34	22	042	&#34;	<b>"</b>	66	42	102	&#66;	<b>B</b>	98	62	142	&#98;	<b>b</b>
3	3	003	<b>ETX</b> (end of text)	35	23	043	&#35;	<b>#</b>	67	43	103	&#67;	<b>C</b>	99	63	143	&#99;	<b>c</b>
4	4	004	<b>EOT</b> (end of transmission)	36	24	044	&#36;	<b>\$</b>	68	44	104	&#68;	<b>D</b>	100	64	144	&#100;	<b>d</b>
5	5	005	<b>ENQ</b> (enquiry)	37	25	045	&#37;	<b>%</b>	69	45	105	&#69;	<b>E</b>	101	65	145	&#101;	<b>e</b>
6	6	006	<b>ACK</b> (acknowledge)	38	26	046	&#38;	<b>&amp;</b>	70	46	106	&#70;	<b>F</b>	102	66	146	&#102;	<b>f</b>
7	7	007	<b>BEL</b> (bell)	39	27	047	&#39;	<b>'</b>	71	47	107	&#71;	<b>G</b>	103	67	147	&#103;	<b>g</b>
8	8	010	<b>BS</b> (backspace)	40	28	050	&#40;	<b>(</b>	72	48	110	&#72;	<b>H</b>	104	68	150	&#104;	<b>h</b>
9	9	011	<b>TAB</b> (horizontal tab)	41	29	051	&#41;	<b>)</b>	73	49	111	&#73;	<b>I</b>	105	69	151	&#105;	<b>i</b>
10	A	012	<b>LF</b> (NL line feed, new line)	42	2A	052	&#42;	<b>*</b>	74	4A	112	&#74;	<b>J</b>	106	6A	152	&#106;	<b>j</b>
11	B	013	<b>VT</b> (vertical tab)	43	2B	053	&#43;	<b>+</b>	75	4B	113	&#75;	<b>K</b>	107	6B	153	&#107;	<b>k</b>
12	C	014	<b>FF</b> (NP form feed, new page)	44	2C	054	&#44;	<b>,</b>	76	4C	114	&#76;	<b>L</b>	108	6C	154	&#108;	<b>l</b>
13	D	015	<b>CR</b> (carriage return)	45	2D	055	&#45;	<b>-</b>	77	4D	115	&#77;	<b>M</b>	109	6D	155	&#109;	<b>m</b>
14	E	016	<b>SO</b> (shift out)	46	2E	056	&#46;	<b>.</b>	78	4E	116	&#78;	<b>N</b>	110	6E	156	&#110;	<b>n</b>
15	F	017	<b>SI</b> (shift in)	47	2F	057	&#47;	<b>/</b>	79	4F	117	&#79;	<b>O</b>	111	6F	157	&#111;	<b>o</b>
16	10	020	<b>DLE</b> (data link escape)	48	30	060	&#48;	<b>0</b>	80	50	120	&#80;	<b>P</b>	112	70	160	&#112;	<b>p</b>
17	11	021	<b>DC1</b> (device control 1)	49	31	061	&#49;	<b>1</b>	81	51	121	&#81;	<b>Q</b>	113	71	161	&#113;	<b>q</b>
18	12	022	<b>DC2</b> (device control 2)	50	32	062	&#50;	<b>2</b>	82	52	122	&#82;	<b>R</b>	114	72	162	&#114;	<b>r</b>
19	13	023	<b>DC3</b> (device control 3)	51	33	063	&#51;	<b>3</b>	83	53	123	&#83;	<b>S</b>	115	73	163	&#115;	<b>s</b>
20	14	024	<b>DC4</b> (device control 4)	52	34	064	&#52;	<b>4</b>	84	54	124	&#84;	<b>T</b>	116	74	164	&#116;	<b>t</b>
21	15	025	<b>NAK</b> (negative acknowledge)	53	35	065	&#53;	<b>5</b>	85	55	125	&#85;	<b>U</b>	117	75	165	&#117;	<b>u</b>
22	16	026	<b>SYN</b> (synchronous idle)	54	36	066	&#54;	<b>6</b>	86	56	126	&#86;	<b>V</b>	118	76	166	&#118;	<b>v</b>
23	17	027	<b>ETB</b> (end of trans. block)	55	37	067	&#55;	<b>7</b>	87	57	127	&#87;	<b>W</b>	119	77	167	&#119;	<b>w</b>
24	18	030	<b>CAN</b> (cancel)	56	38	070	&#56;	<b>8</b>	88	58	130	&#88;	<b>X</b>	120	78	170	&#120;	<b>x</b>
25	19	031	<b>EM</b> (end of medium)	57	39	071	&#57;	<b>9</b>	89	59	131	&#89;	<b>Y</b>	121	79	171	&#121;	<b>y</b>
26	1A	032	<b>SUB</b> (substitute)	58	3A	072	&#58;	<b>:</b>	90	5A	132	&#90;	<b>Z</b>	122	7A	172	&#122;	<b>z</b>
27	1B	033	<b>ESC</b> (escape)	59	3B	073	&#59;	<b>;</b>	91	5B	133	&#91;	<b>[</b>	123	7B	173	&#123;	<b>{</b>
28	1C	034	<b>FS</b> (file separator)	60	3C	074	&#60;	<b>&lt;</b>	92	5C	134	&#92;	<b>\</b>	124	7C	174	&#124;	<b> </b>
29	1D	035	<b>GS</b> (group separator)	61	3D	075	&#61;	<b>=</b>	93	5D	135	&#93;	<b>]</b>	125	7D	175	&#125;	<b>}</b>
30	1E	036	<b>RS</b> (record separator)	62	3E	076	&#62;	<b>&gt;</b>	94	5E	136	&#94;	<b>^</b>	126	7E	176	&#126;	<b>~</b>
31	1F	037	<b>US</b> (unit separator)	63	3F	077	&#63;	<b>?</b>	95	5F	137	&#95;	<b>_</b>	127	7F	177	&#127;	<b>DEL</b>

Source: [www.LookupTables.com](http://www.LookupTables.com)


A = 0x41 = 0100 0001    a = 0x61 = 0110 0001    0 = 0x30, 1 = 0x31, ...

## How does it work?



## How does it work?

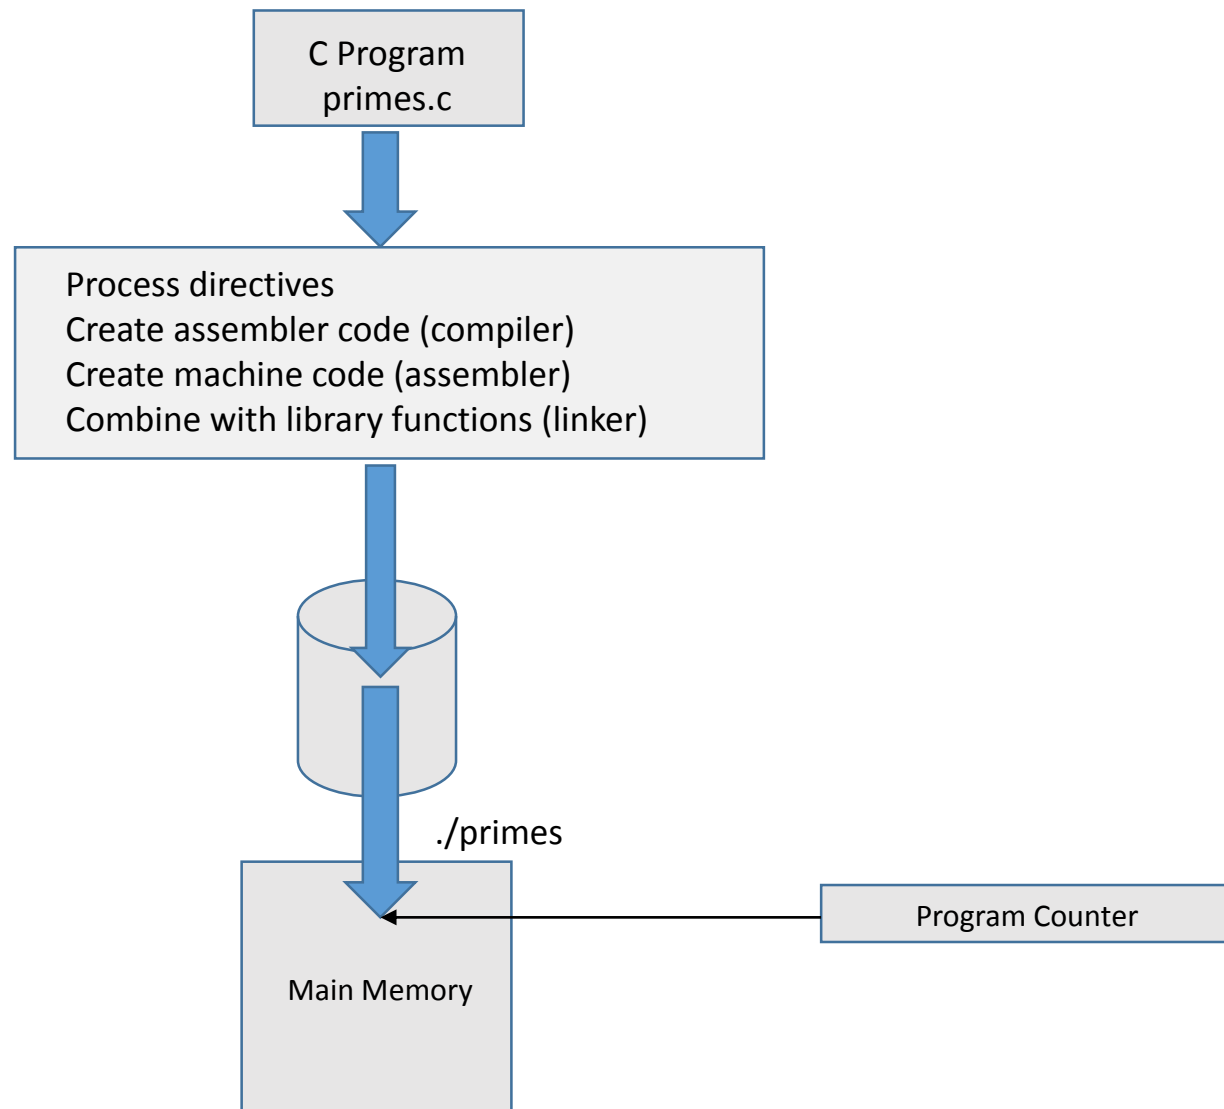
**Power On**



Program Counter = Starting Location

## How does it work?

```
gcc -o primes -lm primes.c
```



## Binary! String of 0's and 1's

01101100  
7 6 5 4 3 2 1 0

$$n \text{ bit value} = \sum_{i=0}^{n-1} x_i * 2^i \quad 2^n \text{ possible values, } -2^{n-1}:2^{n-1} \text{ signed}$$

0110    1100  
3 2 1 0    3 2 1 0

$$4 \text{ bit value} = \sum_{i=0}^3 x_i * 2^i$$

16 possible values: 0:15    0-9, A-F

$$\begin{array}{cc} 0110 = 2^2 + 2^1 = 6 & 1100 = 2^3 + 2^2 = 12 \\ \phantom{0110 = 2^2 + 2^1 = } 6 & \phantom{1100 = 2^3 + 2^2 = } C \end{array}$$

6C  
1 0

$$\text{hexadecimal value} = \sum_{i=0}^1 x_i * 16^i$$

256 possible values: 0:255, -127:127 **signed**

$$6C = 6 \cdot 16 + 12 \cdot 1 = 108$$

**Signed: most significant bit is the sign**

11101100  
7 6 5 4 3 2 1 0

$$n \text{ bit signed value} = -x_{n-1} * 2^{n-1} + \sum_{i=0}^{n-2} x_i * 2^i$$

$$11101100 = -2^7 + 2^6 + 2^5 + 2^3 + 2^2 = -128 + 64 + 32 + 8 + 4 = -20$$

Easy way to make negative: reverse bits and add 1 ... remember to carry

11101100 becomes 00010011 +00000001 becomes 00010100

Which equals  $2^4 + 2^2 = 16 + 4 = 20$



Type	bytes (X68-64)	mnemonic	description
char	1	b	Character (can also be interpreted as integer)
short	2	w	Short integer
int	4	l	Integer
long int	8	q	Long integer
long long int	8	q	Long integer
char *	8	q	Pointer
float	4	s	Single precision floating point
double	8	d	Double precision floating point
long double	16	t	Extended precision floating point

No bit data type

4 sizes: 8, 16, 32 and 64 bits

unsigned/signed

$$\text{n unsigned bit value} = \sum_{i=0}^{n-1} x_i * 2^i$$

$$\text{n bit signed value} = -x_{n-1} * 2^{n-1} + \sum_{i=0}^{n-2} x_i * 2^i$$

**\*\*show binary1\*\***

Command		Effect
	Starting and stopping	
quit		Exit gdb
run		Run your program (give command line arguments here)
kill		Stop your program
	Breakpoints	
break sum		Set breakpoint at entry to function sum
break		*0x8048394 Set breakpoint at address 0x8048394
delete 1		Delete breakpoint 1
delete		Delete all breakpoints
	Execution	
stepi		Execute one instruction
stepi 4		Execute four instructions
nexti		Like stepi, but proceed through function calls
continue		Resume execution
finish		Run until current function returns
	Examining code	
disas		Disassemble current function
disas		sum Disassemble function sum
disas 0x8048397		Disassemble function around address 0x8048397
disas 0x8048394 0x80483a4		Disassemble code within specified address range
print /x		\$eip Print program counter in hex
	Examining data	
print \$eax		Print contents of %eax in decimal
print /x \$eax		Print contents of %eax in hex
print /t \$eax		Print contents of %eax in binary
print 0x100		Print decimal representation of 0x100
print /x 555		Print hex representation of 555
print /x (\$ebp+8)		Print contents of %ebp plus 8 in hex
print *(int *) 0xffff076b0		Print integer at address 0xffff076b0
print *(int *) (\$ebp+8)		Print integer at address %ebp + 8
x/2w 0xffff076b0		Examine two (4-byte) words starting at address 0xffff076b0
x/20b sum		Examine first 20 bytes of function sum
	Useful information	
info frame		Information about current stack frame
info registers		Values of all the registers
help		Get information about gdb s: 8, 16, 32 and 64 bits

```
void main()
{
  int i = 123 ;
  int j = -123 ;
  char c[20] = "Hello World" ;
  int k[10] = { 1,2,3,4,5,6,7,8,9,10 } ;
}
```

## CS33 Fall 2014 Homework 1

Due 6PM Oct 12, 2014

Read and understand the text up to page 41. Alternatively, you can wait to start your homework until after the first lecture. The problem is about the concept of “big endian” and “little endian”, as explained in the text. The assignment is to write a program which has no inputs and its only output is “big” or “little”, depending on whether the host machine that your program is running on uses the “big” or “little” storage method.

The assignment will teach you different ways of examining what is in the computer memory, depending on the data types used.

Estimated solution time: 2 hours.

**CS33 Fall 2014 Lab1****Due Oct 19, 2014 6PM**

This assignment will help you to understand

- 1) how to convert from decimal to binary and back and
- 2) how a computer does addition and multiplication.

The C language does not accommodate a “bit” data type where you can actually input/store/change and output a single bit. For instance, in PL/I, a data type: bit(n) designates a bit string where n is the number of bits in the string. There are mechanisms for accessing any bits in the string (substr and array for example) and a constant type: '0'B and '1'B exists to set and compare bit values. In this assignment, we will work with integer/binary values so we will emulate the “bit” data type by using the C “int” data type and use only values 0 or 1 for values.

**First:** You must write two procedures:

```
void to_binary( int n, int w, int *x, int *o )  
    and  
void from_binary( int *x, int w, int *n )
```

to\_binary shall convert the decimal integer “n” into an array of ints “x” of length “w” where “x” is the binary representation of “n” using zeroes and ones. Negative numbers shall be represented in two’s complement. The error flag “o” shall denote that the integer “n” is larger than can fit into a binary representation of length “w”.

from\_binary shall convert a “binary” array of length “w” to an integer “n” whose value is the decimal value of the binary number.

**Second:** Write a procedure which takes two arrays of “binary” numbers produced by to\_binary and adds them together using **only** Boolean functions: AND, OR, XOR and NOT and outputs a “binary” array:

```
void adder( int *x, int *y, int *z, int *o, int w )
```

“o” is the overflow flag, indicating that the result will not fit in an array of length “w”, where “w” is the length of the binary arrays.

You must allow for either or both numbers to be negative according to the rules of 2's complement arithmetic.. Also, the program must set an error flag, but compute the added result if addition results in an overflow. Of course, you can use FOR loops and assignments.

**Third:** using your adder and conversion routines, write a procedure which takes two numbers and multiplies them using only Boolean functions and assignments. Set an error flag when the multiplication overflows. (This is a bit more difficult than for addition.)

```
void mult( int *x, int *y, int *z, int *o, int w )
```

Overall Note: It does not matter, in `to_binary`, whether the most or least significant bit comes out in `x[0]` or `x[w-1]` just as long as you are consistent in all of your procedures.

Hint: for both the adder and multiplier, go back to your grammar school days and recall how you learned how to do it and emulate that. Also, multiplication tables are not required since multiplying by 0 is 0 and multiplying by 1 is 1 or 0. This may not be the most sophisticated adder and multiplier but it will show you how things work.

Use the `lab1.c` template. Do not change anything in the `main()` procedure!

Estimated solution time: 8 hours.



## Convert from smaller to larger

Two ways:

```
unsigned char i ; i is 8 bits unsigned  
unsigned short j ; j is 16 bits unsigned  
j = i ;
```

assign to unsigned: zero propagation

```
1100 0001 (193) becomes 0000 0000 1100 0001 (193)  
0100 0001 (65) becomes 0000 0000 0100 0001 (65)
```

---

```
signed char i ; i is 8 bits signed  
unsigned short j ; j is 16 bits unsigned  
j = i ;
```

assign to signed: sign propagation

```
1100 0001 (-63) becomes 1111 1111 1100 0001 (65473)  
0100 0001 (65) becomes 0000 0000 0100 0001 (65)
```

no difference if j is signed

## Convert from larger to smaller

Proceed with caution! Value is truncated

convert  $x$  to  $y$  (  $y = x ;$  ) where  $x$  is  $s$  bits and  $y$  is  $t$  bits and  $t < s$

for unsigned  $x$  must be  $\leq 2^t - 1$

for signed  $y$ ,  $\text{abs}(x)$  must be  $\leq 2^{t-1} - 1$ ,  $x \geq -2^{t-1}$

signed to unsigned for width  $w$ , value changes to  $2^w + x$  if  $x < 0$

unsigned char  $j$  ;  $j$  is 8 bits signed

signed short  $i$  ;  $i$  is 16 bits unsigned

$j = i ;$

1111 1111 1100 0001 (-63) becomes 1100 0001 (193)

1111 1111 0100 0001 (-191) becomes 0100 0001 (65)

# Integer Size Conversion Summary

Examples of type conversion: Smaller to larger: signed propagates first bit, unsigned propagates zero, go to smaller: chop off. Comparisons are 32 bits.

From		To		Compare?	top : from, bottom : to
int	-16777216	uint	4278190080	1	11111111000000000000000000000000 11111111000000000000000000000000
short	-12345	int	-12345	1	1100111111000111 11111111111111111111001111111000111
short	-12345	uint	4294954951	1	1100111111000111 11111111111111111111001111111000111
short	-12345	ushort	53191	0	1100111111000111 1100111111000111
short	2345	int	2345	1	0000100100101001 000000000000000000000000100100101001
short	2345	uint	2345	1	0000100100101001 000000000000000000000000100100101001
short	2345	ushort	2345	1	0000100100101001 0000100100101001
uint	40000	int	40000	1	000000000000000000001001110001000000 000000000000000000001001110001000000

# Integer Size Conversion Summary

Examples of type conversion: Smaller to larger: signed propagates first bit, unsigned propagates zero, go to smaller: chop off. Comparisons are 32 bits.

From		To		Compare?	top : from, bottom : to
uint	40000	short	-25536	0	00000000000000001001110001000000 1001110001000000
uint	40000	ushort	40000	1	00000000000000001001110001000000 1001110001000000
ushort	53191	int	53191	1	1100111111000111 00000000000000001100111111000111
ushort	53191	short	-12345	0	1100111111000111 1100111111000111
ushort	53191	uint	53191	1	1100111111000111 00000000000000001100111111000111
uint	118727	short	-12345	0	000000000000000011100111111000111 1100111111000111
uint	118727	ushort	53191	0	000000000000000011100111111000111 1100111111000111
uint	4278190080	int	-16777216	1	11111111000000000000000000000000 11111111000000000000000000000000

# Integer Size Conversion Summary

Examples of type conversion: Smaller to larger: signed propagates first bit, unsigned propagates zero, go to smaller: chop off. Comparisons are 32 bits.

From		To		Compare?	top : from, bottom : to
uint	4278190080	short	0	0	11111111000000000000000000000000 0000000000000000
uint	4278190081	short	1	0	11111111000000000000000000000001 0000000000000001
uint	4294954951	int	-12345	1	1111111111111111111100111111000111 1111111111111111111100111111000111

# Programmatic Comparison Summary

```

int main()
{
    unsigned short k = 53191 ;           Bit pattern assigned 11001111111000111
    short l = -12345 ;                   11001111111000111
    unsigned int m = k ;                 00000000000000001100111111000111
    int n = k ;                         00000000000000001100111111000111
    unsigned int o = l ;                 11111111111111111100111111000111
    int p = l ;                         11111111111111111100111111000111

    printf( "k l ushort short %d\n", k == l ) ;   movzwl -0x14(%rbp),%edx k
                                                    movswl -0x12(%rbp),%eax l
    compare: 0                                     cmp %eax,%edx

    printf( "k m ushort uint %d\n", k == m ) ;   movzwl -0x14(%rbp),%eax k
                                                    cmp -0x10(%rbp),%eax m
    compare: 1

    printf( "k n ushort int %d\n", k == n ) ;   movzwl -0x14(%rbp),%eax k
                                                    cmp -0xc(%rbp),%eax n
    compare: 1

    printf( "k o ushort uint %d\n", k == o ) ;   movzwl -0x14(%rbp),%eax k
                                                    cmp -0x8(%rbp),%eax o
    compare: 0

    printf( "k p ushort int %d\n", k == p ) ;   movzwl -0x14(%rbp),%eax k
                                                    cmp -0x4(%rbp),%eax p
    compare: 0

```

## Programmatic Comparison Summary

```
printf( "l m short uint  %d\n", l == m );   movswl -0x12(%rbp),%eax l
                                              cmp   -0x10(%rbp),%eax m
compare: 0

printf( "l n short int   %d\n", l == n );   movswl -0x12(%rbp),%eax l
                                              cmp   -0xc(%rbp),%eax  n
compare: 0

printf( "l o short uint  %d\n", l == o );   movswl -0x12(%rbp),%eax l
                                              cmp   -0x8(%rbp),%eax  o
compare: 1

printf( "l p short int   %d\n", l == p );   movswl -0x12(%rbp),%eax l
                                              cmp   -0x4(%rbp),%eax  p
compare: 1

printf( "m n uint  int   %d\n", m == n );   mov   -0xc(%rbp),%eax  m
                                              cmp   -0x10(%rbp),%eax n
compare: 1

printf( "m o uint  uint  %d\n", m == o );   mov   -0x10(%rbp),%eax m
                                              cmp   -0x8(%rbp),%eax  o
compare: 0

printf( "m p uint  int   %d\n", m == p );   mov   -0x4(%rbp),%eax  m
                                              cmp   -0x10(%rbp),%eax p
compare: 0
}
```

## Programmatic Comparison Summary

```
printf( "n o int  uint  %d\n", n == o );      mov  -0xc(%rbp),%eax  n
                                              cmp  -0x8(%rbp),%eax  o
      compare: 0

printf( "n p int  int  %d\n", n == p );      mov  -0xc(%rbp),%eax  n
                                              cmp  -0x4(%rbp),%eax  p
      compare: 0

printf( "o p uint  int  %d\n", o == p );      mov  -0x4(%rbp),%eax  p
                                              cmp  -0x8(%rbp),%eax  o
      compare: 1
}
```

When moving from a word (16 bits) to any kind of long (32 bits) (signed or unsigned), sign propagation occurs when the short is signed.

Zero propagation occurs when the short is unsigned.

When comparing a short to a long, it first converts the short to a long with the same sign propagation rules.

Note that when two shorts are compared, it converts them both to longs ( `k == l` ).



## Unsigned

Adding x and y when they are w bits long: Size of x and y : 0 to  $2^w - 1$

$2 * (2^w - 1) = (2^{w+1} - 2)$  is w+1 bits: overflow!

Let's set w = 12 for example

	11	10	9	8	7	6	5	4	3	2	1	0	
x	0	0	0	0	0	1	1	1	1	0	1	1	123
y	0	0	0	1	0	1	0	0	0	0	0	1	321
sum	0	0	0	1	1	0	1	1	1	1	0	0	444
carry	0	0	0	0	0	1	0	0	0	0	1	1	

## How to do it?

$$\text{sum}_i = \text{mod}(x_i + y_i + \text{carry}_{i-1}, 2)$$

$$\text{carry}_i = \text{floor}((x_i + y_i + \text{carry}_{i-1})/2, 1)$$

$$\text{carry}_{i-1} = 0 \text{ when } i-1 < 0$$

## Unsigned

let's set  $w = 8$  for example,

	7	6	5	4	3	2	1	0	
x	1	1	1	0	1	0	1	0	234
y	1	1	1	0	1	0	1	0	234
sum	1	1	0	1	0	1	0	0	212
carry	1	1	1	0	1	0	1	0	

What is wrong with this picture?

What is  $2^w - 1$  ? 255.  $234 + 234 = 468 > 255$ .

$\text{carry}_7 = 1$  ! Overflow!

The actual wrong answer is  $468 - 2^w$  because we lost a bit.

## Signed

Same as unsigned except:

Adding x and y when they are w bits long: Size of x and y :  $-2^{w-1}$  to  $2^{w-1}-1$

$2^* - 2^{w-1} = -2^w$  is w bits: overflow!

	11	10	9	8	7	6	5	4	3	2	1	0	
x	0	0	0	0	0	0	0	0	0	0	0	0	0
y	1	1	1	1	1	1	1	1	1	1	1	1	-1
sum	1	1	1	1	1	1	1	1	1	1	1	1	-1
carry	0	0	0	0	0	0	0	0	0	0	0	0	

## How to do it?

$$\text{sum}_i = \text{mod}(x_i + y_i + \text{carry}_{i-1}, 2)$$

$$\text{carry}_i = \text{floor}((x_i + y_i + \text{carry}_{i-1})/2, 1)$$

$$\text{carry}_{i-1} = 0 \text{ when } i-1 < 0$$

**\*\*show binary2\*\***

How do we do *decimal* multiplication?

$$123 * 321 =$$

$$123 * (3 * 10^2 + 2 * 10^1 + 1 * 10^0) =$$

$$123 * 3 * 100 + 123 * 2 * 10 + 123 * 1 =$$

$$12300 * 3 + 1230 * 2 + 123 * 1 =$$

$$36900 + 2460 + 123 = 39483$$

$$\begin{array}{r} 123 \\ \times 321 \\ \hline 123 \\ 2460 \\ +36900 \\ \hline 39483 \end{array}$$

How do we do binary multiplication: unsigned?

$$123 * 321 = 0111\ 1011 * 1\ 0100\ 0001$$

$$0111\ 1011 * (1 * 2^8 + 1 * 2^6 + 1 * 2^0 + (\text{lots of times zero terms})) =$$

$$0111\ 1011\ 0000\ 0000 + 01\ 1110\ 1100\ 0000 + 0111\ 1011 =$$

$$1001\ 1010\ 0011\ 1011$$

$$\begin{array}{r} 0111\ 1011 \\ \times 1\ 0100\ 0001 \\ \hline 0111\ 1011 \\ 01\ 1110\ 1100\ 0000 \\ 0111\ 1011\ 0000\ 0000 \\ \hline 1001\ 1010\ 0011\ 1011 \end{array} \quad (39483)$$

Overflow problems?

$x * y$  :  $\log_2 x + \log_2 y$  is the size of the result

How do we do binary multiplication: signed?

If either multiplicand is negative, take the two's complement and proceed

$$-23 * 21$$

$$-23 = 1111\ 1110\ 1001$$

$$\text{Two's complement} = 0000\ 0001\ 0111 = 23$$

$$23 * 21 = 0001\ 0111 * 0001\ 0101$$

$$0001\ 0111 * (1 * 2^4 + 1 * 2^2 + 1 * 2^0 + (\text{lots of zero terms}) =$$

$$0001\ 0111\ 0000 + 0000\ 0101\ 1100 + 0000\ 0001\ 0111 =$$

$$1001\ 1010\ 0011\ 1011$$

$$\begin{array}{r} 0001\ 0111 \\ x\ 0001\ 0101 \\ \hline 0001\ 0111 \\ 0101\ 1100 \\ \hline 0001\ 0111\ 0000 \\ 0001\ 1110\ 0011\ (483) \end{array}$$

But must take two's complement of result: 1110 0001 1101 (-483)

Overflow problems?

**\*\*show binary3\*\***

Homework/lab submissions

Oct 22, 2104 CS33 Mid Term Topics

Closed book, open notes.

1. Integer Binary data
  - encode/decode
  - size conversion
  - operations
  - fractional fixed point
2. Floating point
  - encode/decode normalized
  - denormalized
  - operations
3. Boolean operations
  - truth tables
4. Assembly Language
  - operands
5. Stack operation



## Review decimal

$$abc.def_{10} = 10^2a + 10^1b + 10^0c + 10^{-1}d + 10^{-2}e + 10^{-3}f$$

$$123.321 = 1 \cdot 10^2 + 2 \cdot 10^1 + 3 \cdot 10^0 + 3 \cdot 10^{-1} + 2 \cdot 10^{-2} + 1 \cdot 10^{-3}$$

## Could do the same in binary

$$abc.def_2 = 2^2a + 2^1b + 2^0c + 2^{-1}d + 2^{-2}e + 2^{-3}f$$

$$101.101_2 = 2^2 \cdot 1 + 2^1 \cdot 0 + 2^0 \cdot 1 + 2^{-1} \cdot 1 + 2^{-2} \cdot 0 + 2^{-3} \cdot 1 = \\ 4 + 0 + 1 + .5 + 0 + 0.125 = 5.625_{10}$$

## Multiply by 2

$$2^3 \cdot 1 + 2^2 \cdot 0 + 2^1 \cdot 1 + 2^0 \cdot 1 + 2^{-1} \cdot 0 + 2^{-2} \cdot 1 = 8 + 0 + 2 + 1 + 0 + 0.25 = 11.25_{10}$$

\*\*\*Show binary4\*\*\*

## Problems:

Where is the binary point?

Arithmetic

Limited range of values

Rounding ?

to even

towards zero

down

up

Create “floating binary point” designated as part of the binary string

Several floating point formats IEEE-754 most accepted.

IBM “excess 64 hexadecimal”

A p bit floating point number has 3 elements: sign s: 1 bit , significand M (fractional part): n bits, exponent E: p-n-1 = k bits

M and E are n and k bit unsigned binary numbers

The value represented is  $V = -1^s \times M \times 2^E$

s	e	e	..	..	E	f	f	..	..	f
	k-1	k-2			0	n-1	n-2			0
p	p-1	p-2	..	..	p-k	p-k-1	p-k-2	..	..	0

Several floating point formats **IEEE-754** most accepted.  
IBM “excess 64 hexadecimal”

## Normalized:

$e$  is not all zeroes and not all ones. i.e.  $0 > e < 2^k - 1$

$E = e - \text{Bias}$  where  $\text{bias} = 2^{k-1} - 1$  and  $e = e_{k-1}e_{k-2} \dots e_0$

$M = 1 + f$  where  $f = f_{n-1}f_{n-2} \dots f_0$  (note that  $M = 1.f$  in binary and the 1 is implied in  $f$ )

The value represented is  $V = -1^s \times M \times 2^E$

## Denormalized:

exponent field is all zeroes

$E = 1 - \text{Bias}$

$M = f$  (no implied 1)

## Special Values:

$e$ 's are all ones and  $f$ 's are all zeroes: infinity positive or negative depending on  $s$ .

$e$ 's are all ones and  $f$ 's are not all zeroes: invalid value or Not a Number (NaN)

# Floating Point Binary Numbers

**Focus on single precision:** 32 bit: word with bits 31, 30, ... , 1, 0.

	7	6	5	4	3	2	1	0	2	2	2	1	1	1	1	1	1	1	1	1	0	9	8	7	6	5	4	3	2	1	0
3	3	2	2	2	2	2	2	2	2	2	2	1	1	1	1	1	1	1	1	1	0	9	8	7	6	5	4	3	2	1	0
1	0	9	8	7	6	5	4	3	2	1	0	9	8	7	6	5	4	3	2	1	0	9	8	7	6	5	4	3	2	1	0
s	e	e	e	e	e	e	e	e	m	m	m	m	m	m	m	m	m	m	m	m	m	m	m	m	m	m	m	m	m	m	m

Bit 31 = s = 1 when negative, 0 when positive: s or sign

Bits 30-23 = 8 (k) bit unsigned integer “excess 127”: E or exponent

Subtract 127 (bias  $2^8 - 1$ ) from actual value to get the value of e

Values 255 (infinity) and 0 (denormalized number) have reserved meaning

Bits 22-0 = 23 (m) bit unsigned fractional number: significand, mantissa or fraction:

Mantissa is actually 24 bits, supplemented by a “phantom” bit which is always 1, as if there were 24 bits 23-0 and bit 23 is 1 interpreted as:

$$M = 1.b_{22}b_{21} \dots b_1b_0$$

$$\text{which is equal to } 1 + 2^{-1}b_{22} + 2^{-2}b_{21} + \dots + 2^{-22}b_1 + 2^{-23}b_0$$

$2^e$  multiplies M (shifts the binary point + = right, - = left ), attach the sign value  $-1^s$   $2^e$  M

**\*\*show binary5\*\***

For double precision k = 12, m = 52

**Example for 16 bit:** word with bits 15, ... , 1, 0. 5 bit sign

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
	4	3	2	1	0	9	8	7	6	5	4	3	2	1	0
s	e	e	e	e	e	m	m	m	m	m	m	m	m	m	m

Bias =  $2^4 - 1 = 15$  for normalized:  $0 > \text{exponent} < 31$  ( 00001 to 11110 ),  $-14 > E < 15$

Consider bits 9-0 as a 10 bit unsigned binary number: f. Remember  $M = 1.f$  or  $m = 1+f/1024$

Example: 0 10001 10 1100 000 normalized!

exponent = 17 so  $E = 2$  (17-15)

$f = 512+128+64 = 704$  so the fractional part of f is  $512/1024+128/1024+64/1024 = 11/16$

$M = 1.6875$  (phantom) =  $16/16 + 11/16 = 27/16$

$V = 4 \times 27/16 = 27/4 = 6.75$

Example: 0 00000 10 1100 000 Denormalized!

exponent = 0 so  $E = 1 - \text{bias} = -14$

$f = 512 + 128 + 64 = 704$  so the fractional part of  $f$  is  $512/1024 + 128/1024 + 64/1024 = 11/16$

$M = .6875$  due to Denormalized rule (no phantom)

$$V = 2^{-14} \times 11/16 = 11/2^{10} = 11/1024$$

Example: 0 11111 00 0000 0000 +Infinity!

1 11111 00 0000 0000 -Infinity!

Example: 0 11111 00 1000 0000 Nan (Not a Number)

**Focus on single precision:** 32 bit: word with bits 31, 30, ... , 1, 0.

How do we add decimal? Pretend we have 7 digit limit on number of digits. Align the decimal points and add.

$$123 + .321 = 123.0000 + 000.3210$$

123.0000  
000.3210

$$1234 + .321 = 1234.000 + 0000.321$$

1234.000  
0000.321

$$12345 + .321 = 12345.00 = 00000.32$$

12345.00  
00000.321 becomes

12345.00  
00000.32      due to 7 digit limit      for  $A + B$ ,  $\text{abs}(\log_{10}A - \log_{10}B)$  digits lost from smaller addend

Same for binary. Align the binary point and add.

for  $A + B$ ,  $\text{abs}(\log_2A - \log_2B)$  bits lost from smaller addend

**\*\*show binary6\*\***

**Focus on single precision:** 32 bit: word with bits 31, 30, ... , 1, 0.

How do we multiply decimal? Pretend we have 7 digit limit on number of digits. Align the decimal points, multiply and move the decimal point

$$10 * .01 = 10.00 * 00.01$$

Binary: multiply the mantissas and add the exponents

Sum of exponents, e: must be  $\leq 254$ ,  $\geq -126$

**\*\*show binary7\*\***

Due to rounding problems, Floating point arithmetic is NOT associative

$$(A+B) + C \text{ not necessarily } = A + (B+C)$$

If A and B have grossly different logs but B and C the same.

$$(12345 + .321) + .338 = 12345.32 + .338 = 12345.68$$

$$12345 + (.321 + .338) = 12345 + .659 = 12345.66$$



## Floating Point Binary Multiplication

Example of multiplication accuracy

1	1.0000000	3f800000	2.000000e+00	40000000
1	0.1000000	3dcccccd	1.100000e+00	3f8ccccd
1	0.0100000	3c23d70a	1.010000e+00	3f8147ae
1	0.0010000	3a83126e	1.001000e+00	3f8020c5
1	0.0001000	38d1b716	1.000100e+00	3f800347
1	0.0000100	3727c5ab	1.000010e+00	3f800054
1	0.0000010	358637bc	1.000001e+00	3f800008
1	0.0000001	33d6bf93	1.000000e+00	3f800001
10	1.0000000	3f800000	1.100000e+01	41300000
10	0.1000000	3dcccccd	1.010000e+01	4121999a
10	0.0100000	3c23d70a	1.001000e+01	412028f6
10	0.0010000	3a83126e	1.000100e+01	41200419
10	0.0001000	38d1b716	1.000010e+01	41200069
10	0.0000100	3727c5ab	1.000001e+01	4120000a
10	0.0000010	358637bc	1.000000e+01	41200001
10	0.0000001	33d6bf93	1.000000e+01	41200000
100	1.0000000	3f800000	1.010000e+02	42ca0000
100	0.1000000	3dcccccd	1.001000e+02	42c83333
100	0.0100000	3c23d70a	1.000100e+02	42c8051f
100	0.0010000	3a83126e	1.000010e+02	42c80083
100	0.0001000	38d1b716	1.000001e+02	42c8000d
100	0.0000100	3727c5ab	1.000000e+02	42c80001
100	0.0000010	358637bc	1.000000e+02	42c80000
100	0.0000001	33d6bf93	1.000000e+02	42c80000

## Floating Point Binary Multiplication

Example of multiplication accuracy

1000	1.0000000	3f800000	1.001000e+03	447a4000
1000	0.1000000	3dcccccd	1.000100e+03	447a0666
1000	0.0100000	3c23d70a	1.000010e+03	447a00a4
1000	0.0010000	3a83126e	1.000001e+03	447a0010
1000	0.0001000	38d1b716	1.000000e+03	447a0002
1000	0.0000100	3727c5ab	1.000000e+03	447a0000
1000	0.0000010	358637bc	1.000000e+03	447a0000
1000	0.0000001	33d6bf93	1.000000e+03	447a0000
10000	1.0000000	3f800000	1.000100e+04	461c4400
10000	0.1000000	3dcccccd	1.000010e+04	461c4066
10000	0.0100000	3c23d70a	1.000001e+04	461c400a
10000	0.0010000	3a83126e	1.000000e+04	461c4001
10000	0.0001000	38d1b716	1.000000e+04	461c4000
10000	0.0000100	3727c5ab	1.000000e+04	461c4000
10000	0.0000010	358637bc	1.000000e+04	461c4000
10000	0.0000001	33d6bf93	1.000000e+04	461c4000
100000	1.0000000	3f800000	1.000010e+05	47c35080
100000	0.1000000	3dcccccd	1.000001e+05	47c3500d
100000	0.0100000	3c23d70a	1.000000e+05	47c35001
100000	0.0010000	3a83126e	1.000000e+05	47c35000
100000	0.0001000	38d1b716	1.000000e+05	47c35000
100000	0.0000100	3727c5ab	1.000000e+05	47c35000
100000	0.0000010	358637bc	1.000000e+05	47c35000
100000	0.0000001	33d6bf93	1.000000e+05	47c35000

# Floating Point Binary Multiplication

Example of multiplication accuracy

1000000	1.0000000	3f800000	1.000001e+06	49742410
1000000	0.1000000	3dcccccd	1.000000e+06	49742402
1000000	0.0100000	3c23d70a	1.000000e+06	49742400
1000000	0.0010000	3a83126e	1.000000e+06	49742400
1000000	0.0001000	38d1b716	1.000000e+06	49742400
1000000	0.0000100	3727c5ab	1.000000e+06	49742400
1000000	0.0000010	358637bc	1.000000e+06	49742400
1000000	0.0000001	33d6bf93	1.000000e+06	49742400
10000000	1.0000000	3f800000	1.000000e+07	4b189681
10000000	0.1000000	3dcccccd	1.000000e+07	4b189680
10000000	0.0100000	3c23d70a	1.000000e+07	4b189680
10000000	0.0010000	3a83126e	1.000000e+07	4b189680
10000000	0.0001000	38d1b716	1.000000e+07	4b189680
10000000	0.0000100	3727c5ab	1.000000e+07	4b189680
10000000	0.0000010	358637bc	1.000000e+07	4b189680
10000000	0.0000001	33d6bf93	1.000000e+07	4b189680

C allows “Boolean” (George Boole, 1847: Mathematical Analysis of Logic) operations on any data types.

## Bitwise or logical

### Truth Tables

AND &, &&

	0	1
0	0	0
1	0	1

OR |, ||

	0	1
0	0	1
1	1	1

Exclusive OR ^

	0	1
0	0	1
1	1	0

NOT ~, !

	0	1
0	1	0
1	0	1

### Examples

```
char a = 0x23 ;
char b = 0xc5 ;
```

00100011 in binary  
11000101

#### Bitwise

```
char c = a & b ;
c = a | b ;
c = a ^ b ;
c = ~a ;
```

00000001 0x01  
11100111 0xe7  
11100110 0xe6  
11011100 0xdc

#### Logical

```
c = a && b ;
c = a || b ;
c = a ^^ b ;
c = !a ;
```

00000001 0x01  
00000001 0x01  
00000000 0x00  
00000000 0x00

~ is unary.

## Precedence:

& then

^ then

|

Same precedence, left to right

**Example:**  $A \mid B \ \& \ C$      $(A \mid B) \ \& \ C$  with  $A = 1$ ,  $B = 0$ ,  $C = 0$

1		0	&	0
1		0		
1				

1		0	&	0
1			&	0
0				

## Truth Tables

**A | B & C**

A	B	C	B&C	A B&C
0	0	0	0	0
0	0	1	0	0
0	1	0	0	0
0	1	1	1	1
1	0	0	0	1
1	0	1	0	1
1	1	0	0	1
1	1	1	1	1

**( A | B ) & C**

A	B	C	A   B	(A   B) & C
0	0	0	0	0
0	0	1	0	0
0	1	0	1	0
0	1	1	1	1
1	0	0	1	0
1	0	1	1	1
1	1	0	1	0
1	1	1	1	1

$A \& (B \mid C) = (A \& B) \mid (A \& C)$       distributivity

$A \mid (B \& C) = (A \mid B) \& (A \mid C)$       associativity

$A \& B = B \& A$       reflexivity

$A \mid B = B \mid A$       reflexivity

$\sim(A \& B) = \sim A \mid \sim B$

$\sim(A \mid B) = \sim A \& \sim B$  Augustus De Morgan's laws

lots of other rules: similar to arithmetic, set theory

Notice that  $A \wedge B = (A \mid B) \& \sim(A \& B)$

**Truth Table**

A	B	$A \mid B$	$A \& B$	$\sim(A \& B)$	$A \wedge B$
0	0	0	0	1	0
0	1	1	0	1	1
1	0	1	0	1	1
1	1	1	1	0	0



## Programs are made up of machine instructions

Instructions primarily deal with registers, memory

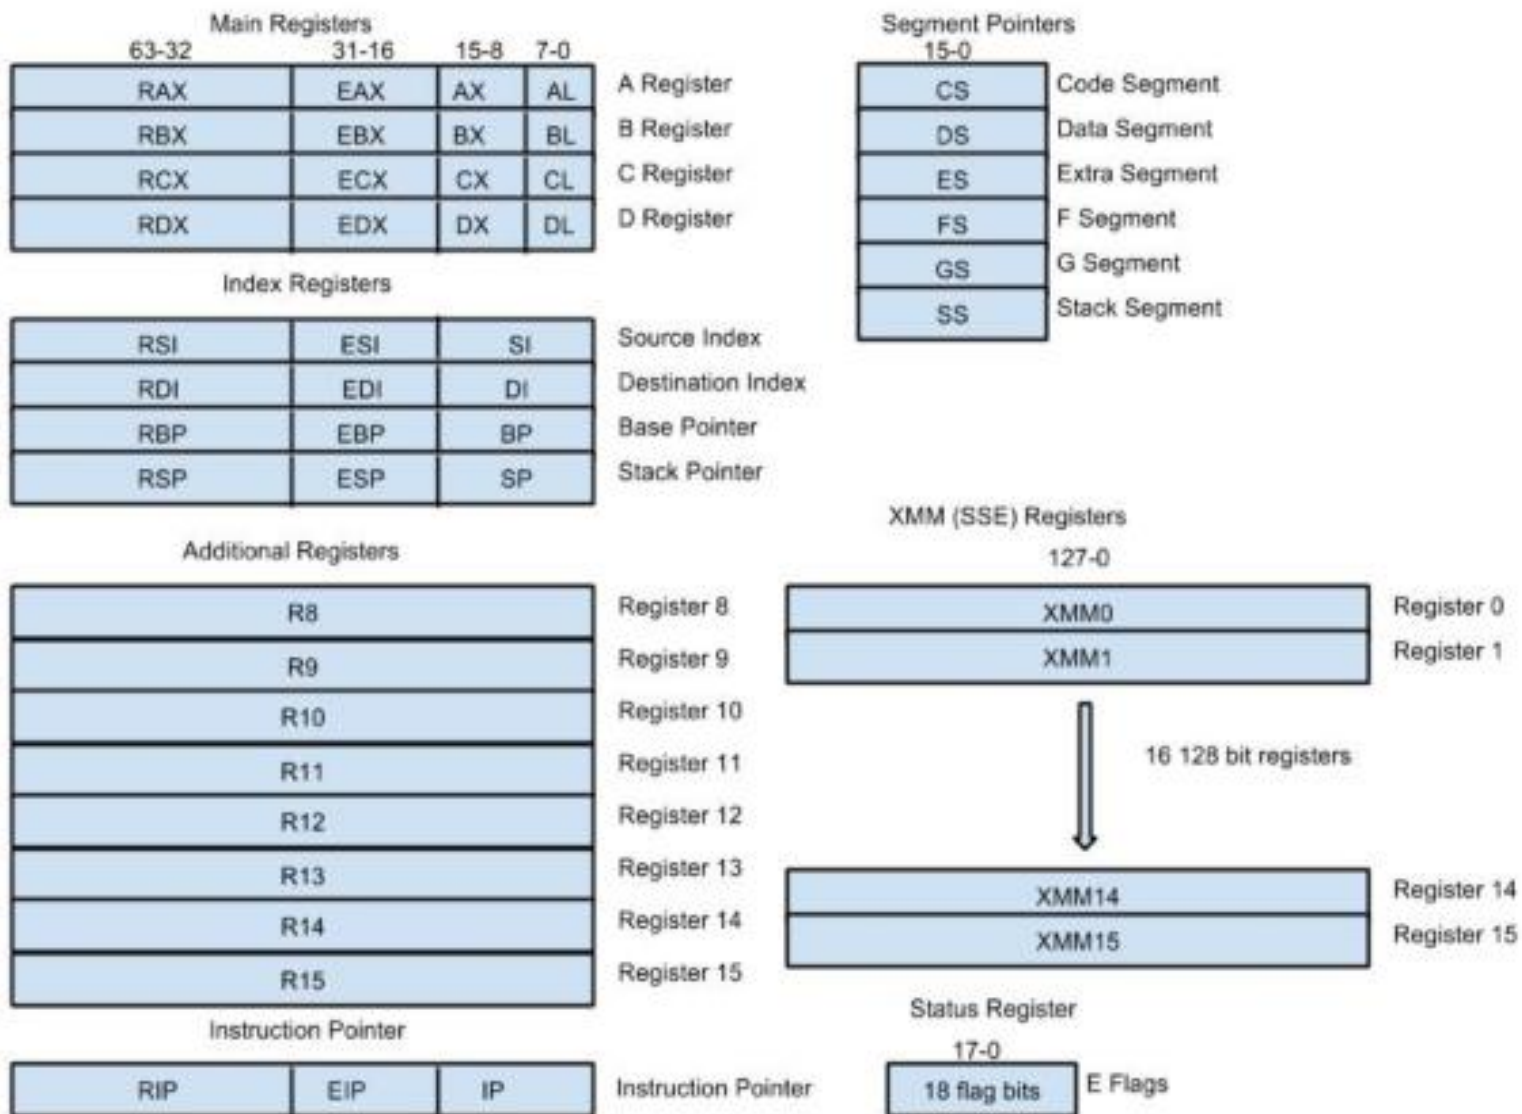
X86-64 Registers: (over IA32: added 32 bits, added general purpose, expanded FP registers)

- 4 main registers
- 4 index registers
- 8 additional registers
- instruction pointer
- 6 segment pointers
- 16 floating point registers
- status register

Memory, an array of  $2^n$  bytes, addressed from 0 to  $2^n - 1$

Aggregates come out of memory all at once.

## X86-64 registers



Huge array of bytes (bits) Each byte has an address. Lets say that x is an int. x has address 0x78, y is a char. Its address is 0x95. z is a short, its address is 0xb4.

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
0																
1																
2																
3																
4																
5																
6																
7									x <sub>3</sub>	x <sub>2</sub>	x <sub>1</sub>	x <sub>0</sub>				
8																
9						y										
A																
B					z <sub>1</sub>	z <sub>0</sub>										
C																
D																
E																
F																

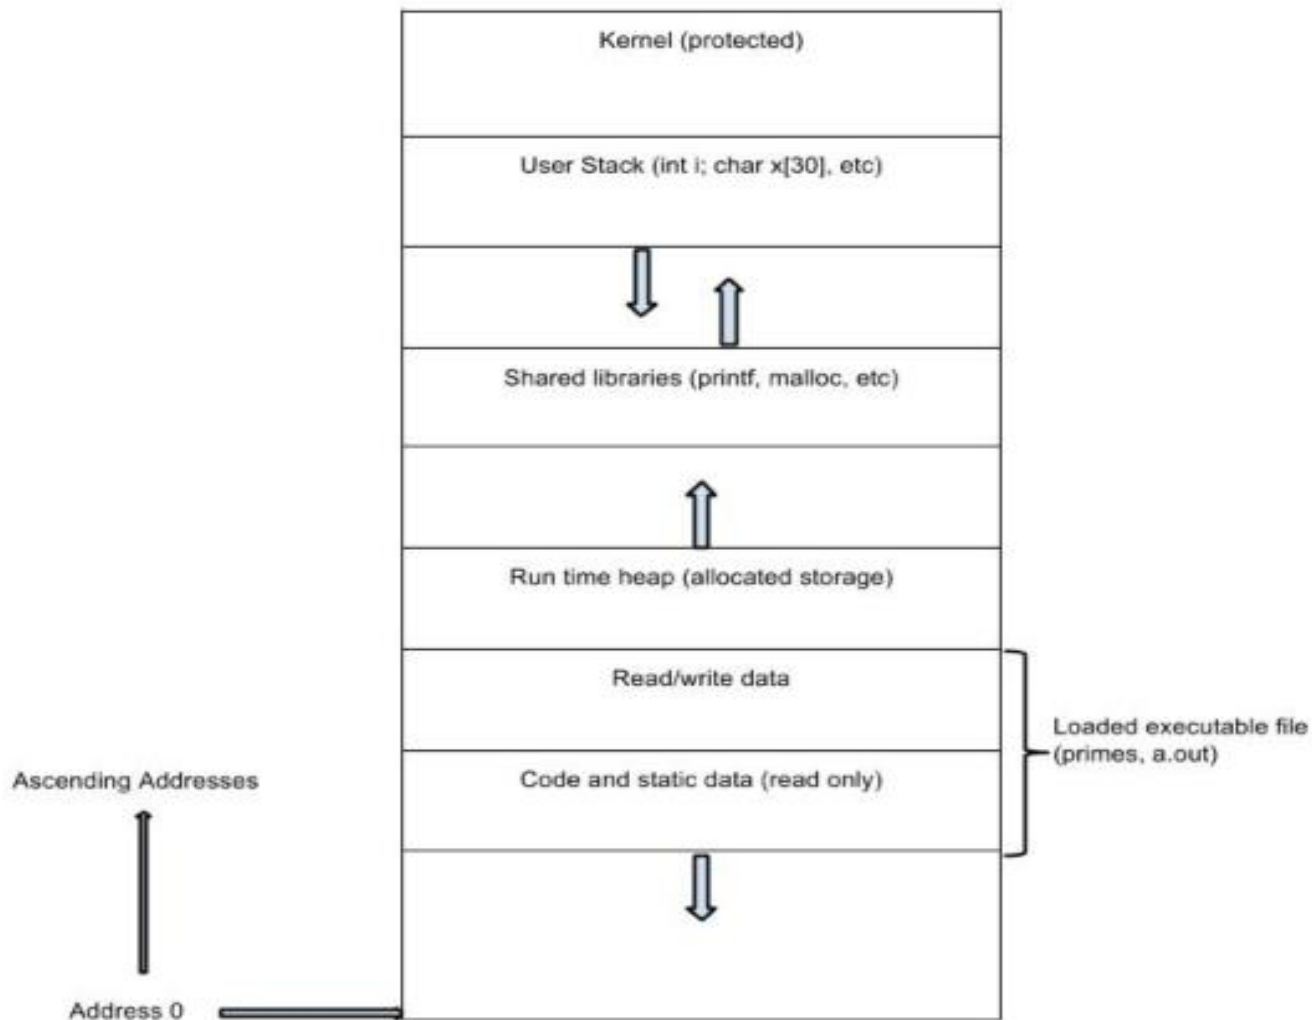
Anything but char considered as an aggregate. Some machines store with least significant on right (big endian), some on the left (little endian). Can retrieve all at once depending on the instruction.

0x12345678 stored as

$[x_3, x_2, x_1, x_0] = [0x12, 0x34, 0x56, 0x78]$  big endian

$[x_3, x_2, x_1, x_0] = [0x78, 0x56, 0x34, 0x12]$  little endian

Virtual memory: each user seems to have exclusive use of the entire memory. To one user, it looks like:



Type	bytes (X68-64)	mnemonic	description
char	1	b	Character (can also be interpreted as integer)
short	2	w	Short integer
int	4	l	Integer
long int	8	q	Long integer
long long int	8	q	Long integer
char *	8	q	Pointer
float	4	s	Single precision floating point
double	8	d	Double precision floating point
long double	16	t	Extended precision floating point

No bit data type

b = byte (8 bits)  
w = word (16 bits)  
l = long (32 bits)  
q = quad (64 bits)

## How does it work from a software point of view?

Program is just a string of operations in memory. Your program is a set of lines in a .txt (.c) file.

```
freq.c:
#include <math.h>    /* sqrt */
int frequency_of_primes (int n) {
    int i,j;
    int freq= n/2+1 ;
    .
    .
    return freq;
}

main.c:
#include <stdio.h>    /* printf */
#include <time.h>     /* clock_t, clock, CLOCKS_PER_SEC */
int main ()
{
    clock_t t;
    int f,i;

    i = 1000000 ;
    printf ("Calculating... %d ", i );
    t = clock();
    f = frequency_of_primes (i);
    printf ( "%d ",f);
    t = clock() - t;
    printf ("It took me %d clicks (%f seconds).\n",t,((float)t)/CLOCKS_PER_SEC);


    return 0;
}
```

# Transform .c to Machine Language

Pre-process, compile, assemble link.

Can stop/start process anywhere, combine any of these files using gcc or g++.  
Pre-process, compile, assemble link.

```
gcc -O1 -lm -o primes freq.c main.c
```



Book is IA32, Inxsrvt at SEAS is X86-64.

## How to specify an operand

% implies a register %rax, %eax, %ax, %al  
(64) (32) (16) (8)

\$ means “immediate” or exactly that value \$0x5: the value 5

parentheses means the value stored at that address memory (%rax)

Type	Form	Operand value	Name
Immediate	\$num	num	Immediate
Register	%rax	%rax	Register
Memory	num	(num)	Absolute
Memory	(\$rax)	(%rax)	Indirect
Memory	num(%rax)	(num+%rax)	Base+displacement
Memory	(%rax,%rbx)	(%rax+%rbx)	Indexed
Memory	num(%rax,%rbx)	(num+%rax+%rbx)	Indexed
Memory	(,%rax,s)	(%rax*s)	Scaled indexed
Memory	num(,%rax,s)	(num+%rax*s)	Scaled indexed
Memory	(%rax,%rbx,s)	(%rax+%rbx*s)	Scaled indexed
Memory	num(%rax,%rbx,s)	(num+%rax+%rbx*s)	Scaled indexed

Note: s may only be 1, 2, 4 or 8



## Operand Examples

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
0	ff	fe	fd	fc	fb	fa	f9	f8	f7	f6	f5	f4	f3	f2	f1	f0
1	fe	fd	fc	fb	fa	f9	f8	f7	f6	f5	f4	f3	f2	f1	f0	ef
2	fd	fc	fb	fa	f9	f8	f7	f6	f5	f4	f3	f2	f1	f0	ef	ee
3	fc	fb	fa	f9	f8	f7	f6	f5	f4	f3	f2	f1	f0	ef	ee	ed
4	fb	fa	f9	f8	f7	f6	f5	f4	f3	f2	f1	f0	ef	ee	ed	ec
5	fa	f9	f8	f7	f6	f5	f4	f3	f2	f1	f0	ef	ee	ed	ec	eb
6	f9	f8	f7	f6	f5	f4	f3	f2	f1	f0	ef	ee	ed	ec	eb	ea
7	f8	f7	f6	f5	f4	f3	f2	f1	f0	ef	ee	ed	ec	eb	ea	e9
8	f7	f6	f5	f4	f3	f2	f1	f0	ef	ee	ed	ec	eb	ea	e9	e8
9	f6	f5	f4	f3	f2	f1	f0	ef	ee	ed	ec	eb	ea	e9	e8	e7

Contents of memory on the left,

Assume %rax contains 0x10,  
%rbx contains 0x40

Specification	Computation	Address	Value
\$0x5	<na>	<na>	5
%rax	<na>	<na>	0x10
0x5	0x05	0x05	0xfa
(%rax)	%rax	0x10	0xfe
0x04(%rax)	0x04+%rax	0x14	0xfa
(%rax,%rbx)	%rax+%rbx	0x50	0xfa
0x04(%rax,%rbx)	0x04+%rax+%rbx	0x54	0xf6
(,%rax,4)	%rax*4	0x40	0xfb
0x05 (,%rax,2)	0x05+%rax*2	0x25	0xf8
(%rax,%rbx,2)	%rax+%rbx*2	0x90	0xf6
0x05 (%rax,%rbx,2)	0x05+%rax+%rbx*2	0x95	0xf1

## Move instructions: mov source and destination

Combinations of source and destination implied. Proper operation code determined by compiler.

immediate to register	mov	immediate,register
register to register	mov	register,register
memory to register	mov	memory,register
immediate to memory	mov	immediate,memory
register to memory	mov	register,memory

Obviously cannot move to an immediate: `mov %rax,$0x40`

Cannot move memory to memory

Moving from smaller to larger: sign extension or zero extension

<code>movs</code>	source,dest	sign extension
<code>movz</code>	source,dest	zero extension

Moving from larger to smaller, take only low order.

Move	<code>movb</code>	<code>movw</code>	<code>movl</code>	<code>movq</code>
Move, sign extension	<code>movsbw</code>	<code>movsbl</code>	<code>movswl</code>	<code>movslq</code>
Move, zero extension	<code>movzbw</code>	<code>movzbl</code>	<code>movzwl</code>	<code>movxql</code>


## Push, pop instructions: stack manipulation

push      operand

short for     $\%rsp = \%rsp - 8$       (IA32?)  
              mov operand, (%rsp)


pop      operand

short for    mov (%rsp), operand  
               $\%rsp = \%rsp + 8$



## Call instruction: stack manipulation


call            xyz  
                short for     $\text{\%rsp} = \text{\%rsp} - 8$             (IA32?)  
                                mov   $\text{\%rip}, (\text{\%rsp})$   
                                addq  $\$0x05, (\text{\%rsp})$   
                                mov  address of xyz,  $\text{\%rip}$



## Leave instruction: stack manipulation

leave


short for    `mov %rbp,%rsp`  
              `mov (%rbp),%rbp`  
              `%rsp = %rsp+8`



## ret instruction: stack manipulation

ret

short for    `mov (%rsp),%rip`  
              `%rsp = %rsp+8`



```
#include <stdio.h>
void to_sub( int *i )
{
    int j = 0x78563412 ; char y[10] ;

    int k = *i+j ;
    other_sub( &j ) ;
    printf( "%p\n", &k ) ;
}

void other_sub( int *j )
{
    int i ;
    i = *j ;
}

int main()
{
    int i; char x[30] ;

    i = 0x12345678 ;
    to_sub( &i ) ;
    return 0 ;
}

**linux interlude** stacktest.c
```

```
4004c4 <+0>: push %rbp
4004c5 <+1>: mov  %rsp,%rbp
4004c8 <+4>: sub  $0x30,%rsp
4004cc <+8>: movl $0x12345678,-0x4(%rbp)
4004d3 <+15>: lea  -0x4(%rbp),%rax
4004d7 <+19>: mov  %rax,%rdi
4004da <+22>: callq 0x4004e6 <to_sub>
4004df <+27>: mov  $0x0,%eax
4004e4 <+32>: leaveq
4004e5 <+33>: retq
```

Dump of assembler code for function main:

```
4004e6 <+0>: push %rbp
4004e7 <+1>: mov  %rsp,%rbp
4004ea <+4>: sub  $0x30,%rsp
4004ee <+8>: mov  %rdi,-0x28(%rbp)
4004f2 <+12>: movl $0x78563412,-0x4(%rbp)
4004f9 <+19>: mov  -0x28(%rbp),%rax
4004fd <+23>: mov  (%rax),%edx
4004ff <+25>: mov  -0x4(%rbp),%eax
400502 <+28>: lea  (%rdx,%rax,1),%eax
400505 <+31>: mov  %eax,-0x14(%rbp)
400508 <+34>: lea  -0x4(%rbp),%rax
40050c <+38>: mov  %rax,%rdi
40050f <+41>: callq 0x40052f <other_sub>
400514 <+46>: mov  $0x400648,%eax
400519 <+51>: lea  -0x14(%rbp),%rdx
40051d <+55>: mov  %rdx,%rsi
400520 <+58>: mov  %rax,%rdi
400523 <+61>: mov  $0x0,%eax
400528 <+66>: callq 0x4003b8 <printf@plt>
40052d <+71>: leaveq
40052e <+72>: retq
```

Dump of assembler code for function to\_sub:



## Before Call

Breakpoint 1, 4004da in main ()

(gdb) x/32w 0x7fffffffef3e0

	0x7fffffffef3e0:	0x00000000	0x00000000	0xd08908ed	0x00000032
	0x7fffffffef3f0:	0x00000000	0x00000000	0x00400560	0x00000000
	0x7fffffffef400:	0x00000000	0x00000000	0x004003a3	0x00000000
	0x7fffffffef410:	0xffffe548	0x00007fff	0x004005a5	0x00000000
%rsp →	0x7fffffffef420:	0xd080fba0	0x00000032	0x00400560	0x00000000
	0x7fffffffef430:	0x00000000	0x00000000	0x004003e0	0x00000000
	0x7fffffffef440:	0xffffe530	0x00007fff	0x00000000	0x12345678
%rbp →	0x7fffffffef450:	0x00000000	0x00000000	0xd081ed1d	0x00000032
	rip	rbp	rsp		
	4da	e450	e420		
					i

## After Call

```

4004e6 in to_sub ()
(gdb) x/32w 0x7fffffffef3e0
0x7fffffffef3e0: 0x00000000    0x00000000    0xd08908ed    0x00000032
0x7fffffffef3f0: 0x00000000    0x00000000    0x00400560    0x00000000
0x7fffffffef400: 0x00000000    0x00000000    0x004003a3    0x00000000
0x7fffffffef410: 0xffffe548    0x00007fff    0x004004df    0x00000000
0x7fffffffef420: 0xd080fba0    0x00000032    0x00400560    0x00000000
0x7fffffffef430: 0x00000000    0x00000000    0x004003e0    0x00000000
0x7fffffffef440: 0xffffe530    0x00007fff    0x00000000    0x12345678
%rbp → 0x7fffffffef450: 0x00000000    0x00000000    0xd081ed1d    0x00000032

rip      rbp      rsp
4e6      e450     e418

```

Return address, %rsp

## After push %rbp

```

                                4004e7 in to_sub ()
(gdb) x/32w 0x7fffffffef3e0
0x7fffffffef3e0: 0x00000000      0x00000000      0xd08908ed      0x00000032
0x7fffffffef3f0: 0x00000000      0x00000000      0x00400560      0x00000000
0x7fffffffef400: 0x00000000      0x00000000      0x004003a3      0x00000000
%rsp → 0x7fffffffef410: 0xffffe450      0x00007fff      0x004004df      0x00000000
0x7fffffffef420: 0xd080fba0      0x00000032      0x00400560      0x00000000
0x7fffffffef430: 0x00000000      0x00000000      0x004003e0      0x00000000
0x7fffffffef440: 0xffffe530      0x00007fff      0x00000000      0x12345678
%rbp → 0x7fffffffef450: 0x00000000      0x00000000      0xd081ed1d      0x00000032

rip      rbp      rsp      (%rbp)
4e7      e450     e410

```

## Mov %rsp,%rbp

```

                                4004ea in to_sub ()
(gdb) x/32w 0x7fffffffef3e0
0x7fffffffef3e0: 0x00000000      0x00000000      0xd08908ed      0x00000032
0x7fffffffef3f0: 0x00000000      0x00000000      0x00400560      0x00000000
%rsp  0x7fffffffef400: 0x00000000      0x00000000      0x004003a3      0x00000000
      → 0x7fffffffef410: 0xffffe450      0x00007fff      0x004004df      0x00000000
%rbp  0x7fffffffef420: 0xd080fba0      0x00000032      0x00400560      0x00000000
0x7fffffffef430: 0x00000000      0x00000000      0x004003e0      0x00000000
0x7fffffffef440: 0xffffe530      0x00007fff      0x00000000      0x12345678
0x7fffffffef450: 0x00000000      0x00000000      0xd081ed1d      0x00000032

rip      rbp      rsp
4ea      e410     e410

```

## Sub 0x30,%rsp

```

                                4004ee in to_sub ()
                                (gdb) x/32w 0x7fffffffef3e0
%rsp → 0x7fffffffef3e0: 0x00000000      0x00000000      0xd08908ed      0x00000032
                                0x7fffffffef3f0: 0x00000000      0x00000000      0x00400560      0x00000000
                                0x7fffffffef400: 0x00000000      0x00000000      0x004003a3      0x00000000
%rbp → 0x7fffffffef410: 0xffffe450      0x00007fff      0x004004df      0x00000000
                                0x7fffffffef420: 0xd080fba0      0x00000032      0x00400560      0x00000000
                                0x7fffffffef430: 0x00000000      0x00000000      0x004003e0      0x00000000
                                0x7fffffffef440: 0xffffe530      0x00007fff      0x00000000      0x12345678
                                0x7fffffffef450: 0x00000000      0x00000000      0xd081ed1d      0x00000032

                                rip      rbp      rsp
                                4ee      e410     e3e0

```

## Before leave

```

Breakpoint 2, 400528 in to_sub ()
(gdb) x/32w 0x7fffffffef3e0
%rsp → 0x7fffffffef3e0: 0x00000000    0x00000000    0xd08908ed    0x00000032
      0x7fffffffef3f0: 0x00000000    0x00000000    0x00400560    0x8a8a8a8a
      0x7fffffffef400: 0x00000000    0x00000000    0x004003a3    0x78563412
%rbp → 0x7fffffffef410: 0xffffe450    0x00007fff    0x004004df    0x00000000
      0x7fffffffef420: 0xd080fba0    0x00000032    0x00400560    0x00000000
      0x7fffffffef430: 0x00000000    0x00000000    0x004003e0    0x00000000
      0x7fffffffef440: 0xffffe530    0x00007fff    0x00000000    0x12345678
      0x7fffffffef450: 0x00000000    0x00000000    0xd081ed1d    0x00000032

      rip      rbp      rsp
      528      e410      e3e0

```

## Before ret

40052d in to\_sub ()

```
(gdb) x/32w 0x7fffffffef3e0
```

0x7fffffffef3e0:	0x00000000	0x00000000	0xd08908ed	0x00000032
0x7fffffffef3f0:	0x00000000	0x00000000	0x00400560	0x8a8a8a8a
0x7fffffffef400:	0x00000000	0x00000000	0x004003a3	0x78563412
0x7fffffffef410:	0xffffe450	0x00007fff	0x004004df	0x00000000
0x7fffffffef420:	0xd080fba0	0x00000032	0x00400560	0x00000000
0x7fffffffef430:	0x00000000	0x00000000	0x004003e0	0x00000000
0x7fffffffef440:	0xffffe530	0x00007fff	0x00000000	0x12345678
0x7fffffffef450:	0x00000000	0x00000000	0xd081ed1d	0x00000032

%rbp →

rip	rbp	rsp
52d	e450	e418

%rsp →

## After ret

```

                                4004df in main ()
(gdb) x/32w 0x7fffffffef3e0
0x7fffffffef3e0: 0x00000000      0x00000000      0xd08908ed      0x00000032
0x7fffffffef3f0: 0x00000000      0x00000000      0x00400560      0x8a8a8a8a
0x7fffffffef400: 0x00000000      0x00000000      0x004003a3      0x78563412
0x7fffffffef410: 0xffffef450      0x00007fff      0x004004df      0x00000000
%rsp → 0x7fffffffef420: 0xd080fba0      0x00000032      0x00400560      0x00000000
0x7fffffffef430: 0x00000000      0x00000000      0x004003e0      0x00000000
0x7fffffffef440: 0xffffef530      0x00007fff      0x00000000      0x12345678
%rbp → 0x7fffffffef450: 0x00000000      0x00000000      0xd081ed1d      0x00000032

rip      rbp      rsp
4df      e450     e420

```



Summary of 4 special move instructions:

call xyz:

3 steps:  $\text{\%rsp} = \text{\%rsp} - 8$ ,  $(\text{\%rsp}) = \text{\%rip} + 5$ ,  $\text{\%rip} = \text{address}(\text{xyz})$

push %rbp

2 steps:  $\text{\%rsp} = \text{\%rsp} - 8$ ,  $(\text{\%rsp}) = \text{\%rbp}$

leave

3 steps:  $\text{\%rsp} = \text{\%rbp}$ ,  $\text{\%rbp} = (\text{\%rbp})$ ,  $\text{\%rsp} = \text{\%rsp} + 8$

ret

2 steps:  $\text{\%rip} = (\text{\%rsp})$ ,  $\text{\%rsp} = \text{\%rsp} + 8$

An aside:  $\text{\%rbx}$  and  $\text{\%r12}$  through  $\text{\%r15}$  must be preserved across a function call. If they are used, they should be saved in the stack and restored before leave.

# Tentative Lecture Schedule

Lecture	Date	Topic
1	10/6/2014	Introduction/Tour of Systems
2	10/8/2014	Data Representation
3	10/13/2014	Data Representation
4	10/15/2014	Machine Language
5	10/20/2014	Assembly Language
6	10/22/2014	Midterm*
7	10/27/2014	Code Optimization
8	10/29/2014	Code Optimization
9	11/3/2014	Memory Hierarchy
10	11/5/2014	Caching
11	11/10/2014	Caching
12	11/12/2014	Midterm*
13	11/17/2014	Linking
14	11/19/2014	Exception Control Flow
15	11/24/2014	Virtual Memory
16	11/26/2014	Virtual Memory
17	12/1/2014	Concurrent Programming
18	12/3/2014	Concurrent Programming
19	12/8/2014	Review
20	12/10/2014	Review
	12/16/2014	Final Exam 8AM-11AM
	11/28/2014	No Discussion - Thanksgiving Holiday
		* No office hour

# Tentative Lecture Schedule

Lecture	Date	Topic
1	10/6/2014	Introduction/Tour of Systems
2	10/8/2014	Data Representation
3	10/13/2014	Data Representation
4	10/15/2014	Machine Language
5	10/20/2014	Assembly Language
6	10/22/2014	Midterm*
7	10/27/2014	Assembly Language
8	10/29/2014	Code Optimization
9	11/3/2014	Code Optimization
10	11/5/2014	Memory Hierarchy
11	11/10/2014	Caching
12	11/12/2014	Caching
13	11/17/2014	Midterm*
14	11/19/2014	Linking
15	11/24/2014	Exception Control Flow
16	11/26/2014	Virtual Memory
17	12/1/2014	Virtual Memory
18	12/3/2014	Concurrent Programming
19	12/8/2014	Concurrent Programming
20	12/10/2014	Review
	12/16/2014	Final Exam 8AM-11AM
	11/28/2014	No Discussion - Thanksgiving Holiday
		* No office hour

**Homework 2**

**Lab 2**

## Arithmetic and Logical Operations

address	leal	memory,register	load effective address (address arithmetic,destination only a register)
unary	inc	register or memory	increment
	dec	register or memory	decrement
	neg	register or memory	negate
	not	register or memory	complement
arithmetic	add	memory or register,register	add
	sub	memory or register,register	subtract
	imul	memory or register,register	integer multiply
	idiv	memory or register	integer divide (divides RDX:RAX by source)
logical	xor	memory or register,register	bitwise exclusive or
	or	memory or register,register	bitwise or
	and	memory or register,register	bitwise and
shift	sal	immediate or one byte register,memory or register	left arithmetic shift (fill right with zeroes)
	shl	immediate or one byte register,memory or register	left logical shift (sal) (fill right with zeroes)
	sar	immediate or one byte register,memory or register	right arithmetic shift (fill left with sign bit)
	shr	immediate or one byte register,memory or register	right logical shift (fill left with zeroes)

only register %cl allowed for register operand

## Practice Problem -- leal

assume %rax contains x, %rbx contains y

	result
leal 6(%rax), %rcx	$x+6$
leal (%rax,%rbx), %rcx	$x+y$
leal (%rax,%rbx,4), %rcx	$x+4y$
leal 7(%rax,%rax,8), %rcx	$x+8x+7$
leal 0x0a(%rbx,4), %rcx	$4y+10$
leal 9(%rax,%rbx,2), %rcx	$x+2y+9$

## Practice Problem -- shifts

assume %rax contains 0x800000000000000f (8 bytes) = 1000 0000 0000 0000 ... 0000 0000 0000 1111 (64 bits)

	result
sal \$2,%rax	0x0000000000000003c
shl \$2,%rax	0x0000000000000003c
sar \$2,%rax	0xe0000000000000003
shr \$2,%rax	0x20000000000000003
sal \$63,%rax	0x80000000000000000
sar \$63,%rax	0xFFFFFFFFFFFFFFFFF
leal \$0x1,%rax	
sal \$63,%rax	
sar \$63,%rax	0xFFFFFFFFFFFFFFFFF
and 0xffffffe,%rax	
sal \$63,%rax	0x00000000000000000 // no matter whar is in %rax

using %eax allows only shifts from 0-31 positions

## Practice Problem -- convert multiply to shifts

assume x is 10

$x * 17$	$17 = 16 + 1$	$(x \ll 4) + x$	$10 + 160 = 170$
$x * -7$	$-7 = -4 - 2 - 1$	$-(x \ll 2) - (x \ll 1) - x$	$-40 - 20 - 10 = -70$
	also	$-(x \ll 3) + x$	$-80 + 10$
$x * 60$	$60 = 32 + 16 + 8 + 4$	$(x \ll 5) + (x \ll 4) + (x \ll 3) + (x \ll 2)$	$320 + 160 + 80 + 40 = 600$
	also	$(x \ll 6) - (x \ll 2)$	$640 - 40$
$x * -112$	$-112 = -64 - 32 - 16$	$-(x \ll 6) - (x \ll 5) - (x \ll 4)$	$-640 - 320 - 160 = -1120$
	also	$-(x \ll 7) + (x \ll 4)$	$-1280 + 160$

What is  $x \ll 4 + x$ ? =  $x \ll (4 + x) = x \ll 14 = 163840$  ! shifts have lowest precedence



## Practice Problem -- arithmetic

assume	Address	Value	Register	Value
	0x100	0xFF	%eax	0x100
	0x104	0xAB	%ecx	0x1
	0x108	0x13	%edx	0x3
	0x10C	0x11		

Instruction	Destination	Value	
addl %ecx,%eax)	0x100	0x101	
subl %edx,4(%eax)	0x104	0xA8	
imull \$16,(%eax,%edx,4)	0x10c	0x110	0x11 = 0001 0001 * 16 = 0001 0001 0000
incl 8(%eax)	0x108	0x14	
decl %ecx	%ecx	0x0	
subl %edx,%eax	%eax	0xFD	

## Status register, set after arithmetic instructions.

### Status codes:

for unsigned int t, a, b and after say,  $t = a + b$

CF: Carry Flag. The most recent operation generated a carry out of the most significant bit. Used to detect overflow for unsigned operations.  
 $(\text{unsigned})\ t < (\text{unsigned})\ a$

ZF: Zero Flag. The most recent operation yielded zero.  
 $t == 0$

for signed int t, a, b and after  $t = a + b$

SF: Sign Flag. The most recent operation yielded a negative value.  
 $t < 0$

OF: Overflow Flag. The most recent operation caused a two's-complement overflow—either negative or positive.  
 $(a < 0 == b < 0) \ \&\& \ (t < 0 != a < 0)$

Show condtest.c

## Compare and test.

all sizes but must be the same. In C, comparing two different sizes, they must first be made the same (larger) size and this depends on signed/unsigned.

arithmetic    `cmp`            memory or register (S2), memory or register (S1)    set code depending on  $S_1 - S_2$

logical        `test`            memory or register (S2), memory or register (S1)    set code depending on  $S_1 \& S_2$

`test %eax,%eax`             $S_1 \& S_1 = S_1$  ! But status is set depending on  $<, =, > 0$ .

`set`            saves various Boolean combinations of the condition register

## Combinations of tests in one instruction.

Destination is one byte (low order byte if register)

			Condition	
-	sete	memory or register	ZF	equal / zero
-	setne	memory or register	$\sim ZF$	not equal / not zero
-	sets	memory or register	SF	negative
-	setns	memory or register	$\sim SF$	not negative
-	setg	memory or register	$\sim (SF \wedge OF) \ \& \ \sim ZF$	greater (signed > )
-	setge	memory or register	$\sim (SF \wedge OF)$	greater or equal (signed $\geq$ )
-	setl	memory or register	$SF \wedge OF$	less (signed < )
-	setle	memory or register	$\sim (SF \wedge OF) \mid ZF$	less or equal (signed $\leq$ )
-	seta	memory or register	$\sim CF \ \& \ \sim ZF$	above (unsigned > )
-	setae	memory or register	!CF	above or equal (unsigned $\geq$ )
-	setb	memory or register	CF	below (unsigned < )
-	setbe	memory or register	$CF \mid ZF$	below or equal (unsigned $\leq$ )

## Status register, set after arithmetic instructions.

### Status codes:

logical

CF: Carry Flag. The most recent operation generated a carry out of the most significant bit. Used to detect overflow for unsigned operations.  
(unsigned)  $t < (\text{unsigned}) a$  when doing  $t = a + b$

ZF: Zero Flag. The most recent operation yielded zero.  
 $t == 0$

arithmetic

SF: Sign Flag. The most recent operation yielded a negative value.  
 $t < 0$

OF: Overflow Flag. The most recent operation caused a two's-complement overflow—either negative or positive.  
 $(a < 0 == b < 0) \ \&\& \ (t < 0 != a < 0)$  when doing  $t = a + b$

Show condtest.c

## Compare and test.

Jump (replaces %rip) goes to the location in the instruction or memory or register (\*operand) if the condition matches. Goes to the next instruction, if not.

		Synonym	Condition	
-	jmp		1	always
-	je	jz	ZF	equal / zero
-	jne	jnz	$\sim$ ZF	not equal / not zero
-	js		SF	negative
-	jns		$\sim$ SF	nonnegative
-	jg	jnle	$\sim$ (SF ^ OF) & $\sim$ ZF	greater (signed >)
-	jge	jnl	$\sim$ (SF ^ OF)	greater or equal (signed >=)
-	jl	jnge	SF ^ OF	less (signed <)
-	jle	jng	(SF ^ OF)   ZF	less or equal (signed <=)
-	ja	jnbe	$\sim$ CF & $\sim$ ZF	above (unsigned >)
-	jae	jnb	$\sim$ CF	above or equal (unsigned >=)
-	jb	jnae	CF	below (unsigned <)
-	jbe	jna	CF   ZF	below or equal (unsigned <=)

**IF**  
**then**  
**else**

GoToLess Programming?

```
    if ( <condition> )  
        statement if true ;  
    else  
        statement if false ;
```

can leave off the else

translates to in !GoToLess programming

```
    if ( !<condition> )  
        goto false_part ;  
true_part:  
    statement if true ;  
    goto do done ;  
false_part:  
    statement if false ;  
done: ;
```

It compiles into exactly the !GoToLess way. (jmp is the assembly language equivalent of goto)

## Compilation?

```

int main()
{
    int i = 1, j = 2, a = 3, b = 4, result = 0;

    if ( i < j )
        result = a;
    else
        result = b;

    return 0;
}

```

```

400478 <+4>:  movl  $0x1,-0x14(%rbp) // i
40047f <+11>: movl  $0x2,-0x10(%rbp) // j
400486 <+18>: movl  $0x3,-0xc(%rbp) // a
40048d <+25>: movl  $0x4,-0x8(%rbp) // b
400494 <+32>: movl  $0x0,-0x4(%rbp) // result
40049b <+39>: mov   -0x14(%rbp),%eax // i
40049e <+42>: cmp   -0x10(%rbp),%eax // compare to j
4004a1 <+45>: jge   0x4004ab <main+55> // if >= goto false_part
4004a3 <+47>: mov   -0xc(%rbp),%eax // true_part: a
4004a6 <+50>: mov   %eax,-0x4(%rbp) // result
4004a9 <+53>: jmp   0x4004b1 <main+61> // goto done
4004ab <+55>: mov   -0x8(%rbp),%eax // false_part: b
4004ae <+58>: mov   %eax,-0x4(%rbp) // result
4004b1 <+61>:          // done:

```



## Do While

GoToLess

```
do {  
    statement(s) }  
while ( <condition> );
```

!GoToLess

loop:

```
    statement(s)  
    If ( <condition> )  
        goto loop ;
```

Test is after statement(s) so goes through at least once.

## Compilation

```
int main()
{
    int n = 20, result = 1 ;

    do {
        result = result*n ;
        n = n-1 ;
    } while ( n > 1 ) ;

    return 0 ;
}
```

```
400478 <+4>:  movl  $0x14,-0x8(%rbp) // n
40047f <+11>:  movl  $0x1,-0x4(%rbp) // result
400486 <+18>:  mov   -0x4(%rbp),%eax // result
400489 <+21>:  imul  -0x8(%rbp),%eax // times n
40048d <+25>:  mov   %eax,-0x4(%rbp) // result
400490 <+28>:  subl  $0x1,-0x8(%rbp) // minus 1
400494 <+32>:  cmpl  $0x1,-0x8(%rbp) // <condition>
400498 <+36>:  jg    0x400486 <main+18>
```

## While

While statement

GoToLess

```
while ( <condition> ) {  
    statement(s) }
```

!GoToLess

loop:

```
    If (! <condition> )  
        goto done ;  
    statement(s)  
    goto loop ;
```

done ;

Test is before statement(s) so possibly goes through not at all.

## Compilation

```
int main()
{
    int n = 20, result = 1 ;

    while ( n > 1 ) {
        result = result*n ;
        n = n-1 ;
    }

    return 0 ;
}
```

```
400478 <+4>:  movl  $0x14,-0x8(%rbp)
40047f <+11>:  movl  $0x1,-0x4(%rbp)
400486 <+18>:  jmp   0x400496 <main+34>
400488 <+20>:  mov   -0x4(%rbp),%eax
40048b <+23>:  imul  -0x8(%rbp),%eax
40048f <+27>:  mov   %eax,-0x4(%rbp)
400492 <+30>:  subl  $0x1,-0x8(%rbp)
400496 <+34>:  cmpl  $0x1,-0x8(%rbp)
40049a <+38>:  jg    0x400488 <main+20>
```

## For

for statement is the same as a while statement

### GoToLess

```
for( init-expr; test-expr; update-expr ) {  
    statement(s) }
```

can be rewritten

```
init-expr ;  
while ( test-expr ) {  
    statement(s)  
    update-expr ;  
}
```

### !GoToLess

```
init-expr ;  
loop:  
    if ( !test-expr )  
        goto done ;  
    statements(s)  
    update-expr ;  
    goto loop ;  
done:
```

but we will see that it is actually implemented as:

```
init-expr ;  
goto test ;  
loop:  
    statements(s)  
    update-expr ;  
test:  
    if ( test-expr )  
        goto loop ;
```

## Compilation

```

int main()
{
    int n = 20, result = 1, i ;

    for( i=n; i>1; i-- )
    {
        result = result*n ;
        n = n-1 ;
    }

    return 0 ;
}

```

```

400478 <+4>:  movl  $0x14,-0xc(%rbp) // n
40047f <+11>:  movl  $0x1,-0x8(%rbp) // result
400486 <+18>:  mov   -0xc(%rbp),%eax // init-expr
400489 <+21>:  mov   %eax,-0x4(%rbp) // i
40048c <+24>:  jmp   0x4004a0 <main+44>
40048e <+26>:  mov   -0x8(%rbp),%eax // result
400491 <+29>:  imul  -0xc(%rbp),%eax // n
400495 <+33>:  mov   %eax,-0x8(%rbp) // result
400498 <+36>:  subl  $0x1,-0xc(%rbp) // update-expr
40049c <+40>:  subl  $0x1,-0x4(%rbp) // i-1
4004a0 <+44>:  cmpl  $0x1,-0x4(%rbp) // test-expr
4004a4 <+48>:  jg    0x40048e <main+26>

```

## cmov

Added to avoid branch prediction errors. Move only if true

	Instruction	Synonym		
-	cmove	cmovz	ZF	Equal / zero
-	cmovne	cmovnz	$\sim$ ZF	Not equal / not zero
-	cmovs		SF	Negative
-	cmovns		$\sim$ SF	Nonnegative
-	cmovg	cmovnle	$\sim$ (SF ^ OF) & $\sim$ ZF	Greater (signed >)
-	cmovge	cmovnl	$\sim$ (SF ^ OF)	Greater or equal (signed >=)
-	cmovl	cmovnge	SF ^ OF	Less (signed <)
-	cmovle	cmovng	(SF ^ OF)   ZF	Less or equal (signed <=)
-	cmova	cmovnbe	$\sim$ CF & $\sim$ ZF	Above (unsigned >)
-	cmovae	cmovnb	$\sim$ CF	Above or equal (Unsigned >=)
-	cmovb	cmovnae	CF	Below (unsigned <)
-	cmovbe	cmovna	CF	ZF below or equal (unsigned <=)

Carefully read pages 208-212.

comparison expression ? true statement : false statement

$x < y ? y - x : x - y$

Two ways to do this

```
if ( x < y )
    ret = y-x ;
else
    ret = x-y ;
return ret ;
```

Or:

```
reta = y-x ;
retb = x-y ;
test = x<y ;
if (test) retb = reta ; // conditional move undoes retb = x-y ;
return retb ;
```



## Compilation

```
int absdiff( int x, int y )
```

```
{
  return x < y ? y-x : x-y ;
}
```

```
400493 <+10>:  mov  -0x4(%rbp),%eax
400496 <+13>:  cmp   -0x8(%rbp),%eax      // x:y
400499 <+16>:  jge   0x4004a9 <absdiff+32>
40049b <+18>:  mov   -0x4(%rbp),%eax      // true: x
40049e <+21>:  mov   -0x8(%rbp),%edx      // y
4004a1 <+24>:  mov   %edx,%ecx
4004a3 <+26>:  sub   %eax,%ecx
4004a5 <+28>:  mov   %ecx,%eax
4004a7 <+30>:  jmp   0x4004b5 <absdiff+44>
4004a9 <+32>:  mov   -0x8(%rbp),%eax      // false: y
4004ac <+35>:  mov   -0x4(%rbp),%edx      // x
4004af <+38>:  mov   %edx,%ecx
4004b1 <+40>:  sub   %eax,%ecx
4004b3 <+42>:  mov   %ecx,%eax
4004b5 <+44>:  leaveq                // done
4004b6 <+45>:  retq
```

-O

```
400474 <+0>:  mov   %esi,%eax
400476 <+2>:  sub   %edi,%eax  // x-y
400478 <+4>:  mov   %edi,%edx
40047a <+6>:  sub   %esi,%edx  // y-x
40047c <+8>:  cmp   %esi,%edi  // compare x:y
40047e <+10>:  cmovge %edx,%eax
400481 <+13>:  retq
```

## Explicit list of cases

The formal definition is:

```
switch ( <expression> )
{
  case <constant1> :
    statments1    // executed when <expression> = <constant1>
                  // if statements end with break ; goto the end
                  // otherwise, execute the next set of statements
  case <constant2> :
    statements2    // statements can be null, in which case, the next set of statements apply
    .
    .
  default :
    default statements
}
```

## Explicit list of cases

```

int switch_eg(int x, int n) {
    int result = x;

    switch (n) {

    case 100:
        result *= 13;    // x * 13
        break;

    case 102:
        result += 10;    // x + 10
        /* Fall through */

    case 103:
        result += 11;
        break;

    case 104:
    case 106:
        result *= result;
        break;

    default:
        result = 0;
    }

    return result;
}

```

Case by case outcome?

```

<100
100
101
102
103
104
105
106
>106

```

## Explicit list of cases

```

4004f3 <+4>:  mov  %edi,-0x14(%rbp) // x
4004f6 <+7>:  mov  %esi,-0x18(%rbp) // n
4004f9 <+10>: mov  -0x14(%rbp),%eax
4004fc <+13>: mov  %eax,-0x4(%rbp) // result
4004ff <+16>: mov  -0x18(%rbp),%eax
400502 <+19>: sub  $0x64,%eax          // n-100
400505 <+22>: cmp  $0x6,%eax
400508 <+25>: ja   0x40053f <switch_eg+80> // jump above: unsigned
40050a <+27>: mov  %eax,%eax          // ?
40050c <+29>: mov  0x400650(,%rax,8),%rax // 0x400650+%rax*8
400514 <+37>: jmpq  *%rax             // indirect unconditional
400516 <+39>: mov  -0x4(%rbp),%edx     // 100
400519 <+42>: mov  %edx,%eax
40051b <+44>: add  %eax,%eax
40051d <+46>: add  %edx,%eax          // shift and add to
40051f <+48>: shl  $0x2,%eax          // multiply by 13
400522 <+51>: add  %edx,%eax
400524 <+53>: mov  %eax,-0x4(%rbp)
400527 <+56>: jmp  0x400546 <switch_eg+87> // break
400529 <+58>: addl $0xa,-0x4(%rbp)     // 102
40052d <+62>: addl $0xb,-0x4(%rbp)     // 103
400531 <+66>: jmp  0x400546 <switch_eg+87> // break
400533 <+68>: mov  -0x4(%rbp),%eax     // 104,106
400536 <+71>: imul -0x4(%rbp),%eax
40053a <+75>: mov  %eax,-0x4(%rbp)
40053d <+78>: jmp  0x400546 <switch_eg+87> // break
40053f <+80>: movl $0x0,-0x4(%rbp)     // default:
400546 <+87>: mov  -0x4(%rbp),%eax

```

## Branch Table

```

400650 <+39> // 100
400658 <+80> // 101
400660 <+58> // 102
400668 <+62> // 103
400670 <+68> // 104
400678 <+80> // 105
400680 <+68> // 106

```

## Example

```
int called(int a, int b, int *c, int d, int e, int f) ;
void call2() ;
main()
{
    int a, b, c, d, e, f ;
    called( a,b,&c,d,e,f ) ;
}
```

```
int called(int a, int b, int *c, int d, int e, int f)
{
    char x[30] ;
    int result = a+b+*c+d+e+f ;
    x[0] = 'a' ;
    call2() ;
    return result;
}
```

```
void call2()
{
    int a,b,c ;

    c = a+b ;
}
```

## Assembly language

Dump of assembler code for function main:

```

400474 <+0>:  push %rbp
400475 <+1>:  mov  %rsp,%rbp
400478 <+4>:  push %rbx
400479 <+5>:  sub  $0x28,%rsp
40047d <+9>:  mov  -0x14(%rbp),%edi
400480 <+12>: mov  -0x18(%rbp),%esi
400483 <+15>: mov  -0x1c(%rbp),%ecx
400486 <+18>: lea  -0x28(%rbp),%rdx
40048a <+22>: mov  -0x20(%rbp),%ebx
40048d <+25>: mov  -0x24(%rbp),%eax
400490 <+28>: mov  %edi,%r9d
400493 <+31>: mov  %esi,%r8d
400496 <+34>: mov  %ebx,%esi
400498 <+36>: mov  %eax,%edi
40049a <+38>: callq 0x4004a6 <called>
40049f <+43>: add  $0x28,%rsp
4004a3 <+47>: pop  %rbx
4004a4 <+48>: leaveq
4004a5 <+49>: retq

```

## Assembly language

Dump of assembler code for function called:

```

4004a6 <+0>:  push %rbp
4004a7 <+1>:  mov  %rsp,%rbp
4004aa <+4>:  sub  $0x50,%rsp
4004ae <+8>:  mov  %edi,-0x34(%rbp)
4004b1 <+11>: mov  %esi,-0x38(%rbp)
4004b4 <+14>: mov  %rdx,-0x40(%rbp)
4004b8 <+18>: mov  %ecx,-0x44(%rbp)
4004bb <+21>: mov  %r8d,-0x48(%rbp)
4004bf <+25>: mov  %r9d,-0x4c(%rbp)
4004c3 <+29>: mov  -0x38(%rbp),%eax
4004c6 <+32>: mov  -0x34(%rbp),%edx
4004c9 <+35>: add  %eax,%edx
4004cb <+37>: mov  -0x40(%rbp),%rax
4004cf <+41>: mov  (%rax),%eax
4004d1 <+43>: lea  (%rdx,%rax,1),%eax
4004d4 <+46>: add  -0x44(%rbp),%eax
4004d7 <+49>: add  -0x48(%rbp),%eax
4004da <+52>: add  -0x4c(%rbp),%eax
4004dd <+55>: mov  %eax,-0x4(%rbp)
4004e0 <+58>: movb $0x61,-0x30(%rbp)
4004e4 <+62>: mov  $0x0,%eax
4004e9 <+67>: callq 0x4004f3 <call2>
4004ee <+72>: mov  -0x4(%rbp),%eax
4004f1 <+75>: leaveq
4004f2 <+76>: retq

```

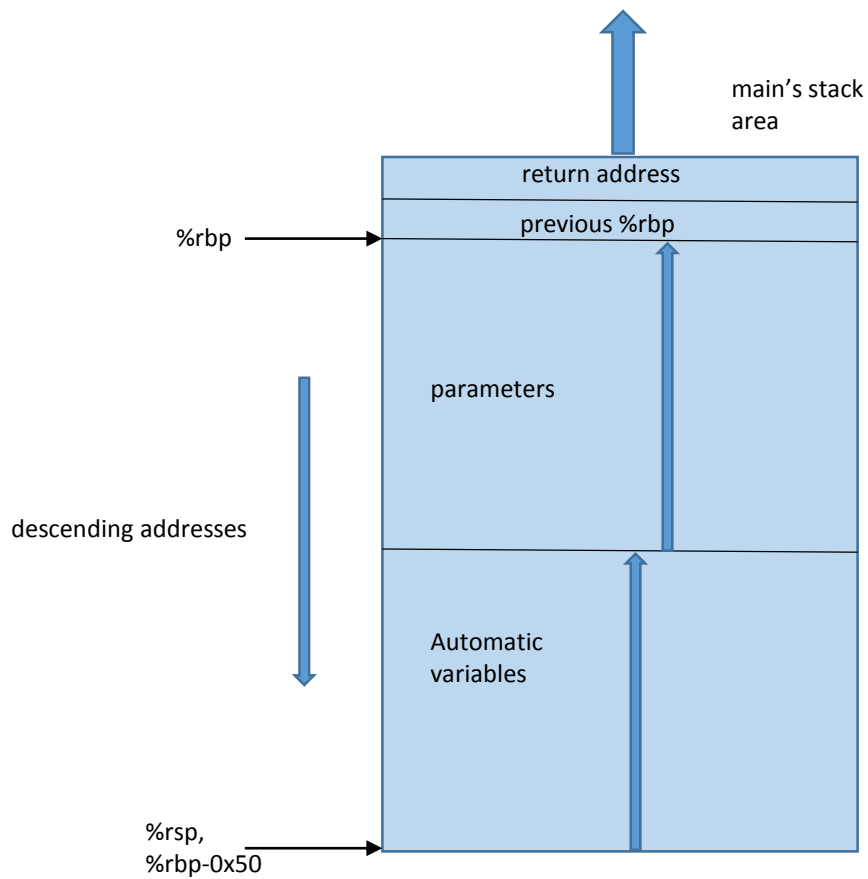
## Assembly language

Dump of assembler code for function call2:

```
40051d <+0>:  push %rbp
40051e <+1>:  mov  %rsp,%rbp
400521 <+4>:  mov  -0x8(%rbp),%eax
400524 <+7>:  mov  -0xc(%rbp),%edx
400527 <+10>: lea  (%rdx,%rax,1),%eax
40052a <+13>:  mov  %eax,-0x4(%rbp)
40052d <+16>:  leaveq
40052e <+17>:  retq
```



Let's have a look at the stack after we have entered "called":



## Procedures which call themselves (recursive)

$N!$  = factorial defined as  $\prod_{i=1}^n i$  but also  $N! = N \times (N-1)!$  And  $1! = 1$

```
#include <stdio.h>
int factorial( int n ) ;
Int kkk = 0 ;

main()
{
    int m,n=6 ;
    m = n ;
    printf( "%d factorial is %d\n", m,factorial(n) ) ;
    return 0 ;
}

int factorial(int n)
{
    int result ;
    kkk = kkk+1
    if ( n<=1 )
        result = 1 ;
    else
        result = n*factorial(n-1) ;
    printf( "n= %2d kkk= %2d exit factorial result= %d\n", n,kkk,result ) ;
    return result ;
}
```

## Procedures which call themselves (recursive)

\*\*\*build factorial machine\*\*\*

```
kkk= 1 in factorial n= 6 result addr 0x7fff201be53c
kkk= 2 in factorial n= 5 result addr 0x7fff201be4fc
kkk= 3 in factorial n= 4 result addr 0x7fff201be4bc
kkk= 4 in factorial n= 3 result addr 0x7fff201be47c
kkk= 5 in factorial n= 2 result addr 0x7fff201be43c
kkk= 6 in factorial n= 1 result addr 0x7fff201be3fc
n= 1 kkk= 6 exit factorial result= 1
n= 2 kkk= 6 exit factorial result= 2
n= 3 kkk= 6 exit factorial result= 6
n= 4 kkk= 6 exit factorial result= 24
n= 5 kkk= 6 exit factorial result= 120
n= 6 kkk= 6 exit factorial result= 720
6 factorial is 720
```

Arrays: they have a name and elements numbered from 0 to  $n-1$  where  $n$  is declared

type  $x[20]$  where type is a C/C++ data type.

for  $\text{char } x[20]$ , each box is 1 byte, total size is 20 bytes

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19

The top row doesn't exist but is just labeling the "address" of each element.

For different types, the total size is  $n * (\text{sizeof type})$ . Remember that in X86-64 pointers are 8 bytes.

## Pointer Arithmetic

In C, an array name can be treated as a pointer:  $*(x+i)$  is the value in  $x[i]$ . The address computation retains the type of  $x$ . That is  $x+i$  has the value  $\&x[0]+i*(\text{sizeof type of } x)$ .

$\&x[i]$  gives the address of  $x[i]$ , automatically adjusted for the size of  $x$ .

$*(x+i)$  gives the value stored at  $x[i]$  (called dereferencing)

$\&x[i]-x$  gives the value  $i$ . (not  $i*\text{sizeof}$ )

The type gives a scale to pointer arithmetic:

<u>type</u>	<u>scale</u>
void	1
char	1
short	2
int, float	4
double	8
pointer	8

```
int *p, *q ;
int x ;
p = &x ;
q = p+1 ; // adds 4
```

The NULL pointer is 0.

Pointers can point to functions:

$(\text{int}) (*\text{fp})(\text{int}, \text{int} *)$  ; declares  $\text{fp}$  as a function pointer.

Can have multiple dimensions:  $x[3][4]$  corresponds to:

	0	1	2	3
0				
1				
2				

So, the number of elements for  $x[n][m]$  is  $n*m$ . The size of the array is  $n*m*(\text{sizeof type})$ .

In linear memory. The array looks like:

0,0	0,1	0,2	0,3	1,0	1,1	1,2	1,3	2,0	2,1	2,2	2,3

So, the right most index runs the fastest.

To go from  $x[i][j]$  to  $x[k]$   $k = j * \max_i + i$  and the memory address is  $k \times \text{sizeof}(\text{type } x)$

## Generalizations

```
#define N 3  
#define M 4
```

```
int x[N][M] ;
```

The address of  $x[i][j]$  is  $\&x[0][0] + (i * M + j) * \text{sizeof}(\text{int})$ . Note that you can still address the array as  $*(x + l)$  but you have to know what you are doing.

For instance, you could

```
for( i=0; i<N; i++)  
    for( j=0; j<M, j++ )  
        {  
            k = i*M+j ;  
            l = *(x+k) ; // compiler multiplies k by sizeof(int)  
//  
// is the same as  
//    l = x[i][j] ;  
//  
        }
```

## Loop compile

```

400481 40>:    push %rbp
400482 <+1>:    mov  %rsp,%rbp
400485 <+4>:    sub  $0x60,%rsp
400489 <+8>:    movl $0x1,-0x18(%rbp) // i
400490 <+15>:   movl $0x2,-0x14(%rbp) // j
400497 <+22>:   movl $0x3,-0x10(%rbp) // k
40049e <+29>:   movl $0x4,-0xc(%rbp) // l
4004a5 <+36>:   movl $0x63,-0x50(%rbp) // x
4004ac <+43>:   movl $0x62,-0x60(%rbp) // y
4004b3 <+50>:   movl $0x0,-0xc(%rbp) // i
4004ba <+57>:   movl $0x0,-0x18(%rbp) // i
4004c1 <+64>:   jmp  0x400507 <main+134> //
4004c3 <+66>:   movl $0x0,-0x14(%rbp) // j
4004ca <+73>:   jmp  0x4004f9 <main+120> //
4004cc <+75>:   mov  -0x14(%rbp),%eax // j
4004cf <+78>:   cltq
4004d1 <+80>:   mov  -0x14(%rbp),%edx // j
4004d4 <+83>:   mov  %edx,-0x60(%rbp,%rax,4) // %rax*4
4004d8 <+87>:   mov  -0x18(%rbp),%edx // i
4004db <+90>:   mov  -0x14(%rbp),%eax // j
4004de <+93>:   cltq
4004e0 <+95>:   movslq %edx,%rdx // i
4004e3 <+98>:   shl  $0x2,%rdx // i*4
4004e7 <+102>:  add  %rax,%rdx // i*4+j
4004ea <+105>:  mov  -0xc(%rbp),%eax // l
4004ed <+108>:  mov  %eax,-0x50(%rbp,%rdx,4) (i*4+j)*4
4004f1 <+112>:  addl $0x1,-0x14(%rbp) // j++
4004f5 <+116>:  addl $0x1,-0xc(%rbp) // l++
4004f9 <+120>:  cmpl $0x3,-0x14(%rbp) // j ... inner loop
4004fd <+124>:  jle  0x4004cc <main+75> //
4004ff <+126>:  addl $0x1,-0x18(%rbp) // i++
400503 <+130>:  addl $0x1,-0xc(%rbp) // l++
400507 <+134>:  cmpl $0x2,-0x18(%rbp) // i ... outer loop
40050b <+138>:  jle  0x4004c3 <main+66> //

```

```

#include <stdio.h>
#define N 3
#define M 4
#define TYP int

void sub1()
{
    int i;
    i = 100;
}

int main()
{
    TYP x[N][M];
    TYP y[M];
    int i = 1, j = 2, k = 3, l = 4, *m;

    x[0][0] = 99;
    y[0] = 98;
    l = 0;
    for( i=0; i<N; i++,l++ )
        for( j=0; j<M; j++,l++ )
        {
            y[j] = j;
            x[i][j] = l;
        }

    for( i=0; i<N; i++ )
        for( j=0; j<M; j++ )
        {
            k = M*i+j;
            m = *(x+k);
        }

    sub1();

    return 0;
}

```



## A road map of a memory area

Structures: a way of aggregating data into one place. Outcome of declaration is a new data type and creates a map of a contiguous area.

History: first widely used in COBOL (1959) (Common Business Oriented Language) supposedly self documenting. “Amazing” Grace Murray Hopper.

Fortran (1953) scientific language John Backus.

PL/I (1965) combines features of COBOL and Fortran, adds elements of LISP, very strong with structures, pointers first introduced, dynamic memory allocation. IBM

```
struct <new data type name>
{
    <type> <var1> ;
    <type> <var2> ;
    .
    .
} <optional name> ;
```

After declaration, <new data type name> is a data type just like char, int, etc.

## Examples

```
struct movies
{
    char title[50];
    int year;
};
```

After this declaration, movies is a data type. If the <optional name> is present, it creates an instance of the data type. You can now declare:

```
movies yours, mine ;
```

Or even

```
movies favorites[50] ;
```

Memory layout for yours, mine is 50 character title followed by year. For favorites:

```
title[0] title[1] title[2] ... title[48] title[49] year
```

For favorites:

```
title[0][0] title[0][1] title[0][2] ... title[0][48] title[0][49] year[0]
title[1][0] title[1][1] title[1][2] ... title[1][48] title[1][49] year[1]
```

## Data alignment

Old days: mandatory

Now: instructions work but C aligns in the interest of speed.

malloc always passes back address on 16 byte boundary

## Compiled?

```

void main()
{
    struct newtype1 {
        char a ;
        float b[10] ;
        char c ;
        int d ;
    } x ;
    struct newtype2 {
        char a ;
        double b[10] ;
        char c ;
        int d ;
    } y ;
    struct newtype3 {
        int a ;
        double b[10] ;
        short c ;
        int d ;
    } z ;

    x.a = 0xff ;
    x.b[0] = 10 ;
    x.c = 0x44 ;
    x.d = 25 ;

    y.a = 0xfe ;
    y.b[0] = 9 ;
    y.c = 0x43 ;
    y.d = 24 ;

    z.a = 0xfd ;
    z.b[0] = 8 ;
    z.c = 0x42 ;
    z.d = 23 ;
}

```

```

400474 <+0>:  push %rbp
400475 <+1>:  mov  %rsp,%rbp
400478 <+4>:  sub  $0x88,%rsp
40047f <+11>: movb  $0xff,-0x40(%rbp)  // x.a   offset -0x40
400483 <+15>: mov  $0x4120,%eax
400488 <+20>: mov  %eax,-0x3c(%rbp)  // x.b[0] offset -0x3c  difference of 4
40048b <+23>: movb  $0x44,-0x14(%rbp)  // x.c   offset -0x14  difference of 40 = 10 * 4
40048f <+27>: movl  $0x19,-0x10(%rbp)  // x.d   offset -0x10  difference of 4
400496 <+34>: movb  $0xfe,-0xa0(%rbp)  // y.a   offset -0xa0
40049d <+41>: movabs $0x4022,%rax
4004a7 <+51>: mov  %rax,-0x98(%rbp)  // y.b[0] offset -0x98  difference of 8
4004ae <+58>: movb  $0x43,-0x48(%rbp)  // y.c   offset -0x48  difference of 80 = 10 * 8
4004b2 <+62>: movl  $0x18,-0x44(%rbp)  // y.d   offset -0x44  difference of 4
4004b9 <+69>: movl  $0xfd,-0x100(%rbp)  // z.a   offset -0x100
4004c3 <+79>: movabs $0x4020,%rax
4004cd <+89>: mov  %rax,-0xf8(%rbp)  // z.b[0] offset -0xf8  difference of 4
4004d4 <+96>: movw  $0x42,-0xa8(%rbp)  // z.c   offset -0xa8  difference of 80 = 10 * 8
4004dd <+105>: movl  $0x17,-0xa4(%rbp)  // z.d   offset -0xa4  difference of 4
4004e7 <+115>: leaveq
4004e8 <+116>: retq

```

## Many Variations

```
struct newtype2
{
    int a ;
    struct inner
    {
        float b ;
        int c[10] ;
    } y ;
    int *d ;
} x ;
```

The scope of the variable name lies inside of the structure. That is the name is not known outside of the structure unless you refer to the variable with its qualification:

`x.a, x.y.c[5]`

You cannot refer to `c[5]` without the qualification unless `char c[]` exists outside of the structure. This means that you can have `c` both inside and outside of the structure. Confusing!

## Passing as parameters

```
struct newtype1
{
    int a ;
    float b[10] ;
    int c ;
} x ;
```

If you want to pass a structure variable as a pointer, you can.

```
void to_sub1( struct newtype1 *x ) ;
```

```
then in to_sub1( struct newtype1 *x )
{
```

```
    x->a = 10.5 ;
```

or

```
    (*x).a = 10.5
```

and

```
    (*x).y.b = 11.5
}
```

but x.a no longer works.

## Compiled code

```
#include <stdio.h>
struct rect
{
    int i;
    int j;
    struct inner
    {
        int i;
        float l;
    } b;
    int c[3];
    int *p;
} x,y;

void to_sub1( struct rect x );
void to_sub2( struct rect *x );

int main()
{
    int i;

    printf( "%d\n", sizeof(x) );

    i = 100;

    x.i = 10;
    to_sub1( x );
    to_sub2( &x );
    return 0;
}
```

Dump of assembler code for function main:

```
4004c4 <+0>:      push    %rbp
4004c5 <+1>:      mov     %rsp,%rbp
4004c8 <+4>:      sub     $0x40,%rsp
4004cc <+8>:      mov     $0x4006a8,%eax
4004d1 <+13>:     mov     $0x28,%esi
4004d6 <+18>:     mov     %rax,%rdi
4004d9 <+21>:     mov     $0x0,%eax
4004de <+26>:     callq   0x4003b8 <printf@plt>
4004e3 <+31>:     movl    $0x64,-0x4(%rbp)
4004ea <+38>:     movl    $0xa,0x2004cc(%rip) # 0x6009c0 <x>
4004f4 <+48>:     mov     0x2004c5(%rip),%rax # 0x6009c0 <x>
4004fb <+55>:     mov     %rax,(%rsp)
4004ff <+59>:     mov     0x2004c2(%rip),%rax # 0x6009c8 <x+8>
400506 <+66>:     mov     %rax,0x8(%rsp)
40050b <+71>:     mov     0x2004be(%rip),%rax # 0x6009d0 <x+16>
400512 <+78>:     mov     %rax,0x10(%rsp)
400517 <+83>:     mov     0x2004ba(%rip),%rax # 0x6009d8 <x+24>
40051e <+90>:     mov     %rax,0x18(%rsp)
400523 <+95>:     mov     0x2004b6(%rip),%rax # 0x6009e0 <x+32>
40052a <+102>:    mov     %rax,0x20(%rsp)
40052f <+107>:    callq   0x400545 <to_sub1>
400534 <+112>:    mov     $0x6009c0,%edi
400539 <+117>:    callq   0x40056c <to_sub2>
40053e <+122>:    mov     $0x0,%eax
400543 <+127>:    leaveq  %rsp
400544 <+128>:    retq
```

## Usage in functions

```
void to_sub1( struct rect x )
{
    int i ;

    i = 100 ;
    x.i = 10 ;
    x.b.i = 5 ;
    i = x.j ;
    i = x.b.i ;
}
```

```
void to_sub2( struct rect *x )
{
    int i ;

    i = 100 ;
    x -> i = 10 ;
    (*x).i = 10 ;

    x -> b.i = 5 ;
    (*x).b.i = 5 ;

}
```

When passed by name, x.i and x.i.b no longer work because x is a pointer. Must use (\*x). instead.



Look like structures but..

In unions, the offset is always 0. This means that each variable overlays or occupies the same storage as the other variables:

```
union u
{
    int i ;
    unsigned char c[4] ;
    float a ;
} examine_endian ;
```

Sound familiar? Pointers are not needed here! But it is dangerous.

Accessing `examine_endian.a` overwrites what is in `examine_endian.i`

## Example

```
#include <stdio.h>

void main()
{
    union endian
    {
        int i;
        unsigned char c[4];
        float a;
    } hw1;

    int j;

    hw1.i = 19088743; /* should be
0x01234567 */

    printf( "\ninteger forward= " );
    for ( j=0; j<4; j++ ) /* prints c[4] */
        printf( "%02x", hw1.c[j] );
    printf( "\n" );

    printf( "integer backward= " );
    for ( j=3; j>=0; j-- ) /* prints c[4] */
        printf( "%02x", hw1.c[j] );
    printf( "\n\n" );

    hw1.a = 10.01;

    printf( "float forward= " );
    for ( j=0; j<4; j++ ) /* prints c[4] */
        printf( "%02x", hw1.c[j] );
    printf( "\n" );

    printf( "float backward= " );
    for ( j=3; j>=0; j-- ) /* prints c[4] */
        printf( "%02x", hw1.c[j] );
    printf( "\n\n" );

}
```

Dump of assembler code for function main:

```

400554 <+0>:      push    %rbp
400555 <+1>:      mov     %rsp,%rbp
400558 <+4>:      sub     $0x10,%rsp
40055c <+8>:      movl    $0x1234567,-0x10(%rbp)

400603 <+175>:     mov     $0x412028f6,%eax
400608 <+180>:     mov     %eax,-0x10(%rbp)
```

result of ./a.out

```
integer forward=  67452301
integer backward= 01234567
```

```
float forward=    f6282041
float backward=   412028f6
```

Many ways to do it

buffer overflow  
pointer goes wild  
array index goes wild

gets/puts example  
pointer error example

```
/* Corrupt1.c Stack corruption with gets */
```

```
#include <stdio.h>
```

```
void echo() ;
```

```
int main()
```

```
{  
    echo() ;  
    printf( "%x\n", EOF ) ;  
}
```

```
void echo()
```

```
{  
    char inp[8] = "012345678901234567890" ;
```

```
    while ( inp != NULL )
```

```
    {  
        gets(inp) ;  
        puts(inp) ;  
    }  
}
```

## Out of bounds subscript

```
/* corrupt2.c Stack corruption with array
overflow */
#include <stdio.h>
void echo() ;
void sub2() ;

int main()
{
    echo() ;
    printf( "%x\n", EOF ) ;
}

void echo()
{
    int i[2] ;
    int j ;
    int k ;

    i[0] = 4 ;
    i[1] = 3 ;
    i[2] = 2 ;
    i[3] = 1 ;
    j = i[4] ;
    i[4] = 0 ; /* destroys return address */
    i[5] = 0 ; /* destroys previous base pointer */
    *(i+4) = -1 ;
```

```
for( k=-4; k>-20; k-- )
    i[k] = k ;

    sub2() ;
}

void sub2()
{
    int i = 5 ;
    int j ;

    j = i ;
}
```

## Out of bounds subscript

```
/* corrupt3.c  stack corruption storing outside of frame */
#include <stdio.h>
void echo() ;
void sub2() ;

int main() {
    echo() ;
    printf( "%x\n", EOF ) ;
}
void echo() {
    int i[2] ;
    int j ;
    int k ;

    for( k=-4; k>-500; k-- )
        i[k] = k ;

    sub2() ;
}
void sub2() {
    int i = 5 ;
    int j ;

    j = i ;
}
```

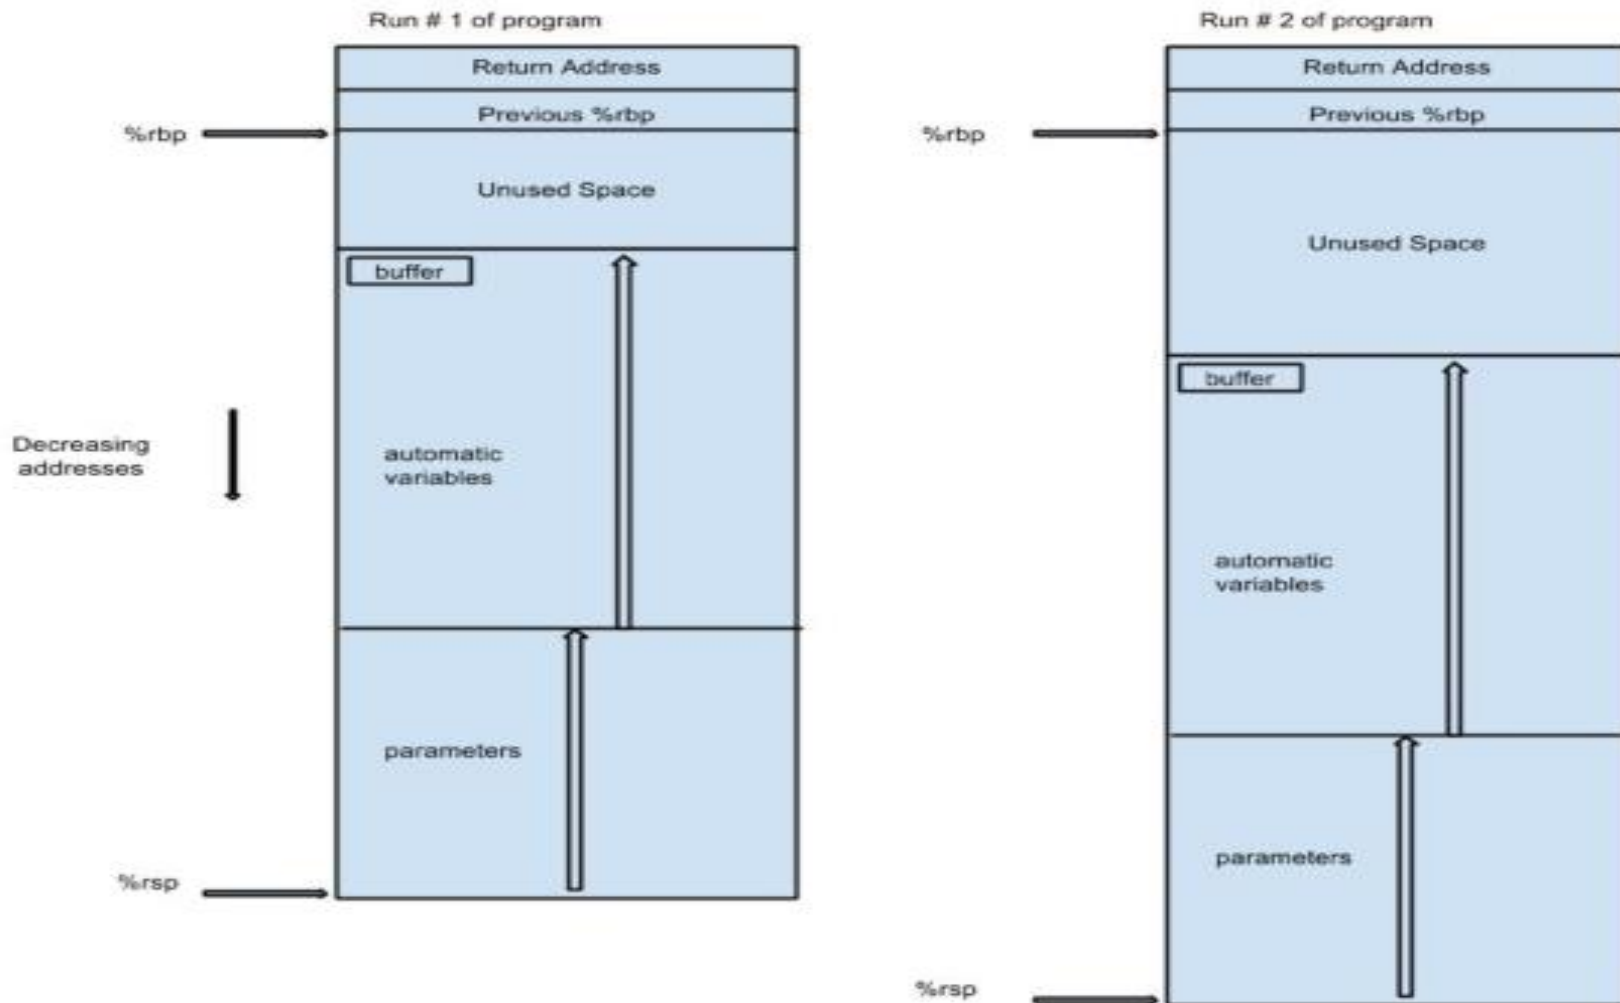
Insert code somewhere (beyond stack). Overlay the return address in the stack. When program returns, it jumps to code.

Stack randomization: after saving the return address and the previous base pointer, allocate a random amount of space in the stack. Then place the automatic variables. This way, the address of the automatic variables and the base pointer has a different offset.

Corruption detection: Store a random value somewhere in stack at the beginning of the program. Store that value in a protected area of memory. At the end of the program compare the values. If changed, raise the red flag.

Hardware which prevents pages from executing code. Memory is divided into 2K or 4K byte “pages”. Each page can be set with read/write/execute bits when in supervisory mode.

## Sample Stack Randomization



## Several ways to do it

The algorithm: smart and mindless ways

- optimize loops

  - procedure calls

  - recomputing items which do not change

  - unrolling

  - blocking

The optimizer

Take advantage of architecture: parallelism, caching

Algorithms: time as a function of set size

    - linear or polynomial time

    - $n^2$  (sort1.c)

    - $n \log n$  (sort2.c)

    - accuracy desired? (bin packing)

Algorithms: time as a function of granularity (sort3, sort4)

The importance of measurement

  - upper, lower bounds

  - average behavior



## Speeding up your program

Converting your program from  $N^2$  to  $N \log N$ : compare the two.

For small  $N$  it does not make much difference.

Amdahl's law.

Say  $T_{\text{old}}$  is the total time a program takes. Let's say part of the program takes a fraction  $f$  of the time. Let's speed up that part by a factor of  $k$ . then

$$T_{\text{new}} = (1-f) * T_{\text{old}} + (f * T_{\text{old}}) / k = T_{\text{old}} * ((1-f) + f/k)$$

then

$$T_{\text{old}} / T_{\text{new}} = 1 / ((1-f) + f/k) = S = \text{speedup factor}$$

try  $f = 0.6$ ,  $k = 3$ .  $S = 1.67$

Say  $k$  is very large, then  $S$  is approximately  $1/(1-f)$  and with  $f = 0.6$ ,  $S = 2.5$

## Compiler vs Programmer

Compiler must analyze code to see where to optimize

- reduce memory references
- take redundant code out of loops
- inline functions

Programmer can do things to allow optimization

- loop unrolling
- taking procedure calls out of loops
- reduce the use of functions: inlining
- reduce memory references
- avoid variable aliasing

## Blocks to optimization

### Memory aliasing

```
void func1(int *xp, int *yp)
{
1  *xp += *yp;
2  *xp += *yp;
}
```

after line 1 in func1, i = 20, after line 2, i = 40.

after line 1 in func2, i = 30

```
void func2(int *xp, int *yp)
{
1  *xp += 2 * *yp;
}
```

in both functions, it "looks like" there are two distinct variables but these names are just aliases for the argument.

the compiler could try to optimize func1 to func2 and the compiled code will give different results.

```
int main
{
  int i = 10 ;

  func1( &i, &i ) ;

  i = 10 ;
  func2( &i, &i ) ;

  return 0 ;
}
```

## Blocks to optimization

How about:

```
x = 1000; y = 3000;  
*q = y;      /* 3000 */  
*p = x;      /* 1000 */  
t1 = *q;     /* 1000 or 3000 */
```

if  $p == q$ , result is  $t1 = 1000$  ; if not  $t1 = 3000$

Can this happen with pass by value variables?

## Blocks to optimization

Consider:

```
void swap(int *xp, int *yp)
{
    *xp = *xp + *yp; /* x+y */
    *yp = *xp - *yp; /* x+y-y = x */
    *xp = *xp - *yp; /* x+y-x = y */
}
```

```
int main()
{
    int x,y ;

    x = 10 ;
    y = 20 ;
    swap( &x, &y ) ;

    x = 10 ;
    y = 10 ;
    swap( &x, &y ) ;
```

```
    x = 10 ;
    y = 10 ;
    swap( &x, &x ) ;
```

```
    return 0 ;
}
```

Start swap

```
*xp *yp
10 20
30 20
30 10
20 10
```

Start swap

```
*xp *yp
10 10
20 10
20 10
10 10
```

Start swap

```
*xp *yp
10 10
20 20
0 0
0 0
```

When  $x \neq y$ , (first two cases), everything works normally.  
Even when  $*xp = *yp$ .

But when  $x == y$ , problems.

## Blocks to optimization

Consider when a function operates on global variables.

```
int counter = 0;
```

```
int f()
{
    printf( "in f() counter= %d\n", counter );
    return counter++;
}
```

```
int main()
{
    printf( "result of f()+f()+f()+f() = %d \n", f()+f()+f()+f() );

    counter = 0 ;

    printf( "result of 4*f() = %d \n", 4*f() );

    return 0 ;
}
```

```
in f() counter= 0
in f() counter= 1
in f() counter= 2
in f() counter= 3
result of f()+f()+f()+f() = 6
```

```
in f() counter= 0
result of 4*f() = 0
End value of counter = 1
```

“Optimized” program gives different results.

## Measuring performance

### Cycles per element

How do we measure performance? Stop watch? Program running time as a function of number of input elements. Also may be a function of the distribution of input data (coarseness).

Run the program many times with different number of input elements and data types. Plot a curve of number of elements versus running time. Fit a line to the data using least squares fit (regression) and you arrive at a formula which is

$$\text{run time} = \text{constant} + \text{coefficient} * N$$

where N is the number of input elements. So, the coefficient expresses the rate of increase in run time per additional data element. The constant expresses the overhead to start the program (run time when  $N = 0$ ).

The coefficient is known as the CPE: cycles per element. Its units are run time per element. In all cases, it is relative to the speed of the computer but we think of it as cycles

## Loop unrolling

```
/* Compute prefix sum of vector a */
void psum1(float a[], float p[], long int n)
{
    long int i;
    p[0] = a[0];
    for (i = 1; i < n; i++)
        p[i] = p[i-1] + a[i];
}
```

This function computes the “prefix sum” of an array of elements: a. It is defined as:

$$p[0] = a[0];$$

$$p[i] = p[i-1] + a[i]$$

So,  $p[i]$  = the sum of all  $a[j]$  where  $j \leq i$

```
void psum2(float a[], float p[], long int n)
{
    long int i;
    p[0] = a[0];
    for (i = 1; i < n-1; i+=2) {
        p[i] = p[i-1] + a[i];
        p[i+1] = p[i] + a[i+1];
    }
    /* For odd n, finish remaining element */
    if (i < n)
        p[i] = p[i-1] + a[i];
}
```

This is 1x unrolling



## Loop unrolling

**\*Caveat\*** The increase in the number of lines in the code affects the savings for loop unrolling.

Lets say that it takes  $x$  units of time to execute the line of code in the loop in psum1,  $y$  units to execute the loop overhead and  $z$  time units are used when the loop is unrolled in psum2.

So, the time to execute the loop in psum1 is

$$a = x*n + y*n \quad \text{execute the code plus the loop overhead}$$

an unrolled loop program would take

$$b = z*n/2 + y*n/2 \quad \text{execute the code half as much and the loop overhead half as much}$$

for it to be faster, we want  $b < a$  or

$$z*n/2 + y*n/2 < x*n + y*n$$

This is the same as

$$0 < x*n + y*n/2 - z*n/2$$

dividing by  $n$  it becomes

$$0 < x + y/2 - z/2 = c = \text{difference in run times old - new}$$

## Loop unrolling

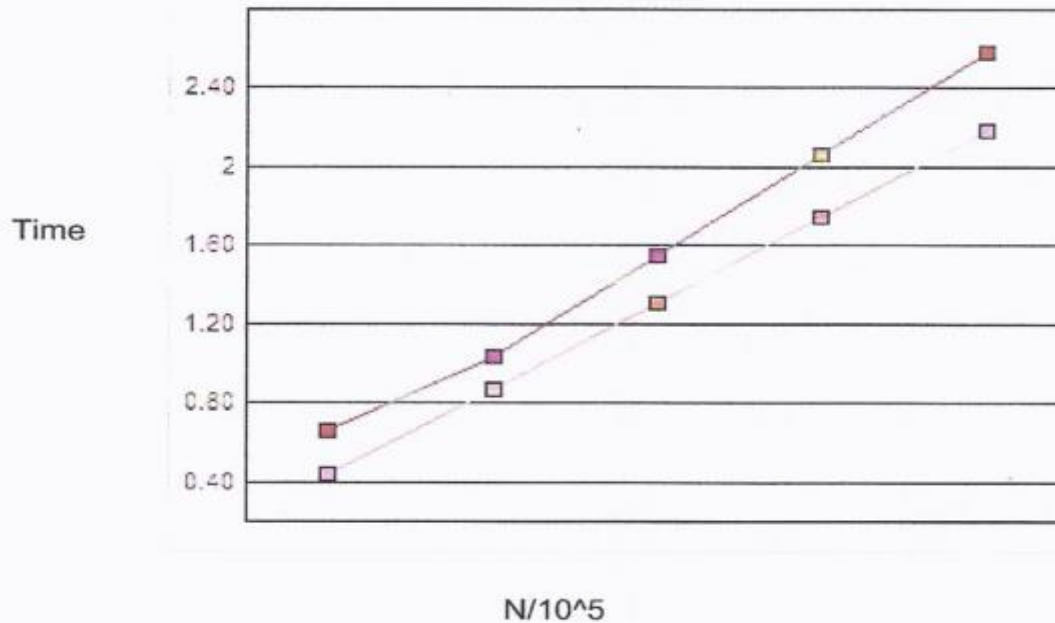
So, depending on the relative values of x, y and z, there is a diminishing rate of return!

x	y	z	c
1	1	2	0.5
1	1	3	0
1	1	4	-0.5
1	1	5	-1
2	1	3	1
2	1	4	0.5
2	1	5	0
2	1	6	-0.5

For the first line in the table, we increase the statement executions by 1 and we get a .5 improvement in the run time, by 2 and we get zero. In line 5, we increase it by 1 from 2 to 3, we get an improvement of 1.

## Loop unrolling

unoptimized psum timing		psum1	psum2
N/10^5	1	0.657	0.438
	2	1.034	0.870
	3	1.550	1.311
	4	2.069	1.751
	5	2.585	2.184
constant		0.112	-0.001
coefficient		0.489	0.437



In the table, the entries are the time to run psum1(upper line) and psum2 (lower line) given the number of elements. Constant and coefficient are the terms which fit:

$$\text{run time} = \text{constant} + \text{coefficient} * n$$

Combine example. A program to add or multiply an array of values.

```
typedef struct
{
    int len;
    data_t *data;
} vec_rec, *vec_ptr;

/*
 * Retrieve vector element and store at dest.
 * Return 0 (out of bounds) or 1 (successful)
 */
int get_vec_element(vec_ptr v, int index, data_t *dest)
{
    if (index < 0 || index >= v->len)
        return 0;
    *dest = v->data[index];
    return 1;
}

/* Return length of vector */
int vec_length(vec_ptr v)
{
    return v->len;
}
```

data\_t is a data type: int or float or double.  
IDENT is 0 or 1  
OP is \* or +

```
void combine1(vec_ptr v, data_t *dest)
{
    int i;

    *dest = IDENT;
    for (i = 0; i < vec_length(v); i++)
    {
        data_t val;
        get_vec_element(v, i, &val);
        *dest = *dest OP val;
    }
}
```

Disassembling the code using GDB shows that both function calls are in the loop

Combine example. A program to add or multiply an array of values.

Running this without optimization with  $1,2,3,4,5 * 10^5$   
gives the following timings

combine1

Time= 1.0622 n= 100000

Time= 2.1249 n= 200000

Time= 3.1871 n= 300000

Time= 4.2490 n= 400000

Time= 5.3124 n= 500000

Running with g++ -O : optimization level 1, gives:

combine1 N

Time= 0.3437 n= 100000

Time= 0.6847 n= 200000

Time= 1.0270 n= 300000

Time= 1.3692 n= 400000

Time= 1.7116 n= 500000

A factor of 3 improvement.

## Compiled code.

Dump of assembler code for function `_Z8combine1P7vec_recPi`:

```

400b1e <+0>:  push %rbp
400b1f <+1>:  mov  %rsp,%rbp
400b22 <+4>:  sub  $0x20,%rsp
400b26 <+8>:  mov  %rdi,-0x18(%rbp)
400b2a <+12>: mov  %rsi,-0x20(%rbp)
400b2e <+16>: mov  -0x20(%rbp),%rax
400b32 <+20>: movl  $0x0,(%rax)
400b38 <+26>: movl  $0x0,-0x4(%rbp)
400b3f <+33>: jmp  0x400b6b <+77>

400b41 <+35>: lea  -0x8(%rbp),%rdx
400b45 <+39>: mov  -0x4(%rbp),%ecx
400b48 <+42>: mov  -0x18(%rbp),%rax
400b4c <+46>: mov  %ecx,%esi
400b4e <+48>: mov  %rax,%rdi
400b51 <+51>: callq 0x400a1e <get_vec_element>
400b56 <+56>: mov  -0x20(%rbp),%rax
400b5a <+60>: mov  (%rax),%edx
400b5c <+62>: mov  -0x8(%rbp),%eax
400b5f <+65>: add  %eax,%edx
400b61 <+67>: mov  -0x20(%rbp),%rax
400b65 <+71>: mov  %edx,(%rax)
400b67 <+73>: addl  $0x1,-0x4(%rbp)
400b6b <+77>: mov  -0x18(%rbp),%rax
400b6f <+81>: mov  %rax,%rdi
400b72 <+84>: callq 0x400a69 <vec_length>
400b77 <+89>: cmp  -0x4(%rbp),%eax
400b7a <+92>: setg %al
400b7d <+95>: test %al,%al
400b7f <+97>: jne  0x400b41 <+35>

400b81 <+99>: leaveq
400b82 <+100>: retq

```

Dump of assembler code for function (with `-O`) `_Z8combine1P7vec_recPi`:

```

400a7d <+0>:  movl  $0x0,(%rsi)
400a83 <+6>:  cmpl  $0x0,(%rdi)
400a86 <+9>:  jle  0x400ac2 <+69>
400a88 <+11>: push  %r12
400a8a <+13>: push  %rbp
400a8b <+14>: push  %rbx
400a8c <+15>: sub  $0x10,%rsp
400a90 <+19>: mov  %rsi,%r12
400a93 <+22>: mov  %rdi,%rbp
400a96 <+25>: mov  $0x0,%ebx

400a9b <+30>: lea  0xc(%rsp),%rdx
400aa0 <+35>: mov  %ebx,%esi
400aa2 <+37>: mov  %rbp,%rdi
400aa5 <+40>: callq 0x4009f8 <get_vec_element>
400aaa <+45>: mov  0xc(%rsp),%eax
400aae <+49>: add  %eax,(%r12)
400ab2 <+53>: add  $0x1,%ebx
400ab5 <+56>: cmp  0x0(%rbp),%ebx
400ab8 <+59>: jl   0x400a9b <+30>

400aba <+61>: add  $0x10,%rsp
400abe <+65>: pop  %rbx
400abf <+66>: pop  %rbp
400ac0 <+67>: pop  %r12
400ac2 <+69>: repz retq

```

There is no call to `vec_length` in sight! It has somehow figured out that the length of the vector is stored at the beginning of the vector.

Combine example. A program to add or multiply an array of values.

data\_t is a data type: int or float or double.

IDENT is 0 or 1

OP is \* or +

```
void combine1(vec_ptr v, data_t *dest)
{
    int i;

    *dest = IDENT;
    for (i = 0; i < vec_length(v); i++)
    {
        data_t val;
        get_vec_element(v, i, &val);
        *dest = *dest OP val;
    }
}
```

Disassembling the code using GDB shows that both function calls are in the loop

Combine example. Remove function call from loop.

data\_t is a data type: int or float or double.

IDENT is 0 or 1

OP is \* or +

```
void combine2(vec_ptr v, data_t *dest)
{
    int l, l = vec_length(v);

    *dest = IDENT;
    for (i = 0; i < l ; i++)
    {
        data_t val;
        get_vec_element(v, i, &val);
        *dest = *dest OP val;
    }
}
```

Disassembling the code using GDB shows that both function calls are in the loop



## Comparison after taking a function call out of loop

Running this without optimization with  $1,2,3,4,5 * 10^5$  gives the following timings

combine1	combine2
Time= 1.0622 n= 100000	Time= 0.9520 n= 100000
Time= 2.1249 n= 200000	Time= 1.9029 n= 200000
Time= 3.1871 n= 300000	Time= 2.8540 n= 300000
Time= 4.2490 n= 400000	Time= 3.8053 n= 400000
Time= 5.3124 n= 500000	Time= 4.7566 n= 500000

Running with g++ -O : optimization level 1, gives:

combine1	combine2
Time= 0.3437 n= 100000	Time= 0.3054 n= 100000
Time= 0.6847 n= 200000	Time= 0.6058 n= 200000
Time= 1.0270 n= 300000	Time= 0.9087 n= 300000
Time= 1.3692 n= 400000	Time= 1.2114 n= 400000
Time= 1.7116 n= 500000	Time= 1.5146 n= 500000

A factor of 3 improvement.

Compiled code (with -O).

with vec\_length reference

```

400b1b <+30>: lea  0xc(%rsp),%rdx
400b20 <+35>: mov  %ebx,%esi
400b22 <+37>: mov  %rbp,%rdi
400b25 <+40>: callq 0x400a78 <get_vec_element>
400b2a <+45>: mov  0xc(%rsp),%eax
400b2e <+49>: add  %eax,(%r12)
400b32 <+53>: add  $0x1,%ebx
400b35 <+56>: cmp  0x0(%rbp),%ebx
400b38 <+59>: jl   0x400b1b <+30>
    
```

without vec\_length reference

```

400b20 <+35>: lea  0xc(%rsp),%rdx
400b25 <+40>: mov  %ebx,%esi
400b27 <+42>: mov  %r12,%rdi
400b2a <+45>: callq 0x400a78 <get_vec_element>
400b2f <+50>: mov  0xc(%rsp),%eax
400b33 <+54>: add  %eax,0x0(%rbp)
400b36 <+57>: add  $0x1,%ebx
400b39 <+60>: cmp  %r13d,%ebx
400b3c <+63>: jne  0x400b20 <vec_recPi+35>
    
```

The only real difference is: no memory reference in compare!

Horrible example of function call in a loop.

```
/* Sample implementation of library function strlen */
/* Compute length of string */
size_t strlen(const char *s)
{
    int length = 0;
    while (*s != '\0')
    {
        s++;
        length++;
    }
    return length;
}
```

```
/* Sample implementation of library function strlen */
/* Compute length of string */
size_t strlenfor(const char *s)
{
    int length ;

    for( length=0;;length++ )
        if( *s == '\0' )
            break ;
        else
            s++ ;

    return length ;
}
```

```
/* Sample implementation of library function strlen */
/* Compute length of string */
size_t strlenif(const char *s)
{
    int length = 0 ;

top: if( *s == '\0' )
    goto ret ;
    else
    {
        s++ ;
        length++ ;
        goto top ;
    }
ret:
    return length ;
}
```

```
/* Convert string to lowercase */
void lower1(char *s)
{
    int i;

    for (i = 0; i < strlen(s); i++)
        if (s[i] >= 'A' && s[i] <= 'Z')
            s[i] -= ('A' - 'a');
}
```

```
/* Convert string to lowercase */
void lower2(char *s)
{
    int i;

    for (i = 0; i < strlenfor(s); i++)
        if (s[i] >= 'A' && s[i] <= 'Z')
            s[i] -= ('A' - 'a');
}
```

```
/* Convert string to lowercase */
```

```
void lower3(char *s)
```

```
{
```

```
    int i;
```

```
    for (i = 0; i < strlenif(s); i++)
```

```
        if (s[i] >= 'A' && s[i] <= 'Z')
```

```
            s[i] -= ('A' - 'a');
```

```
}
```

```
/* Convert string to lowercase: fast */
```

```
void lower4(char *s)
```

```
{
```

```
    int i, n = strlenif(s) ;
```

```
    for (i = 0; i < n; i++)
```

```
        if (s[i] >= 'A' && s[i] <= 'Z')
```

```
            s[i] -= ('A' - 'a');
```

```
}
```

Instructive to look at timings: 4 values of N, optimized and not

## Not Optimized

Method = While	Time= 1.0651 n= 700
Method = For	Time= 1.2104 n= 700
Method = If	Time= 1.2203 n= 700
Method = Nofunc	Time= 0.0049 n= 700
Method = While	Time= 1.3874 n= 800
Method = For	Time= 1.5790 n= 800
Method = If	Time= 1.6095 n= 800
Method = Nofunc	Time= 0.0059 n= 800
Method = While	Time= 1.7541 n= 900
Method = For	Time= 1.9961 n= 900
Method = If	Time= 1.9583 n= 900
Method = Nofunc	Time= 0.0067 n= 900
Method = While	Time= 2.1869 n= 1000
Method = For	Time= 2.4659 n= 1000
Method = If	Time= 2.3654 n= 1000
Method = Nofunc	Time= 0.0074 n= 1000

## Optimized

Method = While	Time= 0.3778 n= 700
Method = For	Time= 0.3807 n= 700
Method = If	Time= 0.3799 n= 700
Method = Nofunc	Time= 0.0014 n= 700
Method = While	Time= 0.4919 n= 800
Method = For	Time= 0.4953 n= 800
Method = If	Time= 0.4944 n= 800
Method = Nofunc	Time= 0.0019 n= 800
Method = While	Time= 0.6214 n= 900
Method = For	Time= 0.6250 n= 900
Method = If	Time= 0.6240 n= 900
Method = Nofunc	Time= 0.0017 n= 900
Method = While	Time= 0.7655 n= 1000
Method = For	Time= 0.7698 n= 1000
Method = If	Time= 0.7691 n= 1000
Method = Nofunc	Time= 0.0021 n= 1000

## Compiled

### Strlen (while)

```

4008f8 <+0>:  push %rbp
4008f9 <+1>:  mov  %rsp,%rbp
4008fc <+4>:  mov  %rdi,-0x18(%rbp)
400900 <+8>:  movl $0x0,-0x4(%rbp)
400907 <+15>: jmp  0x400912 <+26>

400909 <+17>:  addq $0x1,-0x18(%rbp)
40090e <+22>:  addl $0x1,-0x4(%rbp)
400912 <+26>:  mov  -0x18(%rbp),%rax
400916 <+30>:  movzbl (%rax),%eax
400919 <+33>:  test %al,%al
40091b <+35>:  jne  0x400909 <+17>

40091d <+37>:  mov  -0x4(%rbp),%eax
400920 <+40>:  cltq
400922 <+42>:  pop  %rbp
400923 <+43>:  retq

```

### Strlen (for)

```

400924 <+0>:  push %rbp
400925 <+1>:  mov  %rsp,%rbp
400928 <+4>:  mov  %rdi,-0x18(%rbp)
40092c <+8>:  movl $0x0,-0x4(%rbp)

400933 <+15>:  mov  -0x18(%rbp),%rax
400937 <+19>:  movzbl (%rax),%eax
40093a <+22>:  test %al,%al
40093c <+24>:  jne  0x400940 <+28>
40093e <+26>:  jmp  0x40094b <+39>
400940 <+28>:  addq $0x1,-0x18(%rbp)
400945 <+33>:  addl $0x1,-0x4(%rbp)
400949 <+37>:  jmp  0x400933 <+15>

40094b <+39>:  mov  -0x4(%rbp),%eax
40094e <+42>:  cltq
400950 <+44>:  pop  %rbp
400951 <+45>:  retq

```



Combine example.

data\_t is a data type: int or float or double.

IDENT is 0 or 1

OP is \* or +

```
void combine2(vec_ptr v, data_t *dest)
{
    int l, l = vec_length(v);

    *dest = IDENT;
    for (i = 0; i < l ; i++)
    {
        data_t val;
        get_vec_element(v, i, &val);
        *dest = *dest OP val;
    }
}
```

Combine example. Remove additional function call from loop.

data\_t is a data type: int or float or double.

IDENT is 0 or 1

OP is \* or +

```
void combine3(vec_ptr v, data_t *dest)
{
    int l, l = vec_length(v);
    data_t *data = get_vec_start(v);
    *dest = IDENT;
    for (i = 0; i < l; i++)
        *dest = *dest OP data[i];
}
```

## Comparison after taking an additional function call out of loop

Running this without optimization with  $1,2,3,4,5 * 10^5$  gives the following timings

combine1	combine2	Combine3
Time= 1.0622 n= 100000	Time= 0.9520 n= 100000	Time= 0.3845 n= 100000
Time= 2.1249 n= 200000	Time= 1.9029 n= 200000	Time= 0.7586 n= 200000
Time= 3.1871 n= 300000	Time= 2.8540 n= 300000	Time= 1.1377 n= 300000
Time= 4.2490 n= 400000	Time= 3.8053 n= 400000	Time= 1.5170 n= 400000
Time= 5.3124 n= 500000	Time= 4.7566 n= 500000	Time= 1.8961 n= 500000

Running with g++ -O : optimization level 1, gives:

combine1	combine2	Combine3
Time= 0.3437 n= 100000	Time= 0.3054 n= 100000	Time= 0.2312 n= 100000
Time= 0.6847 n= 200000	Time= 0.6058 n= 200000	Time= 0.4536 n= 200000
Time= 1.0270 n= 300000	Time= 0.9087 n= 300000	Time= 0.6811 n= 300000
Time= 1.3692 n= 400000	Time= 1.2114 n= 400000	Time= 0.9095 n= 400000
Time= 1.7116 n= 500000	Time= 1.5146 n= 500000	Time= 1.1332 n= 500000

## Compiled code (with -O).

without vec\_length reference

```

400afd <+0>:  push  %r13
400aff <+2>:  push  %r12
400b01 <+4>:  push  %rbp
400b02 <+5>:  push  %rbx
400b03 <+6>:  sub   $0x10,%rsp
400b07 <+10>: mov   (%rdi),%r13d
400b0a <+13>: movl  $0x0,(%rsi)
400b10 <+19>: test  %r13d,%r13d
400b13 <+22>: jle   0x400b3e <+65>
400b15 <+24>: mov   %rsi,%rbp
400b18 <+27>: mov   %rdi,%r12
400b1b <+30>: mov   $0x0,%ebx

400b20 <+35>: lea   0xc(%rsp),%rdx
400b25 <+40>: mov   %ebx,%esi
400b27 <+42>: mov   %r12,%rdi
400b2a <+45>: callq 0x400a78 <_Z15get_vec_elementP7vec_reciPi>
400b2f <+50>: mov   0xc(%rsp),%eax
400b33 <+54>: add   %eax,0x0(%rbp)
400b36 <+57>: add   $0x1,%ebx
400b39 <+60>: cmp   %r13d,%ebx
400b3c <+63>: jne   0x400b20 <+35>

400b3e <+65>: add   $0x10,%rsp
400b42 <+69>: pop   %rbx
400b43 <+70>: pop   %rbp
400b44 <+71>: pop   %r12
400b46 <+73>: pop   %r13
400b48 <+75>: retq

```

without get\_vec\_start reference in loop

```

400afd <+0>:  mov   (%rdi),%edx
400aff <+2>:  mov   0x8(%rdi),%rcx
400b03 <+6>:  movl  $0x0,(%rsi)
400b09 <+12>: test  %edx,%edx
400b0b <+14>: jle   0x400b25 <+40>
400b0d <+16>: mov   %rcx,%rax
400b10 <+19>: sub   $0x1,%edx
400b13 <+22>: lea   0x4(%rcx,%rdx,4),%rcx

400b18 <+27>: mov   (%rax),%edx
400b1a <+29>: add   %edx,(%rsi)
400b1c <+31>: add   $0x4,%rax
400b20 <+35>: cmp   %rcx,%rax
400b23 <+38>: jne   0x400b18 <+27>

400b25 <+40>: repz retq

```

Two memory references in loop

Also eliminates stack and argument preparation

Combine example.

data\_t is a data type: int or float or double.

IDENT is 0 or 1

OP is \* or +

```
void combine3(vec_ptr v, data_t *dest)
{
    int l, l = vec_length(v);
    data_t *data = get_vec_start(v);
    *dest = IDENT;

    for (i = 0; i < l; i++)
        *dest = *dest OP data[i];
}
```

Combine example, eliminate additional memory reference.

data\_t is a data type: int or float or double.  
IDENT is 0 or 1  
OP is \* or +

```
void combine4(vec_ptr v, data_t *dest)
{
    int l, l = vec_length(v);
    data_t *data = get_vec_start(v);
    data_t acc = IDENT;

    for (i = 0; i < l; i++)
        acc = acc OP data[i];

    *dest = acc;
}
```

## Comparison after eliminating additional memory reference

Running this without optimization with  $1,2,3,4,5 * 10^5$  gives the following timings

combine1	combine2	combine3	combine4
Time= 1.0622 n= 100000	Time= 0.9520 n= 100000	Time= 0.3845 n= 100000	Time= 0.3614 n= 100000
Time= 2.1249 n= 200000	Time= 1.9029 n= 200000	Time= 0.7586 n= 200000	Time= 0.7228 n= 200000
Time= 3.1871 n= 300000	Time= 2.8540 n= 300000	Time= 1.1377 n= 300000	Time= 1.0837 n= 300000
Time= 4.2490 n= 400000	Time= 3.8053 n= 400000	Time= 1.5170 n= 400000	Time= 1.4443 n= 400000
Time= 5.3124 n= 500000	Time= 4.7566 n= 500000	Time= 1.8961 n= 500000	Time= 1.8057 n= 500000

Running with g++ -O : optimization level 1, gives:

combine1	combine2	combine3	combine4
Time= 0.3437 n= 100000	Time= 0.3054 n= 100000	Time= 0.2312 n= 100000	Time= 0.0784 n= 100000
Time= 0.6847 n= 200000	Time= 0.6058 n= 200000	Time= 0.4536 n= 200000	Time= 0.1515 n= 200000
Time= 1.0270 n= 300000	Time= 0.9087 n= 300000	Time= 0.6811 n= 300000	Time= 0.2270 n= 300000
Time= 1.3692 n= 400000	Time= 1.2114 n= 400000	Time= 0.9095 n= 400000	Time= 0.3026 n= 400000
Time= 1.7116 n= 500000	Time= 1.5146 n= 500000	Time= 1.1332 n= 500000	Time= 0.3786 n= 500000

## Compiled code (with -O).

without get\_vec\_start reference in loop

```

400afd <+0>:  mov  (%rdi),%edx
400aff <+2>:  mov  0x8(%rdi),%rcx
400b03 <+6>:  movl  $0x0,(%rsi)
400b09 <+12>: test  %edx,%edx
400b0b <+14>: jle   0x400b25 <+40>
400b0d <+16>: mov   %rcx,%rax
400b10 <+19>: sub   $0x1,%edx
400b13 <+22>: lea   0x4(%rcx,%rdx,4),%rcx

400b18 <+27>: mov   (%rax),%edx
400b1a <+29>: add   %edx,(%rsi)
400b1c <+31>: add   $0x4,%rax
400b20 <+35>: cmp   %rcx,%rax
400b23 <+38>: jne   0x400b18 <+27>

400b25 <+40>: repz retq

```

allowing local accumulator

```

400afd <+0>:  mov  (%rdi),%edx
400aff <+2>:  mov  0x8(%rdi),%rcx
400b03 <+6>:  test %edx,%edx
400b05 <+8>:  jle   0x400b24 <+39>
400b07 <+10>: mov   %rcx,%rax
400b0a <+13>: sub   $0x1,%edx
400b0d <+16>: lea   0x4(%rcx,%rdx,4),%rcx
400b12 <+21>: mov   $0x0,%edx

400b17 <+26>: add   (%rax),%edx
400b19 <+28>: add   $0x4,%rax
400b1d <+32>: cmp   %rcx,%rax
400b20 <+35>: jne   0x400b17 <+26>

400b22 <+37>: jmp   0x400b29 <+44>
400b24 <+39>: mov   $0x0,%edx
400b29 <+44>: mov   %edx,(%rsi)
400b2b <+46>: retq

```

One memory references in loop, one less instruction.



## Summary of combine improvements

remove function call whose value does not change from loop: combine2  
 access data directly rather than through function call: combine3  
 reduce memory references by using local variable: combine4

These changes, by themselves can decrease the running time. But also, they assist the compiler in optimizing the program.

We ran each version using 5 different size of input data: 100000, 200000, 300000, 400000 and 500000 elements in the array to add up.

Using regression, we find the CPE (relative values) for each version:

	1	2	3	4
Unoptimized	1.06	0.95	0.38	0.36
Optmized	0.34	0.31	0.23	0.08

## Optimization

The horizontal lines demarcate the for loop. The rest of the code is just overhead and we are mainly interested in the loop.

		memory	arith	reg	func
<0>: push %r12	stack	1			
<2>: push %rbp	"	1			
<3>: push %rbx	"	1			
<4>: sub \$0x10,%rsp	"		1	1	
<8>: mov %rdi,%rbp	v			1	
<11>: mov %rsi,%r12	dest			1	
<14>: movl \$0x0,(%rsi)	*dest = 0 ;	1			
<20>: cmpl \$0x0,(%rdi)	comp 0 to v.len	1	1		
<23>: jle <61>	if <= go to finish				
<25>: mov \$0x0,%ebx	i = 0			1	
<30>:--->leal 0xc(%rsp),%rdx	parameter &val		1	1	
<35>: mov %ebx,%esi	parameter i			1	
<37>: mov %rbp,%rdi	parameter v			1	
<40>: callq <get_vec_element>					1
<45>: mov 0xc(%rsp),%eax	val	1			
<49>: add %eax,(%r12)	+ *dest	1	1		
<53>: add \$0x1,%ebx	i = i+1		1		
<56>: cmp 0x0(%rbp),%ebx	compare to v.len	1	1		
<59>:--->jl <30>	if less go to loop			1	
<61>: add \$0x10,%rsp	stack manipulation		1		
<65>: pop %rbx	"	1			
<66>: pop %rbp	"	1			
<67>: pop %r12	"	1			
<69>: retq	"				
	inside loop	3	4	4	1

## Formal model of the loop

Let's say that a memory operation is  $M$  cycles, an arithmetic operation is  $A$  cycles, a register operation is  $R$  cycles and a function is  $F$  cycles. Let's add 1 cycle for each byte of instruction. The length of the code is 31 bytes.

So, our loop will take  $M*3 + A*4 + R*4 + F + 31$  cycles.

If we guess  $M=7$ ,  $A=5$ ,  $R=1$  and  $F=15$ , the loop will take 91 cycles.

## Removing function call: combine2

<pre> combine1 &lt;30&gt;:---&gt;lea 0xc(%rsp),%rdx &lt;35&gt;: mov %ebx,%esi &lt;37&gt;: mov %rbp,%rdi &lt;40&gt;: callq &lt;get_vec_element&gt; &lt;45&gt;: mov 0xc(%rsp),%eax &lt;49&gt;: add %eax,(%r12) &lt;53&gt;: add \$0x1,%ebx &lt;56&gt;: cmp 0x0(%rbp),%ebx &lt;59&gt;:---&gt;jl &lt;30&gt; </pre>	<pre> combine2 &lt;35&gt;:---&gt;lea 0xc(%rsp),%rdx &lt;40&gt;: mov %ebx,%esi &lt;42&gt;: mov %r13,%rdi &lt;45&gt;: callq &lt;_get_vec_element&gt; &lt;50&gt;: mov 0xc(%rsp),%eax &lt;54&gt;: add %eax,0x0(%rbp) &lt;57&gt;: add \$0x1,%ebx &lt;60&gt;: cmp %r12d,%ebx &lt;63&gt;:---&gt;jne &lt;35&gt; </pre>
---	--

But, we see that combine1 has

add %eax,(%r12)

combine2 has add %eax,0x0(%rbp) : one less memory op.

Also combine1 has

cmp 0x0(%rbp),%ebx combine2 has cmp %r12d,%ebx

combine2 is one byte shorter. Significant?  $3.3\% = (31/30)*100$

$M*2 + A*4 + R*4 + F + 30 \text{ cycles} = 83$

CPE improvement:  $.31/.34 = .912$  Cycles improvement =  $83/91 = .911$

## Removing get\_vec\_element function call: combine3

Combine3

	memory	arith	reg	func
<0>: mov (%rdi),%ecx				
<2>: mov 0x8(%rdi),%rdi				
<6>: movl \$0x0,(%rsi)				
<12>: test %ecx,%ecx				
<14>: jle <34>				
<16>: mov \$0x0,%eax				
<21>:--->mov (%rdi,%rax,4),%edx	1			
<24>: add %edx,(%rsi)	1	1		
<26>: add \$0x1,%rax		1		
<30>: cmp %eax,%ecx		1		
<32>:--->jg <21>			1	
<34>: repz retq				
inside loop	2	3	1	0

Loop code is 13 bytes

$M*2 + A*3 + R*1 + 0*F + 13 \text{ cycles} = 43$

CPE improvement:  $.23/.31 = ..742$  Cycles improvement =  $46/83 = .518$  not as good

## adding local accumulator: combine4

Combine4

	memory	arith	reg	func
<0>: mov (%rdi),%ecx				
<2>: mov 0x8(%rdi),%rdi				
<6>: test %ecx,%ecx				
<8>: jle <i+33>				
<10>: mov \$0x0,%eax				
<15>: mov \$0x0,%edx				
<20>:--->add (%rdi,%rax,4),%edx	1	1		
<23>: add \$0x1,%rax		1		
<27>: cmp %eax,%ecx		1		
<29>:--->jg <20>			1	
<31>: jmp 0x400aa4 <i+38>				
<33>: mov \$0x0,%edx				
<38>: mov %edx,(%rsi)				
<40>: retq				
inside loop	1	3	1	0

Here we have one less memory reference and one less instruction.

$M*1 + A*3 + R*1 + 0*F + 11 \text{ cycles} = 34$

CPE improvement:  $.08/.23 = ..348$  Cycles improvement =  $34/46 = .739$  not as good

## conclusion

Our estimates of M, R, A and F aren't very good.

Suggestion:

Our model is:  $a*M + b*R + c*A + d*F + e*S$

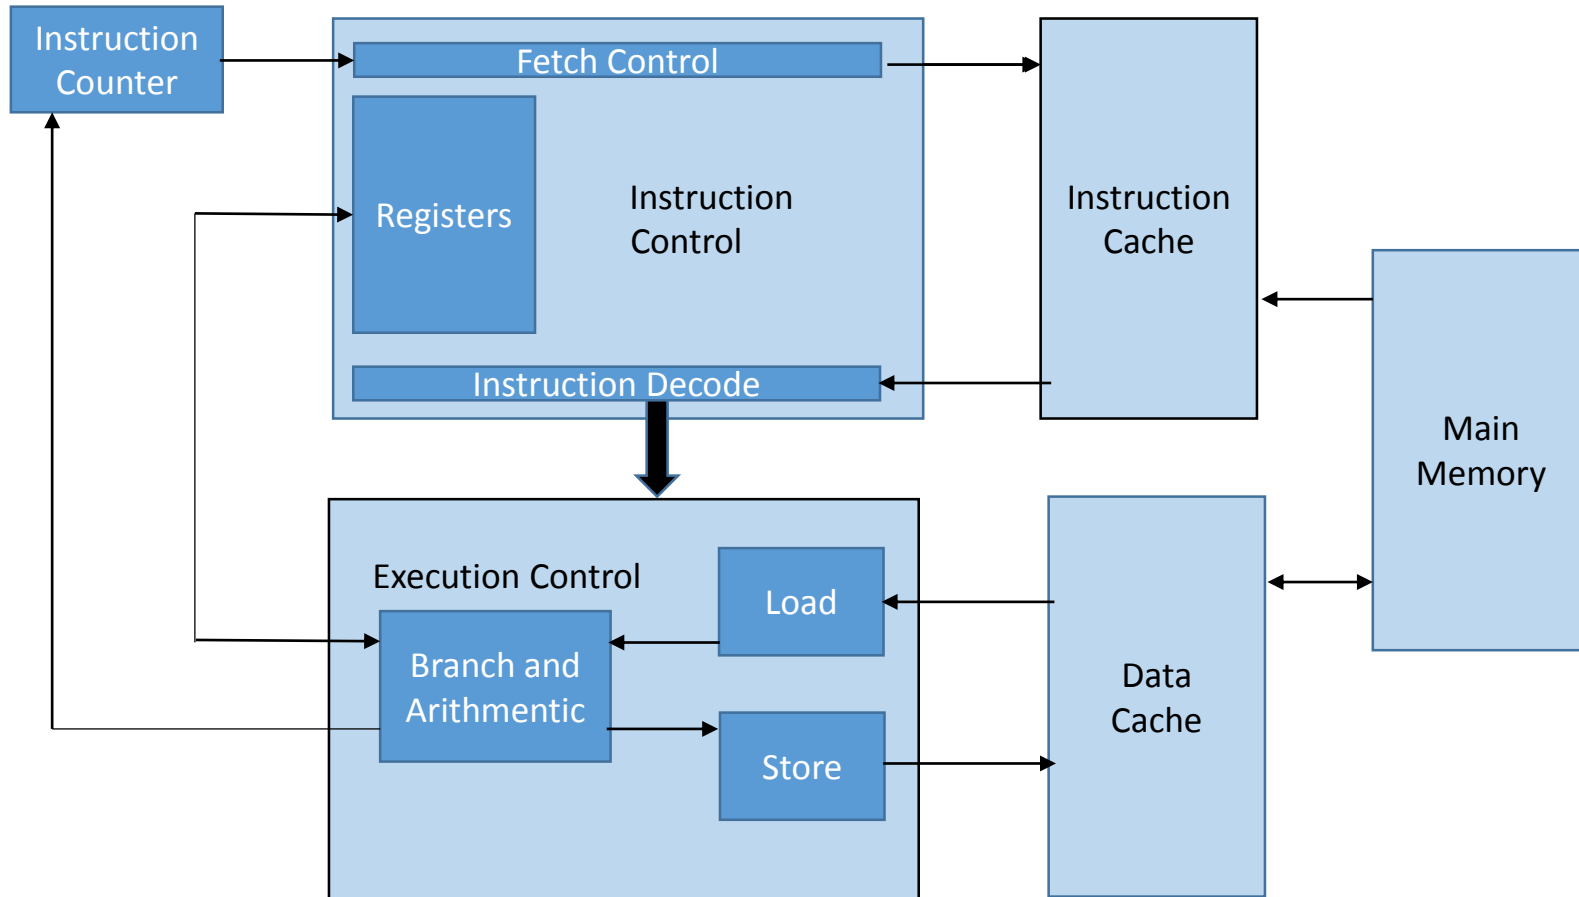
Where

- M is the number of cycles for a memory operation
- R is the number of cycles for a register operation
- A is the number of cycles for an arithmetic operation
- F is the number of cycles for a function.
- S is the number of cycles penalty per byte for instruction fetch

The a, b, c, d are the number of occurrences of each type of operation and s is the size of the code in bytes for a particular loop.

Run many different instances of loops, gather the CPEs, perform multiple regression to solve for the times of each type. But we need a way to compensate for overhead.

## Machine organization






## Steps in CPU operation

Machine Instructions  
in Memory

add (%rdi,%rax,4),%edx
add \$0x1,%rax
cmp %eax,%ecx
jg <20>
mov %edx,(%rsi)



## instruction breakdown

### Fixed point add

```
add %ecx,0x0c(%edx)
```

```
load 0x0c(%edx)
```

```
add %ecx
```

```
store 0x0c(%edx)
```

### Floating point multiply

```
mulss x0c(%edx),%xmm0
```

```
load 0x0c(%edx)
```

```
multiply %ecx
```

```
work with exponents
```


```
work with mantissas
```

```
round off
```

```
replace %xmm0
```

## register renaming or aliasing

During pipelined instructions:

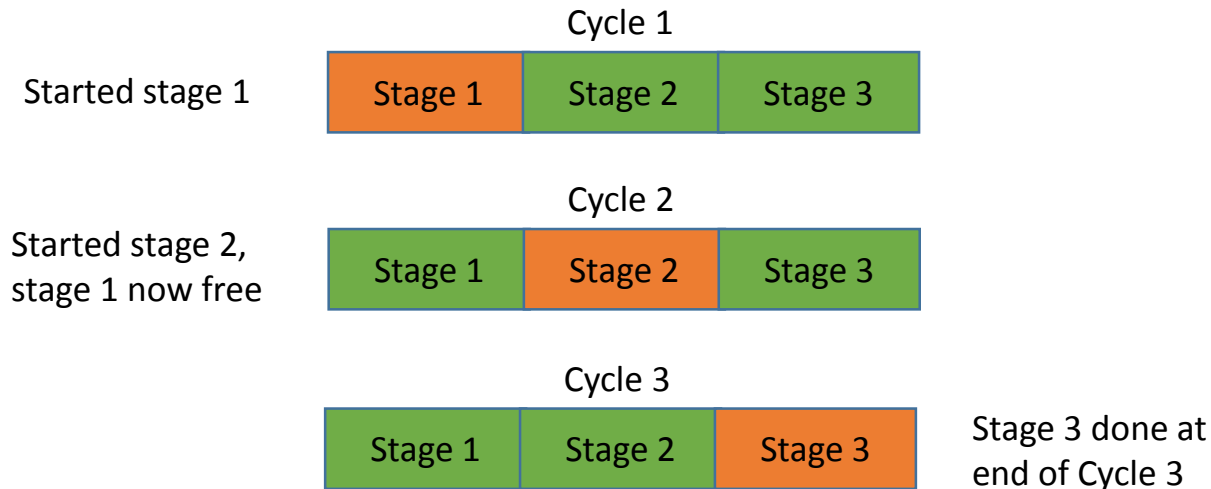


## latency, issue time, throughput

Latency and issue time is measured in cycles.

latency: cycles from start to finish: how fast can you do it, once you have started.

issue time: minimum number of cycles between successive operations. How long do you have to wait before you can do it again? Can be due to multiple units or the operation is done in stages.



Latency: 3, issue time: 1

Throughput is  $1/(\text{issue time}) = \text{number of unrelated operations possible in one cycle}$ .

## Intel I7 measured latencies, issue times

Division is very high, data dependent. Too complicated for our discussions. Focus on add, multiply for integer and floating point.

All issue times for those are 1.

If the issue time is 1, you can do 1 instruction per cycle, when the latency is 1. If the latency is  $> 1$ , the instructions must be “unrelated” if the next operation needs the outcome of the prior operation in order to do 1 instruction per cycle.

	I+	I*	F+	F*	D*
Latency	1	3	3	4	5
Throughput	1	1	1	1	1

I = Integer, F = Floating point single, D = Floating point double,  
+ = add, \* = multiply

Latency gives the maximum time, if you have to wait for the previous result.

Throughput gives the minimum time if you can execute them “maximum pipelined”.


## The trick is to determine whether we are latency or throughput bound

Consider the following loop (combine4 with OP = \*, IDENT = 1)

```
mulss (%rax),%xmm0
add $0x4,%rax
cmp %rdx,%rax
jne loop
```

There are 3 registers involved: how do they change and how are they involved in the loop?


Read-only, write-only, local, loop.



Which operations depend on the outcome of the previous iteration of the loop?

## Eliminate the dependencies which are not loop dependent

```
mulss (%rax),%xmm0
add $0x4,%rax
cmp %rdx,%rax
jne loop
```



Two paths: which path has the highest latency? No-brainer: add path: latency 1, mulss path: latency 4. It is the “critical” path which determines the CPE for the combine function.

The CPE is not additive but you have to figure out the overlapping operations.

## Back to loop unrolling

data\_t is a data type: int or float or double.  
 IDENT is 1  
 OP is \*

```
void combine4(vec_ptr v, data_t *dest)
{
    int l, l = vec_length(v);
    data_t *data = get_vec_start(v);
    data_t acc = IDENT;

    for (i = 0; i < l; i++)
        acc = acc OP data[i];

    *dest = acc;
}
```



Half the number of times, do the OP twice

data\_t is a data type: int or float or double.

IDENT is 1

OP is \*

```
void combine5(vec_ptr v, data_t *dest)
{
    int l, l = vec_length(v);
    data_t *data = get_vec_start(v);
    data_t acc = IDENT;

    for (i = 0; i < l ; i+=2)
        acc = ( acc OP data[i] ) OP data[i+1];

    for( ; i<l ; i++ )
        acc = acc OP data[i] ;

    *dest = acc ;
}
```

## Timings for multiplication ( OP = \*, IDENT = 1, data\_t = float )

Combine4

Time= 0.1527 n= 100000

Time= 0.3016 n= 200000

Time= 0.4523 n= 300000

Time= 0.6029 n= 400000

Time= 0.7537 n= 500000

Combine5 (unroll x 2)

Time= 0.1515 n= 100000

Time= 0.3016 n= 200000

Time= 0.4523 n= 300000

Time= 0.6029 n= 400000

Time= 0.7536 n= 500000

Expected twice the improvement but got none! Why?

## Compiled code

### Combine4

```
400b28 <+29>: mulss (%rax),%xmm0
400b2c <+33>: add  $0x4,%rax
400b30 <+37>: cmp  %rdx,%rax
400b33 <+40>: jne  0x400b28 <+29>
```

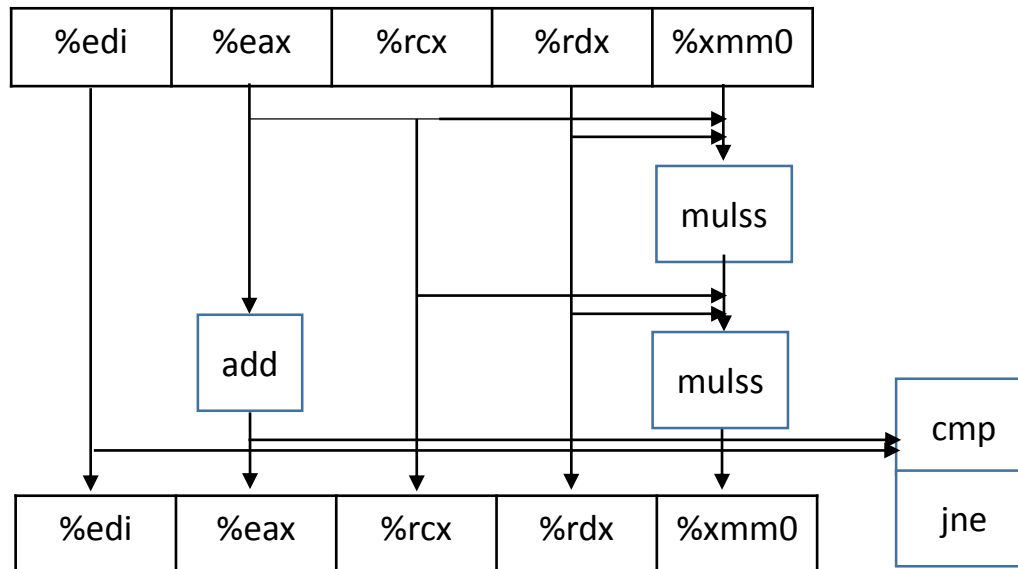
### Combine5

```
400b27 <+28>: movslq %eax,%rcx
400b2a <+31>: mulss (%rdx,%rcx,4),%xmm0
400b2f <+36>: mulss 0x4(%rdx,%rcx,4),%xmm0
400b35 <+42>: add  $0x2,%eax
400b38 <+45>: cmp  %edi,%eax
400b3a <+47>: jl   0x400b27 <+28>
```

## Flow diagram

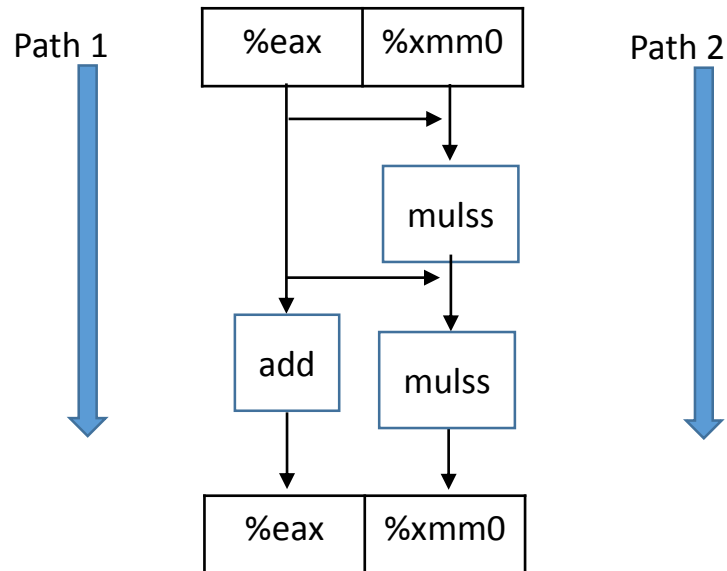
combine5 with OP = \*, IDENT = 1

```
movslq %eax,%rcx
mulss (%rdx,%rcx,4),%xmm0
mulss 0x4(%rdx,%rcx,4),%xmm0
add $0x2,%eax
cmp %edi,%eax
jl 0x400b27 <+28>
```



Which operations depend on the outcome of the previous iteration of the loop?

Eliminate read only and local registers



Path 2 is the critical path with latency of 8: done half of the time

## Increase parallelism by using two accumulators

data\_t is a data type: int or float or double.

IDENT is 1

OP is \*

```
void combine6(vec_ptr v, data_t *dest)
{
    int l, l = vec_length(v);
    data_t *data = get_vec_start(v);
    data_t acc1 = IDENT;
    data_t acc2 = IDENT;

    for (i = 0; i < l ; i+=2)
    {
        acc1 = acc1 OP data[i];
        acc2 = acc2 OP data[i+1];
    }

    for( ; i<l ; i++)
        acc1 = acc1 OP data[i];

    *dest = acc1 OP acc2;
}
```

## Timings for multiplication ( OP = \*, IDENT = 1, data\_t = float )

Combine5 (unroll x 2)

Time= 0.1515 n= 100000

Time= 0.3016 n= 200000

Time= 0.4523 n= 300000

Time= 0.6029 n= 400000

Time= 0.7536 n= 500000

Combine6(unroll x 2), multiple accumulators

Time= 0.0783 n= 100000

Time= 0.1511 n= 200000

Time= 0.2264 n= 300000

Time= 0.3019 n= 400000

Time= 0.3775 n= 500000

Success! Why?

## Compiled code

### Combine5

```
400b27 <+28>: movslq %eax,%rcx
400b2a <+31>: mulss (%rdx,%rcx,4),%xmm0
400b2f <+36>: mulss 0x4(%rdx,%rcx,4),%xmm0
400b35 <+42>: add  $0x2,%eax
400b38 <+45>: cmp  %edi,%eax
400b3a <+47>: jl   0x400b27 <+28>
```

### Combine6

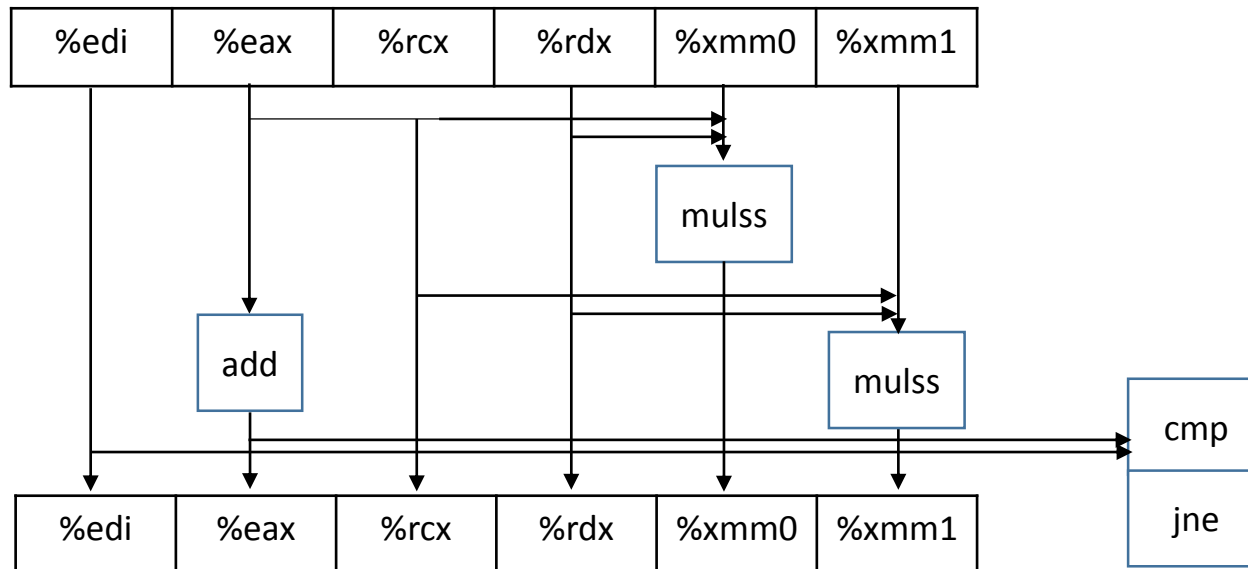
```
400b27 <+28>: movslq %eax,%rdx
400b2a <+31>: mulss (%rcx,%rdx,4),%xmm1
400b2f <+36>: mulss 0x4(%rcx,%rdx,4),%xmm0
400b35 <+42>: add  $0x2,%eax
400b38 <+45>: cmp  %r8d,%eax
400b3b <+48>: jl   0x400b27 <+28>
```



## Flow diagram

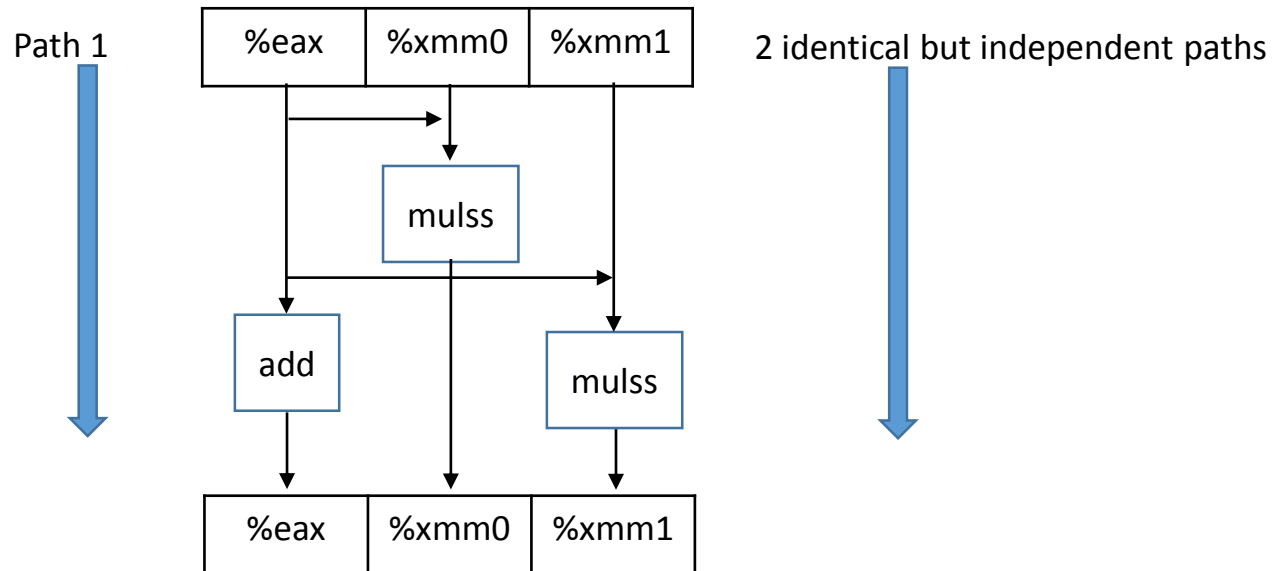
combine6 with OP = \*, IDENT = 1

```
movslq %eax,%rdx
mulss (%rcx,%rdx,4),%xmm1
mulss 0x4(%rcx,%rdx,4),%xmm0
add $0x2,%eax
cmp %r8d,%eax
jl 0x400b27 <+28>
```



Which operations depend on the outcome of the previous iteration of the loop?

Eliminate read only and local registers



Right two paths are critical with latency of 4: done half of the time

## Reassociation

data\_t is a data type: int or float or double.  
 IDENT is 1  
 OP is \*

```
void combine7(vec_ptr v, data_t *dest)
{
    int l, l = vec_length(v);
    data_t *data = get_vec_start(v);
    data_t acc = IDENT;

    for (i = 0; i < l ; i+=2)
        acc = acc OP ( data[i] OP data[i+1] );

    for( ; i<l ; i++ )
        acc = acc OP data[i];

    *dest = acc;
}
```

## Timings for multiplication ( OP = \*, IDENT = 1, data\_t = float )

Combine5 (unroll x 2)

Time= 0.1515 n= 100000

Time= 0.3016 n= 200000

Time= 0.4523 n= 300000

Time= 0.6029 n= 400000

Time= 0.7536 n= 500000

Combine7 (unroll x 2 with reassociation )

Time= 0.0771 n= 100000

Time= 0.1511 n= 200000

Time= 0.2266 n= 300000

Time= 0.3019 n= 400000

Time= 0.3777 n= 500000

Twice the speed. Similar to Combine6. Why?

## Compiled code

## Combine5

```
400b27 <+28>: movslq %eax,%rcx
400b2a <+31>: mulss (%rdx,%rcx,4),%xmm0
400b2f <+36>: mulss 0x4(%rdx,%rcx,4),%xmm0
400b35 <+42>: add  $0x2,%eax
400b38 <+45>: cmp  %edi,%eax
400b3a <+47>: jl   0x400b27 <+28>
```

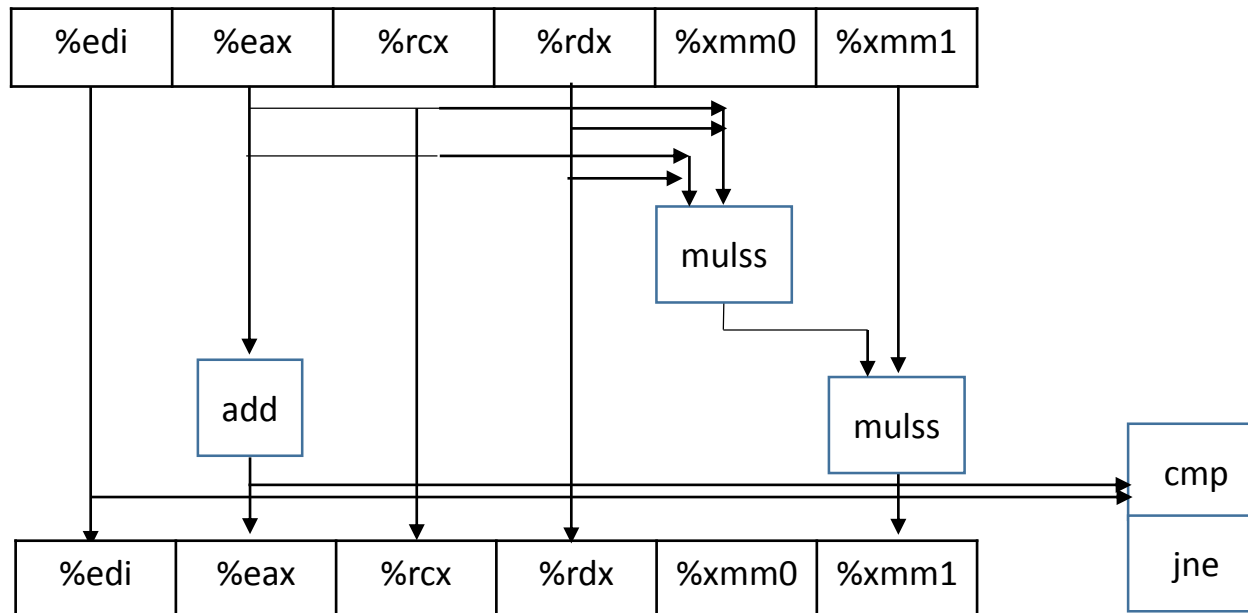
## Combine7

```
400b27 <+28>: movslq %eax,%rcx
400b2a <+31>: movss (%rdx,%rcx,4),%xmm1
400b2f <+36>: mulss 0x4(%rdx,%rcx,4),%xmm1
400b35 <+42>: mulss %xmm1,%xmm0
400b39 <+46>: add  $0x2,%eax
400b3c <+49>: cmp  %edi,%eax
400b3e <+51>: jl   0x400b27 <+28>
```

## Flow diagram

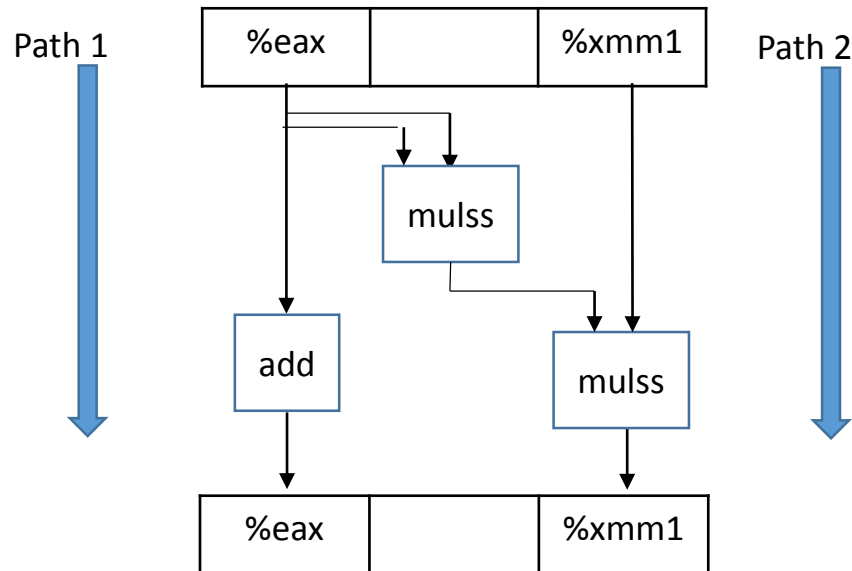
Combine7 with OP = \*, IDENT = 1

```
movslq %eax,%rcx
movss (%rdx,%rcx,4),%xmm1
mulss 0x4(%rdx,%rcx,4),%xmm1
mulss %xmm1,%xmm0
add $0x2,%eax
cmp %edi,%eax
jl 0x400b27 <+28>
```



Which operations depend on the outcome of the previous iteration of the loop?

Eliminate read only and local registers



Path 2 is the critical path with latency of 4: done half of the time. The mulss in the middle can be done in parallel for the next iteration while the one on the right finishes for the previous.

## Additional selected topics

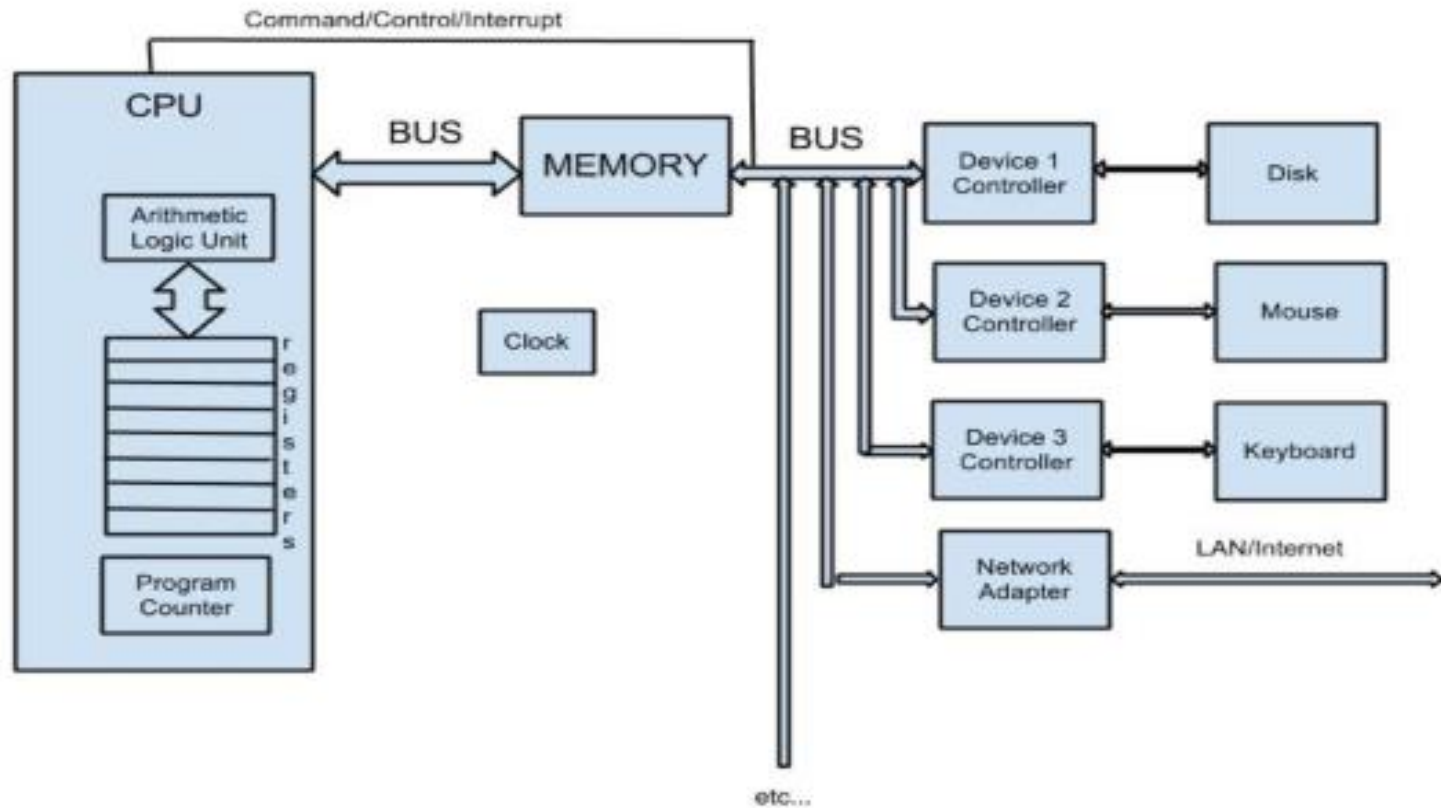
- Multiple accumulators limited by the number of registers
- Branch prediction errors can be minimized by using conditional moves
- Branch prediction errors can be affected by the data
- Reading from memory is usually slower than writing to memory (cache)
- Locate the fat by using a profiler




## Technologies

- Random Access Memories
  - Disk Memories
  - Solid State Disks
- 
- Random Access Memories
    - Read only
      - Programmable
      - Erasable
      - Flash
    - SRAM Static -- fastest, more expensive, stable: caches
    - DRAM Dynamic -- slower, less expensive: main memory speed up by storing parts of words in parallel.
      - Lots of enhanced types

# Overview of a Digital Computer



## Disk Drive



Address on drive: cylinder (0-X), track (0-5), sector (0-Y)

## Disk Drive Terminology

Recording density bytes/inch

Track density tracks/diameter

Areal density = ( Track \* Recording )

Capacity:

Bytes/sector

Sectors/track

Tracks/surface

Surfaces/platter

Platters/disk

Disk capacity is the product of these.

Speed:

Seek time

Rotational latency

Transfer time

Time to read is the sum of these.

## Memory locality

Spatial

Temporal

```
int sumvec(int v[N])
{
    int i, sum = 0;

    for (i = 0; i < N; i++)
        sum += v[i];
    return sum;
}
```

Address	0	4	8	12	16	20	24	28
Contents	v0	v1	v2	v3	v4	v5	v6	v7
Access order	1	2	3	4	5	6	7	8

```
int sumarrayrows(int a[M][N])
{
    int i, j, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < N; j++)
            sum += a[i][j];
    return sum;
}
```

Address	0	4	8	12	16	20	24	28
Contents	a00	a01	a02	a10	a11	a12	a20	a21
Access order	1	2	3	4	5	6	7	8

For N = 3

## Memory locality

```
int sumarraycols(int a[M][N])  
{  
    int i, j, sum = 0;  
  
    for (j = 0; j < N; j++)  
        for (i = 0; i < M; i++)  
            sum += a[i][j];  
    return sum;  
}
```

Accesses every Nth location in memory, M times. Bad spatial locality.

Program instruction display good spatial and temporal locality. Large loops interrupt the spatial and temporal locality.

## Ways to improve locality

Consider `src[n][n]` and `dst[n][n]`. Both functions “rotate” the array by 90 degrees.

```
void rotate1( int *src, int *dst, int n )
{
    int i, j;

    for (i = 0; i < n; i++)
        for (j = 0; j < n; j++)
            dst[(n-1-j)*n+i] = src[(i*n+j)] ;
}
```

```
void rotate2(int *src, int *dst, int n) {
    int i, j;
    int ii, jj;

    for(ii=0; ii < n; ii+=8)
        for(jj=0; jj < n; jj+=8)
            for(i=ii; i < ii+8; i++)
                for(j=jj; j < jj+8; j++)
                    dst[(n-1-j)*n+i] = src[(i)*n+j];
}
```

This is called “blocking” to improve spatial locality.

## Timings

### rotate1

Time= 0.0004 n= 16  
Time= 0.0010 n= 32  
Time= 0.0041 n= 64  
Time= 0.0448 n= 128  
Time= 0.2497 n= 256  
Time= 1.3555 n= 512  
Time= 5.5756 n= 1024  
Time= 76.0182 n= 2048

### rotate2

Time= 0.0009 n= 16  
Time= 0.0032 n= 32  
Time= 0.0099 n= 64  
Time= 0.0338 n= 128  
Time= 0.0876 n= 256  
Time= 0.3815 n= 512  
Time= 2.9030 n= 1024  
Time= 26.7598 n= 2048



## Haber's locality figure of merit

min = moving min of 8 items, same for max

min-max = absolute value min-max

Rotate1

Partial list for n = 16

Rotate2

i	j	$(n-1-j)*n+i$	$(i*n+j)$	min	max	min-max
0	0	240	0			0
0	1	224	1			0
0	2	208	2			0
0	3	192	3			0
0	4	176	4			0
0	5	160	5			0
0	6	144	6			0
0	7	128	7	128	240	112
0	8	112	8	112	224	112
0	9	96	9	96	208	112
0	10	80	10	80	192	112
0	11	64	11	64	176	112
0	12	48	12	48	160	112
0	13	32	13	32	144	112
0	14	16	14	16	128	112
0	15	0	15	0	112	112
1	0	241	16	0	241	241
1	1	225	17	0	241	241
1	2	209	18	0	241	241
1	3	193	19	0	241	241
1	4	177	20	0	241	241
1	5	161	21	0	241	241
1	6	145	22	0	241	241
1	7	129	23	129	241	112
1	8	113	24	113	225	112
1	9	97	25	97	209	112
1	10	81	26	81	193	112
1	11	65	27	65	177	112
1	12	49	28	49	161	112
1	13	33	29	33	145	112

i	j	$(n-1-j)*n+i$	$(i*n+j)$	min	max	min-max
0	0	240	0			
0	1	224	1			
0	2	208	2			
0	3	192	3			
0	4	176	4			
0	5	160	5			
0	6	144	6			
0	7	128	7	128	240	112
1	0	241	16	128	241	113
1	1	225	17	128	241	113
1	2	209	18	128	241	113
1	3	193	19	128	241	113
1	4	177	20	128	241	113
1	5	161	21	128	241	113
1	6	145	22	128	241	113
1	7	129	23	129	241	112
2	0	242	32	129	242	113
2	1	226	33	129	242	113
2	2	210	34	129	242	113
2	3	194	35	129	242	113
2	4	178	36	129	242	113
2	5	162	37	129	242	113
2	6	146	38	129	242	113
2	7	130	39	130	242	112
3	0	243	48	130	243	113
3	1	227	49	130	243	113
3	2	211	50	130	243	113
3	3	195	51	130	243	113
3	4	179	52	130	243	113
3	5	163	53	130	243	113

Haber's locality figure of merit  
Average of min – max

rotate1

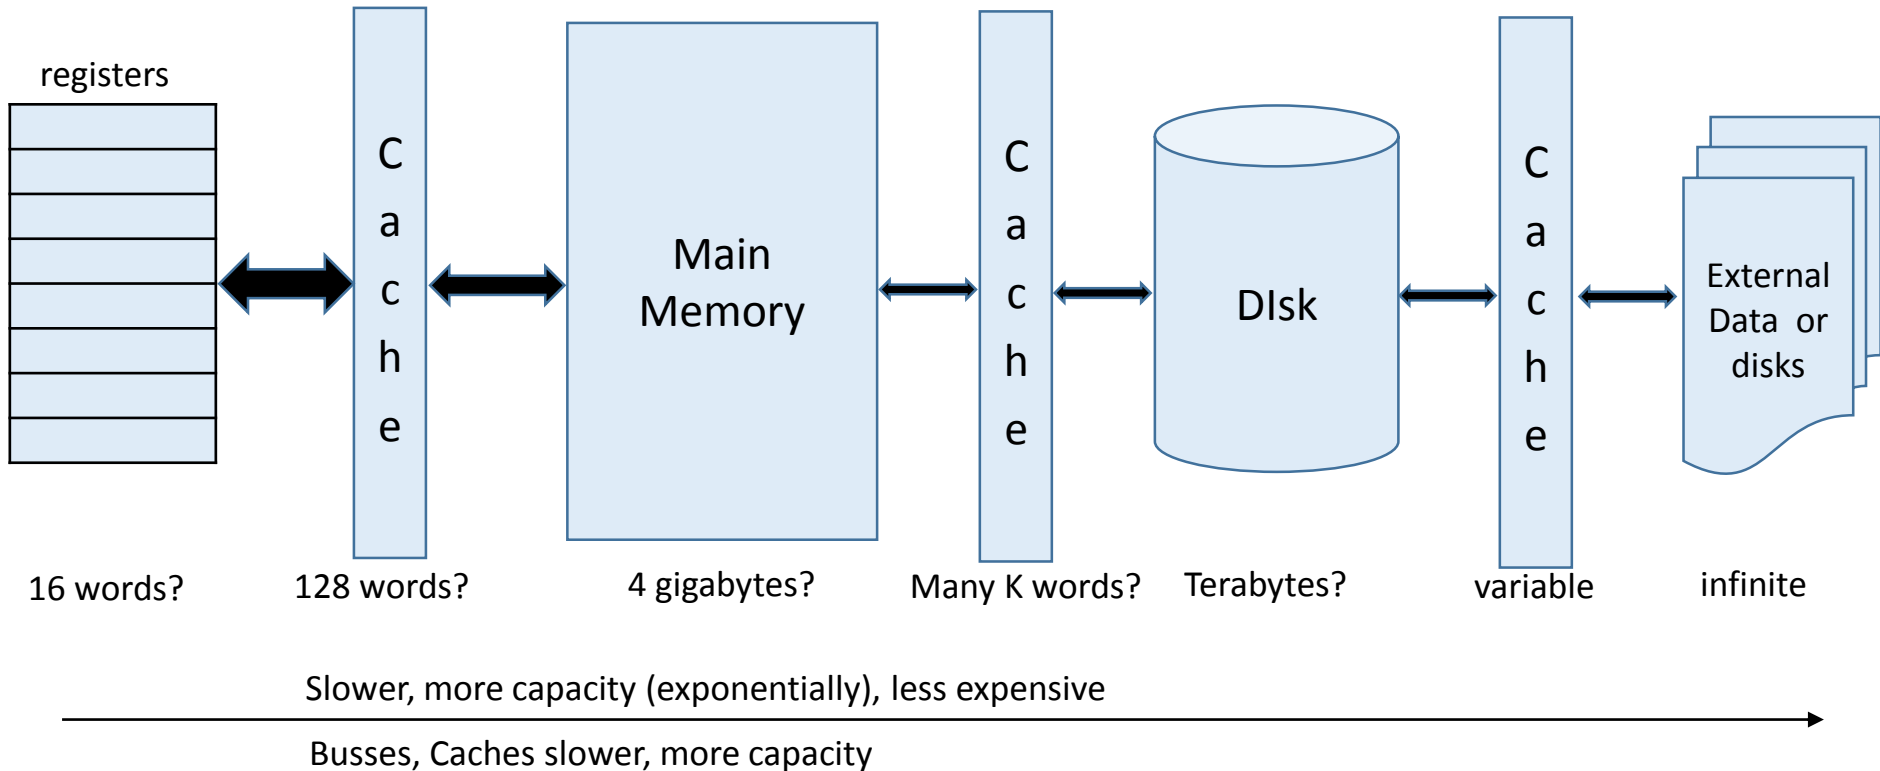
n= 16	166.40
n= 32	388.08
n= 64	834.64
n= 128	1729.83
n= 256	3521.40
n= 512	7105.20
n= 1024	14272.93
n= 2048	30051.08

rotate2

n= 16	116.81
n= 32	241.23
n= 64	492.40
n= 128	995.97
n= 256	2003.17
n= 512	4018.24
n= 1024	8050.08
n= 2048	17160.27

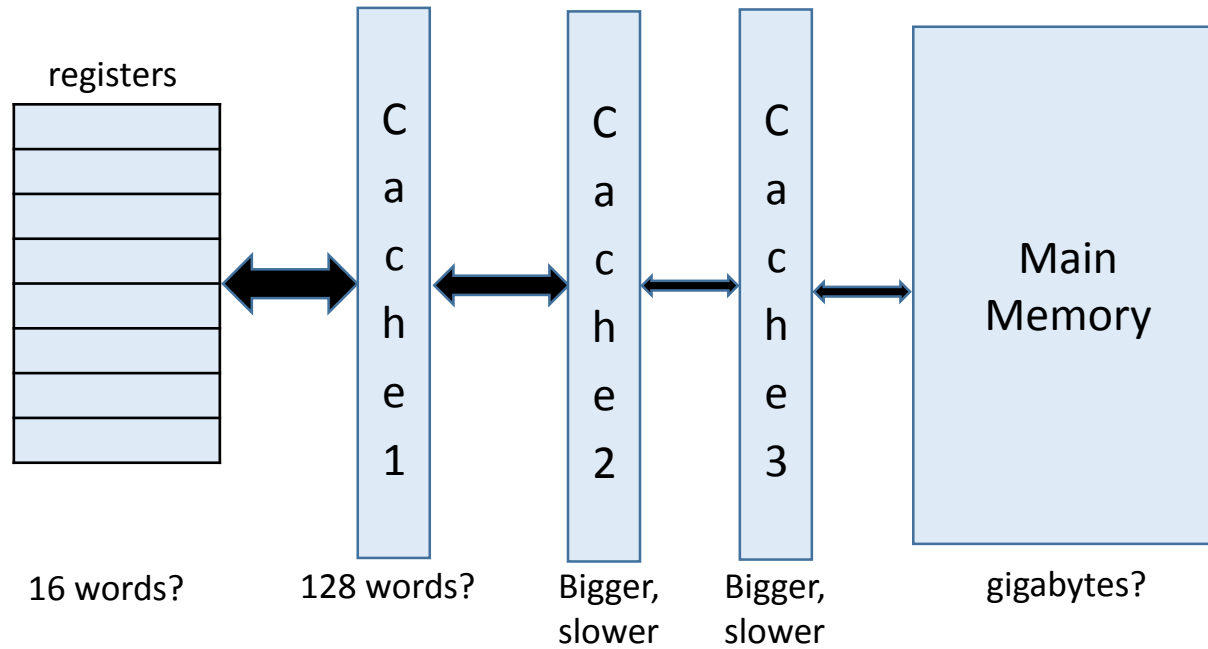
# Memory Hierarchy

Faster to slower, more expensive to cheaper, less capacity to greater



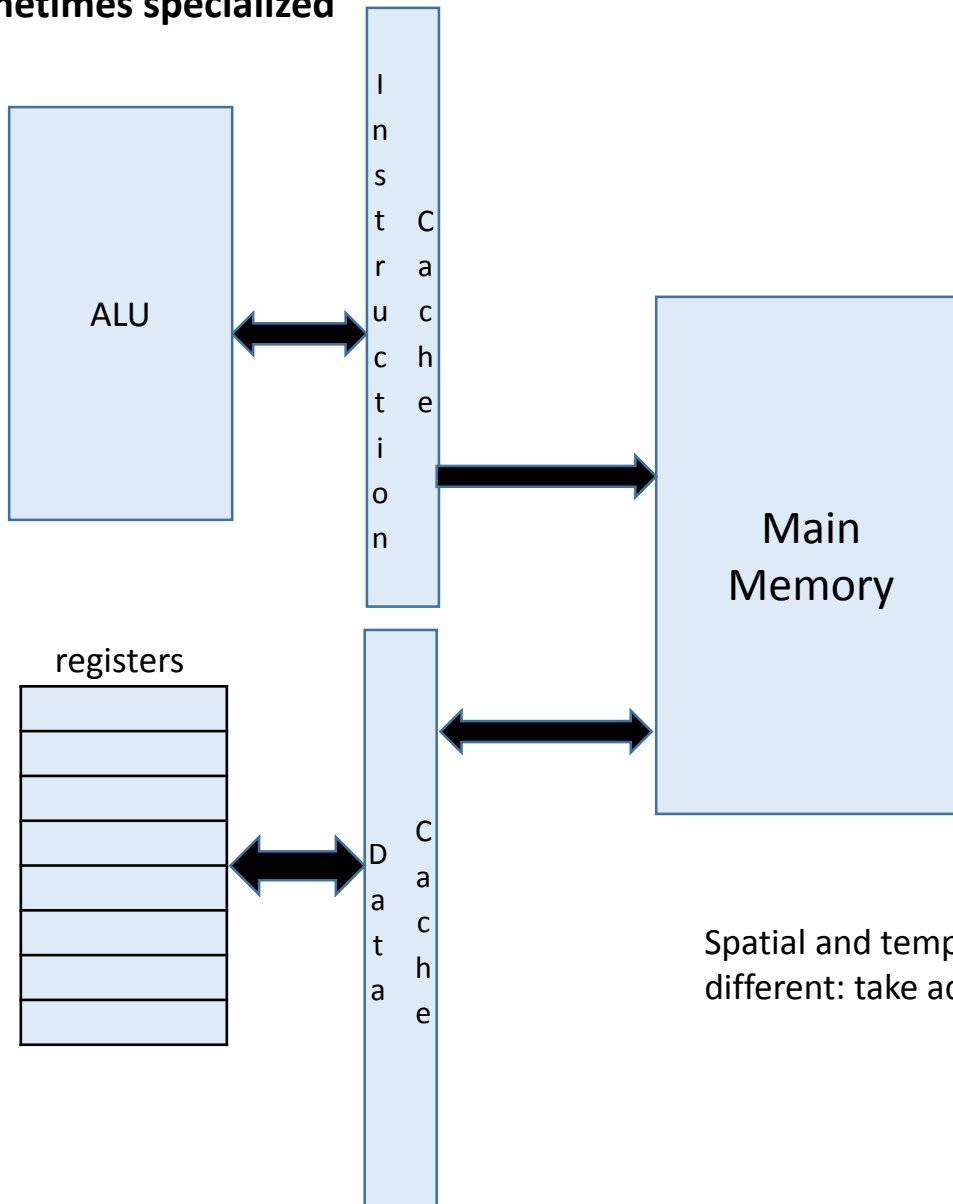
# Memory Hierarchy

Caches are sometime cascaded



Slower, more capacity (exponential), less expensive

Caches are sometimes specialized



Spatial and temporal locality are grossly different: take advantage

## The cache principle

- An intermediary between a slower and faster memory
- A device with memory and an autonomous program
- Speed up the memory access

$M = 2^m$  bytes is the main memory size (m bit unsigned integer address)

Divide slower memory into fixed size blocks (usually small)  $B = 2^b$

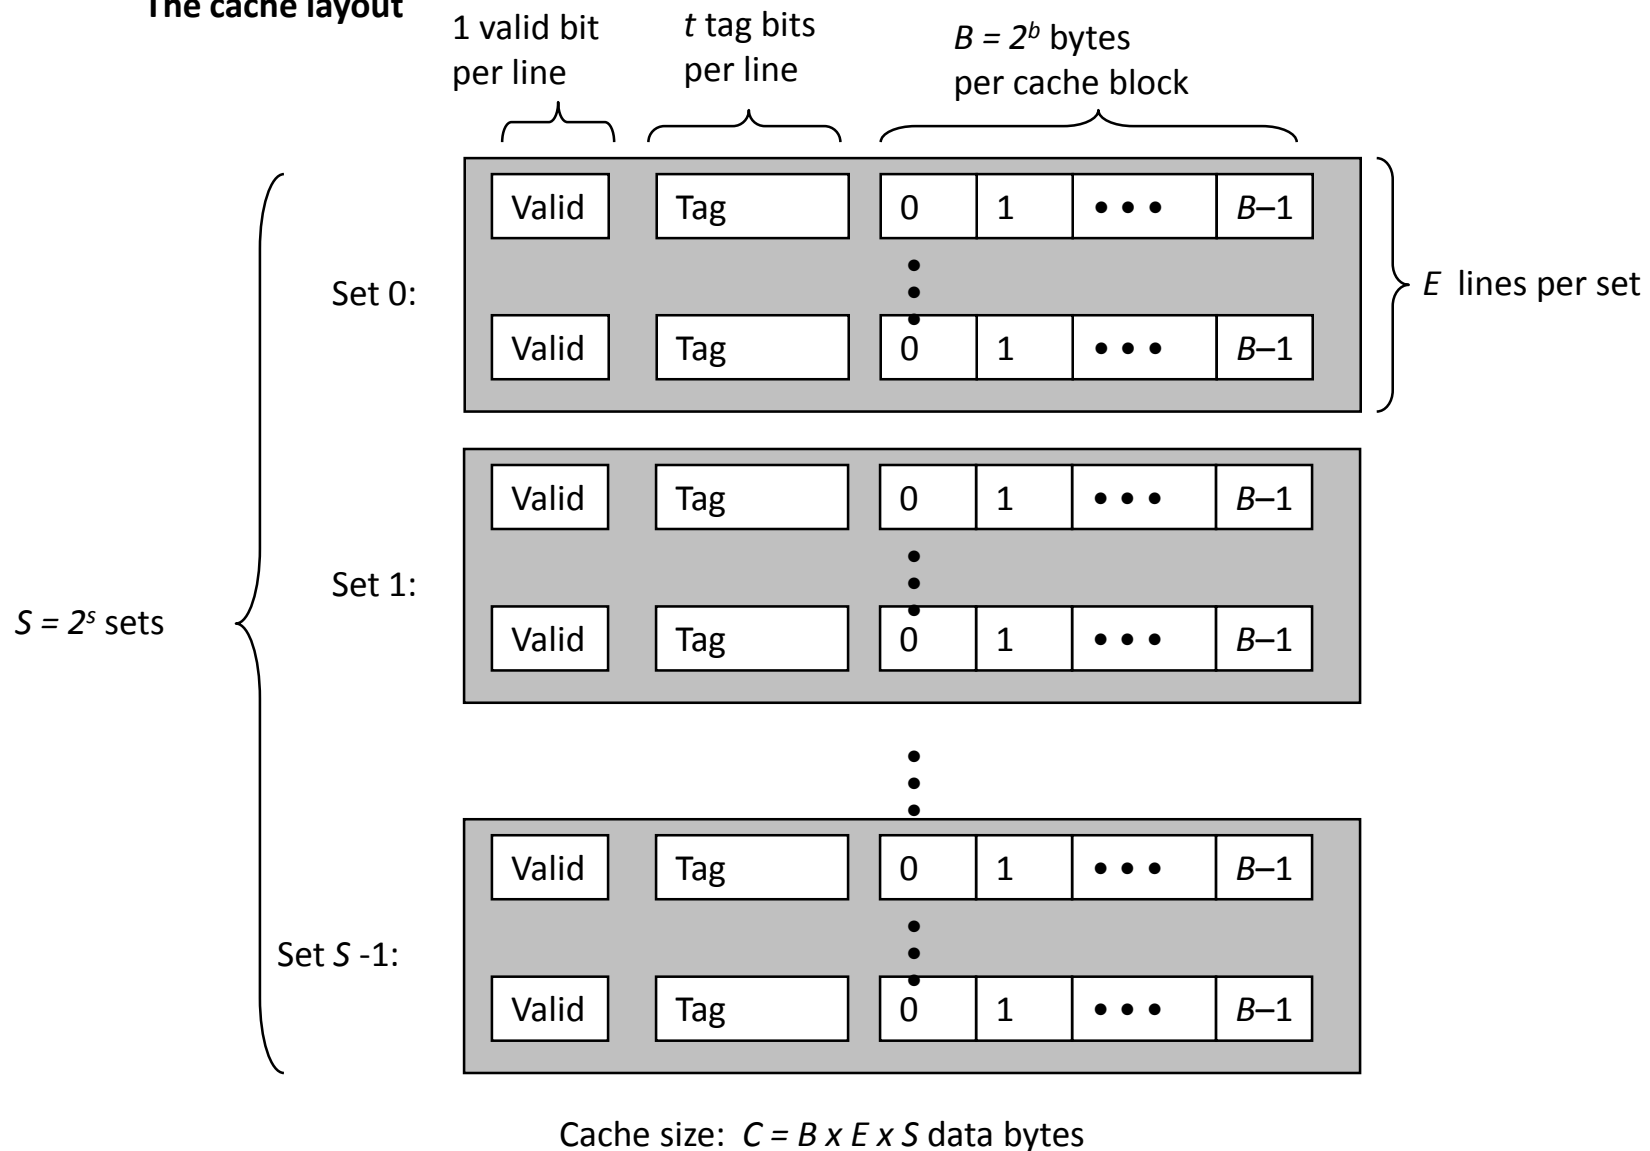
Build cache with a certain number of sets  $S = 2^s$  each containing a B sized block

Some caches have a certain number of lines per set: E. A special case is where  $E = 1$ , call direct mapped cache.

The cache size is  $B \times E \times S = C$  bytes. The cache “covers”  $C/M$  of the memory

Each set has a tag which is t bits.

## The cache layout



## Summary of cache layout

$M = 2^m$  bytes is the main memory size (m bit unsigned integer address)

The cache has

$S = 2^s$  sets

$E$  lines per set


$B = 2^b$  block size in bytes

Each set has a tag which is m bits long

The tag is divided into 3 parts: the block offset, the set index and the tag.  
Since the tag is a total of m bits then:

$$t = m - s - b$$


A m bit memory address is divided into 3 parts the same way:



This means that each  $2^t$  size memory block is mapped into the same cache area.



## A 32 bit memory address example



and in this case, we have  $t = 8$ ,  $s = 8$  and  $b = 16$ . This is a huge cache!

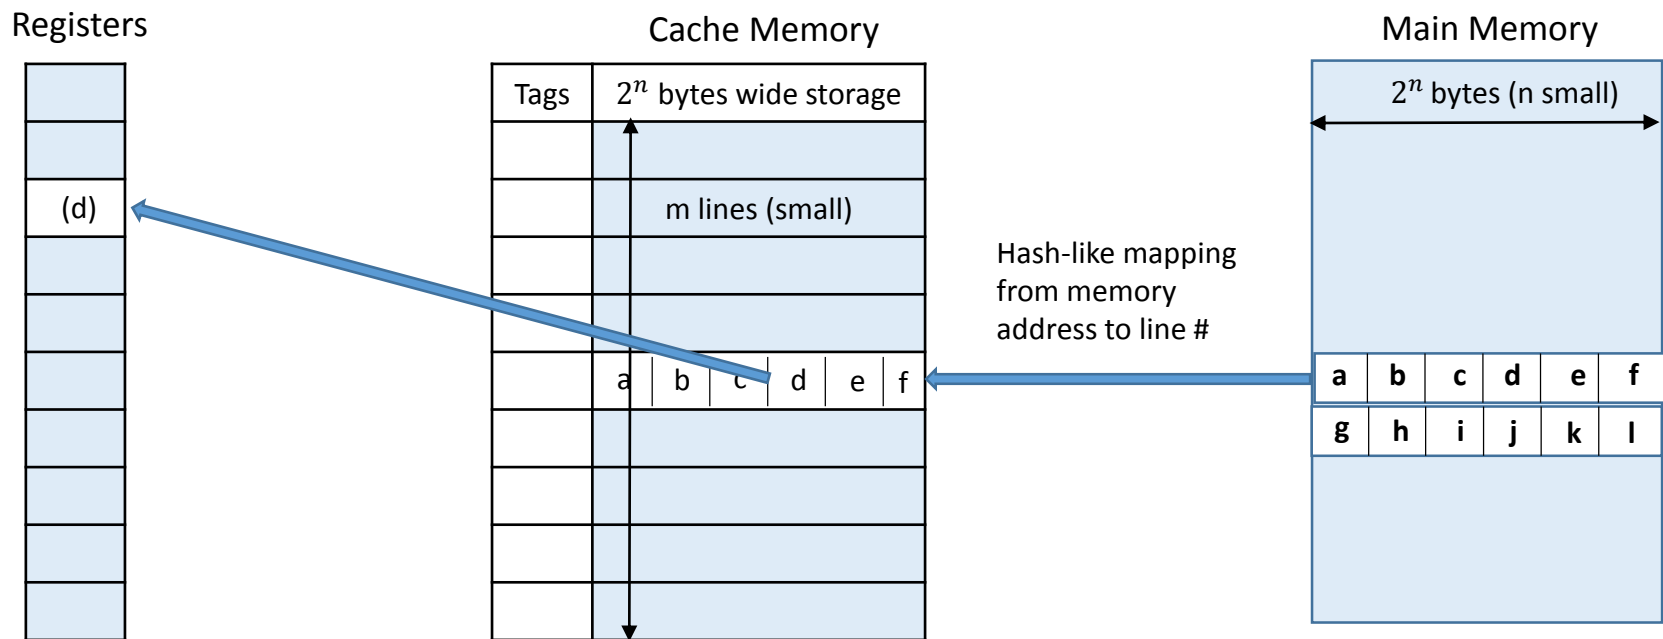
$$B = 65,536$$

$$S = 256$$

## The cache principle

- An intermediary between a slower and faster memory
- A device with memory and an autonomous program
- Speed up the memory access


Clean start : retrieve( address (d), bytes (1, 2, 4 or 8))




## The cache principle

Next Instruction: retrieve( address (e), bytes (1, 2, 4 or 8))


Registers



Cache Memory

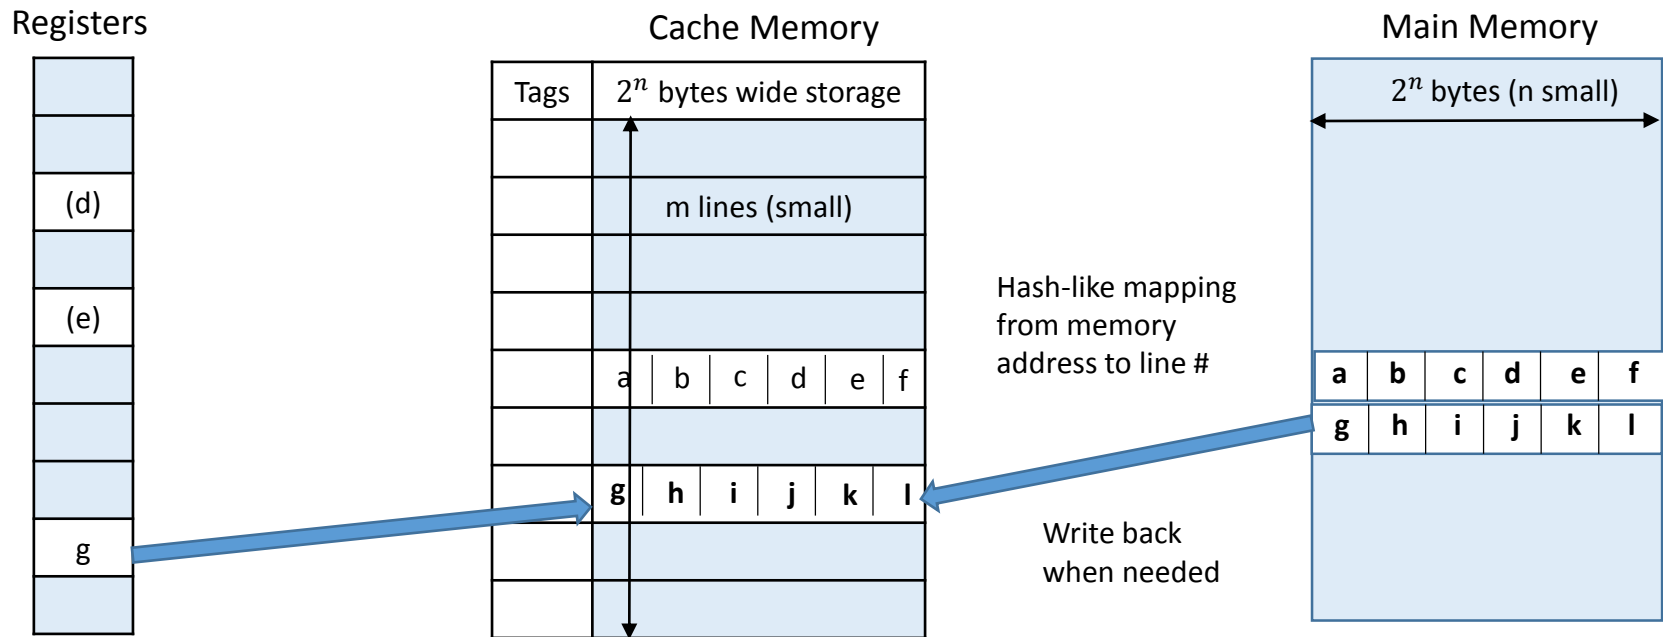


Main Memory



## The cache principle

Next Instruction:      write( address (g), bytes (1, 2, 4 or 8))



## Summary of Cache Operation

### Read from Memory

Item in Cache ?

Read Hit :

Return item

Read Miss :

Select cache line

Read item and surroundings into cache

Return item

### Write to memory

Item in Cache ?

Write Hit:

Write item into cache

Write Miss :

Select cache line

Read item and surroundings into cache

Write item into cache

### Select Cache Line:

Map memory address to cache line number using hash-like formula

Cache line in use?

No: Finished

Yes:

Correct memory segment? Yes: finished

No: Dirty bit set? : Write line from cache to memory

**Cache sample: direct mapped -> E = 1**

lets try a cache like this:

```
#define S    4  // (4 cache sets)
#define E    1  // (direct mapped cache)
#define B   16 // (16 elements in each block)
#define T    2  // (2 tag bits)
#define M   256 // (256 byte memory)

struct cache_t
{
    char valid ; // simulate a bit with this
    char dirty ; // an element was stored but not written
    int tag ;
    char *block ;
} cache[S] ;

char memory[M] ;
```

We would initialize the cache by:

```
for( i=0; i<S; i++ )
{
    cache[i].block = (int *) malloc(B) ;
    cache[i].valid = 0 ;
}
and
int s = log2(S) ;
int b = log2(B) ;
int m = log2(M) ;
```

## Memory retrieval

Retrieve something from memory address:  $a$ .

$\text{int } si = (a \gg m-T-s) \% S;$                       the set index ( $m-T-s = b$ )

$\text{int } ta = a \gg m-T;$                                   the tag

$\text{int } bo = a \% B;$                                       the block offset

$\text{char hit};$

Example: retrieve from address 152

ta		si		bo			
1	0	0	1	1	0	0	0

$ta = 2;$

$si = 1;$

$bo = 8;$

## Memory retrieval

```

If ( cache.valid[si] )
    if ( ta == cache.tag[si] >> m-T )    // the tag
        hit = 1 ;
    else
        {
            if ( cache.dirty[si] )
                for( i = (cache.tag[si] >> b) << b, j = 0; j < B; i++, j++ ) // writes the cache to memory
                    memory[i] = cache.block[si][j] ;
            hit = 0 ;
        }
    else
        hit = 0 ;

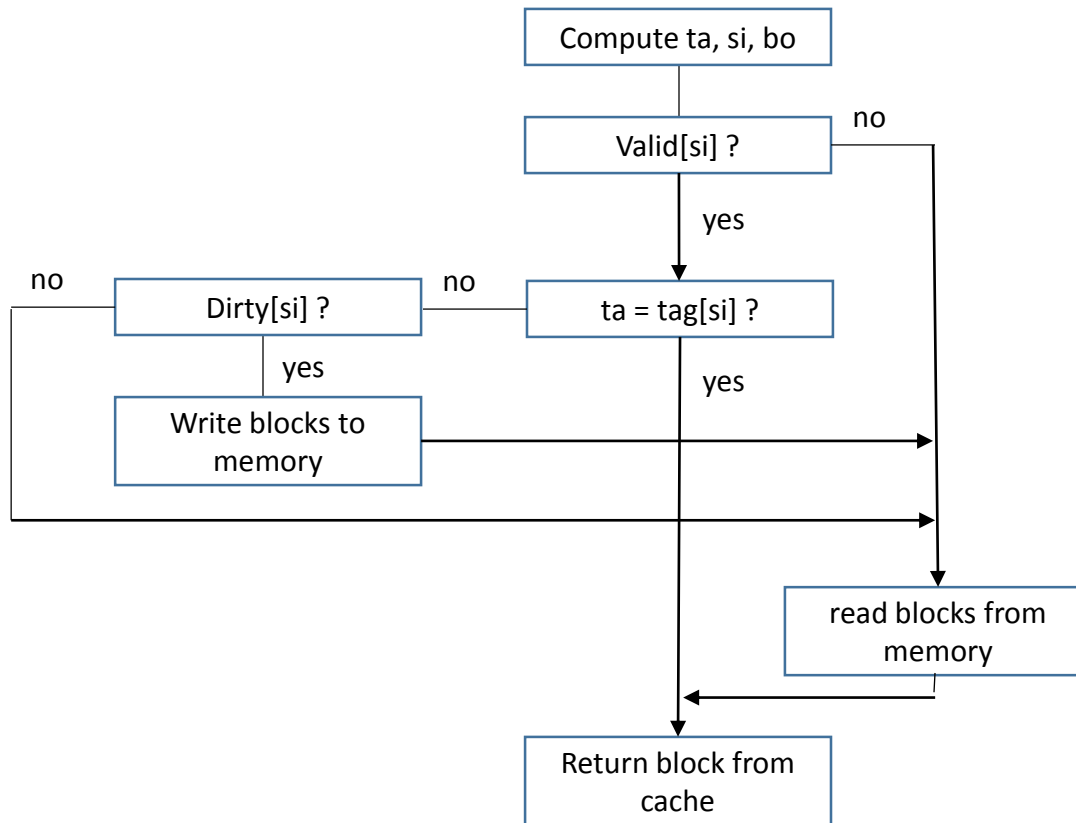
if ( ~hit )
    {
        for( i = ( a >> b ) << b, j = 0; j < B; i++, j++ )
            cache.block[si][j] = memory[i] ;
        cache.valid[si] = 1 ;
        cache.dirty[si] = 0 ;
        cache.tag[si] = a ;
    }

return cache.block[si][bo] ;

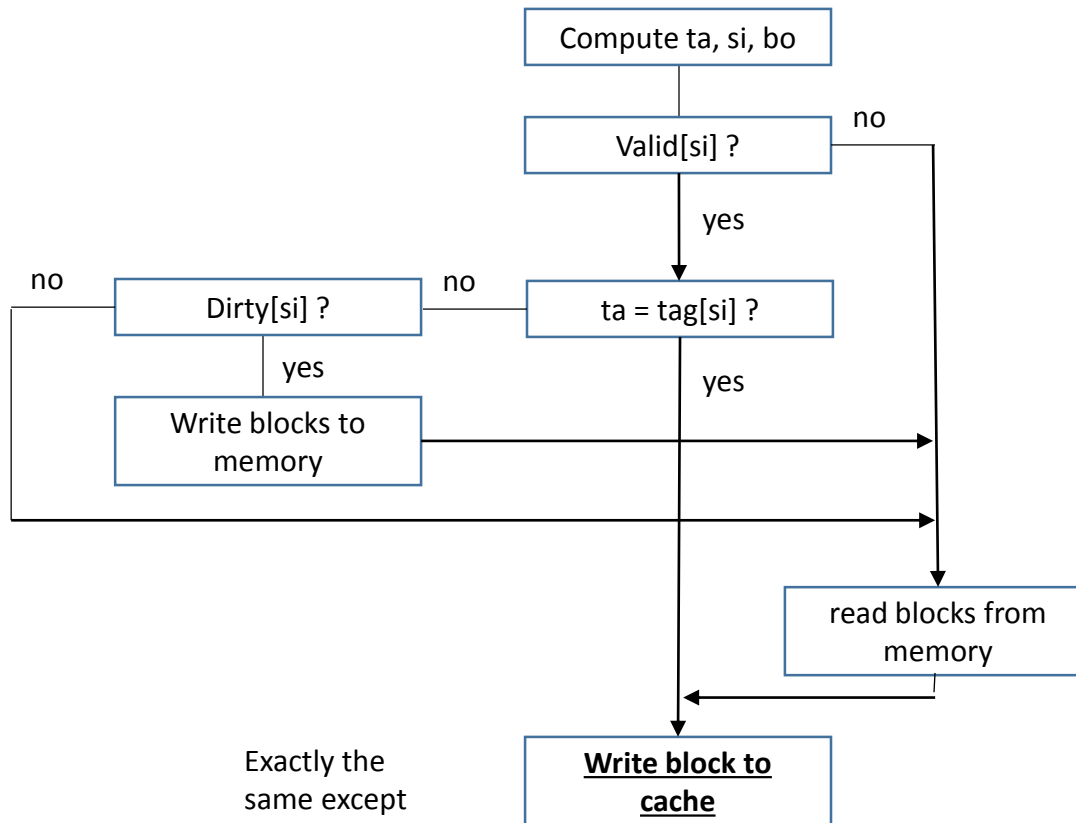
```



## Memory retrieval flowchart




## Write to memory flowchart



## Memory retrieval

If hit = 0, we have to go to memory

get from  $a - a \% B$  to  $a - a \% B + B$



```

If ( ~hit )
{
  for( i = a - a % B, j=0 ; i < a - a % B + B ; i++,j++ )
    cache.block[si][j] = memory[i] ;

  cache.valid[si] = 1 ;
  cache.ta[si] = a ;
}

```

We retrieve block bo from set si

```
*a = cache.block[si][bo] ;
```

cache.c

103

ta		si		bo			
0	1	1	0	0	1	1	1
1		2		7			

Miss! Set 2 now contains 96 – 111 bo of 103 = 7

198

ta		si		bo			
1	1	0	0	0	1	1	0
3		0		6			

Miss! Set 0 now contains 192 – 207 bo of 196 = 6

105

ta		si		bo			
0	1	1	0	1	0	0	1
1		2		9			

Hit! Set 2 now contains 96 – 111 bo = 9

236

ta		si		bo			
1	1	1	0	1	1	0	0
3		2		12			

Line replacement miss!

Set 2 contains 96 – 111, throw it out, replace with 224 - 239

## running cache.c with stride = 1

```
for( i=0; i<256; i++ )
    read from memory
```

one quadrant

We get  $(ta, si, bo) = (0,0,0) (0,0,1) \dots (0,0,15) (0,1,0) (0,1,1) \dots (0,3,15) (1,0,0) \dots (1,3,15) (2,0,0) \dots (3,3,15)$

one set

Looking at  $ta$  and  $si$ , it has mapped all of the possible  $M$  memory addresses onto the  $S$  sets and each set is identified by the tag.  $t$  concatenated with  $s$  is all of the addresses in memory divided by  $B$ .

In our case, the tag denotes which “quadrant” of memory, 0-63, 64-127, 128-191, 192-255. The set index, divides each quadrant into quadrants: 0-15, 16-31, 32-47, 48-63.


# Cache Memory

## Set indexing


## High-order bit indexing

## Middle-order bit indexing


### 4-set cache



0000  
0001  
0010  
0011  
0100  
0101  
0110  
0111  
1000  
1001  
1010  
1011  
1100  
1101  
1110  
1111



0000  
0001  
0010  
0011  
0100  
0101  
0110  
0111  
1000  
1001  
1010  
1011  
1100  
1101  
1110  
1111



Set index bits

## Powers of two problems

Arrays separated by powers of 2 addresses.

In previous example, the memory is in quadrants.

char a[64], b[64] and address of a = 0 and address of b = 0

```
for( i=0; i<64; i++ )
```

```
    b[i] = a[i] ;
```

	i	addr	ta	si	Bo
access a	0	0	0	0	0
access b	0	64	1	0	0
access a	1	1	0	0	1
access b	1	65	1	0	1

Known as cache thrashing

Associative Caches,  $E > 1$ 

```

#define S    4  // (4 cache sets)
#define E    2  // (direct mapped cache)
#define B   16  // (16 elements in each block)
#define T    2  // (2 tag bits)
#define M   256 // (256 byte memory)

struct cache_t
{
    char valid ; // simulate a bit with this
    char dirty ; // an element was stored but not written
    int tag ;
    char *block ;
} cache[S][E] ;

```

Recall that  $C = E \times B \times S$

Two types of associative caches:

$1 < E < C/B$       $E$  way associative

$E = C/B$      fully associative

To convert from a direct cache  $E=1$  to an associative cache of the same size with  $E > 1$ , divide the number of sets by  $2^{E-1}$



## Associative Cache Example, E = 3

aa

set 1	valid	tag	block

set 2	valid	tag	block

⋮

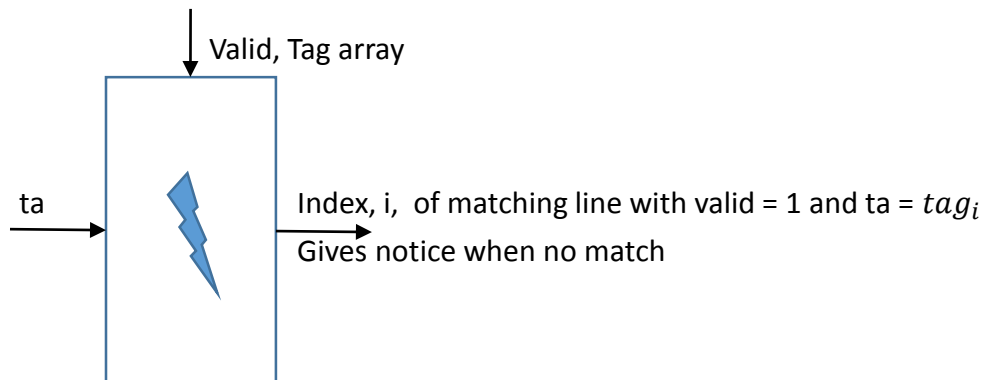
set S-1	valid	tag	block

## Works the same, except

Can store blocks from different quadrant (same offset) in one set. Identified by Tag.

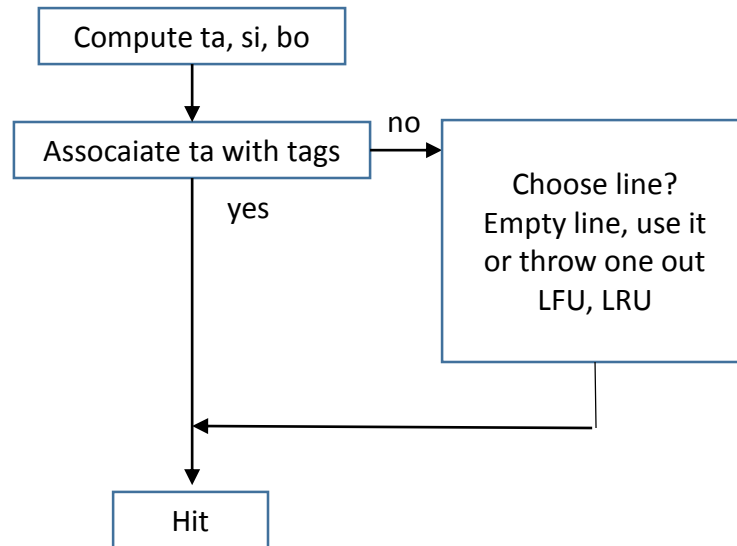
	valid	tag	block
set x	1	$tag_1$	block in quadrant $tag_1$
	0		
	1	$tag_2$	block in quadrant $tag_2$

To find which line: associative hardware:



## Determine if hit?

Extra step in process when no hit



**Fully Associative Caches,  $E = C/B$** 

Only one set,  $C/B$  lines in set.

Recall:

t bits	s bits	b bits
tag	set index	block offset

But now  $s = 0$

t bits	b bits
tag	block offset

Everything works the same

## Write strategies

What to do with a write hit.

- write-through: write the set to memory immediately. High bus traffic can be done in background
- write-back: delay write to memory, requires a dirty bit, complicated replacement algorithm


What to do with a write miss (block not in cache when storing)

- write-allocate: read block from lower level, then update
- write-no allocate

Most caches are write-back and write-allocate

Note that an instruction cash is a no write cache.

## Intel Core i7 cache architecture



## Intel Core i7 cache sizes and speeds

Cache type	L1 i-cache	L1 d-cache	L2 unified cache	L3 unified cache
Access time (cycles)	4	4	11	30-40
Cache size C	32KB	32KB	256KB	8MB
Assoc. (E)	8	8	8	16
Block size (B)	64 b	64 b	64 b	64 b
Sets (S)	64	64	512	8192

## Cache performance metrics:

- Miss rate.  $\frac{\text{\#misses}}{\text{\#references}}$ . Separate into read/write
- Hit rate.  $1 - \text{miss rate}$ .
- Hit time. The time to deliver a word in the cache to the CPU.
- Miss penalty Any additional time required because of a miss. (get from next level)

## Cache performance factors:

- larger cache size: increased hit rate more memory but slower
- larger block size: lessens the impact of large working set. Requires more time to retrieve and store from lower level
- associativity (size of E): reduces thrashing probability but costly in \$ due to complexity and increased record keeping in each line.
- write back strategy: write through is simple, reduces miss penalty



## Cache friendly code

Good spatial locality

Common sense: look for the fat.

Example 1:

```
int sumvec(int v[N])
{
    int i, sum = 0;

    for (i = 0; i < N; i++)
        sum += v[i];

    return sum;
}
```

With B = 4

v[i]		i = 0	1	2	3	4	5	6	7
Access order, [h]it or [m]iss	1	[m]	2	[h]	3	[h]	4	[h]	5
			[m]	6	[h]	7	[h]	8	[h]

## Cache friendly code

Stride-k block size B

$$\text{miss rate} = \min(1, (\text{wordsize} \times k)/B)$$

Example wordsize = 1, k = 1, B = 30. Miss rate = 1/30, one miss every 30 references

Double the word size, double the miss rate. Each access “uses up” wordsize element in the block.

## Another example

```

int sumarrayrows(int a[M][N])
{
    int i, j, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < N; j++)
            sum += a[i][j];
    return sum;
}

```

Still stride 1 because of ordering of arrays in memory

a[i][j]	j = 0	1	2	3	4	5	6	7
i = 0	1 [m]	2 [h]	3 [h]	4 [h]	5 [m]	6 [h]	7 [h]	8 [h]
i = 1	9 [m]	10 [h]	11 [h]	12 [h]	13 [m]	14 [h]	15 [h]	16 [h]
i = 2	17 [m]	18 [h]	19 [h]	20 [h]	21 [m]	22 [h]	23 [h]	24 [h]
i = 3	25 [m]	26 [h]	27 [h]	28 [h]	29 [m]	30 [h]	31 [h]	32 [h]

Follows the miss rate formula.

## Bad example

```

int sumarraycols(int a[M][N])
{
    int i, j, sum = 0;

    for (j = 0; j < N; j++)
        for (i = 0; i < M; i++)
            sum += a[i][j];
    return sum;
}

```

Now stride 8 because of ordering of arrays in memory

a[i][j]	j = 0	1	2	3	4	5	6	7
i = 0	1 [m]	5 [m]	9 [m]	13 [m]	17 [m]	21 [m]	25 [m]	29 [m]
i = 1	2 [m]	6 [m]	10 [m]	14 [m]	18 [m]	22 [m]	26 [m]	30 [m]
i = 2	3 [m]	7 [m]	11 [m]	15 [m]	19 [m]	23 [m]	27 [m]	31 [m]
i = 3	4 [m]	8 [m]	12 [m]	16 [m]	20 [m]	24 [m]	28 [m]	32 [m]

Follows the miss rate formula: = 1

Figure 6.43

```
for (i = 0; i < elems; i += stride)
{
    result += data[i];
}
sink = result; /* So compiler doesn't optimize away the loop */
```

Add up an array of size elems in stride “stride”. Calculate the running time.

Figure 6.43





Figure 6.43 The memory mountain.

Figure 6.44



## November 17, 2014 Exam

### Topics:

- Assembly Language
  - Arithmetic, logical and control instructions
  - Procedure calls, stack frame operations
  - Array structure
  - Structures, data alignment
  - Memory corruption
- Optimization
  - Explaining assembly code: reverse engineering
  - Data Flow graphs
  - Improving code to speed up program
- Memory
  - Locality analysis
  - Direct mapped caches
  - Associative caches
  - Hit/miss analysis



## Stride in matrix multiplication

Matrix multiplication  $A = B \times C$  where all the matrices are  $n \times n$

$$a[i][j] = \text{sum of } (b[i][*] * c[*][j])$$

For  $n = 2$

$$c_{11} = a_{11} * b_{11} + a_{12} * b_{21}$$

$$c_{12} = a_{11} * b_{12} + a_{12} * b_{22}$$

$$c_{21} = a_{21} * b_{11} + a_{22} * b_{21}$$

$$c_{22} = a_{21} * b_{12} + a_{22} * b_{22}$$

For the purposes of this analysis, we make the following assumptions:

Each array is an  $n \times n$  array of double, with  $\text{sizeof}(\text{double}) == 8$ .

There is a single cache with a 32-byte block size ( $B = 32$ ).

The array size  $n$  is so large that a single matrix row does not fit in the L1 cache.

The compiler stores local variables in registers, and thus references to local variables inside loops do not require any load or store instructions.

## Stride in matrix multiplication

(a) Version ijk

```
for (i = 0; i < n; i++)  
  for (j = 0; j < n; j++) {  
    sum = 0.0;  
    for (k = 0; k < n; k++)  
      sum += A[i][k]*B[k][j];  
    C[i][j] += sum;  
  }
```

Look at cache misses per inner loop.

Since block size = 32 and word size = 8, stride 1 gives miss rate = 0.25, according to

$$\text{miss rate} = \min(1, (\text{wordsize} \times k)/B)$$

Stride n yields miss rate 1 because n is very large.

## Program versions

(a) Version ijk

```
for (i = 0; i < n; i++)
  for (j = 0; j < n; j++) {
    sum = 0.0;
    for (k = 0; k < n; k++)
      sum += A[i][k]*B[k][j];
    C[i][j] += sum;
  }
```

(c) Version jki

```
for( j = 0; j < n; j++)
  for (k = 0; k < n; k++) {
    r = B[k][j];
    for (i = 0; i < n; i++)
      C[i][j] += A[i][k]*r ;
  }
```

(e) Version kij

```
for( k = 0; k < n; k++)
  for (i = 0; i < n; i++) {
    r = A[i][k];
    for (j = 0; j < n; j++)
      C[i][j] += r*B[k][j] ;
  }
```

(b) Version jik

```
for (j = 0; j < n; j++)
  for (i = 0; i < n; i++) {
    sum = 0.0;
    for (k = 0; k < n; k++)
      sum += A[i][k]*B[k][j];
    C[i][j] += sum;
  }
```

(d) Version kji

```
for( k = 0; k < n; k++)
  for (j = 0; j < n; j++) {
    r = B[k][j];
    for (i = 0; i < n; i++)
      C[i][j] += A[i][k]*r ;
  }
```

(f) Version ikj

```
for( i = 0; i < n; i++)
  for (k = 0; k < n; k++) {
    r = A[i][k];
    for (j = 0; j < n; j++)
      C[i][j] += r*B[k][j] ;
  }
```

## Miss rates

Stride/miss rate for each array  
in inner loop.

Version	A	B	C	total	memory accesses
a	1/0.25	n/1	0	1.25	2
b	1/0.25	n/1	0	1.25	2
c	n/1	0	n/1	2	3
d	n/1	0	n/1	2	3
e	0	1/0.25	1/0.25	0.5	3
f	0	1/0.25	1/0.25	0.5	3

Shows 3 categories, two version have same total rate.

Versions a,b have less memory accesses but the store  
to  $c[i][j]$  in c,d will have a write hit all the time.  
Miss rate is a better predictor of speed than the  
number of memory accesses.

Versions e,f are the fastest with a miss rate  $\frac{1}{4}$  times the  
slowest. Need to expand the analysis to include the  
middle loop.

## Locality Examples

Consider the following array of structures:

```
#define N 1000
typedef struct {
    double vel[2];
    int acc[2];
} point;
```

```
point p[N];
```

Size of one element in p is:  $2 \times 8 + 2 \times 4 = 24$

Offset to  $p[i].vel[j] = 24 \times i + j \times 8$

Offset to  $p[i].acc[j] = 24 \times i + 16 + j \times 4$

```
void clear1(point *p, int n)
{
    int i, j;

    for (i = 0; i < n; i++) {
        for (j = 0; j < 2; j++)
            p[i].vel[j] = 0;
        for (j = 0; j < 2; j++)
            p[i].acc[j] = 0;
    }
}
```

First 10 memory accesses

i	j	Statement	offset
0	0	$p[0].vel[0]$	0
0	1	$p[0].vel[1]$	8
0	0	$p[0].acc[0]$	16
0	1	$p[0].acc[1]$	20
1	0	$p[1].vel[0]$	24
1	1	$p[1].vel[1]$	32
1	0	$p[1].acc[0]$	40
1	1	$p[1].acc[1]$	44
2	0	$p[2].vel[0]$	48
2	1	$p[2].vel[1]$	56

## Locality Examples

```
void clear2(point *p, int n)
{
    int i, j;

    for (i = 0; i < n; i++) {
        for (j = 0; j < 2; j++) {
            p[i].vel[j] = 0;
            p[i].acc[j] = 0;
        }
    }
}
```

Offset to  $p[i].vel[j] = 24 \times i + j \times 8$

Offset to  $p[i].acc[j] = 24 \times i + 16 + j \times 4$

First 10 memory accesses

i	j	Statement	offset
0	0	$p[0].vel[0]$	0
0	0	$p[0].acc[0]$	16
0	1	$p[0].vel[1]$	8
0	1	$p[0].acc[1]$	20
1	0	$p[1].vel[0]$	24
1	0	$p[1].acc[0]$	40
1	1	$p[1].vel[1]$	32
1	1	$p[1].acc[1]$	44
2	0	$p[2].vel[0]$	48
2	0	$p[2].acc[0]$	64

## Locality Examples

```
void clear3(point *p, int n)
{
    int i, j;

    for (j = 0; j < 2 ; j++) {
        for (i = 0; i < n; i++)
            p[i].vel[j] = 0;
        for (i = 0; i < n; i++)
            p[i].acc[j] = 0;
    }
}
```

Offset to  $p[i].vel[j] = 24 \times i + j \times 8$

Offset to  $p[i].acc[j] = 24 \times i + 16 + j \times 4$

First 10 memory accesses

i	j	Statement	offset
0	0	$p[0].vel[0]$	0
1	0	$p[1].vel[0]$	24
2	0	$p[2].vel[0]$	48
0	0	$p[0].acc[0]$	16
1	0	$p[1].acc[0]$	40
2	0	$p[2].acc[0]$	64
0	1	$p[0].vel[1]$	8
1	1	$p[1].vel[1]$	32
2	1	$p[2].vel[1]$	56
0	1	$p[0].acc[1]$	20

## Locality Examples

What would happen if?

```
#define N 1000
typedef struct {
    double vel[3];
    int acc[3];
} point;
```

```
point p[N];
```

Size of one element in p is:  $3 \times 8 + 3 \times 4 + \underline{4} = 48$

Offset to  $p[i].vel[j] = 48 \times i + j \times 8$

Offset to  $p[i].int[j] = 48 \times i + 24 + j \times 4$

The extra 4 is to align  $vel[0]$  on an 8 byte boundary!



## Cache Configuration examples

cache config

				?	?	?	?
m	C	B	E	S	t	s	b
32	1024	4	1	256	22	8	2
32	1024	8	4	32	24	5	3
32	2048	32	32	1	27	0	5

$$S = C/(B * E)$$

$$b = \log_2 B$$

$$s = \log_2 S$$

$$t = m - b - s$$

## Cache Configuration examples

cache mapping

associative examples

## Cache hit/miss examples

```
typedef int array[2][2];
```

```
void transpose1(array dst, array src)
```

```
{
    int i, j;

    for (i = 0; i < 2; i++) {
        for (j = 0; j < 2; j++) {
            dst[j][i] = src[i][j];
        }
    }
}
```

src at address 0, dst at address 16 (powers of 2 problem?)

M = 32, B = 8, C = 16, S = 2, m = 5

address overlay

t	s	b	b	b
---	---	---	---	---

access sequence:

	a	t	s	b	h/m
read 0	00000	0	0	0	m
write 16	10000	1	0	0	m
read 4	00100	0	0	4	m
write 24	11000	1	1	0	m
read 8	01000	0	1	0	m
write 20	10100	1	0	4	m
read 12	01100	0	1	4	m
write 28	11100	1	1	4	m

## Cache hit/miss examples

Cache size 32

$M = 32, B = 8, C = 32, S = 4, m = 5$

address overlay

s	s	b	b	b
---	---	---	---	---

access sequence:

	a	t	s	b	h/m
read 0	00000	0	0	0	m
write 16	10000	0	2	0	m
read 4	00100	0	0	4	h
write 24	11000	0	3	0	m
read 8	01000	0	1	0	m
write 20	10100	0	2	4	h
read 12	01100	0	1	4	h
write 28	11100	0	3	4	h

4 misses, 8 references rate = 0.5

## Cache hit/miss examples

```
struct algae_position {
    int x;
    int y;
};
```

```
struct algae_position grid[16][16]; // size = 256 x 8 2048 = M, m = 11
int total_x = 0, total_y = 0;
int i, j;
```

```
for (i = 0; i < 16; i++) {
    for (j = 0; j < 16; j++) {
        total_x += grid[i][j].x;
    }
}
```

```
for (i = 0; i < 16; i++) {
    for (j = 0; j < 16; j++) {
        total_y += grid[i][j].y;
    }
}
```

M = 2048, B = 16, C = 1024, S = 64

Address overlay

t	s	s	s	s	s	s	b	b	b	b
---	---	---	---	---	---	---	---	---	---	---

x access sequence:  
0, 8, 16, 24, ...

On x[0] (miss) , set 0 will receive x[0], y[0], x[1], y[1]  
x[1] has address 8, so access 8 is a hit. Alternating hits and misses.

y access sequence:  
4, 12, 20, 28, ...

On y[0] (miss) , set 0 will receive x[0], y[0], x[1], y[1]  
y[1] has address 12, so access 12 is a hit. Alternating hits and misses.

x reads: 256, y reads 256. ½ miss, rate = 0.5

## Cache hit/miss examples

```
for (i = 0; i < 16; i++){
  for (j = 0; j < 16; j++) {
    total_x += grid[j][i].x;
    total_y += grid[j][i].y;
  }
}
```

Access pattern: 0, 4, 64, 68, 128, 132 x is always a miss and y is always a hit: rate = 0.5

Twice the size: double b makes the miss rate 0.25., double s keeps the rate the same.

```
for (i = 0; i < 16; i++){
  for (j = 0; j < 16; j++) {
    total_x += grid[i][j].x;
    total_y += grid[i][j].y;
  }
}
```

Access pattern: 0, 4, 8, 12, 16... best stride. Set holds x[0], y[0], x[1], y[1] so x[0] is a miss and the other three are a hit rate = 0.25

Twice the size: double b makes the miss rate 1/8, double s keeps the rate the same

## Data Flow Graph

```
double poly(double a[], double x, int degree) {  
    long int i;  
    double result = a[0];  
    double xpwr = x; /* Equals x^i at start of loop */  
  
    for (i = 1; i <= degree; i++) {  
        result += a[i] * xpwr;  
        xpwr = x * xpwr;  
    }  
    return result;  
}  
  
void main() {  
    double a[128] ;  
    double x = 3;  
    int degree = 128 ;  
  
    poly( a,x,degree ) ;  
    printf( "%d \n", x ) ;  
}
```

## Data Flow Graph

poly function compiled with -O.

```

0x00000000004004c4 <+0>:  movapd %xmm0,%xmm3
0x00000000004004c8 <+4>:  movsd (%rdi),%xmm0
0x00000000004004cc <+8>:  movslq %esi,%rsi
0x00000000004004cf <+11>: test %rsi,%rsi
0x00000000004004d2 <+14>: jle 0x4004f7 <poly+51>
0x00000000004004d4 <+16>: movapd %xmm3,%xmm1
0x00000000004004d8 <+20>: mov $0x1,%eax
0x00000000004004dd <+25>: movapd %xmm1,%xmm2
0x00000000004004e1 <+29>: mulsd (%rdi,%rax,8),%xmm2
0x00000000004004e6 <+34>: addsd %xmm2,%xmm0
0x00000000004004ea <+38>: mulsd %xmm3,%xmm1
0x00000000004004ee <+42>: add $0x1,%rax
0x00000000004004f2 <+46>: cmp %rsi,%rax
0x00000000004004f5 <+49>: jle 0x4004dd <poly+25>
0x00000000004004f7 <+51>: repz retq

```

End of assembler dump.

From the code, it is clear that upon entry we have:

%xmm0	x
%rdi	addr(a)
%esi	degree
%xmm3	receives x
%xmm0	receives a[0] which is result
%rsi	receives degree
%xmm1	receives x which is xpwr
%eax	receives 1

Main compile

```

0x00000000004004f9 <+0>:  sub $0x408,%rsp
0x0000000000400500 <+7>:  mov %rsp,%rdi
0x0000000000400503 <+10>: mov $0x80,%esi
0x0000000000400508 <+15>: movsd 0x130(%rip),%xmm0
0x0000000000400510 <+23>: callq 0x4004c4 <poly>
0x0000000000400515 <+28>: movsd 0x123(%rip),%xmm0
0x000000000040051d <+36>: mov $0x400638,%edi
0x0000000000400522 <+41>: mov $0x1,%eax
0x0000000000400527 <+46>: callq 0x4003b8 <printf@plt>
0x000000000040052c <+51>: add $0x408,%rsp
0x0000000000400533 <+58>: retq

```



## Data Flow Graph

Here is the assembly code of the poly function compiled with -O.



```

0x00000000004004c4 <+0>:  movapd %xmm0,%xmm3
0x00000000004004c8 <+4>:  movsd (%rdi),%xmm0
0x00000000004004cc <+8>:  movslq %esi,%rsi
0x00000000004004cf <+11>: test  %rsi,%rsi
0x00000000004004d2 <+14>:  jle   0x4004f7 <poly+51>
0x00000000004004d4 <+16>:  movapd %xmm3,%xmm1
0x00000000004004d8 <+20>:  mov   $0x1,%eax
0x00000000004004dd <+25>:  movapd %xmm1,%xmm2
0x00000000004004e1 <+29>:  mulsd (%rdi,%rax,8),%xmm2
0x00000000004004e6 <+34>:  addsd %xmm2,%xmm0
0x00000000004004ea <+38>:  mulsd %xmm3,%xmm1
0x00000000004004ee <+42>:  add   $0x1,%rax
0x00000000004004f2 <+46>:  cmp   %rsi,%rax
0x00000000004004f5 <+49>:  jle   0x4004dd <poly+25>
0x00000000004004f7 <+51>:  repz retq

```

We know that:

%xmm0 is result  
 %xmm1 is xpwr  
 %xmm2 is nothing  
 %xmm3 is x  
 %rax is l  
 %rsi is degree  
 %rdi is address of a



## Data Flow Graph

Here is the assembly code of the poly function compiled with -O.

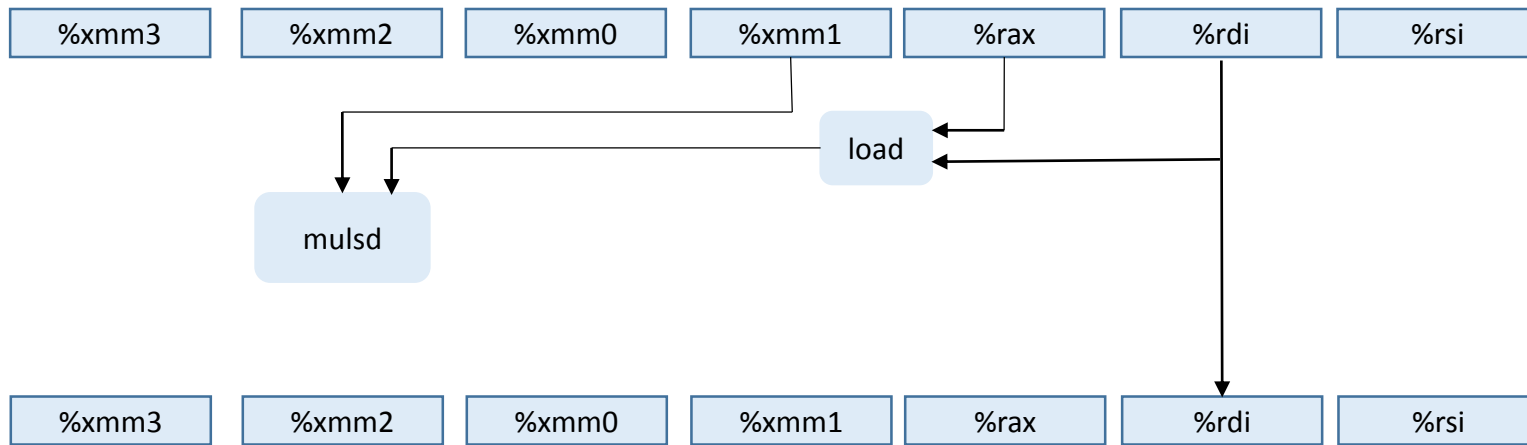
```

0x00000000004004c4 <+0>:  movapd %xmm0,%xmm3
0x00000000004004c8 <+4>:  movsd (%rdi),%xmm0
0x00000000004004cc <+8>:  movslq %esi,%rsi
0x00000000004004cf <+11>: test %rsi,%rsi
0x00000000004004d2 <+14>: jle 0x4004f7 <poly+51>
0x00000000004004d4 <+16>:  movapd %xmm3,%xmm1
0x00000000004004d8 <+20>:  mov $0x1,%eax
0x00000000004004dd <+25>:  movapd %xmm1,%xmm2
0x00000000004004e1 <+29>:  mulsd (%rdi,%rax,8),%xmm2
0x00000000004004e6 <+34>:  addsd %xmm2,%xmm0
0x00000000004004ea <+38>:  mulsd %xmm3,%xmm1
0x00000000004004ee <+42>:  add $0x1,%rax
0x00000000004004f2 <+46>:  cmp %rsi,%rax
0x00000000004004f5 <+49>:  jle 0x4004dd <poly+25>
0x00000000004004f7 <+51>:  repz retq

```

We know that:

%xmm0 is result  
 %xmm1 is xpw  
 %xmm2 is nothing  
 %xmm3 is x  
 %rax is 1  
 %rsi is degree  
 %rdi is address of a



## Data Flow Graph

Here is the assembly code of the poly function compiled with -O.

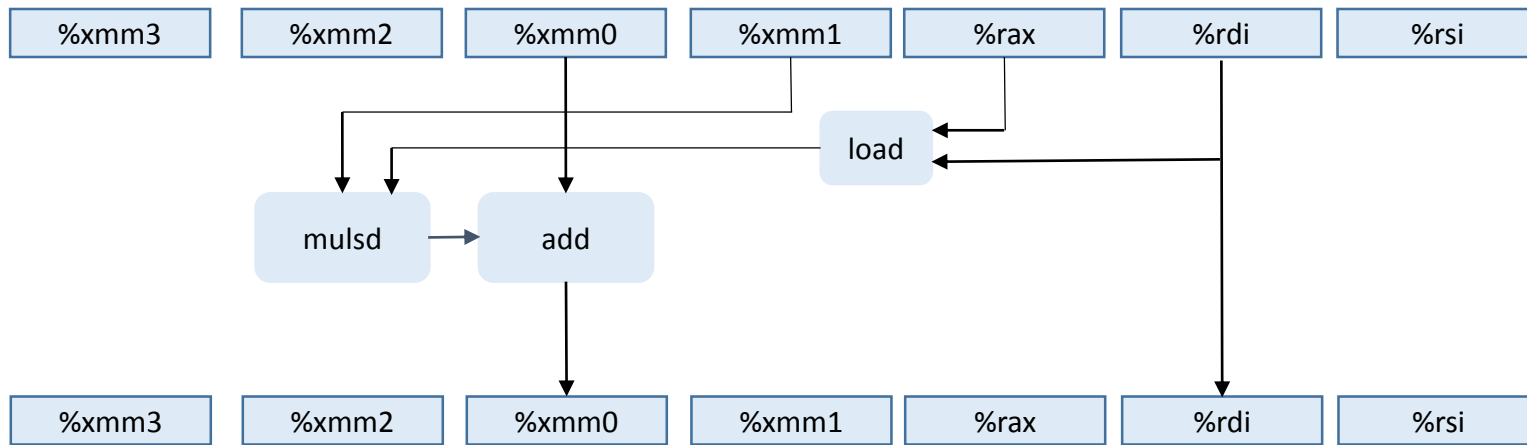
```

0x00000000004004c4 <+0>:  movapd %xmm0,%xmm3
0x00000000004004c8 <+4>:  movsd (%rdi),%xmm0
0x00000000004004cc <+8>:  movslq %esi,%rsi
0x00000000004004cf <+11>: test %rsi,%rsi
0x00000000004004d2 <+14>: jle 0x4004f7 <poly+51>
0x00000000004004d4 <+16>:  movapd %xmm3,%xmm1
0x00000000004004d8 <+20>:  mov $0x1,%eax
0x00000000004004dd <+25>:  movapd %xmm1,%xmm2
0x00000000004004e1 <+29>:  mulsd (%rdi,%rax,8),%xmm2
0x00000000004004e6 <+34>:  addsd %xmm2,%xmm0
0x00000000004004ea <+38>:  mulsd %xmm3,%xmm1
0x00000000004004ee <+42>:  add $0x1,%rax
0x00000000004004f2 <+46>:  cmp %rsi,%rax
0x00000000004004f5 <+49>:  jle 0x4004dd <poly+25>
0x00000000004004f7 <+51>:  repz retq

```

We know that:

%xmm0 is result  
 %xmm1 is xpw  
 %xmm2 is nothing  
 %xmm3 is x  
 %rax is l  
 %rsi is degree  
 %rdi is address of a



## Data Flow Graph

Here is the assembly code of the poly function compiled with -O.

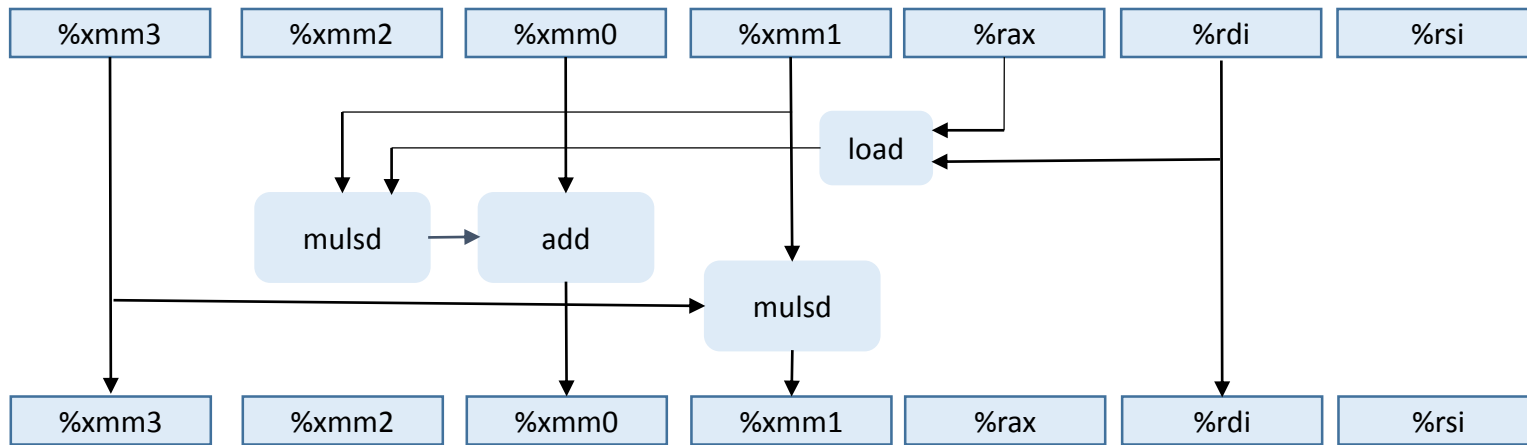
```

0x00000000004004c4 <+0>:  movapd %xmm0,%xmm3
0x00000000004004c8 <+4>:  movsd (%rdi),%xmm0
0x00000000004004cc <+8>:  movslq %esi,%rsi
0x00000000004004cf <+11>: test %rsi,%rsi
0x00000000004004d2 <+14>: jle 0x4004f7 <poly+51>
0x00000000004004d4 <+16>:  movapd %xmm3,%xmm1
0x00000000004004d8 <+20>:  mov $0x1,%rax
0x00000000004004dd <+25>:  movapd %xmm1,%xmm2
0x00000000004004e1 <+29>:  mulsd (%rdi,%rax,8),%xmm2
0x00000000004004e6 <+34>:  addsd %xmm2,%xmm0
0x00000000004004ea <+38>:  mulsd %xmm3,%xmm1
0x00000000004004ee <+42>:  add $0x1,%rax
0x00000000004004f2 <+46>:  cmp %rsi,%rax
0x00000000004004f5 <+49>:  jle 0x4004dd <poly+25>
0x00000000004004f7 <+51>:  repz retq

```

We know that:

%xmm0 is result  
 %xmm1 is x<sup>pwr</sup>  
 %xmm2 is nothing  
 %xmm3 is x  
 %rax is l  
 %rsi is degree  
 %rdi is address of a



## Data Flow Graph

Here is the assembly code of the poly function compiled with -O.

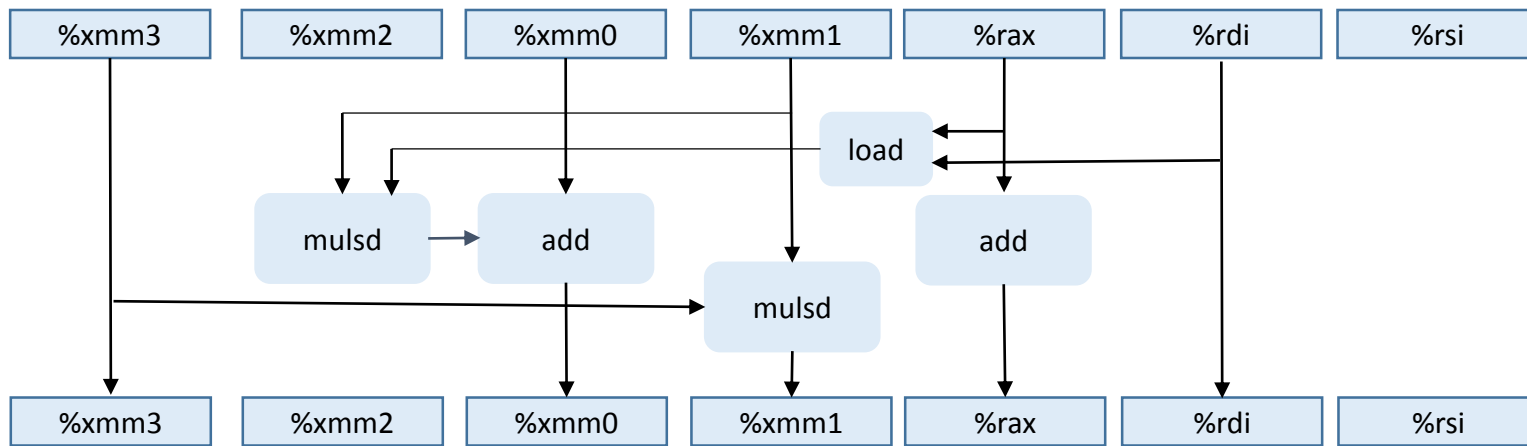
```

0x00000000004004c4 <+0>:  movapd %xmm0,%xmm3
0x00000000004004c8 <+4>:  movsd (%rdi),%xmm0
0x00000000004004cc <+8>:  movslq %esi,%rsi
0x00000000004004cf <+11>: test %rsi,%rsi
0x00000000004004d2 <+14>: jle 0x4004f7 <poly+51>
0x00000000004004d4 <+16>:  movapd %xmm3,%xmm1
0x00000000004004d8 <+20>:  mov $0x1,%rax
0x00000000004004dd <+25>:  movapd %xmm1,%xmm2
0x00000000004004e1 <+29>:  mulsd (%rdi,%rax,8),%xmm2
0x00000000004004e6 <+34>:  addsd %xmm2,%xmm0
0x00000000004004ea <+38>:  mulsd %xmm3,%xmm1
0x00000000004004ee <+42>:  add $0x1,%rax
0x00000000004004f2 <+46>:  cmp %rsi,%rax
0x00000000004004f5 <+49>:  jle 0x4004dd <poly+25>
0x00000000004004f7 <+51>:  repz retq

```

We know that:

%xmm0 is result  
 %xmm1 is x<sup>pwr</sup>  
 %xmm2 is nothing  
 %xmm3 is x  
 %rax is l  
 %rsi is degree  
 %rdi is address of a



## Data Flow Graph

Here is the assembly code of the poly function compiled with -O.

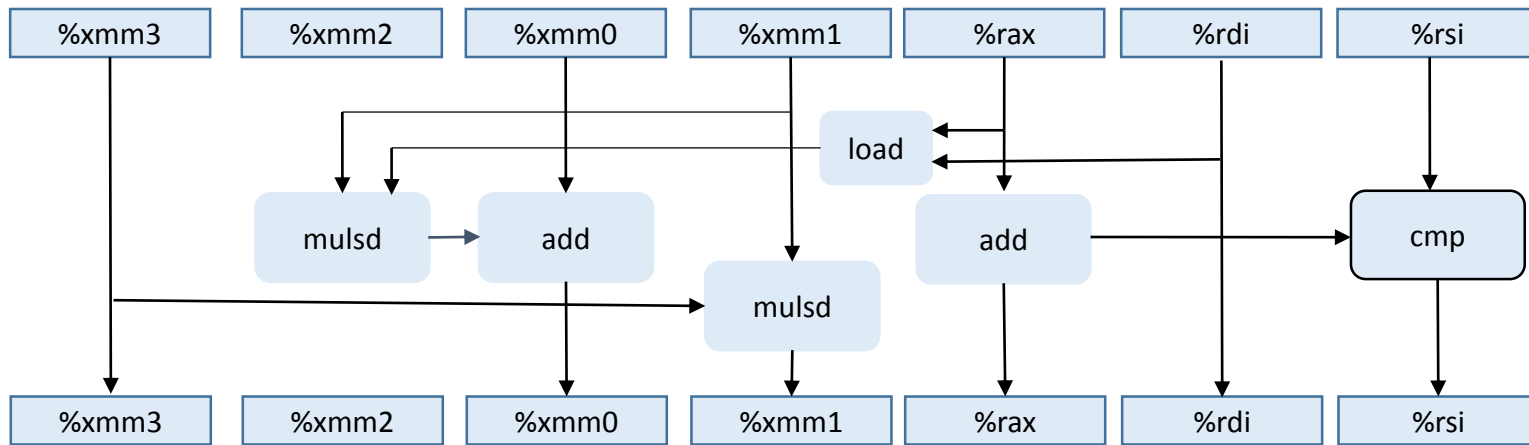
```

0x00000000004004c4 <+0>:  movapd %xmm0,%xmm3
0x00000000004004c8 <+4>:  movsd (%rdi),%xmm0
0x00000000004004cc <+8>:  movslq %esi,%rsi
0x00000000004004cf <+11>: test %rsi,%rsi
0x00000000004004d2 <+14>: jle 0x4004f7 <poly+51>
0x00000000004004d4 <+16>:  movapd %xmm3,%xmm1
0x00000000004004d8 <+20>:  mov $0x1,%rax
0x00000000004004dd <+25>:  movapd %xmm1,%xmm2
0x00000000004004e1 <+29>:  mulsd (%rdi,%rax,8),%xmm2
0x00000000004004e6 <+34>:  addsd %xmm2,%xmm0
0x00000000004004ea <+38>:  mulsd %xmm3,%xmm1
0x00000000004004ee <+42>:  add $0x1,%rax
0x00000000004004f2 <+46>:  cmp %rsi,%rax
0x00000000004004f5 <+49>:  jle 0x4004dd <poly+25>
0x00000000004004f7 <+51>:  repz retq

```

We know that:

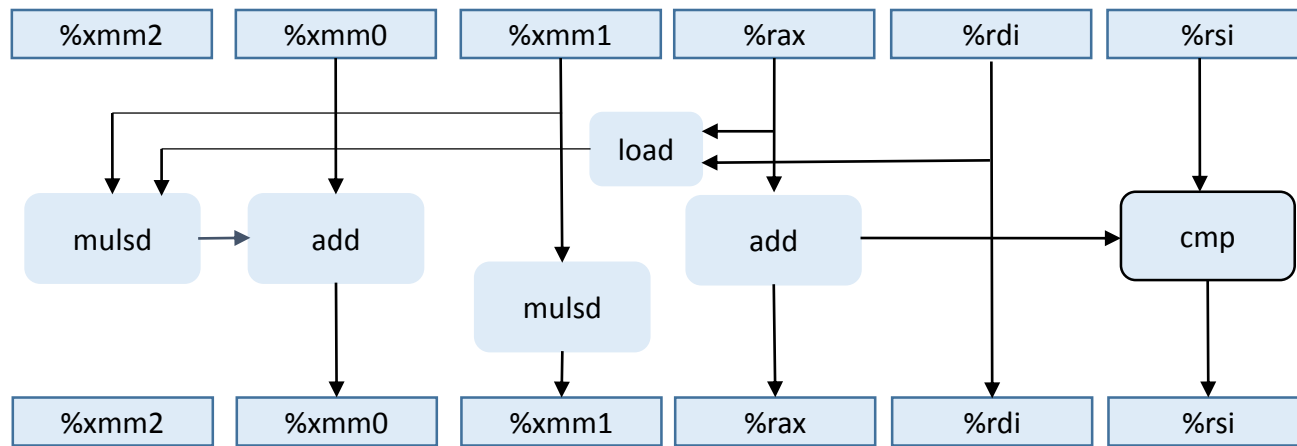
%xmm0 is result  
 %xmm1 is x<sup>pwr</sup>  
 %xmm2 is nothing  
 %xmm3 is x  
 %rax is l  
 %rsi is degree  
 %rdi is address of a



We know that:

%xmm0 is result  
%xmm1 is xpr  
%xmm2 is nothing  
%xmm3 is x  
%rax is l  
%rsi is degree  
%rdi is address of a

Because %xmm3  
does not depend on  
the outcome of the  
loop



We know that:

%xmm0 is result

%xmm1 is xpr

%xmm2 is nothing

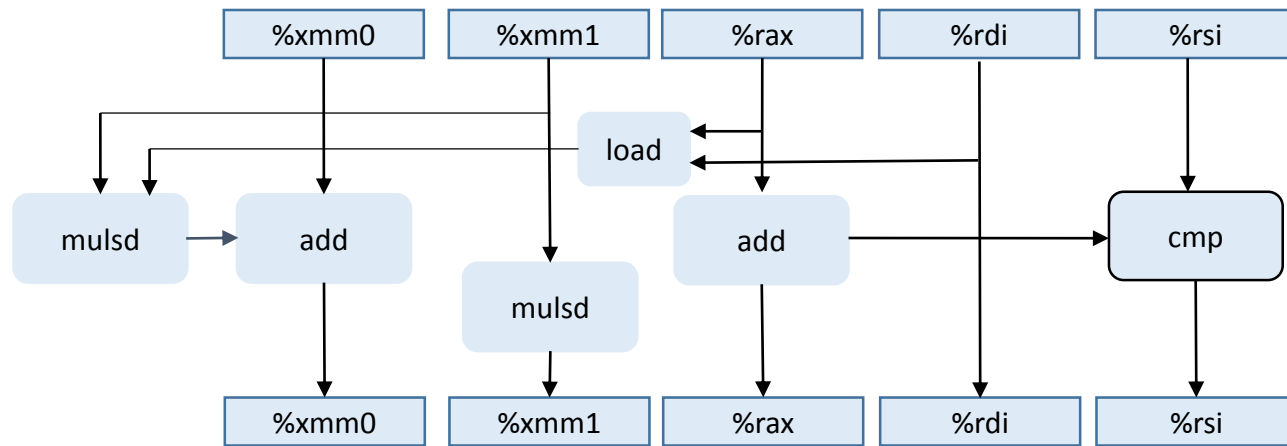
%xmm3 is x

%rax is l

%rsi is degree

%rdi is address of a

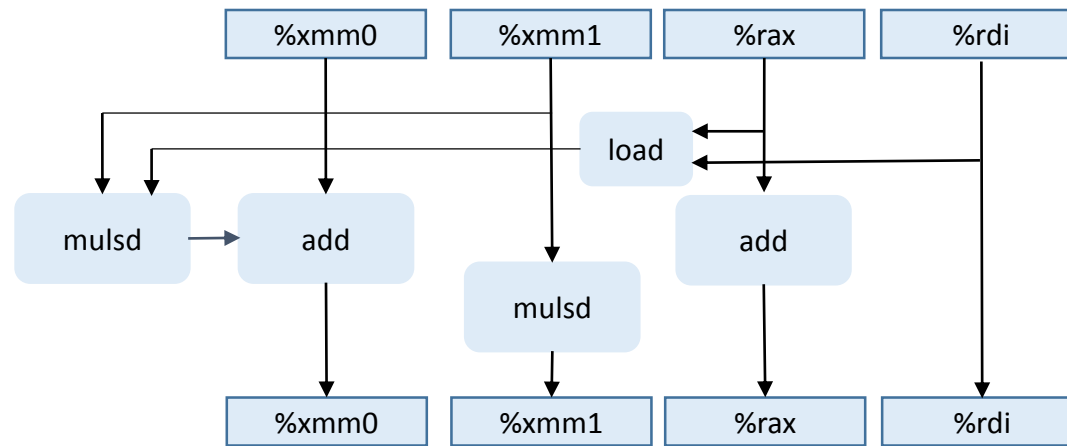
Because %xmm2  
does not depend on  
the outcome of the  
loop





We know that:

%xmm0 is result  
%xmm1 is xpr  
%xmm2 is nothing  
%xmm3 is x  
%rax is l  
%rsi is degree  
%rdi is address of a



Because the compare outcome is done within the loop

We know that:

%xmm0 is result

%xmm1 is xpr

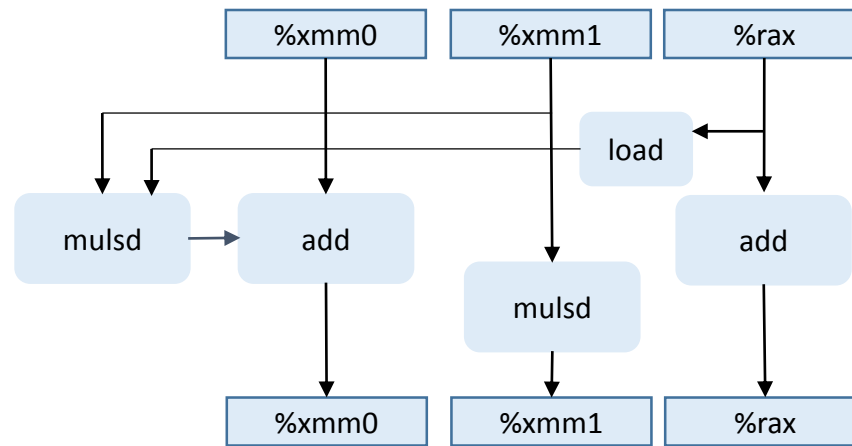
%xmm2 is nothing

%xmm3 is x


%rax is i

%rsi is degree

%rdi is address of a



Because %rdi is read only



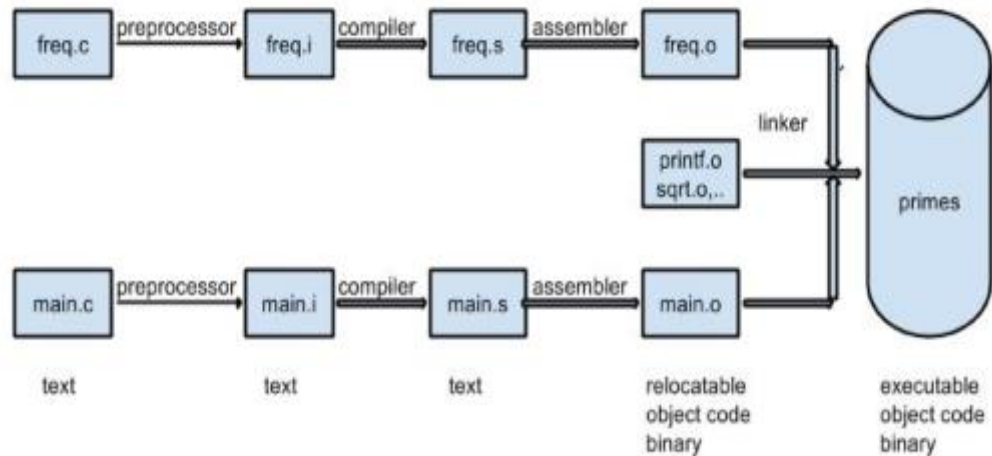
## Tentative Lecture Schedule

Lecture	Date	Topic
14	11/19/2014	Linking
15	11/24/2014	Exception Control Flow
16	11/26/2014	Virtual Memory
	11/30/2014	Lab 3 due
17	12/1/2014	Virtual Memory (Lab 4 issued)
18	12/3/2014	Concurrent Programming
19	12/8/2014	Concurrent Programming
20	12/10/2014	Review (Lab 4 due)
	12/16/2014	Final Exam 8AM-11AM
	11/28/2014	No Discussion - Thanksgiving Holiday
		* No office hour

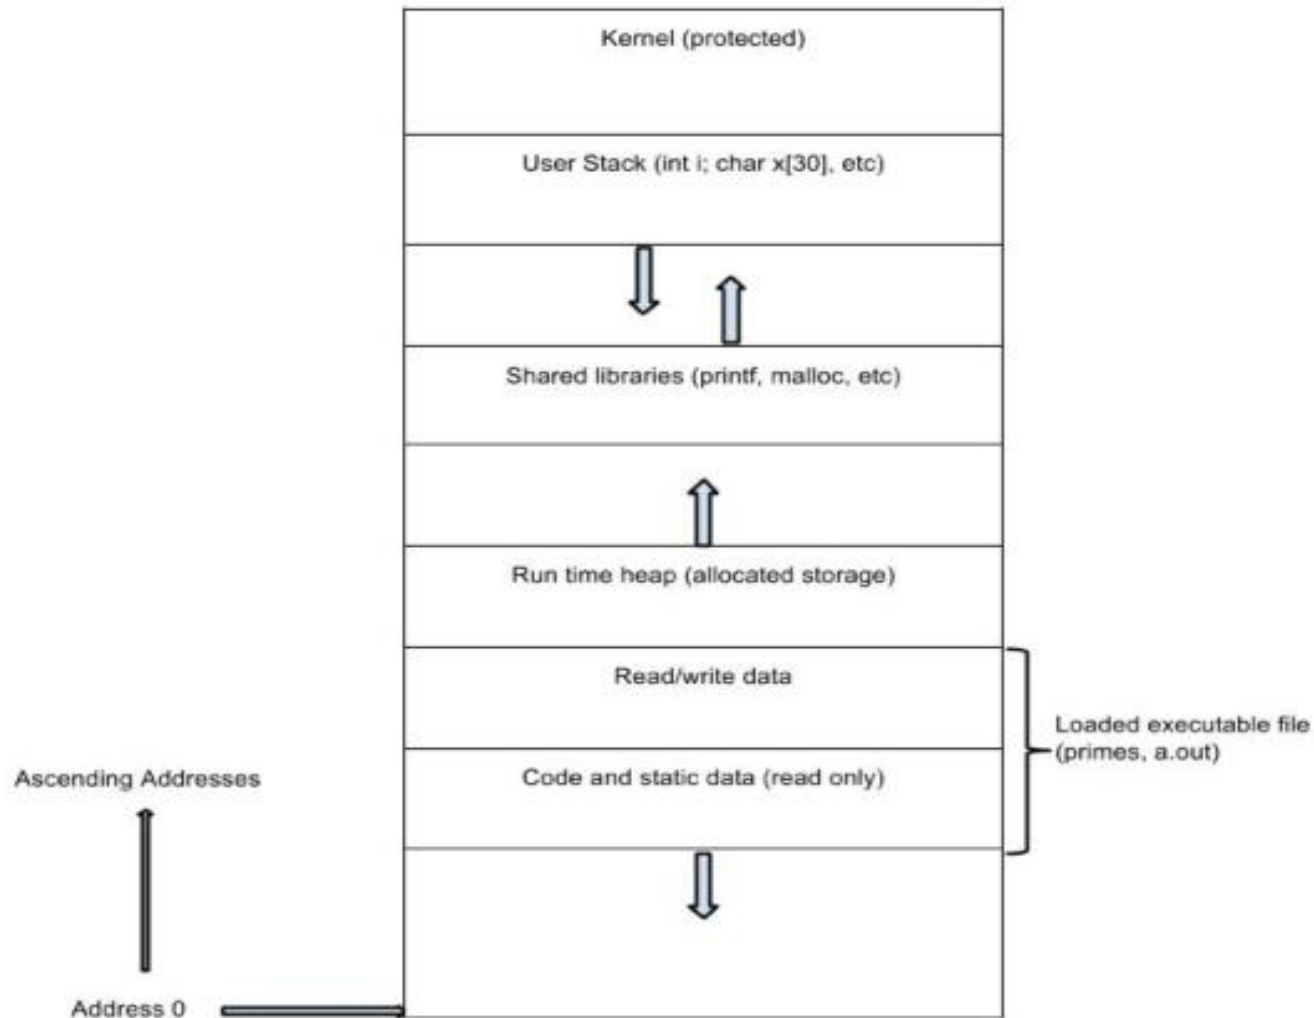
Pre-process, compile, assemble link.

Can stop/start process anywhere, combine any of these files using gcc or g++.  
Pre-process, compile, assemble link.

```
gcc -O1 -lm -o primes freq.c main.c
```



## Layout of program in memory



## Combine/collect code – create executable

1. Combine separately compiled programs.
2. Add library routines (not shared)
3. Relocation

```
#include <stdio.h>
```

```
int a[100];
```

```
int b = 99;
```

```
int poly()
```

```
{
    int i;
```

```
    int sum = 0;
```

```
    for( i=0; i<100; i++ )
        sum = sum+a[i];
```

```
    return sum;
}
```

```
void main()
```

```
{
    printf( "%d \n", poly() );
}
```

```
main()
```

```
0x00000000004004ff <+0>: push %rbp
```

```
0x0000000000400500 <+1>: mov %rsp,%rbp
```

```
0x0000000000400503 <+4>: mov $0x0,%eax
```

```
0x0000000000400508 <+9>: callq 0x4004c4 <poly>
```

```
0x000000000040050d <+14>: mov %eax,%edx
```

```
0x000000000040050f <+16>: mov $0x400628,%eax
```

```
0x0000000000400514 <+21>: mov %edx,%esi
```

```
0x0000000000400516 <+23>: mov %rax,%rdi
```

```
0x0000000000400519 <+26>: mov $0x0,%eax
```

```
0x000000000040051e <+31>: callq 0x4003b8 <printf@plt>
```

```
0x0000000000400523 <+36>: leaveq
```

```
0x0000000000400524 <+37>: retq
```

```
poly()
```

```
0x00000000004004c4 <+0>: push %rbp
```

```
0x00000000004004c5 <+1>: mov %rsp,%rbp
```

```
0x00000000004004c8 <+4>: movl $0x0,-0x4(%rbp)
```

```
0x00000000004004cf <+11>: movl $0x0,-0x8(%rbp)
```

```
0x00000000004004d6 <+18>: jmp 0x4004eb <poly+39>
```

```
0x00000000004004d8 <+20>: mov -0x8(%rbp),%eax
```

```
0x00000000004004db <+23>: cltq
```

```
0x00000000004004dd <+25>: mov 0x600920(%rax,4),%eax
```

```
0x00000000004004e4 <+32>: add %eax,-0x4(%rbp)
```

```
0x00000000004004e7 <+35>: addl $0x1,-0x8(%rbp)
```

```
0x00000000004004eb <+39>: cmpl $0x63,-0x8(%rbp)
```

```
0x00000000004004ef <+43>: jle 0x4004d8 <poly+20>
```

```
0x00000000004004f1 <+45>: mov 0x2003f5(%rip),%eax # 0x6008ec <b>
```

```
0x00000000004004f7 <+51>: add %eax,-0x4(%rbp)
```

```
0x00000000004004fa <+54>: mov -0x4(%rbp),%eax
```


```
0x00000000004004fd <+57>: leaveq
```

```
0x00000000004004fe <+58>: retq
```

Combine/collect code – create executable

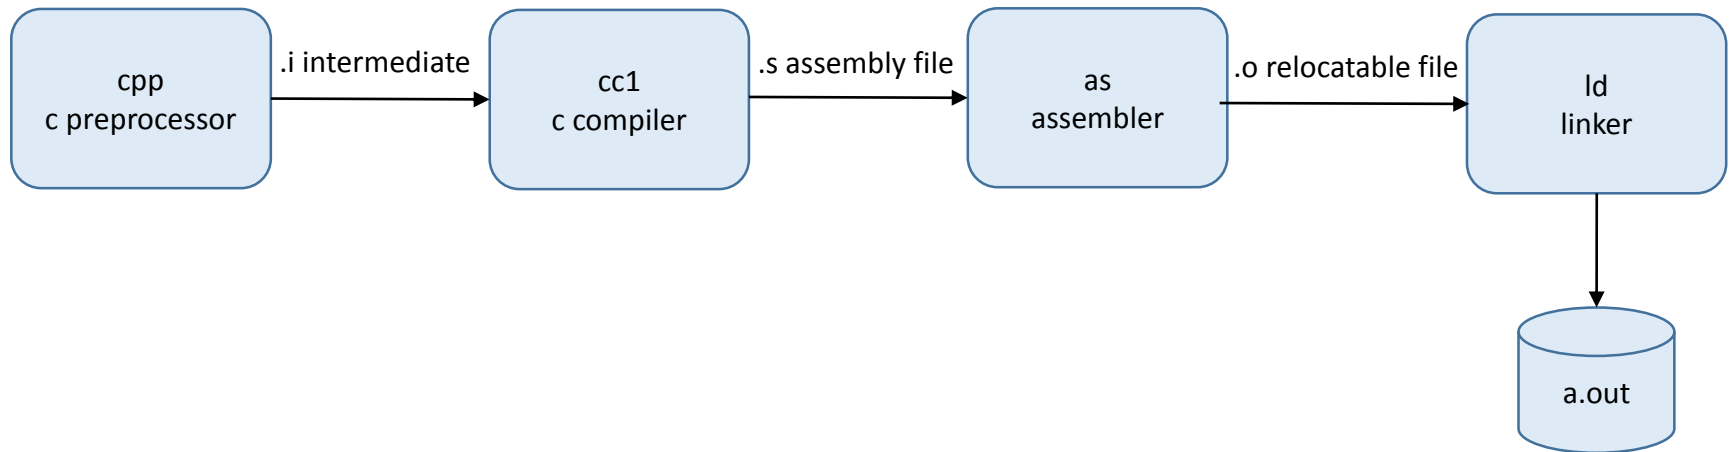
```
/* main.c */  
void swap();  
  
int buf[2] = {1, 2};  
  
int main()  
{  
    swap();  
    return 0;  
}  
  
/* swap.c */  
extern int buf[];  
  
int *bufp0 = &buf[0];  
int *bufp1;  
  
void swap()  
{  
    int temp;  
  
    bufp1 = &buf[1];  
    temp = *bufp0;  
    *bufp0 = *bufp1;  
    *bufp1 = temp;  
}
```

gcc main.c swap.c library.o





gcc is actually



## Static Linking

Programs contain internal and external symbols.

buf is an external symbol

Symbol resolution matches the buf in swap with the buf in main.

```
/* main.c */  
void swap();  
  
int buf[2] = {1, 2};
```

```
int main()  
{  
    swap();  
    return 0;  
}
```

```
/* swap.c */  
extern int buf[];  
  
int *bufp0 = &buf[0];  
int *bufp1;
```

```
void swap()  
{  
    int temp;
```

```
    bufp1 = &buf[1];  
    temp = *bufp0;  
    *bufp0 = *bufp1;  
    *bufp1 = temp;  
}
```

## Relocation

Assembler creates code relative to start of module.

Relocation changes offsets to account for multiple modules.

```
main()
0x00000000004004ff <+0>:  push  %rbp
0x0000000000400500 <+1>:  mov   %rsp,%rbp
0x0000000000400503 <+4>:  mov   $0x0,%eax
0x0000000000400508 <+9>:  callq 0x4004c4 <poly>
0x000000000040050d <+14>: mov   %eax,%edx
0x000000000040050f <+16>: mov   $0x400628,%eax
0x0000000000400514 <+21>: mov   %edx,%esi
0x0000000000400516 <+23>: mov   %rax,%rdi
0x0000000000400519 <+26>: mov   $0x0,%eax
0x000000000040051e <+31>: callq 0x4003b8 <printf@plt>
0x0000000000400523 <+36>: leaveq
0x0000000000400524 <+37>: retq

poly()
0x00000000004004c4 <+0>:  push  %rbp
0x00000000004004c5 <+1>:  mov   %rsp,%rbp
0x00000000004004c8 <+4>:  movl  $0x0,-0x4(%rbp)
0x00000000004004cf <+11>: movl  $0x0,-0x8(%rbp)
0x00000000004004d6 <+18>: jmp   0x4004eb <poly+39>
0x00000000004004d8 <+20>: mov   -0x8(%rbp),%eax
0x00000000004004db <+23>: cltq
0x00000000004004dd <+25>: mov   0x600920(%rax,4),%eax
0x00000000004004e4 <+32>: add   %eax,-0x4(%rbp)
0x00000000004004e7 <+35>: addl  $0x1,-0x8(%rbp)
0x00000000004004eb <+39>: cmpl  $0x63,-0x8(%rbp)
0x00000000004004ef <+43>: jle   0x4004d8 <poly+20>
0x00000000004004f1 <+45>: mov   0x2003f5(%rip),%eax    # 0x6008ec <b>
0x00000000004004f7 <+51>: add   %eax,-0x4(%rbp)
0x00000000004004fa <+54>: mov   -0x4(%rbp),%eax
0x00000000004004fd <+57>: leaveq
0x00000000004004fe <+58>: retq
```

## Object files

Relocatable object file (.o) can be re-linked but not executed.

Executable object file (a.out) can be executed and re-linked (with an option)

Shared object file: dynamically linked.

## Relocatable object file

ELF Header
.text: machine code
.rodata: read only data
.data: initialized globals
.bss: uninitialized globals (description)
.symtab: symbol table (globals and external function info)
.rel.text: relocation information for externals
.rel.data: relocation information for cross referenced data
.debug: -g symbols for gdb
.line: -g line numbers for gdb
.strtab: descriptive strings for .symtab
Section header table: which sections are in the table

```

/* swap.c */
extern int buf[];

int *bufp0 = &buf[0];
int *bufp1;

void swap()
{
    int temp;          // temp local : not in .symtab

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}

```

.symtab contains a binding between it and .strtab called the ELF

## Symbols

Three main types (table constructed by the compiler and sent to the assembler):

Global symbols within a module

Global symbols outside of a module

Local symbols

examp.c

```
extern int z ;  
float y ;  
int f() {
```

y, f() and g() are global but defined in the module

```
    int x = 0 ;  
    return x ;  
}
```

z is global but defined elsewhere

```
int g() {
```

Both x's are local. For gdb, they get unique names.

```
    int x = 1 ;  
    return x ;
```

## ELF descriptors

```

/* swap.c */
extern int buf[];

int *bufp0 = &buf[0];
int *bufp1;

void swap()
{
    int temp;          // temp local : not in .symtab

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}

```

symbol	.symtab?	type	Where	Section
buf	Yes	external	main.o	.data
bufp0	Yes	global	swap.o	.data
bufp1	Yes	global	swap.o	.bss
swap	Yes	global	swap.o	.text
temp	No			

## Symbol resolution

Match up identically named external (global) symbols. Some defined within the module, some outside.

Strong and weak symbols.

Functions and initialized globals are strong.

Uninitialized globals are weak (do not actually take up space but are described).

When the linker encounters two symbols which are the same:

1. Both strong is an error (multiply defined)
2. One strong, one weak, use the strong (the weak references the strong)
3. Both weak, choose either

```
void f(void) ;  
int x = 555 ;  
int main() {  
    f() ;  
}  
int x ;  
void f() {  
    x = 666 ; // refers to the global x  
}
```



## Symbol resolution

```
void f(void) ;  
int x = 555 ;  
int main() {  
    f() ;  
}  
double x ;  
void f() {  
    x = 6.66 ; // refers to the global x but not check that the types match  
}
```

Another example

```
void f(void) ;  
int buf[2] = { 1,2 } ;  
int main() {  
    f() ;  
}  
int buf[] ;  
void f() {  
    buf[5] = 20 ; // bad news  
}
```

## Static libraries

Combine pre-compiled routines into a module from a library. Called static.

`printf`, `rand`, `exp2`, etc

Create a library of common routines.

```
gcc main.c /usr/lib/libc.o
```

But the linker only includes those routines which are referenced in `main.c`. Could create a chain of references.

Static libraries are created using the Unix archive function.

## Relocation

Merging all the modules in a link step:

```
int x = 5, y = 6, z = 7 ;  
void f1() {  
.  
.  
.  
}  
int a = 5, b = 6, c = 7 ;  
void f2() {  
.  
.  
.  
}
```

.data section (initialized globals) becomes something like:

a, b, c, x, y, z

When above is compiled, a and x have offset 0 within f2 and f1. After combining, different offsets.

Same is true for the .bss section (uninitialized globals)

When combined with main(), f1 and f2 will have offsets relative to the module.

## Relocation Tables

Contained in the relocatable object module

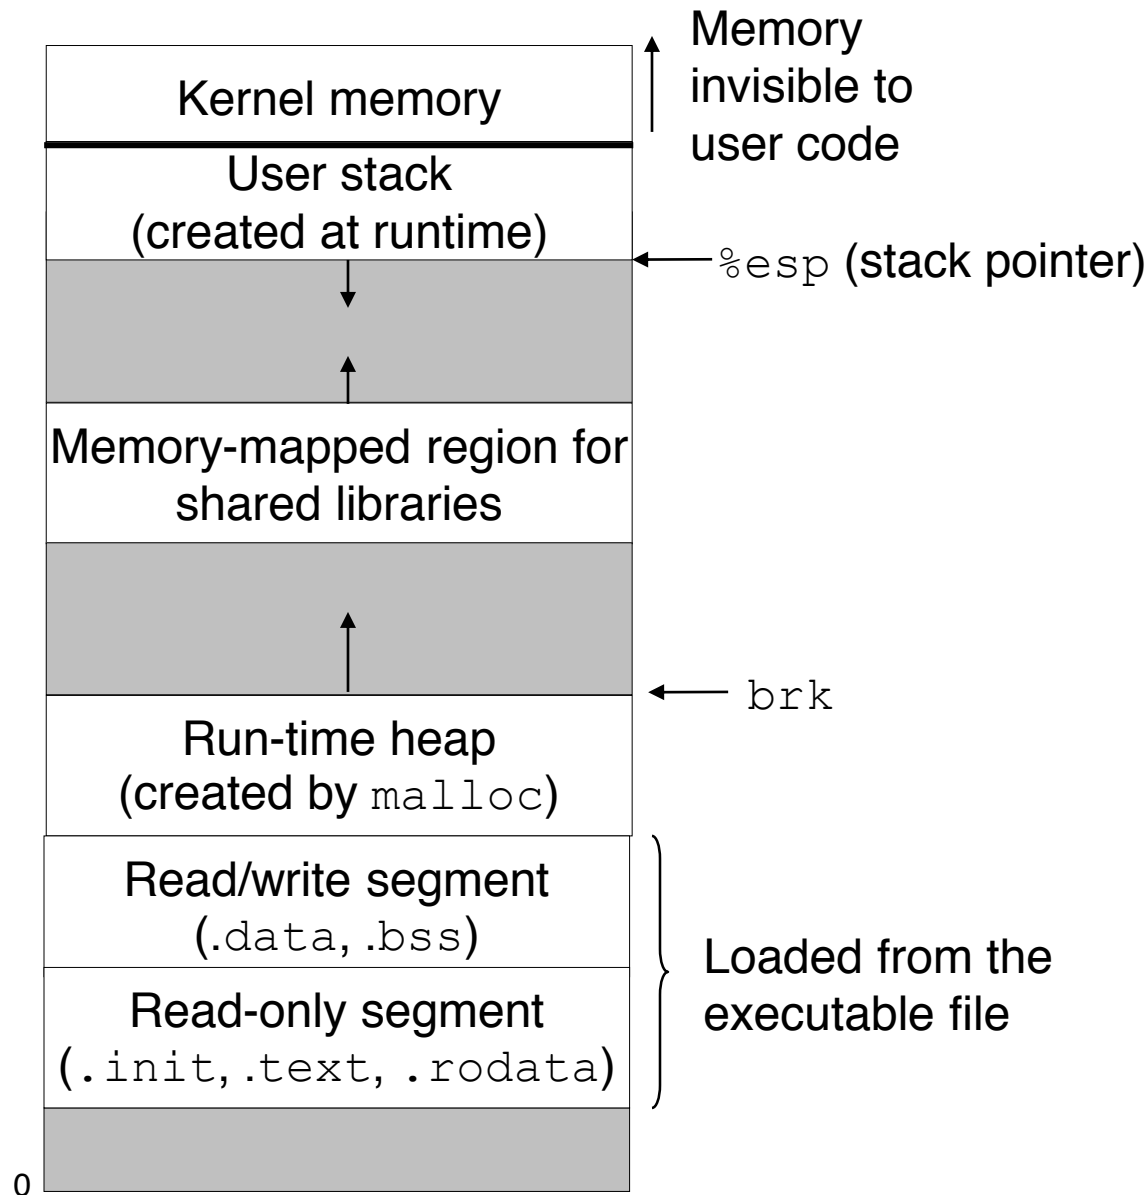
Each entry has the offset within the module of a relocatable object and a type:

2 types are important:

1. relative to the program counter
2. absolute

Depends on how the module was compiled.

## Linux memory map



Use library routines not in the executable

Enhances the maintainability.

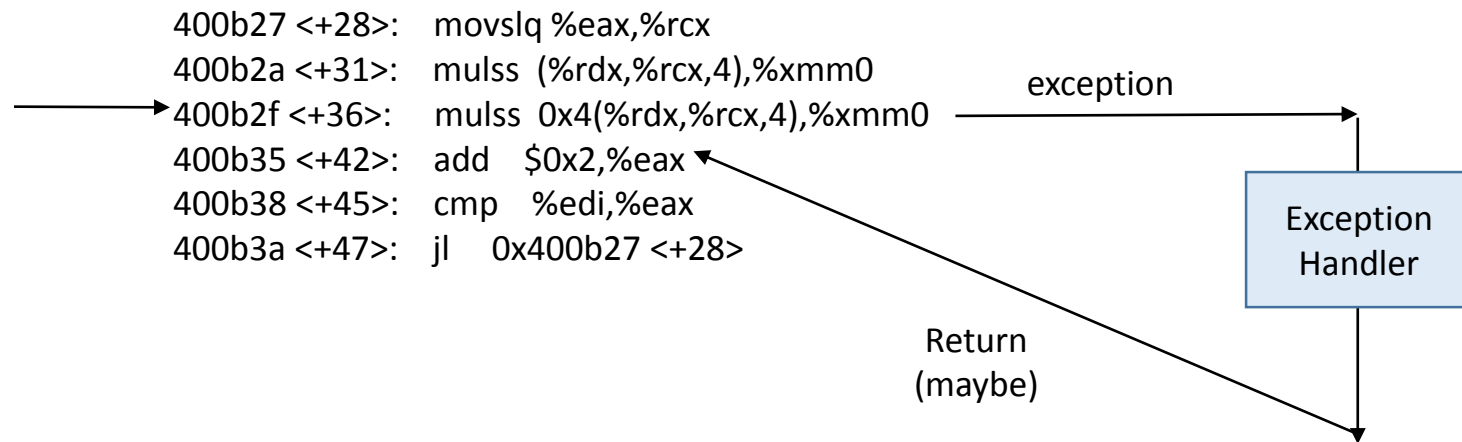
Loads the library into a shared area. Provides tools to locate specific functions within the library.

Functions run in the shared library area.

Functions which contain position independent code can be loaded without relocation. Use a “global offset table” to reference global variables.

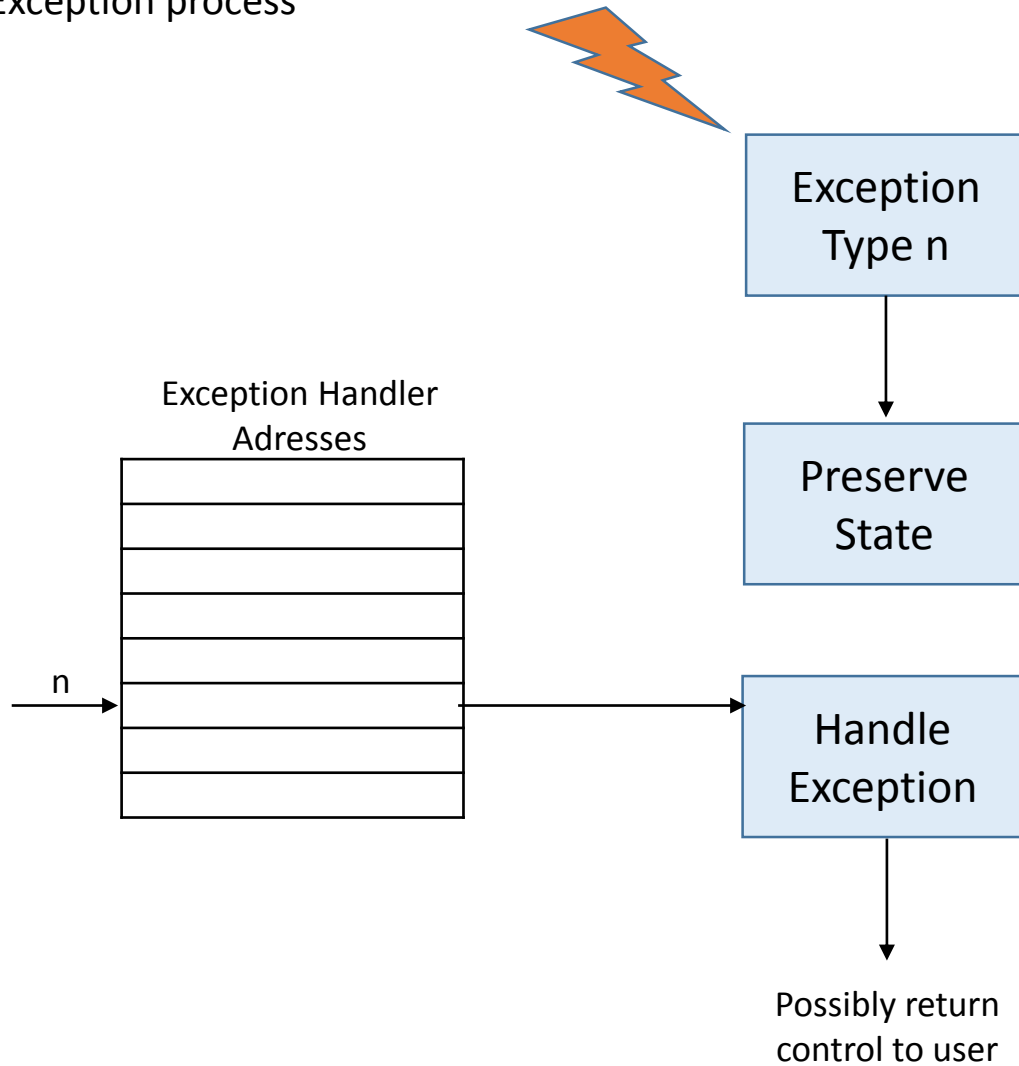
## The key to concurrent execution

Normal execution controlled by the program counter (next instruction or jumps)



# Exceptional Control Flow

Exception process





### Classes of exceptions

Class	Source	Synchronous?	Return?
Interrupt	external	no	Always return to next instruction
Trap	self generated	yes	Always return to next instruction
Fault	Self generated, possibly recoverable	yes	Maybe return to current instruction
Abort	Not recoverable	Yes	Never return

### Types of exceptions

Exception	Class
Divide Error	Fault
Protection Fault	Fault
Page Fault	Fault
Machine Check	Abort
OS-Defined	Interrupt or Trap
System Call	System Call

## System Calls

System Call	Name
Exit	Exit
Fork	Fork
Read file	Read
Write file	Write
Open file	Open
Close file	Close
Wait for child	Waitpid
Load and run	Execve
Go to file offset	Lseek
Get process id	Getpid

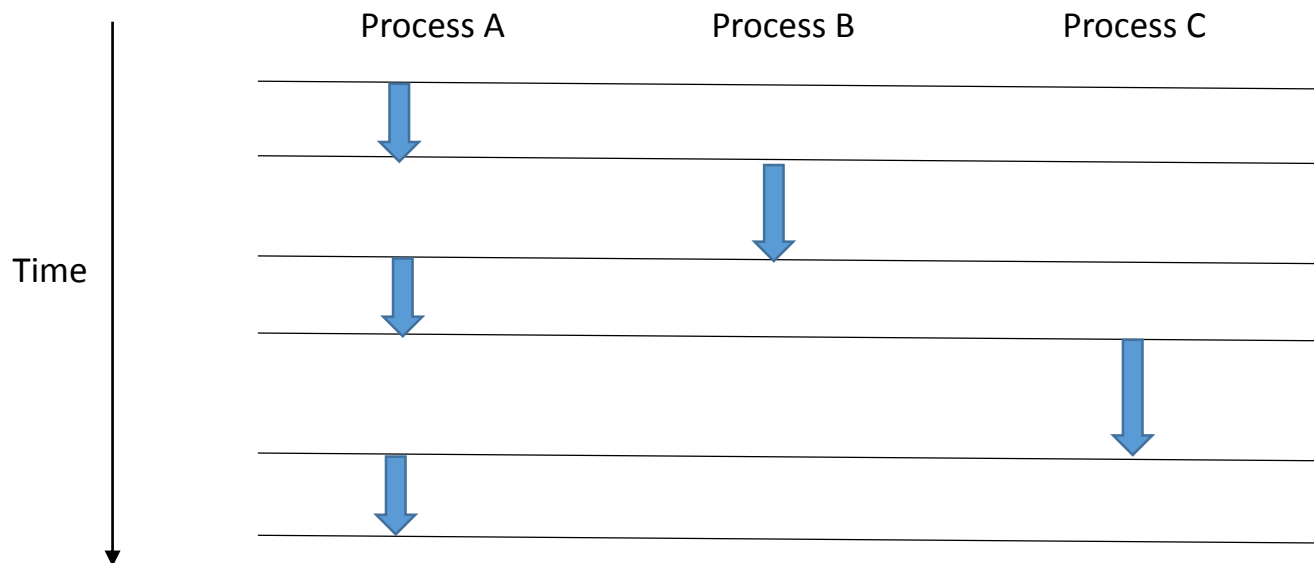
Assembly language instruction is

```
mov $n, %eax // system call number  
int $0x80
```

## Processes

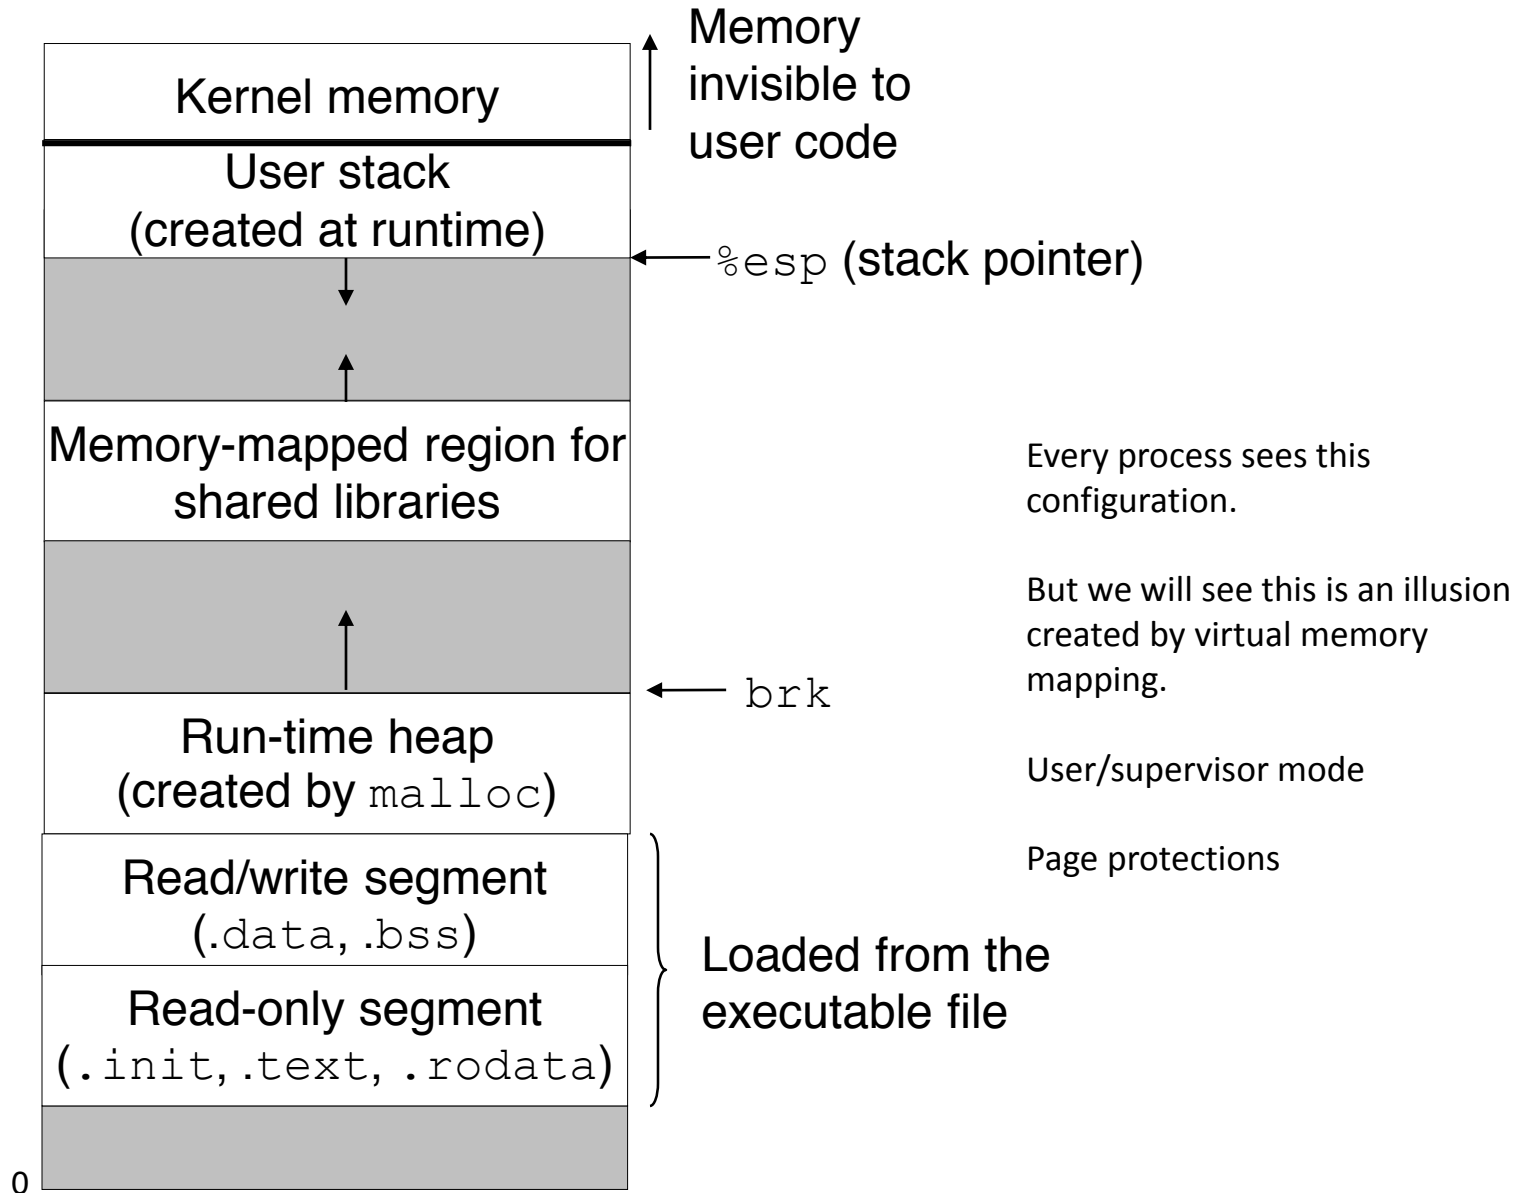
Process: an instance of program execution? Seems to be the only user of the entire resource.

Independent logical control flow.



Concurrency A:B A:C nor B:C

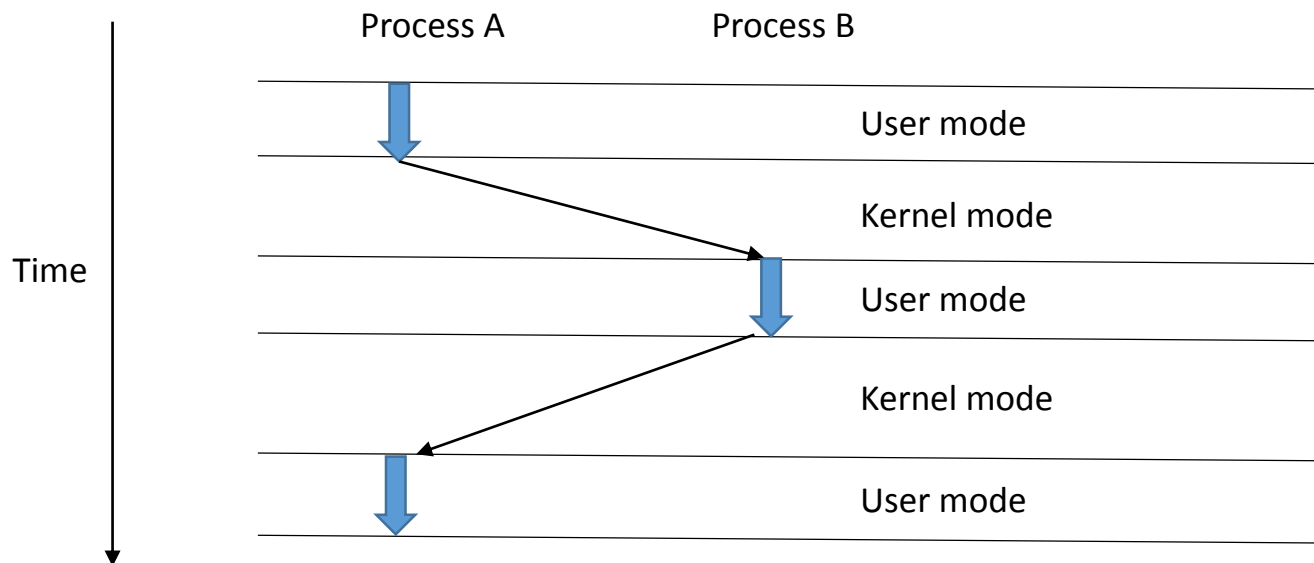
## Linux memory map



## Context Switch

Context: registers, program counter, user stack, files.

Kernel has a scheduler, processes have priority



### Process Control

Unique PID for each process

getpid: get own process ID

getppid: get parent's process ID

Processes are either

- running (could be waiting)

- stopped (stopped by a STOP signal, must be restarted)

- terminated

exit( int status ) terminates with return code

Start new process with fork() ; duplicates the context, creates new PID.

## Fork example

```
#include "csapp.h"
```

```
int main()
```

```
{
```

```
    pid_t pid;
```

```
    int x = 1;
```

```
    pid = fork();
```

```
    if (pid == 0) { /* Child */
```

```
        printf("child : x=%d\n", ++x);
```

```
        exit(0);
```

```
    }
```

```
    /* Parent */
```

```
    printf("parent: x=%d\n", --x);
```

```
    exit(0);
```

```
}
```

Creates new process and returns after the calling point. But it returns twice! But, pid is zero for the child, pid of the child for the parent.

Take out the exit(0) and what happens?

Parent and child now run concurrently.

Address space is duplicated initially but are private (not shared).

Files are shared.



## Exceptional Control Flow

fork example

```
#include "csapp.h"
int main()
{
    fork();
    fork();
    fork();
    printf("hello\n");
    exit(0);
}
```

Main Process

```
fork();
fork();
fork();
printf("hello\n");
exit(0);
```

Child 11

```
fork();
fork();
printf("hello\n");
exit(0);
```

Child 12

```
fork();
printf("hello\n");
exit(0);
```

Child 13

```
printf("hello\n");
exit(0);
```

Child 111

```
fork();
printf("hello\n");
exit(0);
```

Child 112

```
printf("hello\n");
exit(0);
```

Child 121

```
printf("hello\n");
exit(0);
```

Child 1111

```
printf("hello\n");
exit(0);
```

## Process Termination

Process which terminate go into suspended state until reaped.

It is a “zombie” process.

`waitpid( pid_t, int *status, int options )`

when `pid > 0` wait for that child, `pid = -1` wait for any child

`status = -1` when there is no child to wait for

`waitpid` hangs the calling process but can add test options (eg `NOHANG`)

run `fork.c`

sleep and pause

run `fork1.c`

### Run a program

execve loads and runs a program in the current context.

execve does not return: replaces the calling program.

Using execve in conjunction with fork is sort a “shell” behavior.

Example:

```
int main( int argc, char **argv ) {  
  
    execve( “program”, argv, NULL ) ;  
    printf( “hello\n” ) ;  
}
```

hello will never come out. Must do:

```
int main( int argc, char **argv ) {  
  
    If( fork() == 0 )  
        execve( “program”, argv, NULL ) ;  
    printf( “hello\n” ) ;  
}
```

## Signals

High level software implements messaging process. Many types, each has an ID: e.g. kill = ID 9.

Type of interrupt indicating an event has occurred. Similar to exceptions except signal events are handled in user mode.


Requires implementation of signal handlers.

Signal normally sent by kernel but can be sent by a user process to itself.

Signal received by a process by being interrupted and control passing to the signal handler. Process can ignore signals.

Ignored signals are discarded.

Pending signals are queued to a level of 1.



Some examples

#	Name	Default action	Corresponding event
1	SIGHUP	Terminate	Terminal line hangup
2	SIGINT	Terminate	Interrupt from keyboard
3	SIGQUIT	Terminate	Quit from keyboard
4	SIGILL	Terminate	Illegal instruction
5	SIGTRAP	Terminate and dump core	Trace trap
6	SIGABRT	Terminate and dump core	Abort signal from abort function
7	SIGBUS	Terminate	Bus error
8	SIGFPE	Terminate and dump core	Floating point exception
9	SIGKILL	Terminate	Kill program
10	SIGUSR1	Terminate	User-defined signal 1
11	SIGSEGV	Terminate and dump core	Invalid memory eference (seg fault)
12	SIGUSR2	Terminate	User-defined signal 2
13	SIGPIPE	Terminate	Wrote to a pipe with no reader
14	SIGALRM	Terminate	Timer signal from alarm function
15	SIGTERM	Terminate	Software termination signal
16	SIGSTKFLT	Terminate	Stack fault on coprocessor
17	SIGCHLD	Ignore	A child process has stopped or terminated
18	SIGCONT	Ignore	Continue process if stopped
19	SIGSTOP	Stop until next SIGCONT (2)	Stop signal not from terminal
20	SIGTST	Stop until next SIGCONT	Stop signal from terminal
21	SIGTTIN	Stop until next SIGCONT	Background process read from terminal
22	SIGTTOU	Stop until next SIGCONT	Background process wrote to terminal

## Process groups/handlers

A grouping of processes using an ID. Child belongs to parent group by default.

Can change process group ID using `setpgid( pid, pgid ) ;`

Can send the same signal to every process in a group.  
e.g. can kill every process in a group or a single process.

`signal( sig, func )` function “installs” or denotes a func as the handler for signal # sig.

signal can be used to `IGNore`, use default (DFL) or the named function as the handler.

### Example of handler

```
#include <stdio.h>
#include <signal.h>
#include <stdlib.h>

void handler(int sig) /* SIGINT handler */
{
    printf("Caught SIGINT\n");
    exit(0);
}

int main()
{
    /* Install the SIGINT handler */
    if (signal(SIGINT, handler) == SIG_ERR)
        unix_error("signal error");

    pause(); /* Wait for the receipt of a signal */

    exit(0);
}
```

### Pending signals

Blocked signals: same signals are blocked during signal handling. 2nd signal processed after first.

Signal queue: size 1, additional signals are ignored/discarded.

Certain system calls, when interrupted by a signal are not resumed.

Run handler1

### Blocking signals

Signal set contains the signals currently blocked: array of signal numbers

sigemptyset:	unblocks all signals
sigfillset:	blocks all signals
sigaddset:	adds a signal to the set
sigdelset	deletes a signal
sigprocmask:	manipulates the set (blocks/unblocks)
sigismember	tests



### Synchronization

Cannot predict when child or parent runs. Different results.


Run handler3 & 4

## Overview

Illusion of memory.

Actually a version of caching.

Cache line



tt...ttss...ssbb...bb is the address in physical memory. ss....ss is the line in the cache, bb...bb is the block number in the line. tt...tt validates the cache line.

Virtual address is address in “pretend” memory. Physical address is the address in main memory.

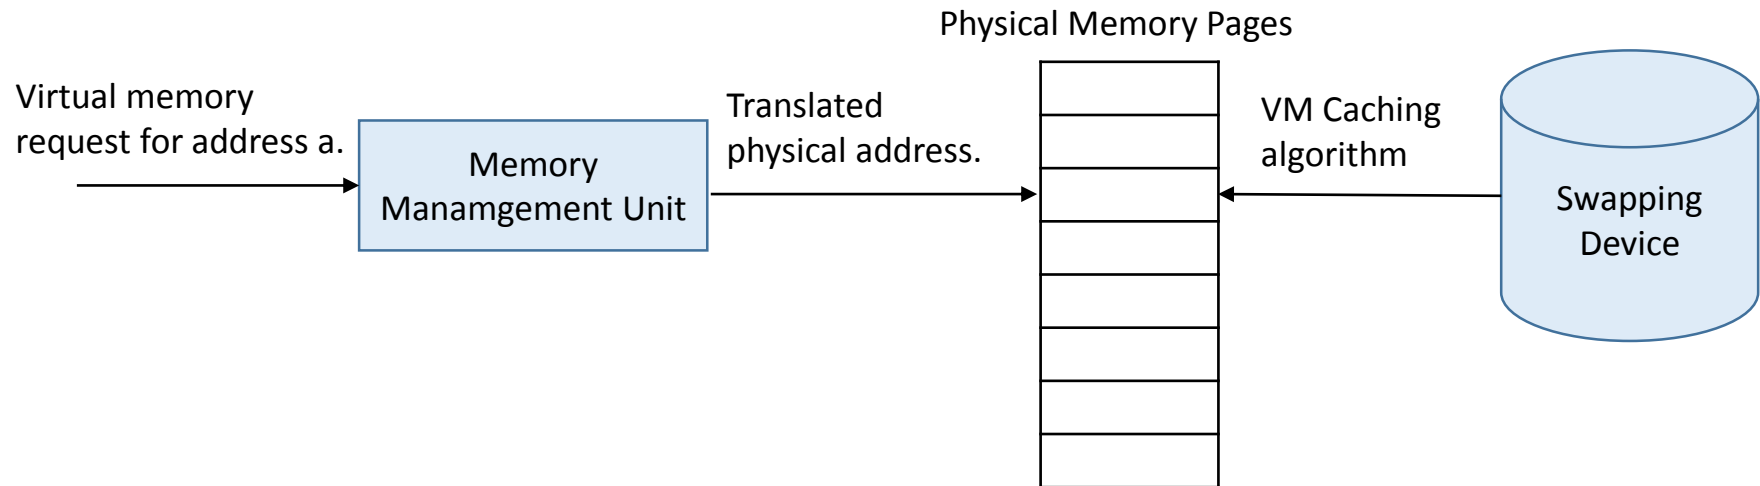
Virtual memory divided into pages. Stored in physical memory using swapping scheme.

Address translation

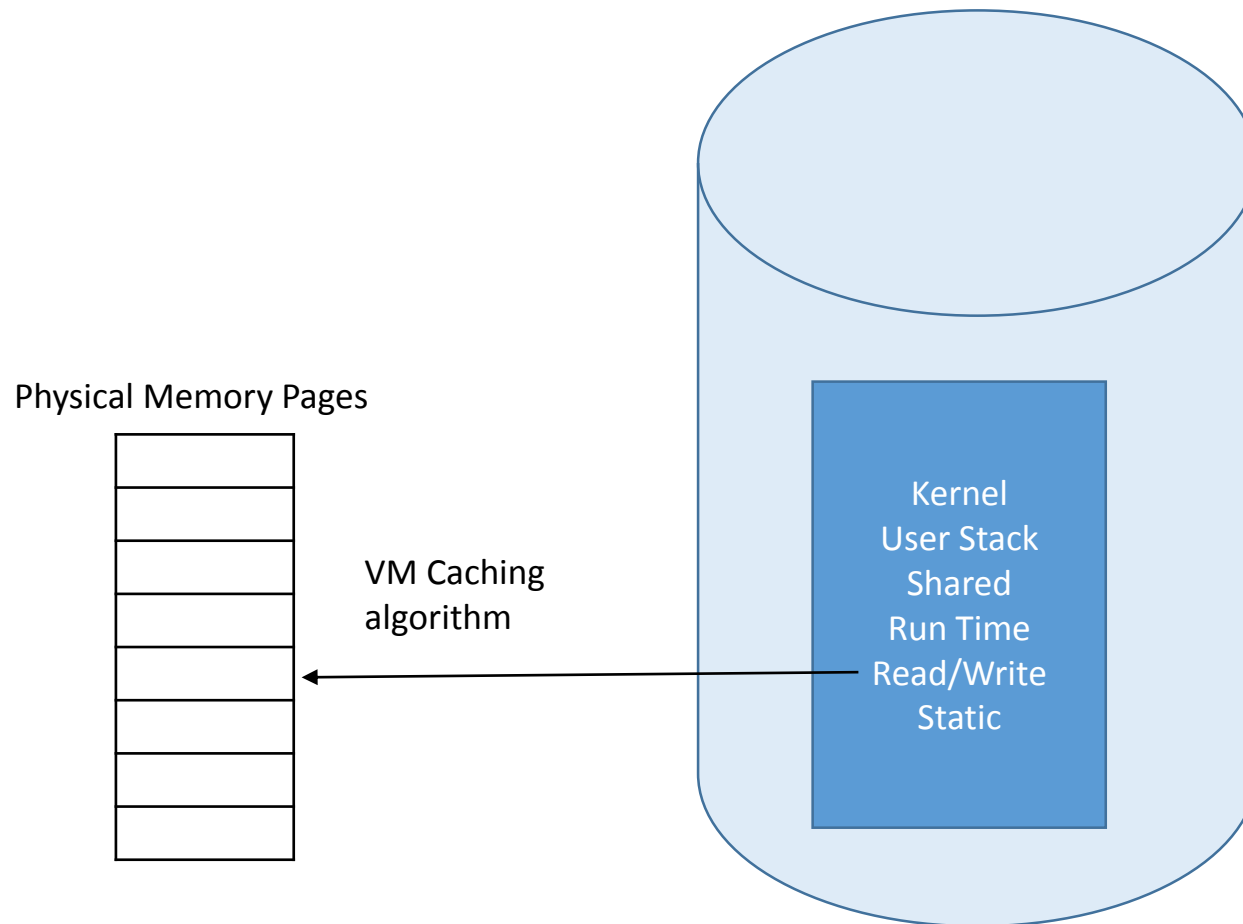
# Virtual Memory

Address translation

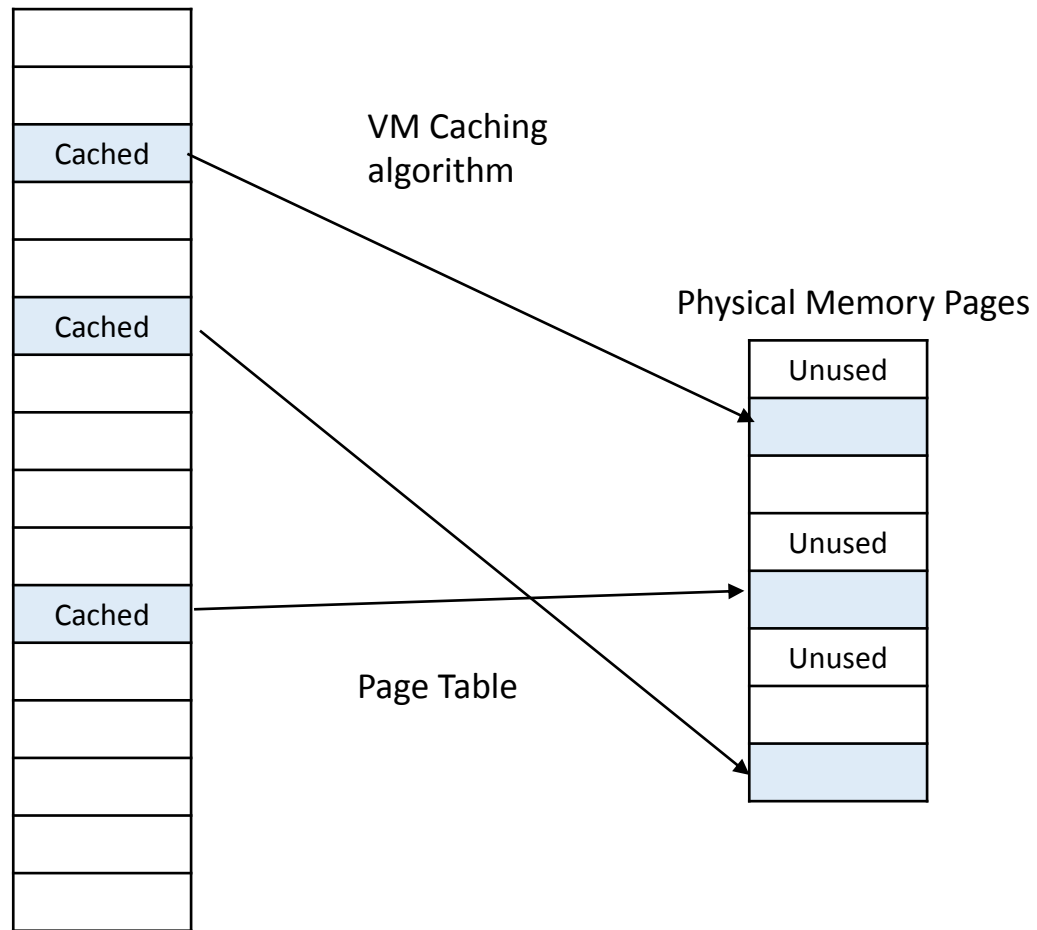
Data fetching



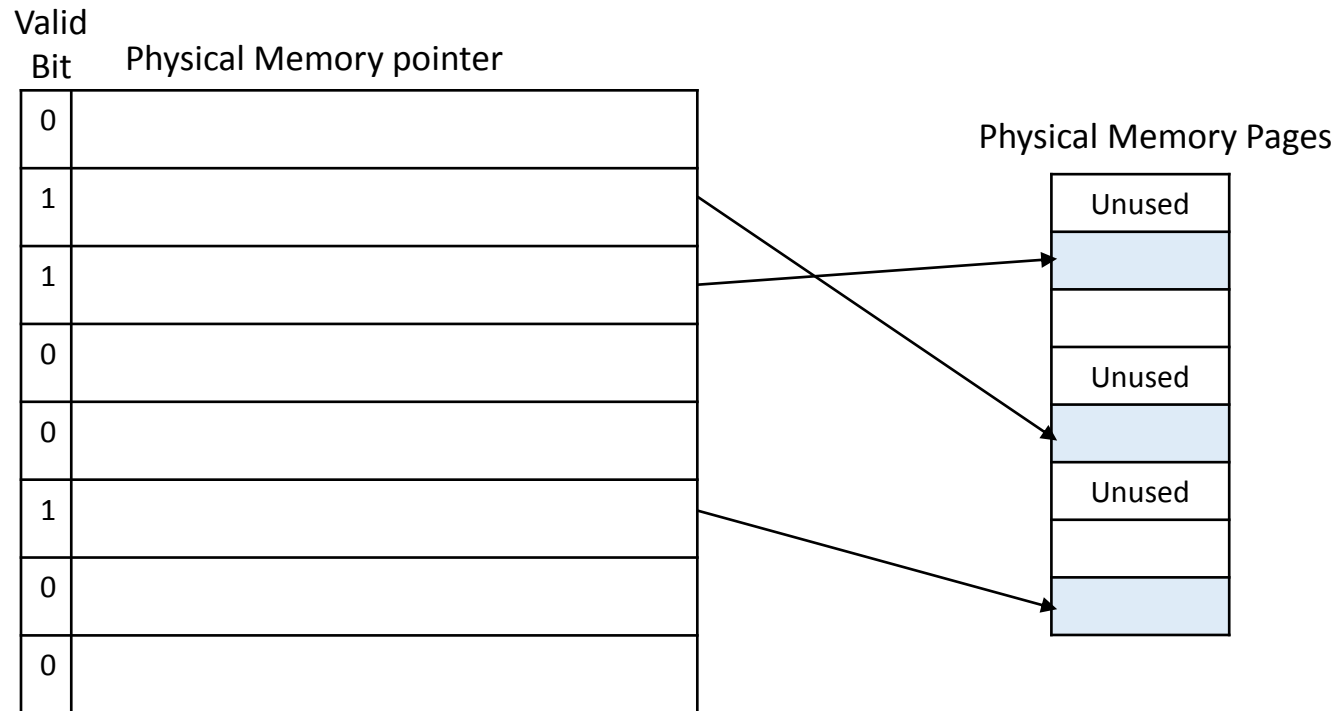
## Physical appearance



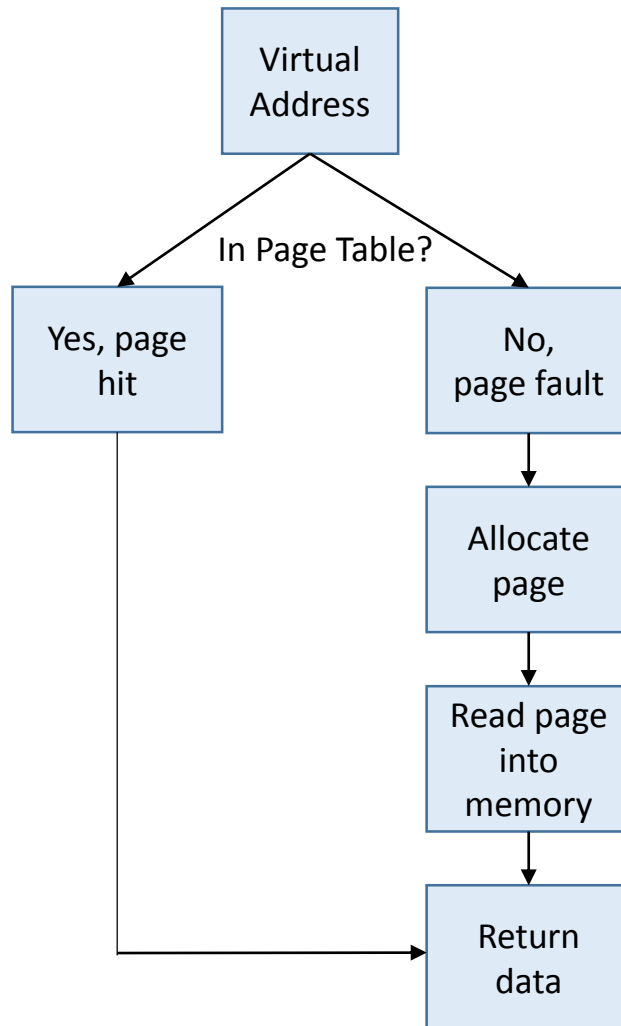
Virtual Memory Pages



## Page Table



Address request




## Multiple processes

Page table for each process.

Switch process, switch page table register.

Possibly shared pages (two table point to the same physical page)





## Memory protection

Add protection bits to page table:

Supervisor mode, read?, write?, execute?

Valid Protection

Bit	Bits				Physical Memory pointer
0					
1					
1					
0					
0					
1					
0					
0					

## Address translation terminology

Similar to caching?

Memory sizes

$N = 2^n$       addresses in virtual memory  
 $M = 2^m$       addresses in physical memory  
 $P = 2^p$       page size

Virtual address made up of

VPN      virtual page number  
 VPO      virtual page offset  
 TLBI      translation lookaside buffer index  
 TLBO      translation lookaside buffer offset

Divide address (n bits) into sections by p

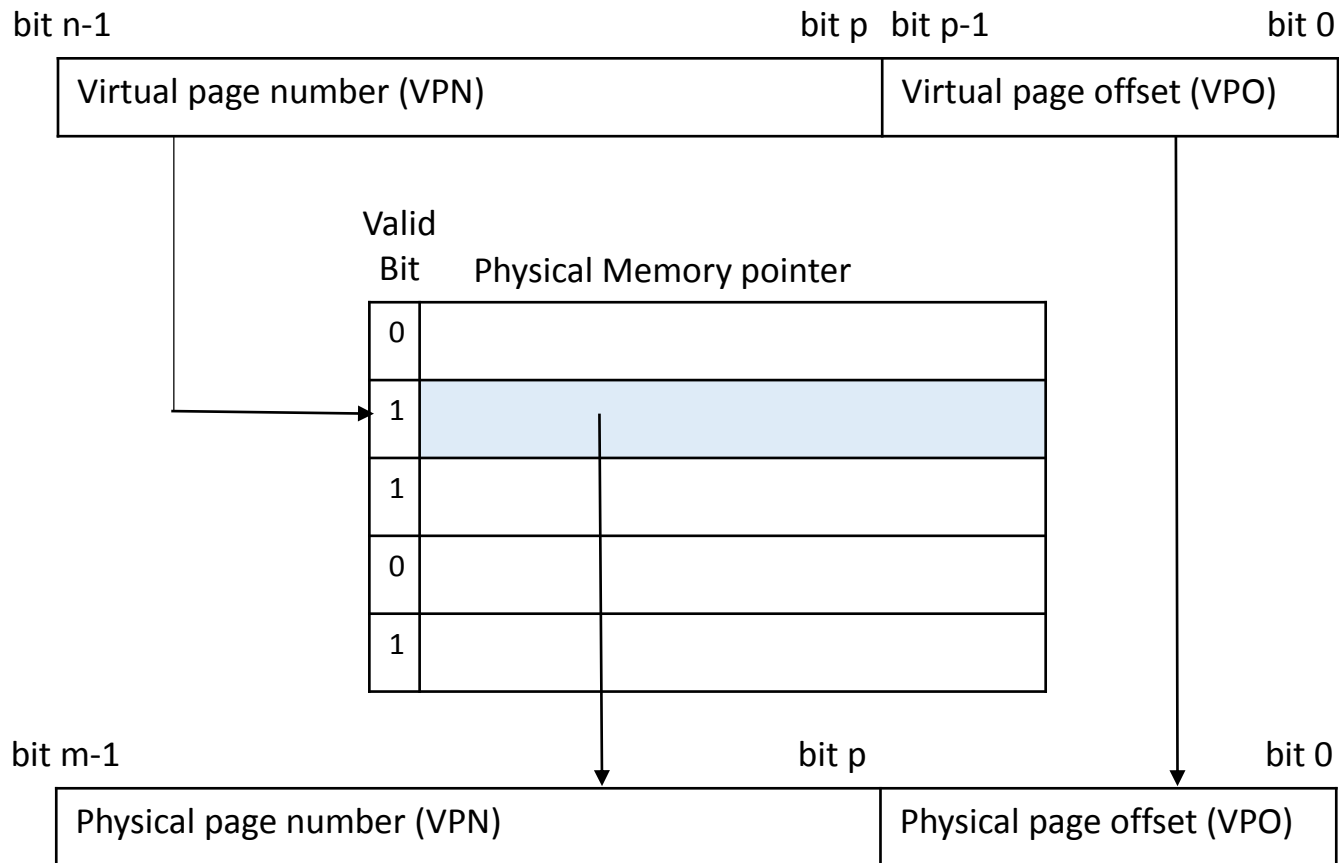
Construct mapping of virtual addresses into physical addresses (VM caching algorithm)

Physical address made up of:

PPN      physical page number  
 PPO      physical page offset  
 CI      cache index      for L1 cache  
 CO      cache offset  
 CT      cache tag

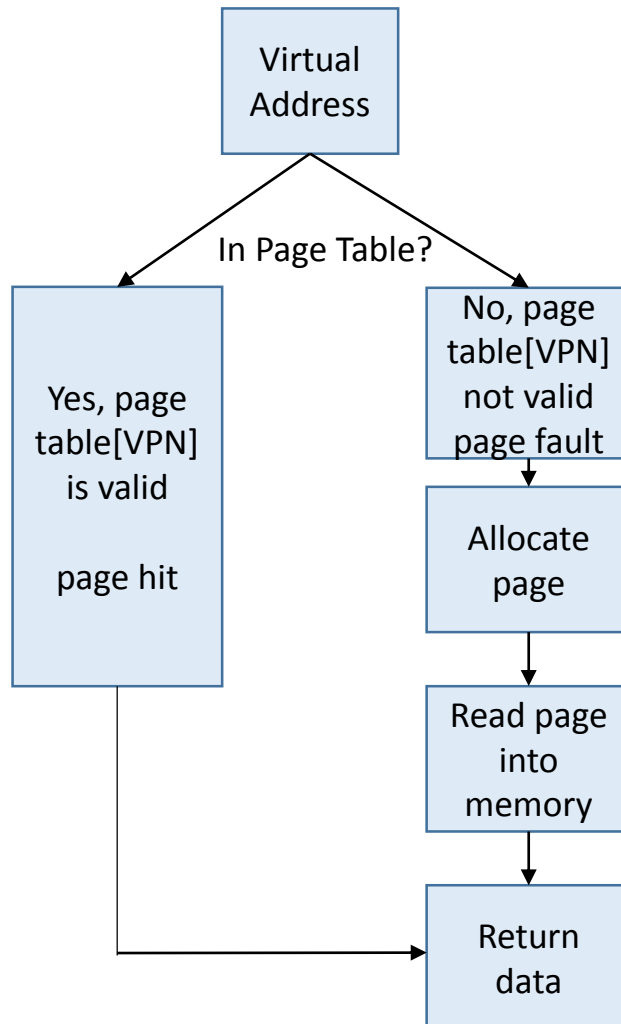
# Virtual Memory

## Address translation methodology



Situation for one process:  
each has its own page table

Direct mapped cache?




## Address translation terminology

N/M virtual pages per physical page.

Number of entries in page table? N/P. Can be quite large and resides in memory!

To reduce memory accesses to the PTE, use a cache in the memory management unit:

Translation lookaside buffer (TLB)



Cache contains one block per line = page table entry for VPN,  $2^t$  lines.

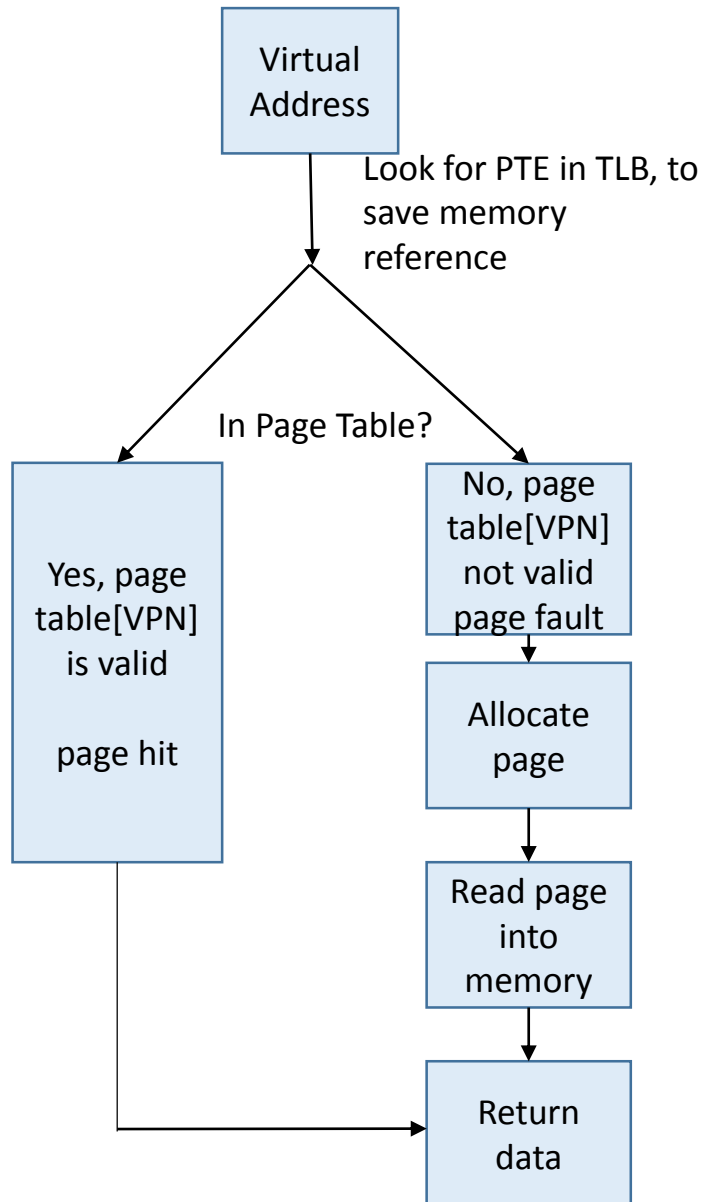
Reduces memory accesses.

## TLB diagram

TLB[0]	TLB Valid	TLB Tag	PTE entry	TLB Valid	TLB Tag	PTE entry	...
TLB[1]							
.				.			
.				.			
.				.			
TLB[2 <sup>t</sup> -2]							
TLB[2 <sup>t</sup> -1]	TLB Valid	TLB Tag	PTE entry	TLB Valid	TLB Tag	PTE entry	...

N-way associative, addressed by TLB index and TLB Tag

Direct mapped cache?



## Hierarchy of page tables

Most processes use only a little bit of the memory at a time.

Most page table entries do not point to anything.

Example:

Use 32 bit address space. Divide up into 4 MB ( 1024, 4096 byte ) sections

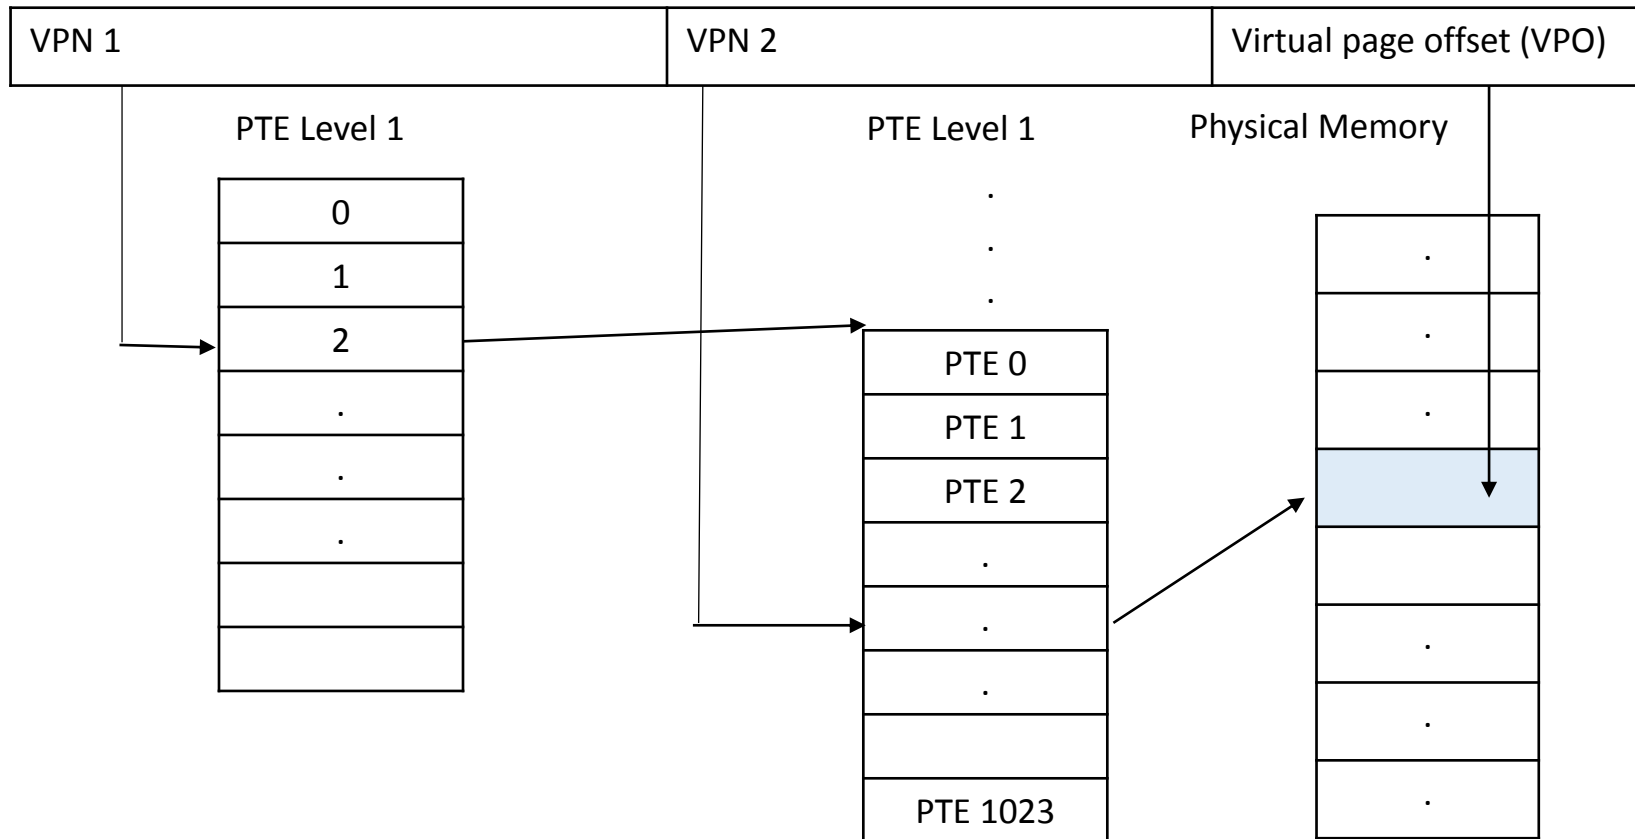
First level points, not to a physical address but the secondary PTE.

Second level corresponds to virtual address in its 4 MB section



## Hierarchy of page tables

corresponds to virtual address in its 4 MB section



## VM Caching example

Assume  $m = 32$  so virtual memory is  $2^{32} = 4,294,967,296$  bytes (4GB)  
 $p = 12$  page size is  $2^{12} = 4096$  bytes (4K)

Divides the memory into  $2^{20} = 1,048,576$  (1MB) pages

For a new process, we would allocate 1MB 4K records on the swapping disk.

Some bookkeeping mechanism to correlate page number to disk location.


In what follows, memory address go up when you go down the page.

## VM Caching example

List of memory addresses on disk after load of program

	Virtual address	Disk address	
read only area	0	0	
	1	.	
	.	.	
	.		
	.	.	
read/write area	x	x	→ Point to 4K records on executable object file (file backed memory)
	x+1	x+1	
	.	NULL	
	.	x+2	
run time heap	y	NULL	→ Uninitialized areas have no records on disk until a dirty write occurs. Afterward, they will point to 4K records in the swap file. (demand zero memory)
		NULL	
	.	NULL	
		NULL	
user stack area		NULL	
	.	NULL	
		NULL	


## Address layout



n's form a 20 bit unsigned number = Virtual Page Number VPN

o's form a 12 bit unsigned number = Virtual Page Offset VPO

Virtual pages from disk will be stored in physical memory, as needed. For our example, let the physical memory be  $2^{20} = 1\text{MB}$  or 256 4K pages



n's form a 8 bit unsigned number = Physical Page Number PPN

o's form a 12 bit unsigned number = Physical Page Offset PPO

## Page Table

Disk pages ( = Virtual Memory pages ) will be stored in physical memory, when being used by the process.

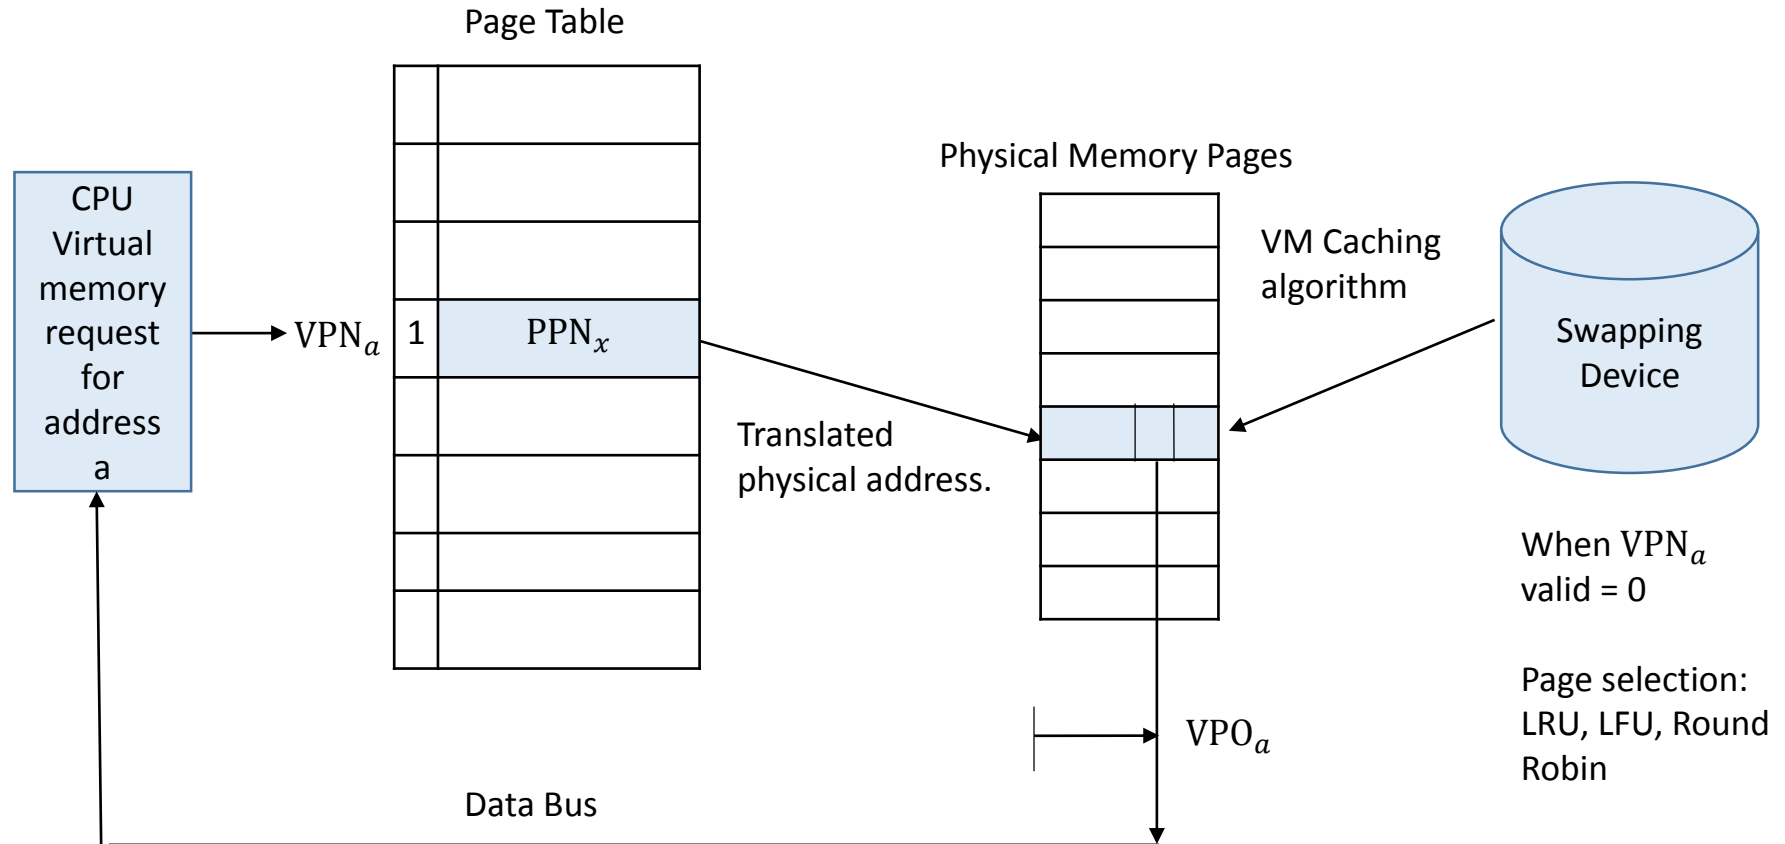
Need a table to translate VPN into PPN. 1MB entries, 4 bytes each, directly addressed by VPN.

	Valid Protection					Physical Memory pointer	
	Bit	Bits					
VPN <sub>0</sub>							Page Table Entry (PTE)
VPN <sub>1</sub>							
.							
.							
.							

When virtual page not in physical memory, valid bit = 0.

# Virtual Memory

## Address translation




## Translation Lookahead Buffer

Page tables are large, stored in memory. 1 entry per virtual page 1MB x 4.

Speed up with a page table entry cache: TLB. n-way associative (fast), with t lines.

Divide VPN into tag and index: for our example, assume TLB has 16 lines, 4 entries:



## TLB diagram

TLB[0]	TLB Valid	TLB Tag	PTE entry	TLB Valid	TLB Tag	PTE entry	TLB Valid	TLB Tag	PTE entry	TLB Valid	TLB Tag	PTE entry
TLB[1]												
.												
.												
.												
TLB[14]												
TLB[15]	TLB Valid	TLB Tag	PTE entry	TLB Valid	TLB Tag	PTE entry						

4-way associative, addressed by TLB index and TLB Tag



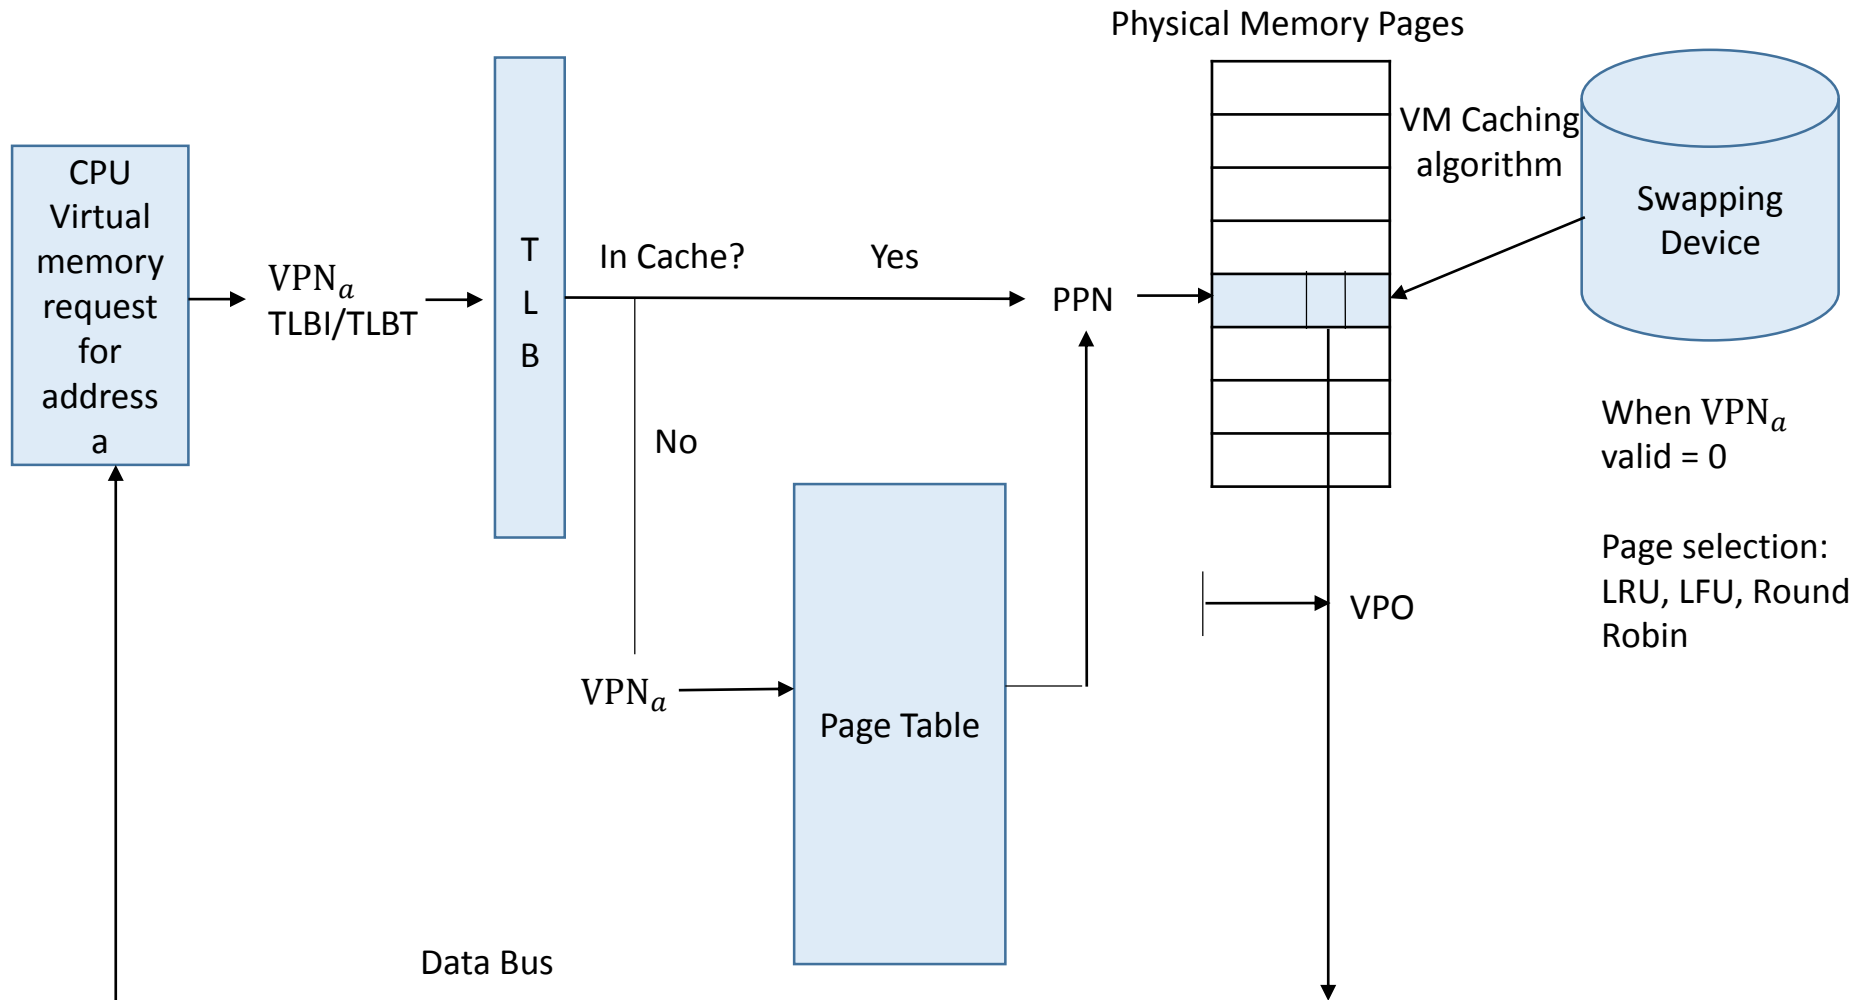
## TLB example

	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
			0			0			0			0
			0	Tag	PTE	1			0			0
			0			0			0			0
			0			0			0			0
			0			0	Tag	PTE	1			0
			0			0			0			0
			0			0			0			0
TLBI →	Tag	PTE	1			0	TLBT	PTE	1			0
			0			0			0			0
			0			0			0			0
			0			0			0			0
			0			0			0			0
			0			0			0			0
			0			0			0			0
			0			0			0			0
			0			0			0	Tag	PTE	1
			0			0			0			0
			0			0			0			0

Found by associativity with valid and tag in line based on TLBI

## TLB example

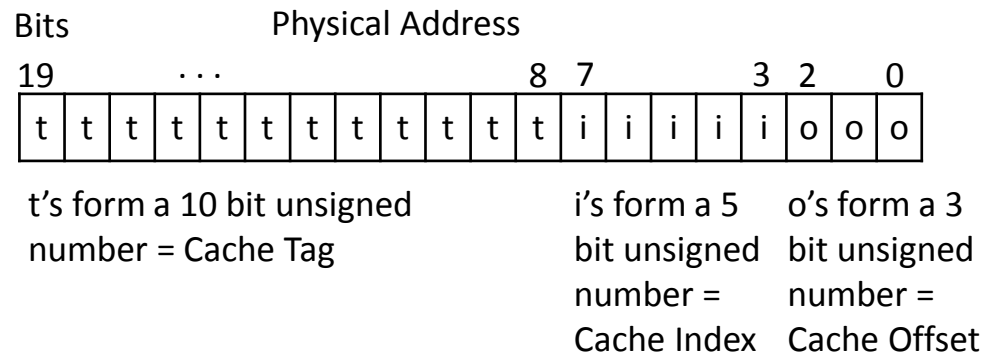
Address retrieval now looks like



## One Final Step

Once the physical address is obtained, need to retrieve from physical memory: L1 cache

Direct mapped. Say 8 blocks per line, 32 lines



# Virtual Memory

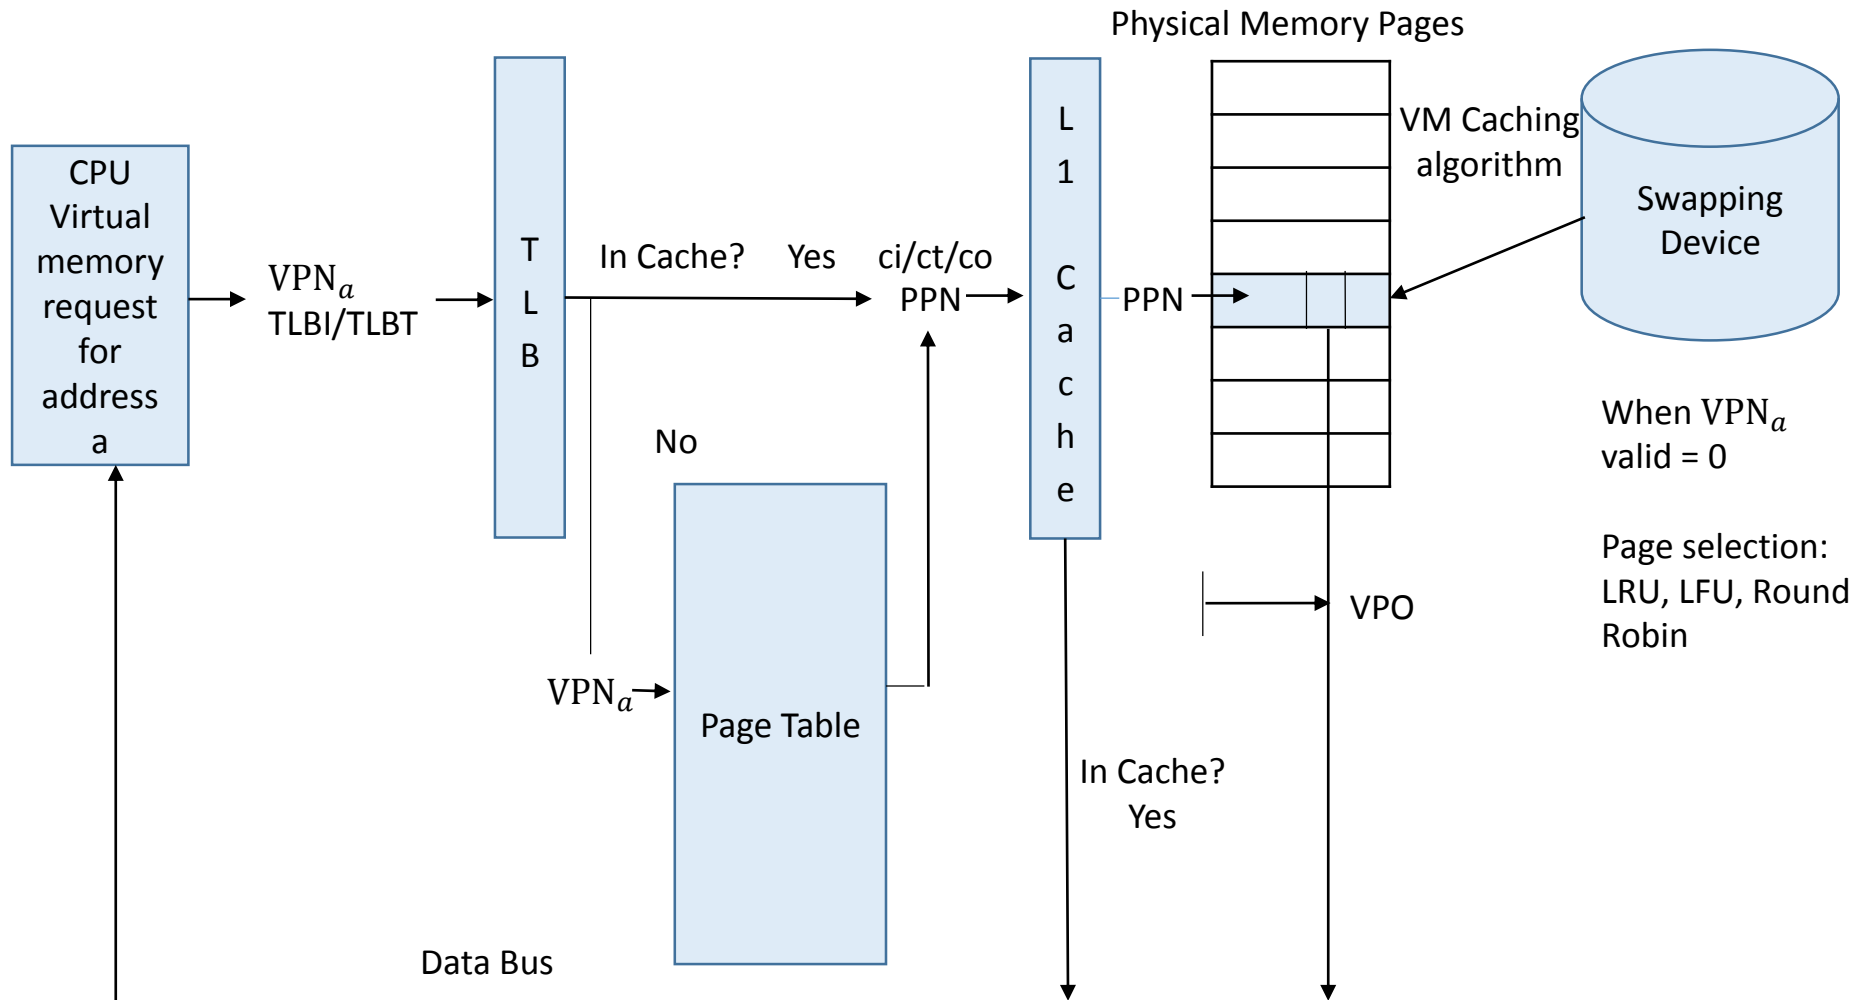
L1 example

	Tag	Valid	Blk 0	Blk 1	Blk 2	Blk 3	Blk 4	Blk 5	Blk 6	Blk 7
		0								
		0								
		0								
	Tag	1	Data <sub>0</sub>	Data <sub>1</sub>	Data <sub>2</sub>	Data <sub>3</sub>	Data <sub>4</sub>	Data <sub>5</sub>	Data <sub>6</sub>	Data <sub>7</sub>
		0								
		0								
		0								
CI →	CT	1	Data <sub>0</sub>	Data <sub>1</sub>	Data <sub>2</sub>	Data <sub>3</sub>	Data <sub>4</sub>	Data <sub>5</sub>	Data <sub>6</sub>	Data <sub>7</sub>
		0								
		0								
		0								
	Tag	1	Data <sub>0</sub>	Data <sub>1</sub>	Data <sub>2</sub>	Data <sub>3</sub>	Data <sub>4</sub>	Data <sub>5</sub>	Data <sub>6</sub>	Data <sub>7</sub>
		0								
		0								
	Tag	1	Data <sub>0</sub>	Data <sub>1</sub>	Data <sub>2</sub>	Data <sub>3</sub>	Data <sub>4</sub>	Data <sub>5</sub>	Data <sub>6</sub>	Data <sub>7</sub>
		0								
		0								

CO ↑

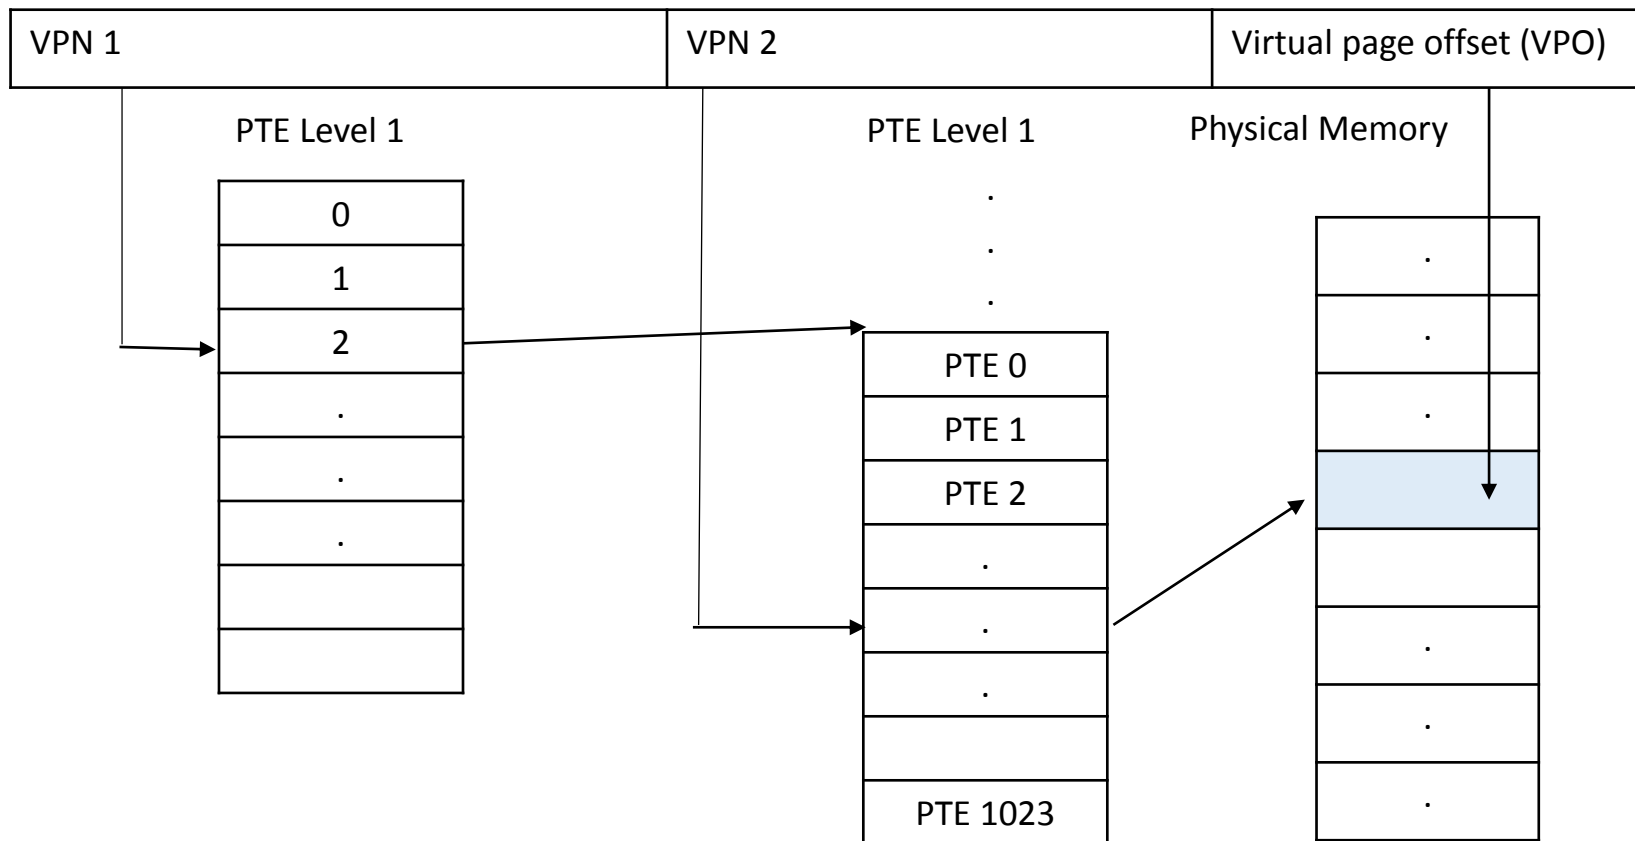
## L1 Cache Example

Address retrieval now looks like



Recall that we could have

corresponds to virtual address in its 4 MB section



Paths through MMU to retrieve address a.

Path 1: hit in TLB, PTE hit in cache

Step 1. Extract VPN and VPO of address

VPO = last  $p$  bits where  $2^p = \text{pages size}$

VPN = first  $n-p$  bits where there are  $2^n$  virtual addresses

$n = 14$   $p = 6$   $t = 2$

Bit	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Virt	vpn	vpn	vpn	vpn	vpn	vpn	vpn	vpn	vpo	vpo	vpo	vpo	vpo	vpo
TLB	tlbt	tlbt	tlbt	tlbt	tlbt	tlbt	tlbi	tlbi						

Step 2. Extract TLBI and TLBT from VPN

TLBI = last  $t$  bits where there are  $2^t$  sets in TLB

TLBT = first  $n-t-p$  bits of VPN

Step 3. TLBT in TLB[TLBI] ? associatively ... valid ?

Paths through MMU to retrieve address a.

Step 4. retrieve PA = PPN concatenated with VPO if hit

Step 5. Extract CO from PA = last  $b$  bits where cache has  $2^b$  blocks per set

Extract CI from PA = middle  $s$  bits where cache has  $2^s$  sets

Extract CT from PA = first  $m-s-b$  bits where  $2^m$  is physical memory size

Bit	13	12	11	10	9	8	7	6	5	4	3	2	1	0
cache			ct	ct	ct	ct	ct	ct	ci	ci	ci	ci	co	co
Phys			pn	pn	pn	pn	pn	pn	po	po	po	po	po	po

Step 6. cache[CI].valid and cache[CI].tag = CT ?

Step 7. return cache[CI].block[CO] as memory value



Paths through MMU to retrieve address a.

Path 2: hit in TLB, PTE miss in cache

Steps 1-5 for Path 1.

Step 6.  $\sim \text{cache}[\text{CI}].\text{valid}$  or  $\text{cache}[\text{CI}] \neq \text{CT}$

Step 7. Retrieve b blocks from memory into cache (direct mapped)

Step 8. return  $\text{cache}[\text{CI}].\text{block}[\text{CO}]$  as memory value

Path 3: miss in TLB

Steps 1-2 for Path 1.

Step 3. TLBT not in  $\text{TLB}[\text{TLBI}]$  ? associatively ... or  $\sim \text{valid}$  ?

Step 4. Fetch PTE from memory

Step 5. Place appropriate PTEs in TLB

Steps 3-7 for path 1. \*\*\*Note\*\*\* cache miss still possible

Paths through MMU to retrieve address a.

Path 2: hit in TLB, PTE miss in cache

Steps 1-5 for Path 1.

Step 6.       $\sim \text{cache}[\text{CI}].\text{valid}$  or  $\text{cache}[\text{CI}] \neq \text{CT}$

Step 7.      Retrieve b blocks from memory into cache (direct mapped)

Step 8.      return  $\text{cache}[\text{CI}].\text{block}[\text{CO}]$  as memory value

Path 3: miss in TLB

Steps 1-2 for Path 1.

Step 3.      TLBT not in  $\text{TLB}[\text{TLBI}]$  ? associatively ... or  $\sim \text{valid}$  ?

Step 4.      Fetch PTE from memory

Step 5.      Place appropriate PTEs in TLB

Steps 3-7 for path 1. \*\*\*Note\*\*\* cache miss still possible

## Example from text

n = 14 bit virtual address

m = 12 bit physical address

p = 6 = 64 byte pages

t = 2 = 4 sets in TLB, 4x associative (4 lines per set)

L1 cache 16 sets, 4 blocks

Looking for data at address 0x03d4

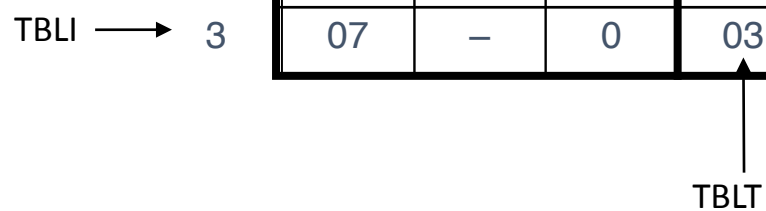
### Worksheet 1

Bit	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Virt	vpn	vpn	vpn	vpn	vpn	vpn	vpn	vpn	vpo	vpo	vpo	vpo	vpo	vpo
TLB	tlbt	tlbt	tlbt	tlbt	tlbt	tlbt	tlbi	tlbi						
0x03d4	0	0	0	0	1	1	1	1	0	1	0	1	0	0

TLB Tag 0x03, TLB Index = 0x3, VPN = 0x0F, VPO = 0x14

## Translation Lookahead Buffer

Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	03	—	0	09	0D	1	00	—	0	07	02	1
1	03	2D	1	02	—	0	04	—	0	0A	—	0
2	02	—	0	08	—	0	06	—	0	03	—	0
TBLI → 3	07	—	0	03	0D	1	0A	34	1	02	—	0


  
 TBLT

PPN = 0x0D =  $1101_2$ , VPO = 0x14 =  $010100_2$

Concatenated =  $1101010100 = 0x354$

PPN to cache

## Worksheet 2

Bit	11	10	9	8	7	6	5	4	3	2	1	0
Phys	ct	ct	ct	ct	ct	ct	ci	ci	ci	ci	co	co
0x0354	0	0	1	1	0	1	0	1	0	1	0	0

CT = 0x0d, CI = 0x05, CO = 0

# Virtual Memory

## L1 Cache

	Idx	Tag	Valid	Blk 0	Blk 1	Blk 2	Blk 3
	0	19	1	99	11	23	11
	1	15	0	—	—	—	—
	2	1B	1	00	02	04	08
	3	36	0	—	—	—	—
	4	32	1	43	6D	8F	09
CI →	5	0D	1	36	72	F0	1D
CT →	6	31	0	—	—	—	—
Valid →	7	16	1	11	C2	DF	03
	8	24	1	3A	00	51	89
	9	2D	0	—	—	—	—
	A	2D	1	93	15	DA	3B
	B	0B	0	—	—	—	—
	C	12	0	—	—	—	—
	D	16	1	04	96	34	15
	E	13	1	83	77	1B	D3
	F	14	0	—	—	—	—

CO

Data Value = 0x36

## Example from text

n = 14 bit virtual address

m = 12 bit physical address

p = 6 = 64 byte pages

t = 2 = 4 sets in TLB, 4x associative (4 lines per set)

L1 cache 16 sets, 4 blocks

Looking for data at address 0x03d7


### Worksheet 1

Bit	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Virt	vpn	vpn	vpn	vpn	vpn	vpn	vpn	vpn	vpo	vpo	vpo	vpo	vpo	vpo
TLB	tlbt	tlbt	tlbt	tlbt	tlbt	tlbt	tlbi	tlbi						
0x03d4	0	0	0	0	1	1	1	1	0	1	0	1	1	1

TLB Tag 0x03, TLB Index = 0x3, VPN = 0x0F, VPO = 0x17

## Translation Lookahead Buffer

Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	03	—	0	09	0D	1	00	—	0	07	02	1
1	03	2D	1	02	—	0	04	—	0	0A	—	0
2	02	—	0	08	—	0	06	—	0	03	—	0
TBLI → 3	07	—	0	03	0D	1	0A	34	1	02	—	0


  
TBLT

PPN = 0x0D =  $1101_2$ , VPO = 0x17 =  $010111_2$

Concatenated =  $1101010111 = 0x357$



PPN to cache

## Worksheet 2

Bit	11	10	9	8	7	6	5	4	3	2	1	0
Phys	ct	ct	ct	ct	ct	ct	ci	ci	ci	ci	co	co
0x0354	0	0	1	1	0	1	0	1	0	1	1	1

CT = 0x0d, CI = 0x05, CO = 3

# Virtual Memory

## L1 Cache

	Idx	Tag	Valid	Blk 0	Blk 1	Blk 2	Blk 3
	0	19	1	99	11	23	11
	1	15	0	—	—	—	—
	2	1B	1	00	02	04	08
	3	36	0	—	—	—	—
	4	32	1	43	6D	8F	09
CI →	5	0D	1	36	72	F0	1D
CT →	6	31	0	—	—	—	↑
	7	16	1	11	C2	DF	03
	8	24	1	3A	00	51	89
	9	2D	0	—	—	—	—
	A	2D	1	93	15	DA	3B
	B	0B	0	—	—	—	—
	C	12	0	—	—	—	—
	D	16	1	04	96	34	15
	E	13	1	83	77	1B	D3
	F	14	0	—	—	—	—

Valid →

CO

Data Value = 0x1D

## Heap Management

Allocate memory when needed. Especially if amount not know at run time

2 types:      explicit: must free unused space  
                implicit: used and unused space recognizable. free by garbage collection

Linux/C/C++ use explicit

`void *malloc( int size in bytes ) ;` returns pointer to block: uninitialized. aligned on 16 byte

`void *calloc( int size of word, int number of words ) ;` same as malloc except area is initialized.

`void *realloc( void *ptr, tin new size in bytes ) ;` changes the size of an existing block


`void *sbrk( int increment the heap size by ) ;`

`void free( void *ptr ) ;` frees unused space. \*ptr must equal a value obtained from malloc/calloc.

## Arbitrary sequences

of malloc, free


p1 = malloc( 4 );




p2 = malloc( 5 );




p3 = malloc( 8 );



free( p3 )



p4 = malloc( 2 )



## Characteristics

1. arbitrary sequences
2. immediate response
3. use the heap
4. alignment
5. only manipulate free area

throughput: speed of the allocator

utilization: estimate of the amount of wasted space as a proportion of the total heap size.

largest total allocated space divided by the heap size.

Fragmentation. If the total free space is  $n$  and a request which is smaller than  $n$  cannot be satisfied, the heap is fragmented (external). If the request size needs to be rounded up, then there is wasted space (internal).

Tuesday, December 16, 2014, 8-11AM LAKRETZ 110: open condensed notes, calculators OK.

## Data representation

- Binary integers (base 10 to binary, etc)
- hexadecimal
- convert from one size to another (zero/sign extension)
- two's complement
- floating point
- integer, floating point addition multiplication, overflow
- boolean ops bitwise vs logical, precedence

## Assembly lang

- operand spec.
- instruction coding: dst, src
- stack frame
- structs, arrays, unions

## Optimization

- calls out of loops
- local vs call by name
- loop unroll
- blocking

## Memory

- Disks

- Spatial, temporal locality

- Cache

  - how works

  - direct mapped

  - associative ( $E > 1$ )

  - write issues

## Exceptional control flow

- exceptions

- concurrent flows

- fork(), wait\_pid, scheduling

- non local jumps

## Virtual Memory

### VM caching

- hits, misses
- page tables
- page allocation
- address translation

### dynamic memory allocation

- allocator performance
- allocation
  - placement (fit algorithms)
  - splitting
- freeing
  - coalescing
- alternative free lists



## Concurrent programming

- processes v threads

  - memory models

  - scheduling

    - races/deadlocks

- reaping/detaching

- synchronization

  - semaphores

- parallelization

- i/o multiplexing

- process/thread scheduling


## Implementation issues

- . how to keep track of free blocks
- . which free block to use
- . how to split into free and used
- . how to coalesce

### Implicit free list

Beginning of each block contains a header with the block size and whether it is free or allocated. Then we can jump through the list using the block size.

$S_0$  = size of block 0,  $F_0$  = 0 if free, 1 if allocated



Use terminator block  $S = 0$ , allocated as end of list

implicit free list requires  $n = \# \text{ free} + \# \text{ allocated}$  steps to search the list

## Implementation issues

Fit strategies:	first fit	large free blocks at the end many small blocks at the beginning
	next fit	search starting from last block split
	best fit	exhaustive search of free list

To split or not. Give the whole block. Could decide based on how bad the fit is.

Do not split: gives poor utilization but better throughput.

## Implementation issues

When blocks are freed, the allocated/free bit is set to free. It might be that there is a free block next to it. Can merge into a larger free block.

Two options:


Immediate coalescing: check whenever a block is freed to see if coalescing is possible. Adds time to the free process.

Deferred coalescing: wait until a request cannot be satisfied.


## Implementation issues

Coalescing cases when freeing the middle block:


Case 1: no adjacent free block




Case 2: previous block is adjacent



Case 3: prior block is adjacent



Case 4: combination of 2 & 3



## Heap simulation: implicit free list

Can speed up coalescing by duplicating header at end of block

Let's make this more concrete by building up our own simulator:

halloc.c:

```
#define size(i) ((struct heap_hdr *) &heap[i]) -> SZE  
#define allt(i) ((struct heap_hdr *) &heap[i]) -> ALLT
```

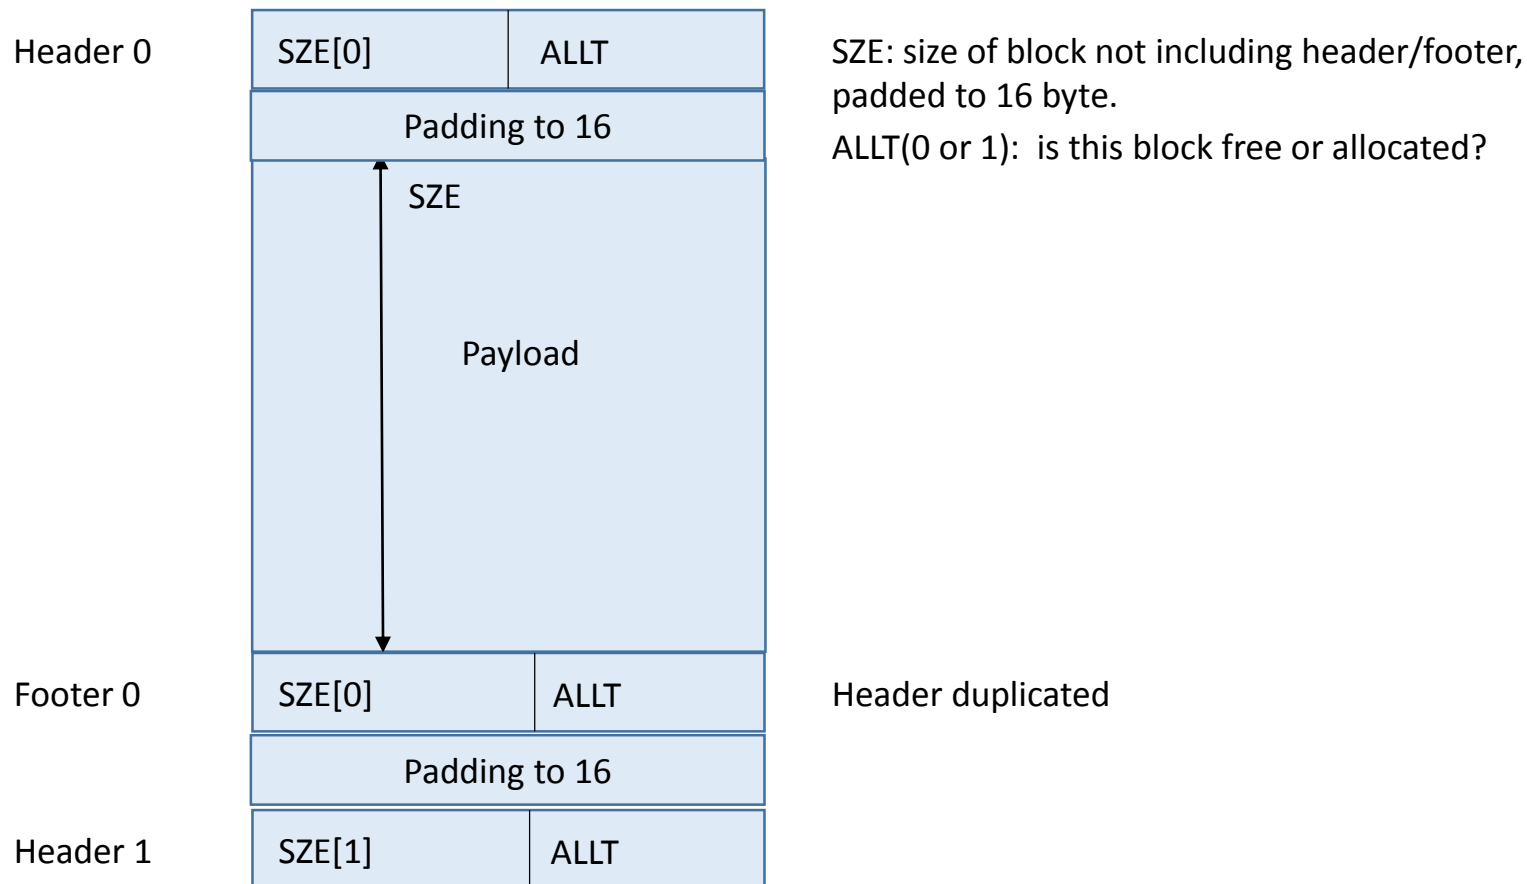
```
#define HEAPSIZE 32768
```

```
struct heap_hdr  
{  
    short SZE ; // size of block (including headers)  
    char ALLT ; // is this block allocated  
};
```

```
char *heap ;
```

# Dynamic Memory Allocation

heap layout



.  
. .  
.

```
init_heap();
dump_heap( "init heap" );
```

```
for( i=0; i<14; i++ )
    blks[j++] = lalloc( rand()%500 );
dump_heap( "first allocs" );
```

```
lfree( &blks[ 9] );
lfree( &blks[10] );
dump_heap( "coalesce with lower (free 9 then 10)" );
```

```
lfree( &blks[ 6] );
lfree( &blks[ 5] );
dump_heap( "coalesce with upper (free 6 then 5)" );
```

```
lfree( &blks[1] );
lfree( &blks[3] );
lfree( &blks[2] );
dump_heap( "coalesce with both (free 1 then 3 then 2)" );
```

```
blks[5] = lalloc( 80000 );
dump_heap( "blow the top off (allocate 80,000)" );
```

```
for( i=0; i<14; i++ )
    if (blks[i] != 0 )
    {
        lfree( &blks[i] );
    }
dump_heap( "free everything" );
```

```
init heap
blk#    addr      size    typ
      0         0  32736     0
      0    32752  32736     0
```



# heap simulation

## Dynamic Memory Allocation

first allocs				coalesce with lower (free 9 then 10)			
blk#	addr	size	typ	blk#	addr	size	typ
0	0	384	1	0	0	384	1
	400	384	1		400	384	1
1	416	400	1	1	416	400	1
	832	400	1		832	400	1
2	848	288	1	2	848	288	1
	1152	288	1		1152	288	1
3	1168	416	1	3	1168	416	1
	1600	416	1		1600	416	1
4	1616	304	1	4	1616	304	1
	1936	304	1		1936	304	1
5	1952	336	1	5	1952	336	1
	2304	336	1		2304	336	1
6	2320	400	1	6	2320	400	1
	2736	400	1		2736	400	1
7	2752	496	1	7	2752	496	1
	3264	496	1		3264	496	1
8	3280	160	1	8	3280	160	1
	3456	160	1		3456	160	1
9	3472	432	1	9	3472	832	0
	3920	432	1		4320	832	0
10	3936	368	1				
	4320	368	1				
11	4336	32	1	11	4336	32	1
	4384	32	1		4384	32	1
12	4400	192	1	12	4400	192	1
	4608	192	1		4608	192	1
13	4624	64	1	13	4624	64	1
	4704	64	1		4704	64	1
14	4720	28016	0	14	4720	28016	0
	32752	28016	0		32752	28016	0

coalesce with lower (free 9 then 10)

coalesce with upper (free 6 then 5)

blk#	addr	size	typ
0	0	384	1
	400	384	1
1	416	400	1
	832	400	1
2	848	288	1
	1152	288	1
3	1168	416	1
	1600	416	1
4	1616	304	1
	1936	304	1
5	1952	336	1
	2304	336	1
6	2320	400	1
	2736	400	1
7	2752	496	1
	3264	496	1
8	3280	160	1
	3456	160	1
9	3472	832	0
	4320	832	0
11	4336	32	1
	4384	32	1
12	4400	192	1
	4608	192	1
13	4624	64	1
	4704	64	1
14	4720	28016	0
	32752	28016	0

blk#	addr	size	typ
0	0	384	1
	400	384	1
1	416	400	1
	832	400	1
2	848	288	1
	1152	288	1
3	1168	416	1
	1600	416	1
4	1616	304	1
	1936	304	1
5	1952	768	0
	2736	768	0
7	2752	496	1
	3264	496	1
8	3280	160	1
	3456	160	1
9	3472	832	0
	4320	832	0
11	4336	32	1
	4384	32	1
12	4400	192	1
	4608	192	1
13	4624	64	1
	4704	64	1
14	4720	28016	0
	32752	28016	0

coalesce with upper (free 6 then 5)

blk#	addr	size	typ
0	0	384	1
	400	384	1
1	416	400	1
	832	400	1
2	848	288	1
	1152	288	1
3	1168	416	1
	1600	416	1
4	1616	304	1
	1936	304	1
5	1952	768	0
	2736	768	0
7	2752	496	1
	3264	496	1
8	3280	160	1
	3456	160	1
9	3472	832	0
	4320	832	0
11	4336	32	1
	4384	32	1
12	4400	192	1
	4608	192	1
13	4624	64	1
	4704	64	1
14	4720	28016	0
	32752	28016	0

coalesce with both (free 1 then 3 then 2)

blk#	addr	size	typ
0	0	384	1
	400	384	1
1	416	1168	0
	1600	1168	0
4	1616	304	1
	1936	304	1
5	1952	768	0
	2736	768	0
7	2752	496	1
	3264	496	1
8	3280	160	1
	3456	160	1
9	3472	832	0
	4320	832	0
11	4336	32	1
	4384	32	1
12	4400	192	1
	4608	192	1
13	4624	64	1
	4704	64	1
14	4720	28016	0
	32752	28016	0

coalesce with both (free 1 then 3 then 2)

blow the top off (allocate 80,000)

blk#	addr	size	typ
0	0	384	1
	400	384	1
1	416	1168	0
	1600	1168	0
4	1616	304	1
	1936	304	1
5	1952	768	0
	2736	768	0
7	2752	496	1
	3264	496	1
8	3280	160	1
	3456	160	1
9	3472	832	0
	4320	832	0
11	4336	32	1
	4384	32	1
12	4400	192	1
	4608	192	1
13	4624	64	1
	4704	64	1
14	4720	28016	0
	32752	28016	0

blk#	addr	size	typ
0	0	384	1
	400	384	1
1	416	1168	0
	1600	1168	0
4	1616	304	1
	1936	304	1
5	1952	768	0
	2736	768	0
7	2752	496	1
	3264	496	1
8	3280	160	1
	3456	160	1
9	3472	832	0
	4320	832	0
11	4336	32	1
	4384	32	1
12	4400	192	1
	4608	192	1
13	4624	64	1
	4704	64	1
14	4720	80000	1
	84736	80000	1
15	84752	13520	0
	98288	13520	0

free everything

blk#	addr	size	typ
0	0	98272	0
	98288	98272	0

## Explicit Free lists

Maintain a linked list of free areas: when looking for a block, do not need to go through the allocated blocks.

```
struct HDR    // free block header/footer/linked list
{
    int  payload ; // size of block (excluding headers)
    char freeall ; // is this block allocated? 0=free/1=allocated
    int  succesr ; // successor free block
    int  previus ; // previous free block
} anchor ;
```

anchor is the head of the free list. The header of each free block contains the same structure.

## Review Lab 4

## Segregated free lists

More than one free list: segregated by size range: many variations.

Simple segregated free lists: all blocks in size range the same size:  
always allocate entire block (no splitting, coalescing)  
request new block, divide it into same size blocks

fast allocation, freeing  
fragmentation danger

Segregated fits: lists segregated by size range  
search list by size range: first fit. split block and put remnant on appropriate list  
free: coalesce and place on size range list

close to best fit search

Buddy system: lists always point to powers of 2 size blocks. always round requests up to power of 2.  
search list. if none available take larger block and split in two, repeatedly.  
when size reached, a “buddy” is left over.  
when free occurs, can coalesce with buddy.

fragmentation

## List of common memory bugs

### Passing value instead of pointer:

```
void f1( int *x ) {  
  
    x = 20 ;  
}  
  
void main() {  
    int z ;  
  
    f1( z ) ; // instead if f1( &z ) ;  
}
```

### Uninitialized values:

```
void main() {  
    int i, *a, b[5], c[5] ;  
  
    int *a = (int *)malloc( 5*sizeof(int) ) ;  
    for( i=0; i<5; i++ )  
        c[i] = a[i]+b[i] ;  
  
}
```



## List of common memory bugs

gets:

```
void f1( int *x ) {  
    char b[10] ;  
  
    gets( b ) ;  
}
```

sizeof( int ) and sizeof( int \* ) not the same:

```
void f1( int n ) ;  
  
    int **a = (int **) malloc( n*sizeof(int) ) ; // instead of sizeof( int * )  
}
```

pointer instead of value referencing:

```
void f1( int *x ) {  
  
    *x-- ; // instead of (*x)-- ;  
}
```

## List of common memory bugs

### pointer arithmetic:

```
void f1( int *x ) {  
  
    x += sizeof(int) // instead of x++ ;  
}
```

### memory pointed to disappears:

```
int *f1( int n ) ;  
    int a ;  
  
    return &a ;  
}
```

### access freed area:

```
void f1( int *x ) {  
    int *y ;  
  
    y = (int *)malloc( 20*sizeof(int) ) ;  
    free y ;  
  
    for( i=0;i<20; i++ )  
        x[i] = y[i] ;  
}
```

## List of common memory bugs

### memory leaks:

```
void f1() {  
    int *x ;  
  
    x = (int *) malloc( 100 ) ;  
    return ;  
}
```

## Items from chapter 11

```
int open_listenfd( int port ) ;
```

Returns a “listening descriptor”, connecting the caller to “port”. -1 if error

```
int accept( int listenfd, struct sockaddr *addr, int *addrlen ) ;
```

Waits for a connection on listenfd ( the port opened ).

## example with processes

```
#include "csapp.h"
void echo(int connfd);

void sigchld_handler(int sig) {
    while (waitpid(-1, 0, WNOHANG) > 0) ;
    return;
}

int main(int argc, char **argv) {
    int listenfd, connfd, port;
    socklen_t clientlen=sizeof(struct sockaddr_in);
    struct sockaddr_in clientaddr;

    if (argc != 2) {
        fprintf(stderr, "usage: %s <port>\n", argv[0]);
        exit(0); }
    port = atoi(argv[1]);

    signal(SIGCHLD, sigchld_handler);
    listenfd = open_listenfd(port);
    while (1) {
        connfd = Accept(listenfd, (SA *) &clientaddr, &clientlen);
        if (fork() == 0) {
            close(listenfd); /* Child closes its listening socket */
            echo(connfd); /* Child services client */
            close(connfd); /* Child closes connection with client */
            exit(0); /* Child exits */
        }
        Close(connfd); /* Parent closes connected socket (important!) */
    } }
```

## processes, descriptor sets

using `fork()` gives separate address space. creates safe environment but inhibits communication.

file descriptor set (remember wait set). An integer: `fd_set`, on bit for each file (port)

```
FD_ZERO( fd_set *fdset );           // zeros descriptor set
FD_CLR( int fd, fd_set *fdset );     // removes fd from descriptor set
FD_SET( int fd, fd_set *fdset );     // enters fd into descriptor set
FD_ISSET(int fd, fd_set *fdset );    // test fd on in descriptor set
```

```
int select( int n, fd_set *fdset, NULL, NULL, NULL );
```

`fdset` is the descriptor set and `n` is the number of descriptors. BLOCKS until a file descriptor is ready. Returns the number of ready descriptors.

## descriptor set example

.

.

.

```

FD_ZERO(&read_set);           /* Clear read set */
FD_SET(STDIN_FILENO, &read_set); /* Add stdin to read set */
FD_SET(listenfd, &read_set);    /* Add listenfd to read set */

while (1) {
    ready_set = read_set;
    select(listenfd+1, &ready_set, NULL, NULL, NULL);
    if (FD_ISSET(STDIN_FILENO, &ready_set))
        command();             /* Read command line from stdin */
    if (FD_ISSET(listenfd, &ready_set)) {
        connfd = accept(listenfd, (SA *)&clientaddr, &clientlen);
        echo(connfd);           /* Echo client input until EOF */
        close(connfd);
    }
}

void command(void) {
    char buf[MAXLINE];
    if (!fgets(buf, MAXLINE, stdin))
        exit(0); /* EOF */
    printf("%s", buf);           /* Process the input command */
}

```

## Threads

Different from processes. No children but “peers” sharing virtual memory.

Thread context: stack, instruction counter, registers and condition code register.

Rest is shared. Same scheduling procedure as processes. But context switch excludes VM.

Main thread: peer threads. less hierarchy.

Creation is by naming a function: `typedef void *(func)(void * ) ; // pointer to a function`

```
int pthread_create( pthread_t *tid, pthread_attr_t *attr, func *f, void *arg ) ;
```

Creates a thread using function f, sets thread ID into tid, returns 0 if OK, -1 if error. Ignore attr, for now.

```
pthread_t pthread_self() ;
```

returns the TID.

```
int pthread_join( pthread_t tid, void **thread_return ) ;
```

Waits for tid to terminate and reaps the thread. can only wait for specific one.



## Example

```
#include <pthread.h>
#include <stdio.h>
void *thread(void *vargp);

int main()
{
    pthread_t tid;
    pthread_create(&tid, NULL, thread, NULL);
    pthread_join(tid, NULL);
    exit(0);
}

void *thread(void *vargp) /* Thread routine */
{
    printf("Hello, world!\n");
    return NULL;
}
```

## Thread control

```
void pthread_exit( void *thread_return ) ;
```

waits for all peer threads (created by the caller) to terminate.

```
void pthread_cancel( pthread_t tid ) ;
```

To end a thread:

- implicit exit by returning

- explicit: call `pthread_exit`, waiting for all peer threads to terminate

- `exit()`: ends the entire process

- call `pthread_cancel` for this TID (by another peer)

```
void pthread_detach( pthread_t tid ) ;
```

“unpeers” the thread. Does not need to be reaped. Can no longer be cancelled. Can detach yourself.

```
pthread_once_t once_control = PTHREAD_ONCE_INIT ;
```

```
int pthread_once( pthread_once_t *once_control, void (*init_routine)(void)) ;
```

Call the init routine the first time and skips all subsequent times.

## Thread example vs process example

```
void echo(int connfd);
void *thread(void *vargp);

int main(int argc, char **argv) {
    int listenfd, *connfdp, port;
    socklen_t clientlen=sizeof(struct sockaddr_in);
    struct sockaddr_in clientaddr;
    pthread_t tid;

    port = atoi(argv[1]);

    listenfd = Open_listenfd(port);
    while (1) {
        connfdp = malloc(sizeof(int));    // create an instance to prevent races
        *connfdp = accept(listenfd, (SA *) &clientaddr, &clientlen);
        pthread_create(&tid, NULL, thread, connfdp);
    }
    /* Thread routine */
    void *thread(void *vargp) {
        int connfd = *((int *)vargp);
        pthread_detach(pthread_self());
        free(vargp);
        echo(connfd);
        close(connfd);
        return NULL;
    }
}
```

## Shared variable issues in threads

```
#include <pthread.h>
#include <stdio.h>
#define N 2
void *thread(void *vargp);
char **ptr; /* Global variable */

int main()
{
    long int i;
    pthread_t tid, pret;
    char *msgs[N] = {
        "Hello from foo",
        "Hello from bar"
    };

    ptr = msgs;
    for (i = 0; i < N; i++)
        pthread_create(&tid, NULL, thread, (void *)i);

    pthread_exit(NULL);
}

void *thread(void *vargp)
{
    long int myid = (long int)vargp;
    static int cnt = 0;
    printf("[%d]: %s (cnt=%d)\n", myid, ptr[myid], ++cnt);
    return NULL;
}
```

C storage classes:

Global variables – one instance

Local automatic variables (stack variables)

Local static variables

Shared variables have only one instance

ptr is a global variable

msgs is local automatic but shared by aliasing

myid is local automatic

cnt is local static

## Counting threads

```
#include <pthread.h>
void *thread(void *vargp); /* Thread routine prototype */

volatile int cnt = 0; /* Counter */

int main(int argc, char **argv)
{
    int niters = 1000000 ;
    pthread_t tid1, tid2;

    /* Create threads and wait for them to finish */
    pthread_create(&tid1, NULL, thread, &niters);
    pthread_create(&tid2, NULL, thread, &niters);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);

    /* Thread routine */
    void *thread(void *vargp)
    {
        int i, niters = *((int *)vargp);

        for (i = 0; i < niters; i++)
            cnt++;
    }
}
```

## Shared variable issues in threads

show threads1.c

				Thread 1	Thread 2
the loop:					
40069f <+11>:	mov	0x2005ab(%rip),%edx	# 0x600c50 <cnt> //	1	4
4006a5 <+17>:	add	\$0x1,%edx	//	2	5
4006a8 <+20>:	mov	%edx,0x2005a2(%rip)	# 0x600c50 <cnt> //	3	6
4006ae <+26>:	add	\$0x1,%eax			
4006b1 <+29>:	cmp	%ecx,%eax			
4006b3 <+31>:	jne	0x40069f <thread+11>			

For thread 1, we have operations 1, 2 and 3, thread 2 , operations 4, 5 and 6

123123	142312	412312	451231	Label %edx with thread #, %edx1 = thread 1's %edx %edx2 = thread 2's %edx	
123124	142315	412315	451236 *		
123142	142351	412351	451263 *		
123145	142356 *	412356 *	451264		
123412	142531	412531	451623 *	cnt = 0	cnt = 0
123415	142536 *	412536 *	451624		
123451	142563 *	412563 *	451642	123456 means:	124356 means:
123456 *	142564	412564	451645	mov cnt,%edx1	mov cnt,%edx1
124312	145231	415231	456123 *	add \$0x1,%edx1	add \$0x1,%edx1
124315	145236 *	415236 *	456124	mov %edx1,cnt	mov cnt,%edx2
124351	145263 *	415263 *	456142	mov cnt,%edx2	mov %edx1,cnt
124356 *	145264	415264	456145	add \$0x1,%edx2	add \$0x1,%edx2
124531	145623 *	415623 *	456412	mov %edx2,cnt	mov %edx2,cnt
124536 *	145624	415624	456415	cnt = 2	cnt = 1
124563 *	145642	415642	456451		
124564	145645	415645	456456		

## Semaphores

Need to “reserve” shared variables to protect.

```
int sem_init( sem_t *x, 0, unsigned int value) ;           // initializes to value (really just an int)

int sem_wait( sem_t x ) ;                                   // reserve resource (decrement x)
                                                            // if 0, wait until ready

int sem_post( sem_t x ) ;                                   // release resource (increment x)
                                                            // if > 0, start first waiting thread
```

While x is shared, incrementing, decrementing, testing done in uninterruptible mode.

Semaphore an integer.

show threads4.c

## Counting Semaphores

can also initialize a semaphore,  $n$ , to something  $> 1$ .

Then  $n$  is a counter. waiting decrements (if not zero, continue, if zero, wait), posting increments. Also, posting a semaphore multiple times

Producer-Consumer: have a display case with product. if not full, a producer can place a product on the shelf, if full, must wait for a space. a consumer can take a product off the shelf, if empty must wait for a product. Need to synchronize because the shelf is a shared object.

show semaphore1.c

Readers:Writers: When reading an item, no changes are made so it is OK to have multiple readers. But when writing, must take exclusive control. Readers over writers: if anyone is reading, the writer must wait, even if more readers come along. Writers over readers, once a writer comes, no new readers can come until the writer is finished.

show semaphore2.c



## Parallelism with threads

Concurrency vs. parallelism revisited

one core vs multiple cores

Create threads to take advantage of cores. But must look out for shared variables.

show cores.c

What is the difference w.r.t process vs threads for multiple cores: context change.

show coref.c


What about competing threads?

show corew.c

show cores.qpw

## Scheduling?

Single core: one queue. What about renegade processes: `while(1) {}` ?? Time slicing



Indifferent whether process or thread. Context switch has more overhead for processes

CPU affinity

## Other Issues

“Thread Safety”: guarantees that a function always gives the correct answer in a concurrent, threaded environment.

- Class 1: does not protect read/write shared variables
- 2: keeps state in non-automatic variables
- 3: return pointers to static variables
- 4: functions which call unsafe functions

Re-entrant functions: shared library routines cannot reference shared read/write data

Races: having the correct value of a shared value depends on who gets there first.

Deadlocks (deadly embrace):

Thread A holds resource x, B holds y, A needs y but B needs x in order to release y.

Thread A:

```
pthread_wait( &x );  
pthread_wait( &y );
```

Thread B:

```
pthread_wait( &y );  
pthread_wait( &x );  
pthread_post( &y );
```

Rule of thumb: all threads must wait and post multiple resources in the same order.

Here: A waits x then y, B waits y then x.

## Threads, cores and semaphores recap

Counting thread adds 1 to a global variable:

```
/* Thread routine */
void *thread(void *vargp)
{
    int i, niters = *((int *)vargp);

    for (i = 0; i < niters; i++)
        cnt++;                // cnt a global variable

    return NULL;
}
```

Run it with random size niter, increasing number of threads

metric:  $\text{cnt}/(\text{niter} \times \text{number of threads})$  is  $\geq 1/(\text{number of threads})$ ,  $\leq 1.000$

metric =  $1/(\text{number of threads})$  means that  $\text{cnt} = \text{niter}$  (addition wrong every time)


metric = 1 means that  $\text{cnt} = \text{niter} \times (\text{number of threads})$  (addition correct every time)

## Threads conclusions

### Conclusions:

- 1) For large niter, the metric is  $1/(\text{number of threads})$ , up to 5.

		Metric
1000000	2 x niter=2000000 cnt=1006491	0.503
1000000	2 x niter=2000000 cnt=1000073	0.500
1000000	2 x niter=2000000 cnt=1003433	0.502
1000000	3 x niter=3000000 cnt=1021719	0.341
1000000	3 x niter=3000000 cnt=1001199	0.334
1000000	3 x niter=3000000 cnt=1006640	0.336
1000000	4 x niter=4000000 cnt=1013463	0.253
1000000	4 x niter=4000000 cnt=1013291	0.253
1000000	4 x niter=4000000 cnt=1092506	0.273
1000000	5 x niter=5000000 cnt=1057547	0.212
1000000	5 x niter=5000000 cnt=1000114	0.200
1000000	5 x niter=5000000 cnt=1022921	0.205
1000000	6 x niter=6000000 cnt=1730981	0.288
1000000	6 x niter=6000000 cnt=1776257	0.296
1000000	6 x niter=6000000 cnt=1588655	0.265



decreasing niter ----->

green = 2 threads  
red = 3 threads  
brown = 4 threads  
blue = 6 threads

- 2) For decreasing niter, metric goes to 1.

## Threads conclusions

Renegade thread:

```
/* Thread routine */
```

```
void *while1(void *vargp)
```

```
{
```

```
while( 1 )
```

```
{
```

```
}
```

```
return NULL;
```

```
}
```

Starting renegade threads in addition to the counting threads has little effect

1000000 4 x niter=4000000 cnt=1524595	0.381
1000000 4 x niter=4000000 cnt=1298805	0.325
1000000 4 x niter=4000000 cnt=1279227	0.320
1000000 4 x niter=4000000 cnt=1266988	0.317
1000000 5 x niter=5000000 cnt=1073560	0.215
1000000 5 x niter=5000000 cnt=1364531	0.273
1000000 5 x niter=5000000 cnt=1142935	0.229
1000000 5 x niter=5000000 cnt=1392169	0.278
1000000 5 x niter=5000000 cnt=1296054	0.259
1000000 6 x niter=6000000 cnt=1724155	0.287
1000000 6 x niter=6000000 cnt=1444135	0.241
1000000 6 x niter=6000000 cnt=1311778	0.219
1000000 6 x niter=6000000 cnt=1241825	0.207
1000000 6 x niter=6000000 cnt=1297965	0.216

## Threads/parallelism recap

Parallel counting: create threads to add up 1-n (n very large) by dividing into sectors and have each thread add up one sector:

```
void *sum(void *vargp)
{
    int i = *((int *)vargp) ;
    long long int j,k,l,a = 0 ;

    k = (i-1)*sector_size ;
    l = i *sector_size-1 ;
    for( j=k; j<l; j++ )
        a = a+j ;

    tot[i] = a ;
}
```

If  $T$  = the time to add up 1-n, then  $T/n$  should be the time to add the numbers with  $n$  threads: up to the number of cores. But we get some surprising results

Threads	Sector Size	T	T/T <sub>0</sub>	1/Threads
1	1073741824	0.807	1.000	1.000
2	536870912	0.414	0.513	0.500
3	357913941	0.272	0.337	0.333
4	268435456	0.205	0.254	0.250
5	214748364	0.171	0.212	0.200
6	178956970	0.143	0.177	0.167
7	153391689	0.125	0.155	0.143
8	134217728	0.107	0.132	0.125
9	119304647	0.152	0.188	0.111

Improvement continues after 4 (number of cores on server)

## Threads vs Processes

Can do the same test using processes instead of threads. get close to the same results.

Creating large working set in the process slows things down a little.

Threads	Sector Size	T	T/T <sub>0</sub>	1/Threads
1	1073741824	0.817	1.000	1.000
2	536870912	0.421	0.516	0.500
3	357913941	0.288	0.353	0.333
4	268435456	0.222	0.272	0.250
5	214748364	0.179	0.219	0.200
6	178956970	0.189	0.231	0.167
7	153391689	0.197	0.241	0.143
8	134217728	0.182	0.223	0.125
9	119304647	0.179	0.220	0.111


Adding 4 renegade threads has an impact but reasons are not explainable.

Threads	Sector Size	T	T/T <sub>0</sub>	1/Threads
1	1073741824	0.850	1.000	1.000
2	536870912	0.426	0.500	0.500
3	357913941	0.481	0.566	0.333
4	268435456	0.383	0.451	0.250
5	214748364	0.313	0.368	0.200
6	178956970	0.282	0.332	0.167
7	153391689	0.243	0.286	0.143
8	134217728	0.237	0.279	0.125
9	119304647	0.230	0.270	0.111



## Renegade threads

Compare with  
renegade threads



## Semaphore programs

### Producer/consumer simulation

```
sem_init( &x, 0, 1 ); // display case lock
sem_init( &n, 0, 5 ); // empty spaces
sem_init( &d, 0, 0 ); // filled spaces

void *produce(void *vargp) {
    sem_wait( &n ); // waits if full ( n = 0 )
    sem_wait( &x ); // locks case
    printf( "Produce %2d %2d %2d\n", i, n.__align, d.__align+1 );
    sem_post( &x ); // unlocks case
    sem_post( &d ); // produces
    return NULL;
}

void *consume(void *vargp) {
    sem_wait( &d ); // waits if empty ( d = 0 )
    sem_wait( &x ); // locks case
    printf( "Consume %2d %2d %2d\n", i, n.__align+1, d.__align );
    sem_post( &x ); // unlocks case
    sem_post( &n ); // consumes
    return NULL;
}
```

CCPPPCCCPPPCPCPCPP

	s	e	f
Produce	1	4	1
Consume	1	5	0
Produce	2	4	1
Produce	3	3	2
Produce	4	2	2
Produce	5	1	2
Consume	5	2	0
Consume	3	3	0
Consume	4	4	0
Consume	2	5	0
Produce	6	4	1
Consume	6	5	0
Produce	7	4	1
Consume	7	5	0
Produce	8	4	1
Consume	8	5	0
Produce	9	4	1
Consume	9	5	0
Produce	10	4	1

s = sequence

e = empty spaces

f = full spaces

## Semaphore programs

### Reader/writer simulation

```
sem_init( &x, 0, 1 ); // update readers count lock
sem_init( &n, 0, 1 ); // writer lock

void *reader(void *vargp) {
    sem_wait( &x ); // lock readers
    readers++;
    if( readers == 1 ) // wait on writer if first reader
        sem_wait( &n );
    sem_post( &x );

    sem_wait( &x ); // lock readers count
    readers--;
    if( readers == 0 ) / if last reader, let writer go
        sem_post( &n );
    sem_post( &x );
    return NULL;
}

void *writer(void *vargp) {
    {
        sem_wait( &n ); // wait for readers
        sem_post( &n ); // let readers go
        return NULL;
    }
}
```

WRRWWWRRRRWRWRWRRRRRW

	s	r	w
Writer start	1	0	1
Reader start	1	1	0
Reader end	1	0	1
Writer start	2	0	1
Reader start	2	1	0
Reader end	2	0	1
Writer start	3	0	1
Reader start	3	1	0
Reader end	3	0	1
Writer start	4	0	1
Reader start	4	1	0
Reader end	4	0	1
Reader start	5	1	0
Reader end	5	0	1
Writer start	5	0	1
Reader start	6	1	0
Reader start	7	2	0
Reader end	6	1	0
Reader end	7	0	1
Reader start	8	1	0

s = sequence

r = # readers

w = writer locked

## Reader/writer observations

Process highly influenced by amount of time reader takes.  
Introduce `sleep(1)` in reader, delays all writers.

## C Data Types

Type	bytes (X68-64)	mnemonic	description
char	1	b	Character (can also be interpreted as integer)
short	2	w	Short integer
int	4	l	Integer
long int	8	q	Long integer
long long int	8	q	Long integer
char *	8	q	Pointer
float	4	s	Single precision floating point
double	8	d	Double precision floating point
long double	16	t	Extended precision floating point

## No bit data type

b	byte (8 bits)
w	word (16 bits)
l	long (32 bits)
q	quad (64 bits)
s	single precision (32 bits)
d	double precisions (64 bits)
t	extended precision (128 bits)

## X86-64 registers

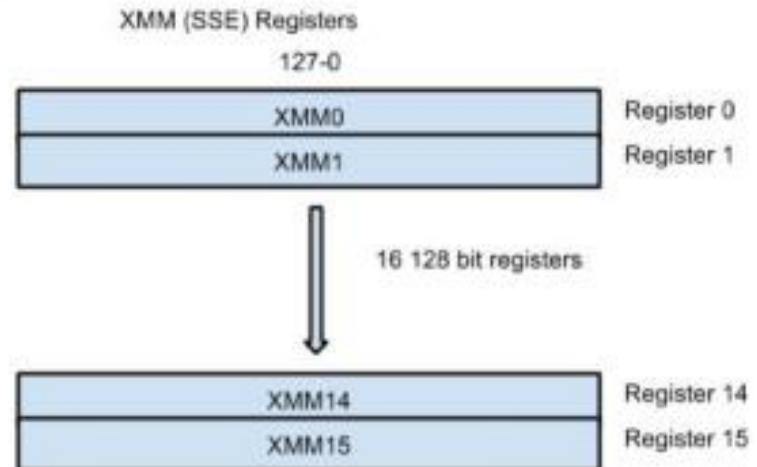
Main Registers				
63-32	31-16	15-8	7-0	
RAX	EAX	AX	AL	A Register
RBX	EBX	BX	BL	B Register
RCX	ECX	CX	CL	C Register
RDX	EDX	DX	DL	D Register

Index Registers			
RSI	ESI	SI	Source Index
RDI	EDI	DI	Destination Index
RBP	EBP	BP	Base Pointer
RSP	ESP	SP	Stack Pointer

Additional Registers	
R8	Register 8
R9	Register 9
R10	Register 10
R11	Register 11
R12	Register 12
R13	Register 13
R14	Register 14
R15	Register 15

Instruction Pointer			
RIP	EIP	IP	Instruction Pointer

Segment Pointers	
15-0	
CS	Code Segment
DS	Data Segment
ES	Extra Segment
FS	F Segment
GS	G Segment
SS	Stack Segment



Status Register	
17-0	
18 flag bits	E Flags

## How to specify an operand

% implies a register

%rax, %eax, %ax, %al  
(64) (32) (16) (8)

\$ means “immediate” or exactly that value

\$0x5: the value 5: 0x means hexadecimal

parentheses means the value stored at that memory address

(%rax)

Type	Form	Operand value	Name
Immediate	\$num	num	Immediate
Register	%rax	%rax	Register
Memory	num	(num)	Absolute
Memory	(%rax)	(%rax)	Indirect
Memory	num(%rax)	(num+%rax)	Base+displacement
Memory	(%rax,%rbx)	(%rax+%rbx)	Indexed
Memory	num(%rax,%rbx)	(num+%rax+%rbx)	Indexed
Memory	(,%rax,s)	(%rax*s)	Scaled indexed
Memory	num(,%rax,s)	(num+%rax*s)	Scaled indexed
Memory	(%rax,%rbx,s)	(%rax+%rbx*s)	Scaled indexed
Memory	num(%rax,%rbx,s)	(num+%rax+%rbx*s)	Scaled indexed

Note: s may only be 1, 2, 4 or 8

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
0	ff	fe	fd	fc	fb	fa	f9	f8	f7	f6	f5	f4	f3	f2	f1	f0
1	fe	fd	fc	fb	fa	f9	f8	f7	f6	f5	f4	f3	f2	f1	f0	ef
2	fd	fc	fb	fa	f9	f8	f7	f6	f5	f4	f3	f2	f1	f0	ef	ee
3	fc	fb	fa	f9	f8	f7	f6	f5	f4	f3	f2	f1	f0	ef	ee	ed
4	fb	fa	f9	f8	f7	f6	f5	f4	f3	f2	f1	f0	ef	ee	ed	ec
5	fa	f9	f8	f7	f6	f5	f4	f3	f2	f1	f0	ef	ee	ed	ec	eb
6	f9	f8	f7	f6	f5	f4	f3	f2	f1	f0	ef	ee	ed	ec	eb	ea
7	f8	f7	f6	f5	f4	f3	f2	f1	f0	ef	ee	ed	ec	eb	ea	e9
8	f7	f6	f5	f4	f3	f2	f1	f0	ef	ee	ed	ec	eb	ea	e9	e8
9	f6	f5	f4	f3	f2	f1	f0	ef	ee	ed	ec	eb	ea	e9	e8	e7

Contents of memory on the left,

Assume %rax contains 0x10,  
%rbx contains 0x40

Specification	Computation	Address	Value
\$0x5	<na>	<na>	5
%rax	<na>	<na>	0x10
0x5	0x05	0x05	0xfa
(%rax)	%rax	0x10	0xfe
0x04(%rax)	0x04+%rax	0x14	0xfa
(%rax,%rbx)	%rax+%rbx	0x50	0xfa
0x04(%rax,%rbx)	0x04+%rax+%rbx	0x54	0xf6
(,%rax,4)	%rax*4	0x40	0xfb
0x05 (,%rax,2)	0x05+%rax*2	0x25	0xf8
(%rax,%rbx,2)	%rax+%rbx*2	0x90	0xf6
0x05 (%rax,%rbx,2)	0x05+%rax+%rbx*2	0x95	0xf1



## Move instructions: mov source and destination

Combinations of source and destination implied. Proper operation code determined by compiler.

immediate to register	mov	immediate,register
register to register	mov	register,register
memory to register	mov	memory,register
immediate to memory	mov	immediate,memory
register to memory	mov	register,memory

Obviously cannot move to an immediate: `mov %rax,$0x40`

Cannot move memory to memory

Moving from smaller to larger: sign extension or zero extension

<code>movs</code>	source,dest	sign extension
<code>movz</code>	source,dest	zero extension

Moving from larger to smaller, take only low order.

Move	<code>movb</code>	<code>movw</code>	<code>movl</code>	<code>movq</code>
Move, sign extension	<code>movsbw</code>	<code>movsbl</code>	<code>movswl</code>	<code>movslq</code>
Move, zero extension	<code>movzbw</code>	<code>movzbl</code>	<code>movzwl</code>	<code>movxqlq</code>

## mov examples

Build up the move instruction:


mov[s|z][l1][l2]

s = sign extension

z = zero extension

l1 = mnemonic of source

l2 = mnemonic of destination



b  
w  
l  
q  
s  
d  
t

byte (8 bits)  
word (16 bits)  
long (32 bits)  
quad (64 bits)  
single precision (32 bits)  
double precisions (64 bits)  
extended precision (128 bits)

s|z omitted if source and destination lengths  
are the same

movslq mov[s][l][q] = move sign extension, long to quad

mov %rsp,%rbp

register to register

64 bits

movl \$0x12345678,-0x4(%rbp)

immediate to memory

32 bits

mov \$0x0,%eax

immediate to register

%eax implies 32 bits

mov %rdi,-0x28(%rbp)

register to memory

%rdi implies 64 bits

mov -0x28(%rbp),%rax

memory to register

%rdi implies 64 bits

mov (%rax),%edx

memory to register

%edx implies 32 bits

mov %eax,-0x14(%rbp)

register to memory

%eax implies 32 bits

movzwl -0x14(%rbp),%edx

memory to register

word to long, zero extension

movswl -0x12(%rbp),%eax

memory to register

word to long, sign extension

## stack instructions

push	operand	enlarges stack, places operand on stack
	short for	leal -0x08(%rsp),%rsp mov operand,(%rsp)
pop	operand	places top of stack in operand, decreases stack
	short for	mov (%rsp),operand leal 0x08(%rsp),%rsp
call	xyz	enlarges stack, places return address on stack, jumps to xyz
	short for	leal -0x08(%rsp),%rsp mov %rip,(%rsp) addq \$0x05,(%rsp) mov address of xyz, %rip
leave		copies base pointer to stack pointer, restores base pointer decrease stack
	short for	mov %rbp,%rsp mov (%rbp),%rbp leal 0x08(%rsp),%rsp
ret		place return address in rip, decrease stack
	short for	mov (%rsp),%rip leal 0x08(%rsp),%rsp

## call/return stack manipulation

```
#include <stdio.h>
void to_sub( int *i )
{
    int j = 0x78563412 ; char y[10] ;

    int k = *i+j ;
    other_sub( &j ) ;
    printf( "%p\n", &k ) ;
}

void other_sub( int *j )
{
    int i ;
    i = *j ;
}

int main()
{
    int i; char x[30] ;

    i = 0x12345678 ;
    to_sub( &i ) ;
    return 0 ;
}
```

```
main      4004c4 <+0>:  push %rbp
          4004c5 <+1>:  mov  %rsp,%rbp
          4004c8 <+4>:  sub  $0x30,%rsp
          4004cc <+8>:  movl $0x12345678,-0x4(%rbp)
          4004d3 <+15>: lea  -0x4(%rbp),%rax
          4004d7 <+19>:  mov  %rax,%rdi
          4004da <+22>:  callq 0x4004e6 <to_sub>
          4004df <+27>:  mov  $0x0,%eax
          4004e4 <+32>:  leaveq
          4004e5 <+33>:  retq

to_sub    4004e6 <+0>:  push %rbp
          4004e7 <+1>:  mov  %rsp,%rbp
          4004ea <+4>:  sub  $0x30,%rsp
          4004ee <+8>:  mov  %rdi,-0x28(%rbp)
          4004f2 <+12>:  movl $0x78563412,-0x4(%rbp)
          4004f9 <+19>:  mov  -0x28(%rbp),%rax
          4004fd <+23>:  mov  (%rax),%edx
          4004ff <+25>:  mov  -0x4(%rbp),%eax
          400502 <+28>:  lea  (%rdx,%rax,1),%eax
          400505 <+31>:  mov  %eax,-0x14(%rbp)
          400508 <+34>:  lea  -0x4(%rbp),%rax
          40050c <+38>:  mov  %rax,%rdi
          40050f <+41>:  callq 0x40052f <other_sub>
          400514 <+46>:  mov  $0x400648,%eax
          400519 <+51>:  lea  -0x14(%rbp),%rdx
          40051d <+55>:  mov  %rdx,%rsi
          400520 <+58>:  mov  %rax,%rdi
          400523 <+61>:  mov  $0x0,%eax
          400528 <+66>:  callq 0x4003b8 <printf@plt>
          40052d <+71>:  leaveq
          40052e <+72>:  retq
```

### Before Call

Breakpoint 1, 4004da in main ()

(gdb) x/32w 0x7fffffffef3e0

	0x7fffffffef3e0:	0x00000000	0x00000000	0xd08908ed	0x00000032
	0x7fffffffef3f0:	0x00000000	0x00000000	0x00400560	0x00000000
	0x7fffffffef400:	0x00000000	0x00000000	0x004003a3	0x00000000
	0x7fffffffef410:	0xffffe548	0x00007fff	0x004005a5	0x00000000
%rsp →	0x7fffffffef420:	0xd080fba0	0x00000032	0x00400560	0x00000000
	0x7fffffffef430:	0x00000000	0x00000000	0x004003e0	0x00000000
	0x7fffffffef440:	0xffffe530	0x00007fff	0x00000000	0x12345678
%rbp →	0x7fffffffef450:	0x00000000	0x00000000	0xd081ed1d	0x00000032
	rip	rbp	rsp		
	4da	e450	e420		

i

### After Call

```

                                4004e6 in to_sub ()
(gdb) x/32w 0x7fffffffef3e0
0x7fffffffef3e0: 0x00000000      0x00000000      0xd08908ed      0x00000032
0x7fffffffef3f0: 0x00000000      0x00000000      0x00400560      0x00000000
0x7fffffffef400: 0x00000000      0x00000000      0x004003a3      0x00000000
0x7fffffffef410: 0xffffe548      0x00007fff      0x004004df      0x00000000
0x7fffffffef420: 0xd080fba0      0x00000032      0x00400560      0x00000000
0x7fffffffef430: 0x00000000      0x00000000      0x004003e0      0x00000000
0x7fffffffef440: 0xffffe530      0x00007fff      0x00000000      0x12345678
%rbp → 0x7fffffffef450: 0x00000000      0x00000000      0xd081ed1d      0x00000032

rip      rbp      rsp
4e6      e450     e418

```

Return address, %rsp

executed

```

leal -0x08(%rsp),%rsp
mov  %rip,%rsp
addq $0x05,%rsp
mov  address of xyz, %rip

```

### After push %rbp

```

                                4004e7 in to_sub ()
(gdb) x/32w 0x7fffffff3e0
0x7fffffff3e0: 0x00000000      0x00000000      0xd08908ed      0x00000032
0x7fffffff3f0: 0x00000000      0x00000000      0x00400560      0x00000000
0x7fffffff400: 0x00000000      0x00000000      0x004003a3      0x00000000
%rsp → 0x7fffffff410: 0xffffe450      0x00007fff      0x004004df      0x00000000
0x7fffffff420: 0xd080fba0      0x00000032      0x00400560      0x00000000
0x7fffffff430: 0x00000000      0x00000000      0x004003e0      0x00000000
0x7fffffff440: 0xffffe530      0x00007fff      0x00000000      0x12345678
%rbp → 0x7fffffff450: 0x00000000      0x00000000      0xd081ed1d      0x00000032

rip      rbp      rsp
4e7      e450     e410      (%rbp)

```

executed

```

leal -0x08(%rsp),%rsp
mov %rbp,(%rsp)

```

**Mov %rsp,%rbp**

```

                                4004ea in to_sub ()
(gdb) x/32w 0x7fffffffef3e0
0x7fffffffef3e0: 0x00000000      0x00000000      0xd08908ed      0x00000032
0x7fffffffef3f0: 0x00000000      0x00000000      0x00400560      0x00000000
%rsp  0x7fffffffef400: 0x00000000      0x00000000      0x004003a3      0x00000000
      → 0x7fffffffef410: 0xffffe450      0x00007fff      0x004004df      0x00000000
%rbp  0x7fffffffef420: 0xd080fba0      0x00000032      0x00400560      0x00000000
0x7fffffffef430: 0x00000000      0x00000000      0x004003e0      0x00000000
0x7fffffffef440: 0xffffe530      0x00007fff      0x00000000      0x12345678
0x7fffffffef450: 0x00000000      0x00000000      0xd081ed1d      0x00000032

rip      rbp      rsp
4ea      e410     e410
executed

mov %rsp,%rbp

```



### Sub 0x30,%rsp

```

                                4004ee in to_sub ()
                                (gdb) x/32w 0x7fffffffef3e0
%rsp → 0x7fffffffef3e0: 0x00000000      0x00000000      0xd08908ed      0x00000032
                                0x7fffffffef3f0: 0x00000000      0x00000000      0x00400560      0x00000000
                                0x7fffffffef400: 0x00000000      0x00000000      0x004003a3      0x00000000
%rbp → 0x7fffffffef410: 0xffffe450      0x00007fff      0x004004df      0x00000000
                                0x7fffffffef420: 0xd080fba0      0x00000032      0x00400560      0x00000000
                                0x7fffffffef430: 0x00000000      0x00000000      0x004003e0      0x00000000
                                0x7fffffffef440: 0xffffe530      0x00007fff      0x00000000      0x12345678
                                0x7fffffffef450: 0x00000000      0x00000000      0xd081ed1d      0x00000032

                                rip      rbp      rsp
                                4ee      e410     e3e0
                                executed

                                sub 0x30,%rsp

```

### Before leave

```

Breakpoint 2, 400528 in to_sub ()
(gdb) x/32w 0x7fffffffef3e0
%rsp → 0x7fffffffef3e0: 0x00000000    0x00000000    0xd08908ed    0x00000032
      0x7fffffffef3f0: 0x00000000    0x00000000    0x00400560    0x8a8a8a8a
      0x7fffffffef400: 0x00000000    0x00000000    0x004003a3    0x78563412
%rbp → 0x7fffffffef410: 0xffffe450    0x00007fff    0x004004df    0x00000000
      0x7fffffffef420: 0xd080fba0    0x00000032    0x00400560    0x00000000
      0x7fffffffef430: 0x00000000    0x00000000    0x004003e0    0x00000000
      0x7fffffffef440: 0xffffe530    0x00007fff    0x00000000    0x12345678
      0x7fffffffef450: 0x00000000    0x00000000    0xd081ed1d    0x00000032

      rip      rbp      rsp
      528      e410      e3e0

```

**After leave, before ret**

```

                                40052d in to_sub ()
(gdb) x/32w 0x7fffffffef3e0
0x7fffffffef3e0: 0x00000000      0x00000000      0xd08908ed      0x00000032
0x7fffffffef3f0: 0x00000000      0x00000000      0x00400560      0x8a8a8a8a
0x7fffffffef400: 0x00000000      0x00000000      0x004003a3      0x78563412
0x7fffffffef410: 0xffffe450      0x00007fff      0x004004df      0x00000000
0x7fffffffef420: 0xd080fba0      0x00000032      0x00400560      0x00000000
0x7fffffffef430: 0x00000000      0x00000000      0x004003e0      0x00000000
0x7fffffffef440: 0xffffe530      0x00007fff      0x00000000      0x12345678
0x7fffffffef450: 0x00000000      0x00000000      0xd081ed1d      0x00000032

```

**%rbp** → 0x7fffffffef450

**%rsp** → 0x004004df

rip	rbp	rsp
52d	e450	e418

executed

```

mov %rbp,%rsp
mov (%rbp),%rbp
leal $0x08(rbp),

```

## After ret

```

                                4004df in main ()
(gdb) x/32w 0x7fffffffef3e0
0x7fffffffef3e0: 0x00000000      0x00000000      0xd08908ed      0x00000032
0x7fffffffef3f0: 0x00000000      0x00000000      0x00400560      0x8a8a8a8a
0x7fffffffef400: 0x00000000      0x00000000      0x004003a3      0x78563412
0x7fffffffef410: 0xffffe450      0x00007fff      0x004004df      0x00000000
%rsp → 0x7fffffffef420: 0xd080fba0      0x00000032      0x00400560      0x00000000
0x7fffffffef430: 0x00000000      0x00000000      0x004003e0      0x00000000
0x7fffffffef440: 0xffffe530      0x00007fff      0x00000000      0x12345678
%rbp → 0x7fffffffef450: 0x00000000      0x00000000      0xd081ed1d      0x00000032

rip      rbp      rsp
4df      e450     e420

```

executed

```

mov (%rsp),%rip
leal 0x08(%rsp),%rsp

```

## Arithmetic and Logical Operations

address	leal	memory,register	load effective address (address arithmetic,destination only a register)
unary	inc	register or memory	increment
	dec	register or memory	decrement
	neg	register or memory	negate
	not	register or memory	complement
arithmetic	add	memory or register,register	add
	sub	memory or register,register	subtract
	imul	memory or register,register	integer multiply
	idiv	memory or register	integer divide (divides RDX:RAX by source)
logical	xor	memory or register,register	bitwise exclusive or
	or	memory or register,register	bitwise or
	and	memory or register,register	bitwise and
shift	sal	immediate or one byte register,memory or register	left arithmetic shift (fill right with zeroes)
	shl	immediate or one byte register,memory or register	left logical shift (sal) (fill right with zeroes)
	sar	immediate or one byte register,memory or register	right arithmetic shift (fill left with sign bit)
	shr	immediate or one byte register,memory or register	right logical shift (fill left with zeroes)

only register %cl allowed for register operand

## Practice Problem -- leal

assume %rax contains x, %rbx contains y

	result
leal 6(%rax), %rcx	$x+6$
leal (%rax,%rbx), %rcx	$x+y$
leal (%rax,%rbx,4), %rcx	$x+4y$
leal 7(%rax,%rax,8), %rcx	$x+8x+7$
leal 0x0a(%rbx,4), %rcx	$10+4y$
leal 9(%rax,%rbx,2), %rcx	$9+x+2y$

## Practice Problem -- shifts

assume %rax contains 0x800000000000000f (8 bytes) = 1000 0000 0000 0000 ... 0000 0000 0000 1111 (64 bits)

	result
sal \$2,%rax	0x0000000000000003c
shl \$2,%rax	0x0000000000000003c
sar \$2,%rax	0xe0000000000000003
shr \$2,%rax	0x20000000000000003
sal \$63,%rax	0x80000000000000000
sar \$63,%rax	0xfffffffffffffffffff
leal \$0x1,%rax	
sal \$63,%rax	
sar \$63,%rax	0xfffffffffffffffffff
and 0xffffffe,%eax	
sal \$63,%rax	0x00000000000000000 // no matter whar is in %rax

using %eax allows only shifts from 0-31 positions (can only shift 0 – length of register -1)

## Practice Problem -- convert multiply to shifts

assume x is 10

$x * 17$	$17 = 16 + 1$	$(x \ll 4) + x$	$160 + 10 = 170$
$x * -7$	$-7 = -4 - 2 - 1$	$-(x \ll 2) - (x \ll 1) - x$	$-40 - 20 - 10 = -70$
	also	$-(x \ll 3) + x$	$-80 + 10$
$x * 60$	$60 = 32 + 16 + 8 + 4$	$(x \ll 5) + (x \ll 4) + (x \ll 3) + (x \ll 2)$	$320 + 160 + 80 + 40 = 600$
	also	$(x \ll 6) - (x \ll 2)$	$640 - 40$
$x * -112$	$-112 = -64 - 32 - 16$	$-(x \ll 6) - (x \ll 5) - (x \ll 4)$	$-640 - 320 - 160 = -1120$
	also	$-(x \ll 7) + (x \ll 4)$	$-1280 + 160$

What is  $x \ll 4 + x$ ?  $= x \ll (4 + x) = x \ll 14 = 163840$  ! shifts have lowest precedence



## Practice Problem -- arithmetic

assume	Address	Value	Register	Value
	0x100	0xFF	%eax	0x100
	0x104	0xAB	%ecx	0x1
	0x108	0x13	%edx	0x3
	0x10C	0x11		

Instruction	Destination	Value	
addl %ecx,%eax)	0x100	0x101	
subl %edx,4(%eax)	0x104	0xA8	
imull \$16,(%eax,%edx,4)	0x10c	0x110	0x11 = 0001 0001 * 16 = 0001 0001 0000
incl 8(%eax)	0x108	0x14	
decl %ecx	%ecx	0x0	
subl %edx,%eax	%eax	0xFD	

## Status register, set after arithmetic instructions.

### Status codes:

logical

CF: Carry Flag. The most recent operation generated a carry out of the most significant bit. Used to detect overflow for unsigned operations.  
(unsigned)  $t < (\text{unsigned}) a$  when doing  $t = a + b$

ZF: Zero Flag. The most recent operation yielded zero.  
 $t == 0$

arithmetic

SF: Sign Flag. The most recent operation yielded a negative value.  
 $t < 0$

OF: Overflow Flag. The most recent operation caused a two's-complement overflow—either negative or positive.  
 $(a < 0 == b < 0) \ \&\& \ (t < 0 != a < 0)$  when doing  $t = a + b$

Show condtest.c

## Compare and test.

all sizes but must be the same. In C, comparing two different sizes, they must first be made the same (larger) size and this depends on signed/unsigned.

arithmetic	cmp	memory or register (S2), memory or register (S1)	set code depending on $S_1 - S_2$
logical	test	memory or register (S2), memory or register (S1)	set code depending on $S_1 \& S_2$
	test	%eax,%eax	$S_1 \& S_1 = S_1$ ! But status is set depending on $<, =, > 0$ .
	set	dest, sets 0 or 1 into dest to store the condition	unsigned and signed suffixes to test for the usual conditions  eg jmpge, jmpne
	jmp	label	
	cmov	memory or register, register conditional move	

suffixes		
signed	unsigned	either
g >	a >	e = 0
ge >=	ae >=	ne != 0
l <	b <	s < 0
le <=	be <=	ns >= 0

## if statements in C

```

if (expression)
    <then statement>
else
    <else statement> } optional
    
```

```

if ( x<y )
    l = 99 ;
else
    l = 100 ;
    
```

assembly language:

```

        if ( !expression )
            goto false ;
        then statement
        goto done ;
false:
        else statement
done:
    
```

assembly language:

```

400486 <+18>: mov  -0xc(%rbp),%eax
400489 <+21>: cmp  -0x8(%rbp),%eax
40048c <+24>: jge  0x400497 <main+35>
40048e <+26>: movl  $0x63,-0xc(%rbp)
400495 <+33>: jmp  0x40049e <main+42>
false:
400497 <+35>: movl  $0x64,-0x4(%rbp)
done:
    
```

## do while

```
do {  
    statements  
}  
while( expression );
```

assembly language:

```
loop:  
    statements  
    if( expression )  
        goto loop ;
```

```
do  
{  
    i = i-1 ;    // statements executed at least once  
    j = j+1 ;  
}  
while ( i>j ) ;
```

assembly language:

```
loop:  
04004cc <+8>:    subl  $0x1,-0x10(%rbp)  
04004d0 <+12>:   addl  $0x1,-0xc(%rbp)  
04004d4 <+16>:   mov   -0x10(%rbp),%eax  
04004d7 <+19>:   cmp   -0xc(%rbp),%eax  
04004da <+22>:   jg    0x4004cc <main+8>  
04004dc <+24>:
```

## while

```
while( expression )  
{  
    statements  
}
```

```
while ( i>j )  
{  
    i = i-1 ;  
    j = j+1 ;  
};
```

assembly language:

```
loop:    go to test ;  
  
        {  
        statements  
        }  
  
test:    if( expression )  
         goto loop ;
```

assembly language:

```
4004cc <+8>:    jmp    0x4004d6 <main+18>  
loop:  
4004ce <+10>:   subl    $0x1,-0x10(%rbp)  
4004d2 <+14>:   addl    $0x1,-0xc(%rbp)  
test:  
4004d6 <+18>:   mov     -0x10(%rbp),%eax  
4004d9 <+21>:   cmp     -0xc(%rbp),%eax  
4004dc <+24>:   jg      0x4004ce <main+10>
```

## for loops

```
for( init; expression, update )  
{  
    statements  
}
```

assembly language:

```
init  
go to test ;  
  
loop:  
{  
    statements  
}  
update  
  
test:  
    if( expression )  
        goto loop ;
```

```
for( k=0; k<l; k++ )  
{  
    i = i-1 ;  
    j = j+1 ;  
}
```

assembly language:

```
04004cc <+8>:    movl  $0x0,-0x8(%rbp) // init  
04004d3 <+15>:   jmp   0x4004e1 <main+29>  
loop:  
04004d5 <+17>:   subl  $0x1,-0x10(%rbp)  
04004d9 <+21>:   addl  $0x1,-0xc(%rbp)  
04004dd <+25>:   addl  $0x1,-0x8(%rbp) // update  
test:  
04004e1 <+29>:   mov   -0x8(%rbp),%eax  
04004e4 <+32>:   cmp   -0x4(%rbp),%eax  
04004e7 <+35>:   jl    0x4004d5 <main+17>
```

## conditional move

var = expression ? true statement : else statement     m = i < j ? k : l ;

assembly language:

```
t = true statement
if( expression )  } implemented
    u = t ;        } using cmov
```

assembly language:

```
04004cc <+8>:  mov    -0x10(%rbp),%eax
04004cf <+11>:  cmovge -0xc(%rbp),%eax
04004d2 <+14>:  mov    %eax,-0x8(%rbp)
```



**switch statements**

```
switch ( <expression> )
{
  case <constant1> :
    statements1 // executed when <expression> = <constant1>
                // if statements end with break ; goto the end
                // otherwise, execute the next set of
  case <constant2> :
    statements2 // executed when <expression> = <constant1>
statements can be null, in which case, the           // next
case applies
.
.
default :
  default statements
}
```

## Switch statement: Explicit list of cases

```
int switch_eg(int x, int n) {
    int result = x;

    switch (n) {

        case 100:
            result *= 13;  // x * 13
            break;

        case 102:
            result += 10;  // x + 10
            /* Fall through */

        case 103:
            result += 11;  // x + 11
            break;

        case 104:
        case 106:
            result *= result; // x * x
            break;

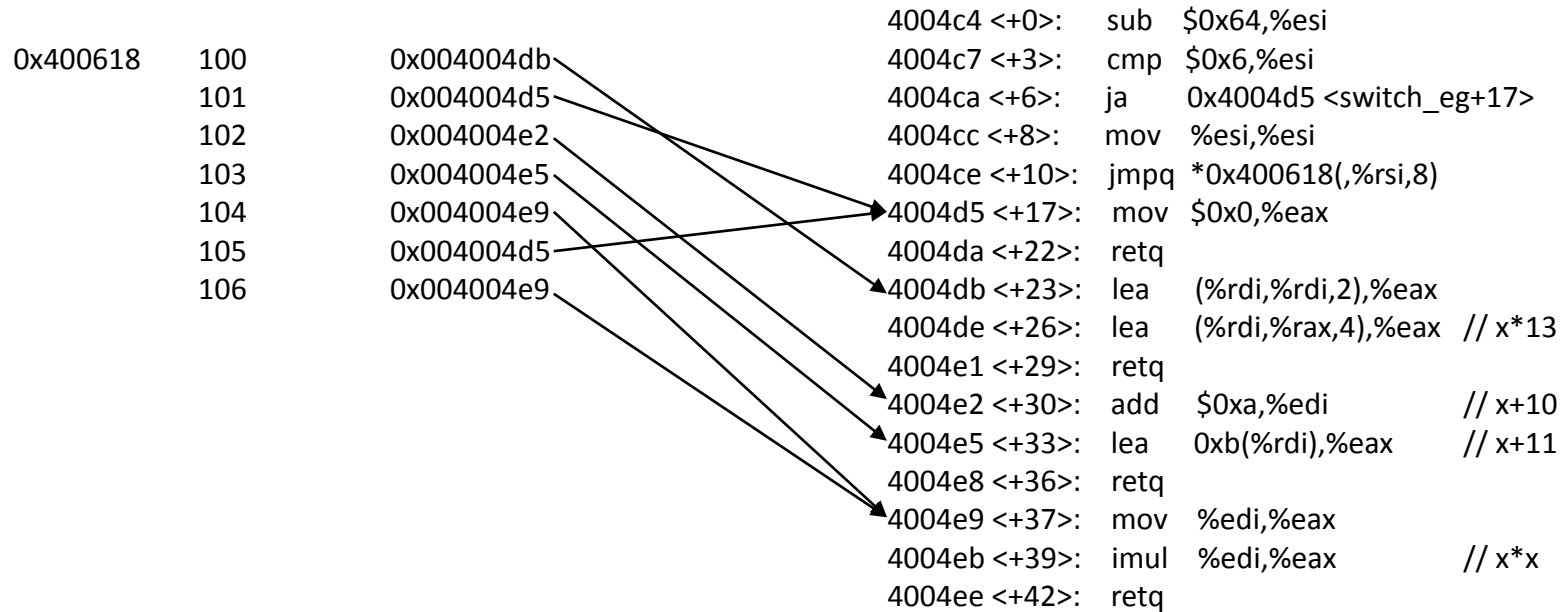
        default:
            result = 0;
    }

    return result;
}
```

Case by case outcome?

	result	x=2
<100	0	0
100	x*13	26
101	0	0
102	x+10+11	23
103	x+11	13
104	x*x	4
105	0	0
106	x*x	4
>106	0	0

## Assembly with jump table



show switchm1.c, switchm2.c

## Procedure calls

Recall stack frame operations:

```
call
push %rbp
sub 0x..,%rsp
leave
ret
```

Registers: caller-save  
callee-save

different on x86-64 because more registers

show caller.c with/without -m32

## Recursive calls

Stack frame operation ideal: each recursive call has its own context.

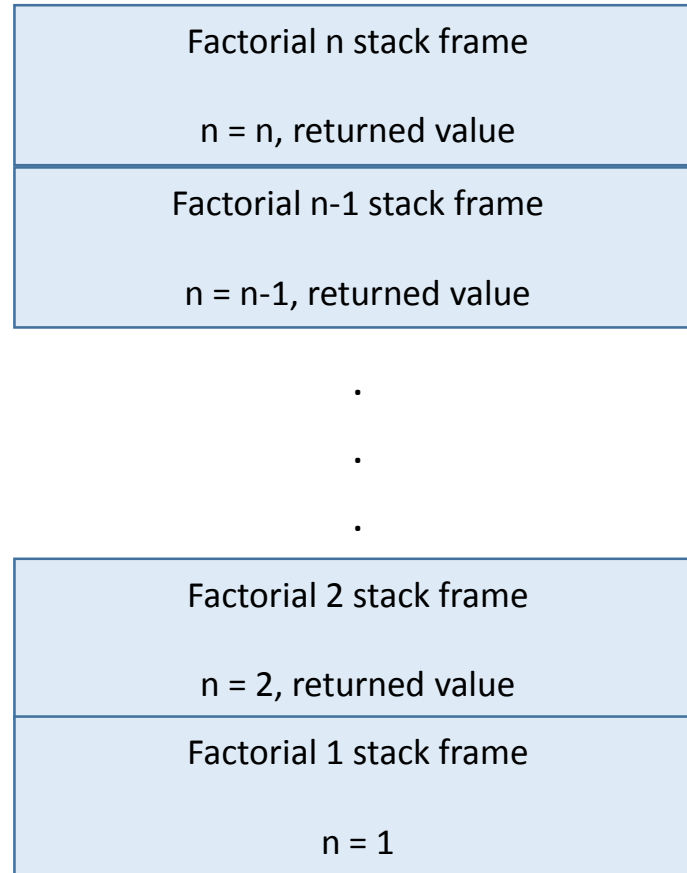
$N!$  = factorial defined as  $\prod_{i=1}^n i$  but also  $N! = N \times (N-1)!$  And  $1! = 1$

```
#include <stdio.h>
int factorial( int n ) ;
Int kkk = 0 ;

main()
{
    int m,n=6 ;
    m = n ;
    printf( "%d factorial is %d\n", m,factorial(n) ) ;
    return 0 ;
}

int factorial(int n)
{
    int result ;
    kkk = kkk+1
    if ( n<=1 )
        result = 1 ;
    else
        result = n*factorial(n-1) ;
    printf( "n= %2d kkk= %2d exit factorial result= %d\n", n,kkk,result ) ;
    return result ;
}
```

## Recursive calls



show factorial.c

## Arrays

Arrays: aggregate of same type values:

```
int x[20]; // 20 integer values of x
```

declarations functions to allocate space and create a pointer

x is a pointer to the start of the array. x[i] is the  $i_{th}$  element.

successive memory address of size int: x in memory

0	1	2	.	.	.	.	.	18	19
---	---	---	---	---	---	---	---	----	----

20 x 4 = 80 bytes total

char y[20] is 20 bytes

int \*z[20] is an array of pointers to integers: 160 bytes total

arithmetic: x+i is the address of element i. Must multiply i by the sizeof(int).

\*(x+i) is the value of element i. (dereferencing)

## Array arithmetic (int)

```
i : -0x14(%rbp)
x : -0x70(%rbp)
y : -0x10(%rbp)
z : -0x4(%rbp)
```

```
typedef int T ;
void main() {
```

```
int i ;
T x[20] ;
T *y ;
T z ;

y = x ;
z = x[0] ;
z = x[i] ;
y = &x[2] ;
y = x+i-1 ;
z = *(x+i-3) ;
z = &x[i]-x ;
}
```

```
400478 <+4>: lea    -0x70(%rbp),%rax
40047c <+8>: mov    %rax,-0x10(%rbp)        // y = x
400480 <+12>: mov    -0x70(%rbp),%eax
400483 <+15>: mov    %eax,-0x4(%rbp)        // z = x[0]
400486 <+18>: mov    -0x14(%rbp),%eax
40048b <+23>: mov    -0x70(%rbp,%rax,4),%eax
40048f <+27>: mov    %eax,-0x4(%rbp)        // z = x[i]
400492 <+30>: lea    -0x70(%rbp),%rax
400496 <+34>: add    $0x8,%rax
40049a <+38>: mov    %rax,-0x10(%rbp)        // y = &x[2]
40049e <+42>: mov    -0x14(%rbp),%eax
4004a3 <+47>: sub    $0x1,%rax
4004a7 <+51>: lea    0x0(,%rax,4),%rdx
4004af <+59>: lea    -0x70(%rbp),%rax
4004b3 <+63>: add    %rdx,%rax
4004b6 <+66>: mov    %rax,-0x10(%rbp)        // y = x+i-1
4004ba <+70>: mov    -0x14(%rbp),%eax
4004bf <+75>: sub    $0x3,%rax
4004c3 <+79>: lea    0x0(,%rax,4),%rdx
4004cb <+87>: lea    -0x70(%rbp),%rax
4004cf <+91>: add    %rdx,%rax
4004d2 <+94>: mov    (%rax),%eax
4004d4 <+96>: mov    %eax,-0x4(%rbp)        // z = *(x+i+3)
4004d7 <+99>: mov    -0x14(%rbp),%eax
4004dc <+104>: shl    $0x2,%rax
4004e0 <+108>: sar    $0x2,%rax
4004e4 <+112>: mov    %eax,-0x4(%rbp)        // z = &x[i]-x
```



## Array arithmetic (char)

i : -0x14(%rbp)  
x : -0x30(%rbp)  
y : -0x10(%rbp)  
z : -0x1(%rbp)

```
typedef int T ;
void main() {
```

```
    int i ;
    T x[20] ;
    T *y ;
    T z ;

    y = x ;
    z = x[0] ;
    z = x[i] ;
    y = &x[2] ;
    y = x+i-1 ;
    z = *(x+i-3) ;
    z = &x[i]-x ;
}
```

```
400478 <+4>: lea -0x30(%rbp),%rax
40047c <+8>: mov %rax,-0x10(%rbp)           // y = x
400480 <+12>: movzbl -0x30(%rbp),%eax
400484 <+16>: mov %al,-0x1(%rbp)             // z = x[0]
400487 <+19>: mov -0x14(%rbp),%eax
40048c <+24>: movzbl -0x30(%rbp,%rax,1),%eax
400491 <+29>: mov %al,-0x1(%rbp)             // z = x[i]
400494 <+32>: lea -0x30(%rbp),%rax
400498 <+36>: add $0x2,%rax
40049c <+40>: mov %rax,-0x10(%rbp)           // y = &x[2]
4004a0 <+44>: mov -0x14(%rbp),%eax
4004a5 <+49>: lea -0x1(%rax),%rdx
4004a9 <+53>: lea -0x30(%rbp),%rax
4004ad <+57>: add %rdx,%rax
4004b0 <+60>: mov %rax,-0x10(%rbp)           // y = x+i-1
4004b4 <+64>: mov -0x14(%rbp),%eax
4004b9 <+69>: lea -0x3(%rax),%rdx
4004bd <+73>: lea -0x30(%rbp),%rax
4004c1 <+77>: add %rdx,%rax
4004c4 <+80>: movzbl (%rax),%eax
4004c7 <+83>: mov %al,-0x1(%rbp)             // z = *(x+i-3)
4004ca <+86>: mov -0x14(%rbp),%eax
4004cd <+89>: mov %al,-0x1(%rbp)             // x = &x[i]-x
```

## Structures

Structures: a way of aggregating data into one place. Outcome of declaration is a new data type and creates a map of a contiguous area.

```
struct <new data type name>
{
    <type> <var1> ;
    <type> <var2> ;
    .
    .
} <optional name> ;
```

After declaration, <new data type name> is a data type just like char, int, etc.

## Variations of Structures

An element can be an array:

```
struct newtype1
{
    int a ;
    float b[10] ;
    int c ;
} x ;
```

4 byte int a followed by 10 float b's followed by in c.

Or another structure

```
struct newtype2
{
    int a ;
    struct inner
    {
        float b ;
        int c[10] ;
    } y ;
    int *d ;
} x ;
```

4 byte a followed by a float b, 10 int c's followed by an 8 byte int pointer

## Data alignment

Old days: mandatory

Now: instructions work but C aligns in the interest of speed.

malloc always passes back address on 16 byte boundary

## Compiled?

```
void main()
```

```
{
  struct newtype1 {
```

```
    char a ;
```

```
    float b[10] ;
```

```
    char c ;
```

```
    int d ;
```

```
  } x ;
```

```
  struct newtype2 {
```

```
    char a ;
```

```
    double b[10] ;
```

```
    char c ;
```

```
    int d ;
```

```
  } y ;
```

```
  struct newtype3 {
```

```
    int a ;
```

```
    double b[10] ;
```

```
    short c ;
```

```
    int d ;
```

```
  } z ;
```

```
x.a = 0xff ;
```

```
x.b[0] = 10 ;
```

```
x.c = 0x44 ;
```

```
x.d = 25 ;
```

```
y.a = 0xfe ;
```

```
y.b[0] = 9 ;
```

```
y.c = 0x43 ;
```

```
y.d = 24 ;
```

```
z.a = 0xfd ;
```

```
z.b[0] = 8 ;
```

```
z.c = 0x42 ;
```

```
z.d = 23 ;
```

```
}
```

```
400474 <+0>:  push %rbp
400475 <+1>:  mov  %rsp,%rbp
400478 <+4>:  sub  $0x88,%rsp
40047f <+11>: movb  $0xff,-0x40(%rbp) // x.a  offset -0x40
400483 <+15>: mov  $0x4120,%eax
400488 <+20>: mov  %eax,-0x3c(%rbp) // x.b[0] offset -0x3c  difference of 4
40048b <+23>: movb  $0x44,-0x14(%rbp) // x.c  offset -0x14  difference of 40 = 10 * 4
40048f <+27>: movl  $0x19,-0x10(%rbp) // x.d  offset -0x10  difference of 4

400496 <+34>: movb  $0xfe,-0xa0(%rbp) // y.a  offset -0xa0
40049d <+41>: movabs $0x4022,%rax
4004a7 <+51>: mov  %rax,-0x98(%rbp) // y.b[0] offset -0x98  difference of 8
4004ae <+58>: movb  $0x43,-0x48(%rbp) // y.c  offset -0x48  difference of 80 = 10 * 8
4004b2 <+62>: movl  $0x18,-0x44(%rbp) // y.d  offset -0x44  difference of 4

4004b9 <+69>: movl  $0xfd,-0x100(%rbp) // z.a  offset -0x100
4004c3 <+79>: movabs $0x4020,%rax
4004cd <+89>: mov  %rax,-0xf8(%rbp) // z.b[0] offset -0xf8  difference of 4
4004d4 <+96>: movw  $0x42,-0xa8(%rbp) // z.c  offset -0xa8  difference of 80 = 10 * 8
4004dd <+105>: movl  $0x17,-0xa4(%rbp) // z.d  offset -0xa4  difference of 4

4004e7 <+115>: leaveq
4004e8 <+116>: retq
```

## Passing as parameters

```
struct newtype1
{
    int a ;
    float b[10] ;
    int c ;
} x ;
```

If you want to pass a structure variable as a pointer, you can.

```
void to_sub1( struct newtype1 *x ) ;
```

```
then in to_sub1( struct newtype1 *x )
{
```

```
    x->a = 10.5 ;
```

or

```
    (*x).a = 10.5
```

and

```
    (*x).y.b = 11.5
}
```

but x.a no longer works.

## Compiled code

```
#include <stdio.h>
struct rect
{
    int i;
    int j;
    struct inner
    {
        int i;
        float l;
    } b;
    int c[3];
    int *p;
} x,y;

void to_sub1( struct rect x );
void to_sub2( struct rect *x );

int main()
{
    int i;

    printf( "%d\n", sizeof(x) );

    i = 100;

    x.i = 10;
    to_sub1( x );
    to_sub2( &x );
    return 0;
}
```

Dump of assembler code for function main:

```
4004c4 <+0>:      push    %rbp
4004c5 <+1>:      mov     %rsp,%rbp
4004c8 <+4>:      sub     $0x40,%rsp
4004cc <+8>:      mov     $0x4006a8,%eax
4004d1 <+13>:     mov     $0x28,%esi
4004d6 <+18>:     mov     %rax,%rdi
4004d9 <+21>:     mov     $0x0,%eax
4004de <+26>:     callq   0x4003b8 <printf@plt>
4004e3 <+31>:     movl    $0x64,-0x4(%rbp)
4004ea <+38>:     movl    $0xa,0x2004cc(%rip) # 0x6009c0 <x>
4004f4 <+48>:     mov     0x2004c5(%rip),%rax # 0x6009c0 <x>
4004fb <+55>:     mov     %rax,(%rsp)
4004ff <+59>:     mov     0x2004c2(%rip),%rax # 0x6009c8 <x+8>
400506 <+66>:     mov     %rax,0x8(%rsp)
40050b <+71>:     mov     0x2004be(%rip),%rax # 0x6009d0 <x+16>
400512 <+78>:     mov     %rax,0x10(%rsp)
400517 <+83>:     mov     0x2004ba(%rip),%rax # 0x6009d8 <x+24>
40051e <+90>:     mov     %rax,0x18(%rsp)
400523 <+95>:     mov     0x2004b6(%rip),%rax # 0x6009e0 <x+32>
40052a <+102>:    mov     %rax,0x20(%rsp)
40052f <+107>:    callq   0x400545 <to_sub1>
400534 <+112>:    mov     $0x6009c0,%edi
400539 <+117>:    callq   0x40056c <to_sub2>
40053e <+122>:    mov     $0x0,%eax
400543 <+127>:    leaveq  %rsp
400544 <+128>:    retq
```

## Usage in functions

```
void to_sub1( struct rect x )  
{  
    int i ;  
  
    i = 100 ;  
    x.i = 10 ;  
    x.b.i = 5 ;  
    i = x.j ;  
    i = x.b.i ;  
}
```

```
void to_sub2( struct rect *x )  
{  
    int i ;  
  
    i = 100 ;  
    x -> i = 10 ;  
    (*x).i = 10 ;  
  
    x -> b.i = 5 ;  
    (*x).b.i = 5 ;  
  
}
```

When passed by name, x.i and x.i.b no longer work because x is a pointer. Must use (\*x). instead.



## Unions

In unions, the offset is always 0. This means that each variable overlays or occupies the same storage as the other variables:

```
union u
{
    int i ;
    unsigned char c[4] ;
    float a ;
} examine_endian ;
```

Sound familiar? Pointers are not needed here! But it is dangerous.

Writing to `examine_endian.a` overwrites what is in `examine_endian.i`

## Example

```
#include <stdio.h>

void main()
{
    union endian
    {
        int i;
        unsigned char c[4];
        float a;
    } hw1;

    int j;

    hw1.i = 19088743 ; /* should be
0x01234567 */

    printf( "\ninteger forward= " );
    for ( j=0; j<4; j++ ) /* prints c[4] */
        printf( "%02x", hw1.c[j] );
    printf( "\n" );

    printf( "integer backward= " );
    for ( j=3; j>=0; j-- ) /* prints c[4] */
        printf( "%02x", hw1.c[j] );
    printf( "\n\n" );

    hw1.a = 10.01 ;

    printf( "float forward= " );
    for ( j=0; j<4; j++ ) /* prints c[4] */
        printf( "%02x", hw1.c[j] );
    printf( "\n" );

    printf( "float backward= " );
    for ( j=3; j>=0; j-- ) /* prints c[4] */
        printf( "%02x", hw1.c[j] );
    printf( "\n\n" );

}
```

Dump of assembler code for function main:

```
400554 <+0>:      push    %rbp
400555 <+1>:      mov     %rsp,%rbp
400558 <+4>:      sub     $0x10,%rsp
40055c <+8>:      movl    $0x1234567,-0x10(%rbp)

400603 <+175>:     mov     $0x412028f6,%eax
400608 <+180>:     mov     %eax,-0x10(%rbp)
```

result of ./a.out

```
integer forward=  67452301
integer backward= 01234567
```

```
float forward=   f6282041
float backward=  412028f6
```

## Memory corruption

buffer overflow  
pointer goes wild  
array index goes wild

gets/puts example  
pointer error example

```
/* Corrupt1.c Stack corruption with gets */
```

```
#include <stdio.h>
```

```
void echo() ;
```

```
int main()
```

```
{  
    echo() ;  
    printf( "%x\n", EOF ) ;  
}
```

```
void echo()
```

```
{  
    char inp[8] = "012345678901234567890" ;
```

```
    while ( inp != NULL )
```

```
    {  
        gets(inp) ;  
        puts(inp) ;  
    }  
}
```

## Out of bounds subscript

```

/* corrupt2.c Stack corruption with array
overflow */
#include <stdio.h>
void echo() ;
void sub2() ;

int main()
{
    echo() ;
    printf( "%x\n", EOF ) ;
}

void echo()
{
    int i[2] ;
    int j ;
    int k ;

    i[0] = 4 ;
    i[1] = 3 ;
    i[2] = 2 ;
    i[3] = 1 ;
    j = i[4] ;
    i[4] = 0 ; /* destroys return address */
    i[5] = 0 ; /* destroys previous base pointer */
    *(i+4) = -1 ;

```

```

for( k=-4; k>-20; k-- )
    i[k] = k ;

    sub2() ;
}

void sub2()
{
    int i = 5 ;
    int j ;

    j = i ;
}

```

## Out of bounds subscript

```
/* corrupt3.c  stack corruption storing outside of frame */
#include <stdio.h>
void echo() ;
void sub2() ;

int main() {
    echo() ;
    printf( "%x\n", EOF ) ;
}
void echo() {
    int i[2] ;
    int j ;
    int k ;

    for( k=-4; k>-500; k-- )
        i[k] = k ;

    sub2() ;
}
void sub2() {
    int i = 5 ;
    int j ;

    j = i ;
}
```

## Machine Takeover

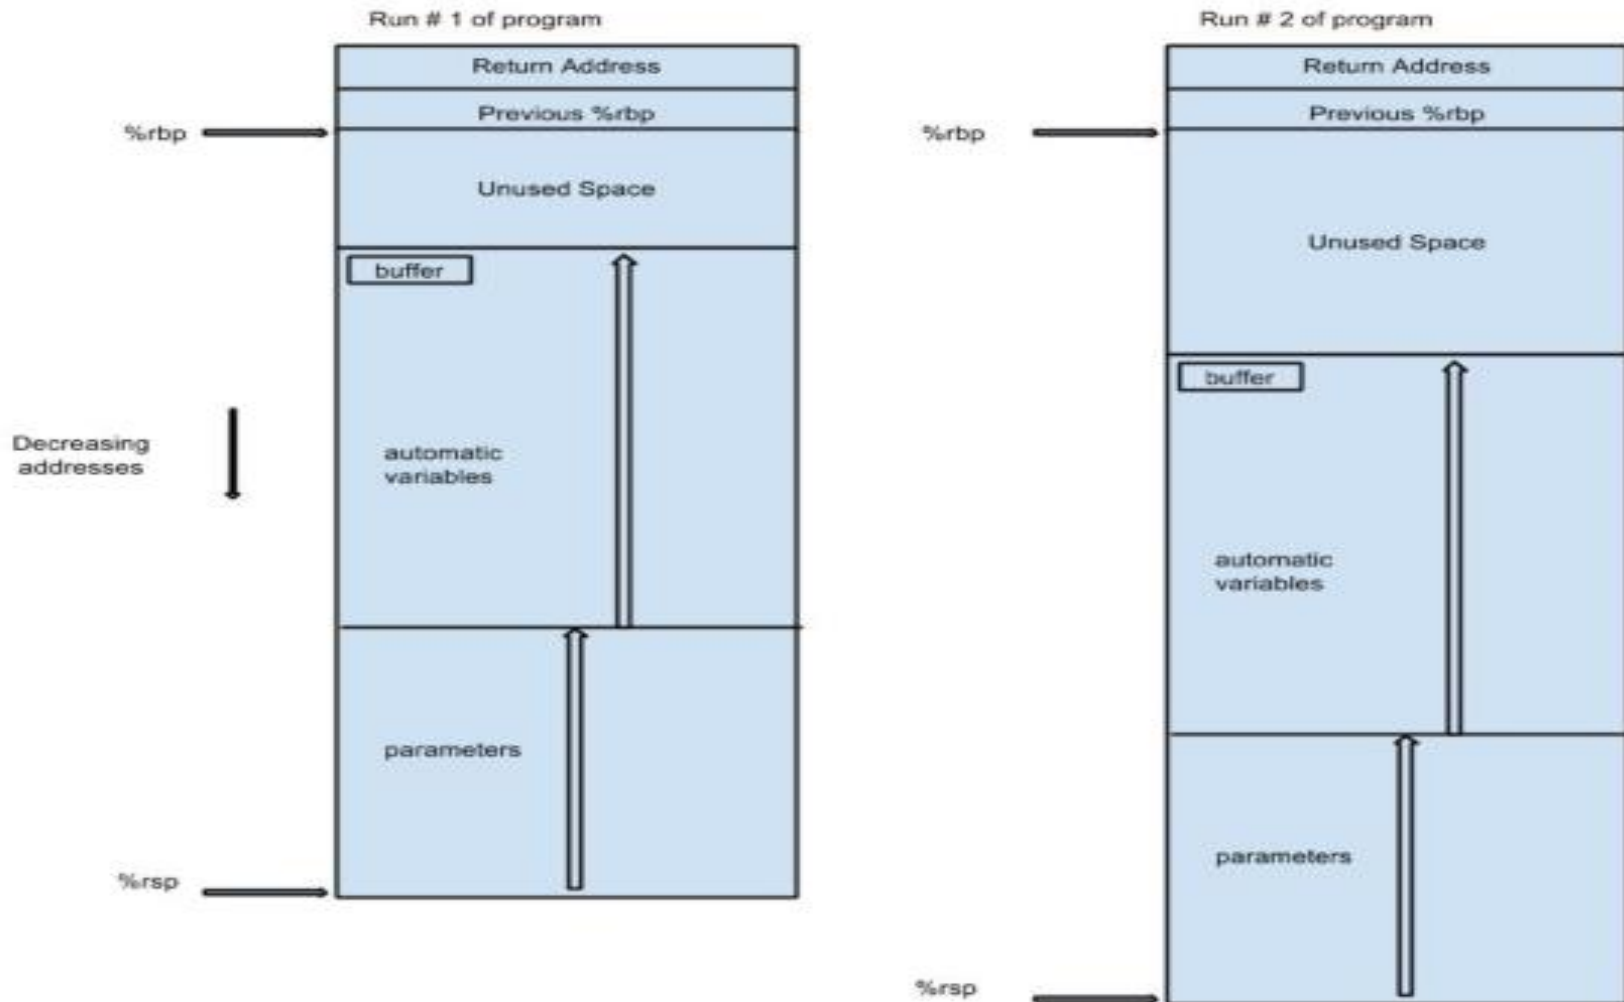
Insert code somewhere (beyond stack). Overlay the return address in the stack. When program returns, it jumps to code.

Stack randomization: after saving the return address and the previous base pointer, allocate a random amount of space in the stack. Then place the automatic variables. This way, the address of the automatic variables and the base pointer has a different offset.

Corruption detection: Store a random value somewhere in stack at the beginning of the program. Store that value in a protected area of memory. At the end of the program compare the values. If changed, raise the red flag.

Hardware which prevents pages from executing code. Memory is divided into 2K or 4K byte “pages”. Each page can be set with read/write/execute bits when in supervisory mode.

## Sample Stack Randomization



## List of common memory bugs

### Passing value instead of pointer:

```
void f1( int *x ) {  
  
    x = 20 ;  
}  
  
void main() {  
    int z ;  
  
    f1( z ) ; // instead if f1( &z ) ;  
}
```

### Uninitialized values:

```
void main() {  
    int i, *a, b[5], c[5] ;  
  
    int *a = (int *)malloc( 5*sizeof(int) ) ;  
    for( i=0; i<5; i++ )  
        c[i] = a[i]+b[i] ;  
  
}
```



## List of common memory bugs

gets:

```
void f1( int *x ) {  
    char b[10] ;  
  
    gets( b ) ;  
}
```

sizeof( int ) and sizeof( int \* ) not the same:

```
void f1( int n ) ;  
  
    int **a = (int **) malloc( n*sizeof(int) ) ; // instead of sizeof( int * )  
}
```

pointer instead of value referencing:

```
void f1( int *x ) {  
  
    *x-- ; // instead of (*x)-- ;  
}
```

## List of common memory bugs

### pointer arithmetic:

```
void f1( int *x ) {  
  
x += sizeof(int) // instead of x++ ;  
}
```

### memory pointed to disappears:

```
int *f1( int n ) ;  
    int a ;  
  
    return &a ;  
}
```

### access freed area:

```
void f1( int *x ) {  
    int *y ;  
  
    y = (int *)malloc( 20*sizeof(int) ) ;  
    free y ;  
  
    for( i=0;i<20; i++ )  
        x[i] = y[i] ;  
}
```

## Several ways to do it

The algorithm: smart and mindless ways

- optimize loops

  - procedure calls

  - recomputing items which do not change

  - unrolling

  - blocking

The optimizer

Take advantage of architecture: parallelism, caching

Algorithms: time as a function of set size

- linear or polynomial time

  - $n^2$  (sort1.c)

  - $n \log n$  (sort2.c)

  - accuracy desired? (bin packing)

Algorithms: time as a function of granularity (sort3, sort4)

The importance of measurement

- upper, lower bounds

- average behavior

## Speeding up your program

Converting your program from  $N^2$  to  $N \log N$ : compare the two.

For small  $N$  it does not make much difference.

Amdahl's law.

Say  $T_{\text{old}}$  is the total time a program takes. Let's say part of the program takes a fraction  $f$  of the time. Let's speed up that part by a factor of  $k$ . then

$$T_{\text{new}} = (1-f) * T_{\text{old}} + (f * T_{\text{old}}) / k = T_{\text{old}} * ((1-f) + f/k)$$

then

$$T_{\text{old}} / T_{\text{new}} = 1 / ((1-f) + f/k) = S = \text{speedup factor}$$

try  $f = 0.6$ ,  $k = 3$ .  $S = 1.67$

Say  $k$  is very large, then  $S$  is approximately  $1/(1-f)$  and with  $f = 0.6$ ,  $S = 2.5$

## Compiler vs Programmer

Compiler must analyze code to see where to optimize

- reduce memory references
- take redundant code out of loops
- inline functions

Programmer can do things to allow optimization

- loop unrolling
- taking procedure calls out of loops
- reduce the use of functions: inlining
- reduce memory references
- avoid variable aliasing

## Blocks to optimization

### Memory aliasing

```
void func1(int *xp, int *yp)
{
1  *xp += *yp;
2  *xp += *yp;
}
```

after line 1 in func1, i = 20, after line 2, i = 40.

after line 1 in func2, i = 30

```
void func2(int *xp, int *yp)
{
1  *xp += 2 * *yp;
}
```

in both functions, it "looks like" there are two distinct variables but these names are just aliases for the argument.

the compiler could try to optimize func1 to func2 and the compiled code will give different results.

```
int main
{
  int i = 10 ;

  func1( &i, &i ) ;

  i = 10 ;
  func2( &i, &i ) ;

  return 0 ;
}
```

## Blocks to optimization

How about:

```
x = 1000; y = 3000;  
*q = y;      /* 3000 */  
*p = x;      /* 1000 */  
t1 = *q;     /* 1000 or 3000 */
```

if  $p == q$ , result is  $t1 = 1000$  ; if not  $t1 = 3000$

Can this happen with pass by value variables?

## Blocks to optimization

Consider:

```
void swap(int *xp, int *yp)
{
    *xp = *xp + *yp; /* x+y */
    *yp = *xp - *yp; /* x+y-y = x */
    *xp = *xp - *yp; /* x+y-x = y */
}
```

```
int main()
{
    int x,y ;

    x = 10 ;
    y = 20 ;
    swap( &x, &y ) ;

    x = 10 ;
    y = 10 ;
    swap( &x, &y ) ;
```

```
    x = 10 ;
    y = 10 ;
    swap( &x, &x ) ;
```

```
    return 0 ;
}
```

Start swap

```
*xp *yp
10 20
30 20
30 10
20 10
```

Start swap

```
*xp *yp
10 10
20 10
20 10
10 10
```

Start swap

```
*xp *yp
10 10
20 20
0 0
0 0
```

When  $x \neq y$ , (first two cases), everything works normally.  
Even when  $*xp = *yp$ .

But when  $x == y$ , problems.



## Blocks to optimization

Consider when a function operates on global variables.

```
int counter = 0;
```

```
int f()
{
    printf( "in f() counter= %d\n", counter );
    return counter++;
}
```

```
int main()
{
    printf( "result of f()+f()+f()+f() = %d \n", f()+f()+f()+f() );

    counter = 0 ;

    printf( "result of 4*f() = %d \n", 4*f() );

    return 0 ;
}
```

```
in f() counter= 0
in f() counter= 1
in f() counter= 2
in f() counter= 3
result of f()+f()+f()+f() = 6
```

```
in f() counter= 0
result of 4*f() = 0
End value of counter = 1
```

“Optimized” program gives different results.

## Measuring performance

### Cycles per element

How do we measure performance? Stop watch? Program running time as a function of number of input elements. Also may be a function of the distribution of input data (coarseness).

Run the program many times with different number of input elements and data types. Plot a curve of number of elements versus running time. Fit a line to the data using least squares fit (regression) and you arrive at a formula which is

$$\text{run time} = \text{constant} + \text{coefficient} * N$$

where N is the number of input elements. So, the coefficient expresses the rate of increase in run time per additional data element. The constant expresses the overhead to start the program (run time when  $N = 0$ ).

The coefficient is known as the CPE: cycles per element. Its units are run time per element. In all cases, it is relative to the speed of the computer but we think of it as cycles

## Loop unrolling

```
/* Compute prefix sum of vector a */
void psum1(float a[], float p[], long int n)
{
    long int i;
    p[0] = a[0];
    for (i = 1; i < n; i++)
        p[i] = p[i-1] + a[i];
}
```

This function computes the “prefix sum” of an array of elements: a. It is defined as:

$$p[0] = a[0];$$

$$p[i] = p[i-1] + a[i]$$

So,  $p[i]$  = the sum of all  $a[j]$  where  $j \leq i$

```
void psum2(float a[], float p[], long int n)
{
    long int i;
    p[0] = a[0];
    for (i = 1; i < n-1; i+=2) {
        p[i] = p[i-1] + a[i];
        p[i+1] = p[i] + a[i+1];
    }
    /* For odd n, finish remaining element */
    if (i < n)
        p[i] = p[i-1] + a[i];
}
```

This is 1x unrolling

## Loop unrolling

**\*Caveat\*** The increase in the number of lines in the code affects the savings for loop unrolling.

Lets say that it takes  $x$  units of time to execute the line of code in the loop in psum1,  $y$  units to execute the loop overhead and  $z$  time units are used when the loop is unrolled in psum2.

So, the time to execute the loop in psum1 is

$$a = x*n + y*n \quad \text{execute the code plus the loop overhead}$$

an unrolled loop program would take

$$b = z*n/2 + y*n/2 \quad \text{execute the code half as much and the loop overhead half as much}$$

for it to be faster, we want  $b < a$  or

$$z*n/2 + y*n/2 < x*n + y*n$$

This is the same as

$$0 < x*n + y*n/2 - z*n/2$$

dividing by  $n$  it becomes

$$0 < x + y/2 - z/2 = c = \text{difference in run times old - new}$$

## Loop unrolling

So, depending on the relative values of x, y and z, there is a diminishing rate of return!

x	y	z	c
1	1	2	0.5
1	1	3	0
1	1	4	-0.5
1	1	5	-1
2	1	3	1
2	1	4	0.5
2	1	5	0
2	1	6	-0.5

For the first line in the table, we increase the statement executions by 1 and we get a .5 improvement in the run time, by 2 and we get zero. In line 5, we increase it by 1 from 2 to 3, we get an improvement of 1.